Statistical Modeling for Engineers and Scientists

Eric M Reyes

Updated: 04 March 2025

Table of contents

Pr	reface	8
	Unit I: Language and Logic of Inference	9
1	The Statistical Process 1.1 Overview of Drawing Inference	11 12 13 16
2	Case Study: Health Effects of the Deepwater Horizon Oil Spill	17
3	Asking the Right Questions 3.1 Characterizing a Variable	20 20 23
4	Gathering the Evidence (Data Collection) 4.1 What Makes a Sample Reliable	27 27 30 32 34
5	Presenting the Evidence (Summarizing Data) 5.1 Characteristics of a Distribution (Summarizing a Single Variable)	40 40 48
6	Assessing the Evidence (Quantifying the Variability in Estimates) 6.1 Conceptualizing the Sampling Distribution	51 53 55 60 63 66 68
7	Quantifying the Evidence (Rejecting Bad Models)	71
	/ 1 30mg 31mtigtigg	/ * /

	7.2 Assuming the Null Hypothesis	73 76 80 82
8	Using the Tools Together 8.1 Framing the Question (Fundamental Idea I)	83 83 84 84 85 87 88
П	Unit II: Inference on the Overall Mean Response	89
9	Case Study: Birth Weights of Babies	91
11	Model for the Data Generating Process 10.1 General Formulation	
Ш	Unit III: Modeling the Average Response as a Function of a Continuous Predictor	114
13	Case Study: Seismic Activity in Greece	116
14	Myriad of Potential Questions	118
15	Nature of Collecting Multivariable Data	120
16		122 122 125
17	Bulding our Statistical Model 17.1 Statistical Model for A Quantitative Response and Quantitative Predictor	128

	17.2	Estimating the Parameters	131
18			134
	18.1	Correctly Specified Model	
		18.1.1 Interpreting the Parameters	
		18.1.2 Embedding our Question in a Statistical Framework	
		Independent Errors	
		Same Degree of Precision	
	18.4	Specific form of the Error Distribution	
		18.4.1 Modeling the Population	
		18.4.2 The Normal Distribution	141
	18.5	Classical Regression Model	145
	18.6	Imposing the Conditions	146
	18.7	Recap	149
19	Qua	ntifying the Quality of a Model Fit	150
	19.1	Partitioning Variability	150
	19.2	R-squared	154
	19.3	Hypothesis Testing	156
20	Asse	essing the Modeling Conditions	164
	20.1	Residuals	164
	20.2	Assessing the Mean-0 Condition	165
	20.3	Assessing the Independence Condition	167
	20.4	Assessing Homoskedasticity	169
	20.5	Assessing Normality	171
	20.6	General Tips for Assessing Assumptions	174
21	Exte	ending the Regression Model	175
	21.1	Using a Categorical Predictor	175
	21.2	Including Multiple Precitors	178
		21.2.1 General Model Formulation	179
		21.2.2 Interpretation of Parameters	180
		21.2.3 Inference for the General Formulation	182
	21.3	Modifying an Effect	185
22	Putt	ing it All Together	190
	22.1	Graphical Summary	190
		Development of Statistical Model	190
		Assessment of Conditions	
			193

IV	Unit IV: Comparing the Average Response Across Groups	196
23	Case Study: Organic Foods and Superior Morals	198
24	Framing the Question 24.1 General Setting	200 202
25	Study Design25.1 Aspects of a Well-Designed Study	209
26	Presenting the Data	211
27	Building the Statistical Model 27.1 Statistical Model for A Quantitative Response and a Categorical Predictor 27.2 Indicator Variable	
28	Conditions on the Error Term of an ANOVA Model 28.1 Independent Errors	221 221 222 223
29	Quantifying the Evidence29.1 Partitioning Variability	229 230 231 231
30	Assessing the Modeling Conditions in ANOVA 30.1 Residual	240 240 243

31	Using the Tools Together	246
	31.1 Framing the Question (Fundamental Idea I)	246
	31.2 Getting Good Data (Fundamental Idea II)	247
	31.3 Presenting the Data (Fundamental Idea III)	247
	31.4 Quantifying the Variability in the Estimate (Fundamental Idea IV)	248
	31.5 Quantifying the Evidence (Fundamental Idea V) $\dots \dots \dots \dots$	250
	31.6 Conclusion	251
V	Unit V: Comparing the Average Response Across Correlated Groups	252
32	Case Study: Paying a Premium for the Experience	254
33	Framing the Question	256
33	33.1 General Setting	256
	ooil denotal sevening	200
34	Correlated Data	258
	34.1 Blocking	258
35	Presenting Correlated Data	262
36	Modeling Correlated Responses	265
	36.1 Statistical Model for Correlated Responses	265
37	Conditions on the Error Term of the Repeated Measures ANOVA Model	271
	37.1 Conditions on the Error Distribution	
	37.1.1 Correctly Specified Model	
	37.1.2 Independent Errors	
	37.1.3 Same Degree of Precision	
	37.1.4 Specific Form of the Error Distribution	
	37.2 Conditions on the Block Effects	
	37.2.1 Independent Block Effects	
	37.3 Classical Repeated Measures ANOVA Model	
	37.4 Imposing the Conditions	
	37.5 Recap	
38	Quantifying the Evidence	279
	Assessing the Modeling Conditions in Repeated Measures ANOVA	285
<i></i>	39.1 Residual	286
	39.2 Assessing the Independence Condition	
	39.3 Time-Series Plot	
	39.4 Assessing Homoskedasticity	288

	39.5	Assessing Normality	288
	39.6	Probability Plot	289
	39.7	Assessing Whether the Model is Correctly Specified	290
	39.8	General Tips for Assessing Assumptions	291
40	Usin	g the Tools Together 2	293
	40.1	Framing the Question (Fundamental Idea I)	293
	40.2	Getting Good Data (Fundamental Idea II)	294
	40.3	Presenting the Data (Fundamental Idea III)	294
	40.4	Quantifying the Variability in the Estimate (Fundamental Idea IV) 2	294
	40.5	Quantifying the Evidence (Fundamental Idea V)	298
	40.6	Conclusion	299
Re	feren	ces 3	800
Α _Ι	open	dices 3	30 :
Α	App	roaches for Modeling Sampling Distributions 3	801
		Residual Bootstrap	
		A.1.1 Bootstrap for Hypothesis Testing	
	A.2	Wild Bootstrap	
		Classical Theory	
	A.4	Asymptotic Theory	806
В	Mat	hematical Results for Standardized Statistics 3	308
	B.1	Single Mean Response, Same Signal	308
	B.2	Single Mean Response, Two Standardized Statistics	309
	B.3	Least Squares Estimate for Intercept-Only Model	310
	B.4	Class of Hypotheses that are Testable	312
	B.5	Equivalence of Hypotheses with Blocking	13
C	Glos	sary 3	15
	C.1	Five Fundamental Ideas of Inference	15
	C.2	Distributional Quartet	316
	C.3	Models for the Data Generating Process	316
	C.4	Glossary	18

Preface

Engineers and scientists are routinely asked to make decisions based on observed data. That data will be subject to variability; that is, measured characteristics will vary from one observation to the next. Learning to characterize that variability and make decisions in its presence is the idea behind statistics.

This text will introduce statistical concepts in the context of engineering, the physical, and the biological sciences. The text emphasizes statistical literacy (interpretation and clear communication of statistical concepts, methods, and results) and statistical reasoning (defining the need for data to address questions, modeling variability in a process, and choosing the appropriate methodology to address a question of interest). We describe approaches to collecting data, summarizing the information contained within the data, building a model to address a question of interest, using the data to estimate the unknowns in the model, assessing the model, and interpreting the results based on the model.

This text is applied, focusing primarily on knowing when various modeling strategies are appropriate and how to interpret their results. Our aim is to provide a strong foundation in statistical ideas enabling readers to engage with research encountered in their field.

Part I

Unit I: Language and Logic of Inference

Children learn the alphabet before tackling *The Odyssey*. Musicians become proficient in scales before playing in a symphony. And, chefs create world-class culinary experiences because they are experts at working with their ingredients. Similarly, working with statistical models benefits from understanding the language and logic of statistical inference.

This first unit introduces the key components of any analysis — asking well-posed questions, collecting useful data, summarizing your data to tell a story, quantifying the strength of a signal. Once we are familiar with these ingredients, we can begin putting them together to address a range of interesting questions.

1 The Statistical Process

Is driving while texting as dangerous as driving while intoxicated? Is there evidence that my measurement device is calibrated inappropriately? How much force, on average, can our concrete blocks withstand before failing? Regardless of your future career path, you will eventually need to answer a question. The discipline of statistics is about using data to address questions by converting that data into valuable information.



Statistics is the discipline of converting data into information.

It might be natural at this point to ask "do I really need an entire class about answering questions with data? Isn't this simple?" Sometimes, it is simple; other times, it can be far from it. Let's illustrate with the following example from Tintle et al. (2015).

Example 1.1 (Organ Donation). Even though organ donations save lives, recruiting organ donors is difficult. Interestingly, surveys show that about 85% of Americans approve of organ donation in principle and many states offer a simple organ donor registration process when people apply for a driver's license. However, only about 38% of licensed drivers in the United States are registered to be organ donors. Some people prefer not to make an active decision about organ donation because the topic can be unpleasant to think about. But perhaps phrasing the question differently could affect a person's willingness to become a donor.

Johnson and Goldstein (2003) recruited 161 participants for a study, published in the journal *Science*, to address the question of organ donor recruitment. The participants were asked to imagine they had moved to a new state and were applying for a driver's license. As part of this application, the participants were to decide whether or not to become an organ donor. Participants were presented with one of three different default choices:

- Some of the participants were forced to make a choice of becoming a donor or not, without being given a default option (the "neutral" group).
- Other participants were told that the default option was not to be a donor but that they could choose to become a donor if they wished (the "opt-in" group).
- The remaining participants were told that the default option was to be a donor but that they could choose not to become a donor if they wished (the "opt-out" group).

The study found that 79% of those in the neutral group, 42% of those in the opt-in group, and 82.0% of those in the opt-out group agreed to become donors.

The results of the study are presented in Figure 1.1. It seems obvious that using the "optin" strategy results in fewer people agreeing to organ donation. However, does the "opt-out" strategy, in which people are by default declared organ donors, result in more people agreeing to organ donation compared to the "neutral" strategy? On the one hand, a higher percentage did agree to organ donation under the "opt-out" (82% compared to 79%). However, since this study involved only a subset of Americans, is this enough evidence to claim the "opt-out" strategy is really superior compared to the "neutral" strategy in the broader population? The discipline of statistics provides a framework for addressing such ambiguity.

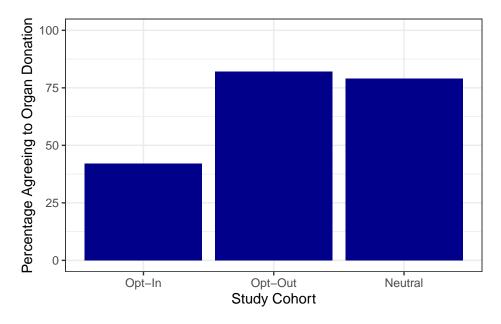


Figure 1.1: Summary of the responses for the Organ Donation Study described in Example 1.1.

1.1 Overview of Drawing Inference

Let's begin by taking a step back and considering the big picture of how data is turned into information. Every research question we pose, at its heart, is trying to characterize a **population**, the group of subjects of ultimate interest.

Definition 1.1 (Population). The collection of subjects we would like to say something about.

In the Organ Donation study (Example 1.1), the researchers would like to say something about Americans who are of the age to consent to organ donation; in particular, they would like to

quantify how likely it is that someone from this group agrees to organ donation. Therefore, the population is all Americans who are of the age to consent to organ donation.

In general, the subjects (or units of observation) in a population need not be people; in some studies, the population could be a collection of screws, cell phones, sheet metal...whatever characterizes the objects from which we would *like to* obtain measurements. We use the phrase "like to" because in reality it is often impossible (or impractical) to observe the entire population. Instead, we make observations on a subset of the population; this smaller group is known as the **sample**.

Definition 1.2 (Sample). The collection of subjects for which we actually obtain measurements (data).

i Note

Some readers may associate "subjects" with people; to avoid this confusion, you may prefer "unit of observation" to subject. In this text, we use subject to mean any unit on which observations could be taken.

For each subject within the sample, we obtain a collection of measurements forming our set of data. The goal of statistical modeling is to use the sample (the group we actually observe) to say something about the population of interest (the group we wish we had observed); this process is known as **statistical inference** (illustrated in Figure 1.2).

Definition 1.3 (Statistical Inference). The process of using a sample to characterize some aspect of the underlying population.

1.2 Anatomy of a Dataset

Once we have our sample, we take measurements on each of the subjects within this sample. These measurements form the data. When we hear the word "data," most of us envision a large spreadsheet. In reality, data can take on many forms — spreadsheets, images, text files, unstructured text from a social media feed, etc. Regardless of the form, all datasets contain information for each subject in the sample; this information, the various measurements, are called **variables**.

Definition 1.4 (Variable). A measurement, or category, describing some aspect of the subject.

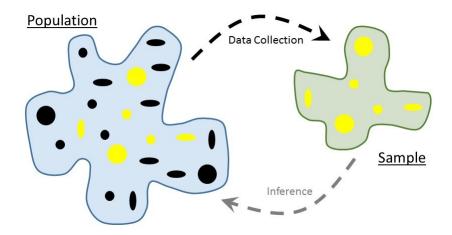


Figure 1.2: Illustration of the statistical process, using a sample to characterize some aspect of the underlying population.

Variables come in one of two flavors. Categorical variables are those which denote a grouping to which the subject belongs. Examples include marital status, manufacturer, and experimental treatment group. Numeric variables are those which take on values for which ordinary arithmetic (e.g., addition and multiplication) makes sense. Examples include height, age of a product, and diameter. Note that sometimes numeric values are used to represent the levels of a categorical variable in a dataset; for example, 0 may indicate "No" and 1 may indicate "Yes" for a variable capturing whether a person is a registered organ donor. Therefore, just because a variable has a numeric value does not make it a numeric variable; the key here is that numeric variables are those for which arithmetic makes sense.

Definition 1.5 (Categorical Variable). Also called a "qualitative variable," a measurement on a subject which denotes a grouping or categorization.

Definition 1.6 (Numeric Variable). Also called a "quantitative variable," a measurement on a subject which takes on a numeric value *and* for which ordinary arithmetic makes sense.

While it may be natural to think of a dataset as a spreadsheet, not all spreadsheets are created equal.

Characteristics of Well-Structured Data

A well-structured dataset should adhere to the following characteristics:

• Each column contains a unique variable.

Table 1.1: Example of a common data structure which does not correspond to the characteristics of well-structured data we recommend. The data is from a hypothetical study comparing battery lifetimes (hours).

Brand A	Brand B
8.3	8.4
5.1	8.6
3.3	3.8
5.3	4.1
5.7	4.5
	4.0

- Each record (row in the dataset) corresponds to a different observation of the variables.
- If you have multiple datasets, they should include a column in the table that allows them to be linked (subject identifier).

These characteristics ensure the data is properly formatted for an analysis. Even unstructured data such as images or text files must be processed prior to performing a statistical analysis.



Warning

We note the above description eliminates a common method of storing data in engineering and scientific disciplines — storing each sample in a different column.

To illustrate the above description, suppose we conduct a study comparing the lifetime (in hours) of two brands of batteries. We measure the lifetime of five batteries of Brand A and six of Brand B. It is common to see a dataset like that in Table 1.1; the problem here is that the first record of the dataset contains information on two different units of observation. We have the lifetime from a battery of Brand A in the same row as the lifetime from a battery of Brand B. This violates the second characteristic of datasets described above.

In order to adhere to the characteristics of well-structured data outlined above, we can reformat the data in Table 1.1 to that shown in Table 1.2. Here, each record represents a unique observation and each column is a different variable. We have also added a unique identifier.

It may take some time to get used to storing data in this format, but it makes analysis easier and avoids time spent managing the data later.

Table 1.2: Example of a well-structured dataset. The data is from a hypothetical study comparing battery lifetimes (hours).

Battery	Brand	Lifetime
1	A	8.3
2	A	5.1
3	A	3.3
4	A	5.3
5	A	5.7
6	В	8.4
7	В	8.6
8	В	3.8
9	В	4.1
10	В	4.5
11	В	4.0

1.3 A Note on Codebooks

A dataset on its own is meaningless if you cannot understand what the values represent. *Before* you access a dataset, you should always review any available **codebooks**.

Definition 1.7 (Codebook). Also called a "data dictionary," these provide complete information regarding the variables contained within a dataset.

Some codebooks are excellent, with detailed descriptions of how the variables were collected alongside appropriate units for the measurements. Other codebooks give only an indication of what each variable represents. Whenever you are working with previously collected data, reviewing a codebook is the first step; and, you should be prepared to revisit the codebook often throughout an analysis. When you are collecting your own dataset, constructing a codebook is essential for others to make use of your data.

2 Case Study: Health Effects of the Deepwater Horizon Oil Spill

On the evening of April 20, 2010, the *Deepwater Horizon*, an oil drilling platform positioned off the coast of Louisiana, was engulfed in flames as the result of an explosion. The drilling rig, leased and operated by BP, had been tasked with drilling an oil well in water nearly 5000 feet deep. Eleven personnel were killed in the explosion. The following screenshot is from the initial coverage by the *New York Times*¹:

The incident is considered the worst oil spill in US history, creating an environmental disaster along the Gulf Coast. In addition to studying the effects on the local environment, researchers have undertaken studies to examine the short and long-term health effects caused by the incident. As an example, we might ask whether volunteers who were directly exposed to oil, such as when cleaning wildlife, are at higher risk of respiratory irritation compared to those volunteers who were helping with administrative tasks (and therefore were not directly exposed to oil). An article appearing in *The New England Journal of Medicine* (Goldstein, Osofsky, and Lichtveld 2011) reported the results from a health symptom survey performed in the Spring and Summer of 2010 by the National Institute for Occupational Safety and Health. Of 54 volunteers assigned to wildlife cleaning and rehabilitation, 15 reported experiencing "nose irritation, sinus problems, or sore throat." Of 103 volunteers who had no exposure to oil, dispersants, cleaners, or other chemicals, 16 reported experiencing "nose irritation, sinus problems, or sore throat."

While a larger fraction of volunteers cleaning wildlife in the study reported respiratory symptoms compared to those who were not directly exposed to irritants, would we expect similar results if we were able to interview all volunteers? What about during a future oil spill? Is there evidence that more than 1 in 5 volunteers who clean wildlife will develop respiratory symptoms? What is a reasonable value for the increased risk of respiratory symptoms for those volunteers with direct exposure compared to those without?

In the first part of this text, we use this motivating example as the context for discussing how research questions should be framed, methods for data collection, summarizing and presenting data clearly, quantifying the variability in an estimate, and quantifying the degree to which the data disagrees with a proposed model. We capture these ideas in what we call the *Five*

¹http://www.nytimes.com/2010/04/22/us/22rig.html?rref=collection%2Ftimestopic%2FOil% 20Spills&action=click&contentCollection=timestopics®ion=stream&module=stream_unit&version=search&contentPlacement=1&pgtype=collection

Search Continues After Oil Rig Blast

By CAMPBELL ROBERTSON APRIL 21, 2010





The rig burned Wednesday about 50 miles southeast of Venice, La. Firefighting efforts were causing it to take on water and list. Gerald Herbert/Associated Press

NEW ORLEANS — An explosion on an \underline{oil} drilling rig off the coast of southeast Louisiana left at least 3 people critically injured and 11 others missing as of Wednesday night.

Figure 2.1: New York Times coverage of the Deepwater Horizon oil spill.

Fundamental Ideas of Inference. We will also see that any statistical analysis moves between the components of what we call the Distributional Quartet. These two frameworks allow us to describe the language and logic of inference, serving as a foundation for the statistical thinking and reasoning needed to address more complex questions encountered later in the text.

3 Asking the Right Questions

The discipline of statistics is about turning data into information in order to address some question. While there may be no such thing as a stupid question, there are ill-posed questions — those which cannot be answered as stated. Consider the Deepwater Horizon Case Study (Chapter 2). It might seem natural to ask "if a volunteer cleans wildlife, will she develop adverse respiratory symptoms?" Let's consider the data. Of the 54 volunteers assigned to wildlife cleaning and rehabilitation, 15 reported experiencing adverse respiratory symptoms ("nose irritation, sinus problems, or sore throat"); while some volunteers developed symptoms, others did not. It seems the answer to our question is then "it depends" or "maybe." This is an example of an *ill-posed question*. Such questions exist because of variability, the fact that every subject in the population does not behave in exactly the same way; that is, the value of a variable potentially differs from one observation to the next. In our example not every volunteer had the same reaction when directly exposed to oil.

It is variability that creates a need for the discipline of statistics; in fact, you could think of statistics as the study and characterization of variability. We must therefore learn to ask the right questions — those which can be answered in the presence of variability.

Definition 3.1 (Variability). The notion that measurements differ from one observation to another.



Big Idea

The presence of variability makes some questions ill-posed; statistics concerns itself with how to address questions in the presence of variability.

3.1 Characterizing a Variable

Recall that the goal of statistical inference is to say something about the population; as a result, any question we ask should then be about on this larger group. The first step to constructing a well-posed question is then to identify the population of interest for the study. For the Deepwater Horizon Case Study, it is unlikely that we are only interested in these 54 observed volunteers assigned to wildlife cleaning. In reality, we probably want to say something about volunteers for any oil spill. The 54 volunteers in our dataset form the sample, a subset from all volunteers who clean wildlife following an oil spill. Our population of interest is comprised of all volunteers who clean wildlife following an oil spill.

Note

When identifying the population of interest for a research question you have, be specific! Suppose you are trying to estimate the average height of trees. Are you really interested in *all* trees? Or, are you interested in Maple trees within the city limits of Terre Haute, Indiana?

Since we expect that the reaction to oil exposure — the primary variable of interest for this study, sometimes called the **response** — to vary from one individual to another, we cannot ask a question about the *value* of the reaction (whether they experienced symptoms or not). Instead, we want to characterize the **distribution** of the response.

Definition 3.2 (Response). The primary variable of interest within a study. This is the variable you would either like to explain or estimate.

Definition 3.3 (Distribution). The pattern of variability corresponding to a set of values.

Notice that in this case, the response is a categorical variable; describing the distribution of such a variable is equivalent to describing how individuals are divided among the possible groups. With a finite number of observations, we could present the number of observations, the **frequency**, within each group. For example, of the 54 volunteers, 15 experienced adverse symptoms and 39 did not. This works well within the sample; however, as our population is infinitely large (all volunteers cleaning wildlife following an oil spill), reporting the frequencies is not appropriate. In this case, we report the fraction of observations, the **relative frequency**, falling within each group; this helps convey information about the distribution of this variable. That is, the relative frequencies give us a sense of which values of the variable are more or less common in the sample.

Definition 3.4 (Frequency). The number of observations in a sample falling into a particular group (level) defined by a categorical variable.

Definition 3.5 (Relative Frequency). Also called the "proportion," the fraction of observations falling into a particular group (level) of a categorical variable.

Numeric quantities, like the proportion, which summarize the distribution of a variable within the population are known as **parameters**.

Definition 3.6 (Parameter). Numeric quantity which summarizes the distribution of a variable within the *population* of interest. Generally denoted by Greek letters in statistical formulas.

While the value of a variable may vary across the population, the parameter is a single fixed constant which summarizes the variable for that population. For example, the grade received on an exam varies from one student to another in a class; but, the average exam grade is a fixed number which summarizes the class as a whole. Well-posed questions can be constructed if we limit ourselves to questions about the parameter. The second step in constructing well-posed questions is then to identify the parameter of interest.

The questions we ask generally fall into one of two categories:

- Estimation: what proportion of volunteers who clean wildlife following an oil spill will experience adverse respiratory symptoms?
- Hypothesis Testing: is it reasonable no more than 1 in 5 volunteers who clean wildlife following an oil spill will experience adverse respiratory symptoms; or, is there evidence more than 1 in 5 volunteers who clean wildlife following an oil spill will experience adverse respiratory symptoms?

Definition 3.7 (Estimation). Using the sample to approximate the value of a parameter from the underlying population.

Definition 3.8 (Hypothesis Testing). Using a sample to determine if the data is consistent with a working theory or if there is evidence to suggest the data is not consistent with the theory.

Since we do not get to observe the population (we only see the sample), we cannot observe the value of the parameter. That is, we will never know the true proportion of volunteers who experience symptoms. However, we can determine what the data suggests about the population (that is what inference is all about).



Big Idea

Parameters are unknown values and can never, in general, be known.

It turns out, the vast majority of research questions can be framed in terms of a parameter. This is the first of what we consider the Five Fundamental Ideas of Inference.

Fundamental Idea I

A research question can often be framed in terms of a parameter that characterizes the population. Framing the question should then guide our analysis.

We now have a way of describing a well-posed question, a question which can be addressed using data. Well posed questions are about the population and can be framed in terms of a parameter which summarizes that population. We now describe how these questions are typically framed.

3.2 Framing the Question

In engineering and scientific applications, many questions fall under the second category of **hypothesis testing**, which is a form of model comparison in which data is collected to help the researcher choose between two competing theories for the parameter of interest. In this section, we consider the terminology surrounding specifying such questions.

For the Deepwater Horizon Case Study suppose we are interested in addressing the following question:

Is there evidence that more than 1 in 5 volunteers who clean wildlife following an oil spill will develop adverse respiratory symptoms?

The question itself is about the population (all volunteers assigned to clean wildlife following an oil spill) and is centered on a parameter (the proportion who develop adverse respiratory symptoms). That is, this is a well-posed question that can be answered with appropriate data. The overall process for addressing these types of questions is similar to conducting a trial in a court of law. In the United States, a trial has the following essential steps:

- 1. Assume the defendant is innocent.
- 2. Present evidence to establish guilt, to the contrary of innocence (prosecution's responsibility).
- 3. Consider the weight of the evidence presented (jury's responsibility).
- 4. Make a decision. If the evidence is "beyond a reasonable doubt," the jury declares the defendant guilty; otherwise, the jury declares the defendant not guilty.

The process of conducting a hypothesis test has similar essential steps:

- 1. Assume the opposite of what we want the data to show (develop a working theory).
- 2. Gather data and compare it to the proposed model from step (1).
- 3. Quantify the likelihood of our data from step (2) under the proposed model.
- 4. If the likelihood is small, conclude the data is not consistent with the working model (there is evidence for what we want to show); otherwise, conclude the data is consistent with the working model (there is no evidence for what we want to show).

Notice that a trial focuses not on proving guilt but on disproving innocence; similarly, in statistics, we are able to establish evidence *against* a specified theory. This is one of several subtle points in hypothesis testing. We will discuss these subtleties at various points throughout the text and revisit the overall concepts often. Here, we focus solely on that first step — developing a working theory that we want to *disprove*.

Note

This process may seem counter-intuitive; it is natural to ask "why can't we prove guilt directly?" However, when you disprove one statement, you are proving that statement's opposite — a technique known in mathematics as "proof by contradiction." So, our approach to proving a statement is to disprove all other possibilities. It is similar to the technique of the fictional detective Sherlock Holmes (Doyle 1890, pg. 92): "Eliminate all other factors, and the one which remains must be the truth."

Consider the above question for the Deepwater Horizon Case Study. We want to find evidence that the proportion experiencing adverse symptoms exceeds 0.20 (1 in 5). Therefore, we would like to disprove (or provide evidence against) the statement that the proportion experiencing adverse symptoms is no more than 0.20. This statement that we would like to disprove is known as the null hypothesis; the opposite of this statement, called the alternative hypothesis, captures what we as the researchers would like to establish.

Definition 3.9 (Null Hypothesis). The statement (or theory) about the parameter that we would like to *disprove*. This is denoted H_0 , read "H-naught" or "H-zero".

Definition 3.10 (Alternative Hypothesis). The statement (or theory) about the parameter capturing what we would like to provide evidence for; this is the opposite of the null hypothesis. This is denoted H_1 or H_a , read "H-one" and "H-A" respectively.

For the Deepwater Horizon Case Study, we write:

 H_0 : The proportion of volunteers assigned to clean wildlife following an oil spill who experience adverse respiratory symptoms is no more than 0.20.

 H_1 : The proportion of volunteers assigned to clean wildlife following an oil spill who experience adverse respiratory symptoms exceeds 0.20.

Each hypothesis is a well-posed statement (about a parameter characterizing the entire population), and the two statements are exactly opposite of one another meaning only one can be a true statement.

Note

When framing your questions, be sure your null hypothesis and alternative hypothesis are exact opposites of one another, and ensure the "equality" component *always* goes in the null hypothesis.

We can now collect data and determine if it is *consistent* with the null hypothesis (a statement similar to "not guilty") or if the data provides *evidence* against the null hypothesis and in favor of the alternative (a statement similar to "guilty").

A Consistent vs. Evidence

The term "consistent" and "reasonable" will be used interchangeably throughout the text; however, these terms differ substantially from the term "evidence." The data is said to be consistent with a statement if the data is aligned with that statement. We have evidence for a statement when the data is aligned with that statement and it is not aligned with the opposite of the statement (when we can disprove something). We will see this develop more in Chapters 6 and 7.

Often these statements are written in a bit more of a mathematical structure in which a Greek letter is used to represent the parameter of interest. For example, we might write

Let θ represent the proportion of volunteers (assigned to clean wildlife following an oil spill) who experience adverse respiratory symptoms.

 $H_0: \theta \le 0.20$ $H_1: \theta > 0.20$

In the above statements, θ represents the parameter of interest; the value 0.20 is known as the null value.

Definition 3.11 (Null Value). The value associated with the equality component of the null hypothesis; it forms the threshold or boundary between the hypotheses. Note: not all questions of interest require a null value be specified.



🥊 Big Idea

Hypothesis testing is a form of statistical inference in which we quantify the evidence against a working theory (captured by the null hypothesis). We essentially argue that the data supports the alternative if it is not consistent with the working theory.

This section has focused on developing the null and alternative hypothesis when our question of interest is best characterized as one of comparing models or evaluating a particular statement. If our goal is estimation, a null and alternative hypothesis are not applicable. For example, we might have the following goal:

Estimate the proportion of volunteers (assigned to clean wildlife following an oil spill) who experience adverse respiratory symptoms.

In this version of our research "question" there is no statement which needs to be evaluated. We are interested in estimation, not hypothesis testing and thus there is no corresponding null and alternative hypothesis.

i Process for Framing a Question

In order to frame a research question, consider the following steps:

- 1. Identify the population of interest.
- 2. Identify the parameter(s) of interest.
- 3. Determine if you are interested in estimating the parameter(s) or quantifying the evidence against some working theory.
- 4. If you are interested in testing a working theory, make the null hypothesis the working theory and the alternative hypothesis the exact opposite statement (capturing what you want to provide evidence for).

4 Gathering the Evidence (Data Collection)

Consider again the goal of statistical inference — to use a sample as a snapshot to say something about the underlying population (see Figure 4.1). This generally provokes unease in people, leading to a distrust of statistical results. In this section we attack that distrust head on.

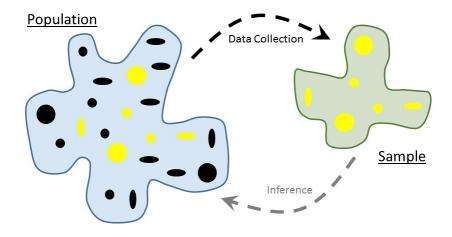


Figure 4.1: Illustration of the statistical process (reprint of Figure 1.2).

4.1 What Makes a Sample Reliable

If we are going to have some amount of faith in the statistical results we produce, we must have data in which we can place our trust. *The Treachery of Images* (Figure 4.2) is a canvas painting depicting a pipe, below which the artist wrote the French phrase "This is not a pipe." Regarding the painting, the artist said

The famous pipe. How people reproached me for it! And yet, could you stuff my pipe? No, it's just a representation, is it not? So if I had written on my picture "This is a pipe," I'd have been lying!



Figure 4.2: The Treachery of Images by René Magritte.

Just as a painting is a representation of the object it depicts, so a sample should be a representation of the population under study. This is the primary requirement if we are to rely on the resulting data.



Big Idea

In order for a statistical analysis to be reliable, the sample must be representative of the population under study.

We need to be careful to not get carried away in our expectations. What constitutes "representative" really depends on the question, just as an artist chooses their depiction based on how they want to represent the object. Let's consider the following example.

Example 4.1 (School Debt). In addition to a degree, college graduates also tend to leave with a large amount of debt due to college loans. In 2012, a graduate with a student loan had an average debt of \$29,400; for graduates from private non-profit institutions, the average debt was $$32,300^{1}$.

Suppose we are interested in determining the average amount of debt in student loans carried by a graduating senior from Rose-Hulman Institute of Technology, a small private non-profit

¹http://ticas.org/sites/default/files/pub_files/Debt_Facts_and_Sources.pdf

engineering school. There are many faculty at Rose-Hulman who choose to send their children to the institute. Suppose we were to ask 25 such faculty members who have a child that attended the institute to report the amount of student loans their children carried upon graduation from Rose-Hulman. Further, suppose we compile the responses and compute the average amount of debt. Using the data, we might report that based on our study, there is significant evidence the average debt carried by a graduate of Rose-Hulman is far below the \$32,300 reported above (great news for this year's graduating class)!

Why should we be hesitant to trust the results from our study?

Many objections to statistical results stem from a distrust of whether the data (the sample) is really representative of the population of interest. Rose-Hulman, like many other universities, has a policy that the children of faculty may attend their university (assuming admittance) tuition-free. We would therefore expect their children to carry much less debt than the typical graduating senior. There is a mismatch between the group we would like to study and the data we have collected.

This example provides a nice backdrop for discussing what it means to be representative. First, let's define our population; in this case, we are interested in graduating seniors from Rose-Hulman. The variable of interest is the amount of debt carried in student loans; the parameter of interest is then the *average* amount of debt in student loans carried by graduating seniors of Rose-Hulman. However, the sample consists of only graduating seniors of Rose-Hulman who have a parent employed by the institute.

With regard to grade point average, the students in our sample are probably similar to all graduating seniors; the starting salary of the students in our sample is probably similar to all graduating seniors; the fraction of mechanical engineering majors versus math majors is probably similar. So, in many regards the sample is representative of the population; however, it fails to be representative with regard to the variable of interest. This is our concern. The amount of debt carried by students in our sample is not representative of that debt carried by all graduating seniors from the university.

Note

When thinking about whether a sample is representative, focus your attention to the characteristics specific to your research question or with regard to how you intend to generalize the results.

Does that mean the sample we collected in Example 4.1 is useless? Yes and no. The sample collected cannot be used to answer our initial question of interest since it is not representative of our population. No statistical method can fix bad data; statistics adheres to the "garbage-in, garbage-out" phenomena. If the data is bad, no analysis will undo that. However, while the sample cannot be used to answer our initial question, it could be used to address a different question:

What is the average amount of debt in student loans carried by graduating seniors from Rose-Hulman whose parent is a faculty member at the institute?

For this revised question, the sample may indeed be representative. If we are working with previously collected data, we must consider the population to which our results will generalize. That is, for what population is the given sample representative? If we are collecting our data, we need to be sure we collect data in such a way that the data is representative of our target population. Let's first look at what *not* to do.

4.2 Poor Methods of Data Collection

Example 4.1 is an example of a "convenience sample," when the subjects in the sample are chosen simply due to ease of collection. Examples include surveying students only in your sorority when you are interested in all students who are part of a sorority on campus; taking soil samples from only your city when you are interested in the soil for the entire state; and, obtaining measurements from only one brand of phone, because it was the only one you could afford on your budget, when you are interested in studying all cell phones on the market. A convenience sample is unlikely to be representative if there is a relationship between the ease of collection and the variable under study. This was true in the School Debt example; the relationship of a student to a faculty member, which is what increased the ease of collection, was directly related to the amount of debt they carried. As a result, the resulting sample was not representative of the population.

When conducting a survey with human subjects, it is common to only illicit responses from volunteers. Such "volunteer samples" tend to draw in those with extreme opinions. Consider product ratings on Amazon. Individual ratings tend to cluster around 5's and 1's. This is because those customers who take time to submit a review (which is voluntary) tend to be those who are really thrilled with their product (and want to encourage others to purchase it) and those who are really disappointed with their purchase (and want to encourage others to avoid it). Such surveys often fail to capture those individuals in the population who have "middle of the road" opinions.

We could not possibly name all the poor methods for collecting a sample; but, poor methods all share something in common — it is much more likely the resulting sample is not representative. Failing to be representative results in **biased** estimates of the parameter.

Definition 4.1 (Bias). A set of measurements is said to be biased if they are *consistently* too high (or too low). Similarly, an estimate of a parameter is said to be biased if it is *consistently* too high (or too low).

To illustrate the concept of bias, consider shooting at a target as in Figure 4.3. We can consider the center of our target to be the parameter we would like to estimate within the population;

in this case, some measure of center. The values in our sample (the strikes on the target) will vary around the parameter; while we do not expect any one value to hit the target precisely, a "representative" sample is one in which the values tend to be clustered about the parameter (unbiased). When the sample is not representative, the values in the sample tend to cluster off the mark (biased). Notice that to be unbiased, it may be that not a single value in the sample is perfect, but aggregated together, they point in the right direction. So, bias is not about an individual measurement being an "outlier," (more on those in Chapter 5) but about consistently shooting in the wrong direction.

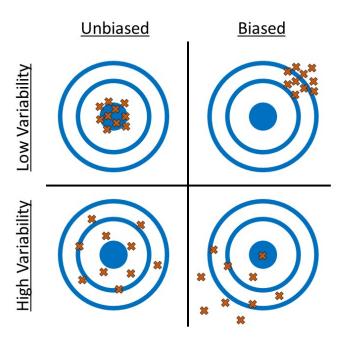


Figure 4.3: Illustration of bias and precision.

Accuracy vs. Precision

There is a difference between accuracy and precision. Generally, accuracy refers to location (and therefore relates to bias); we say a process is accurate when it is unbiased. Precision refers to the variability; data which is more precise has less variability.

Big Idea

Biased results are typically due to poor sampling methods that result in a sample which is not representative of the population of interest.

The catch (there is always a catch) is that we will never *know* with certainty if a sample is actually representative or not. In practice, we critically examine the method in which the sample was collected, and we use summaries of the sample to make educated decisions on whether to generalize the results. Better, however, is to employ methods of data collection that help to minimize the bias in the sample.

4.3 Preferred Methods of Sampling

No method guarantees a perfectly representative sample; but, we can take measures to reduce or eliminate bias. A useful strategy is to employ *randomization*. This is summarized in our second Fundamental Idea.

Fundamental Idea II

If data is to be useful for making conclusions about the population, a process referred to as drawing inference, proper data collection is crucial. Randomization can play an important role ensuring a sample is representative and that inferential conclusions are appropriate.

Consider the School Debt example (Example 4.1) again. Suppose instead of the data collection strategy described there, we had done the following:

We constructed a list of all graduating seniors from the institute. We placed the name of each student on an index card; then, we thoroughly shuffled the cards and chose the top 25 cards. For these 25 individuals, we recorded the amount of debt in student loans each carried.

This essentially describes using a lottery to select the sample. This popular method is known as taking a **simple random sample**. By conducting a lottery, we make it very unlikely that our sample consists of only students with a very small amount of student debt (as occurred when we used a convenience sample).

Definition 4.2 (Simple Random Sample). Often abbreviated SRS, this is a sample of size n such that *every* collection of size n is equally likely to be the resulting sample. This is equivalent to a lottery.

Note

It is convention to use n to represent the sample size.

The primary benefit of a simple random sample is that it removes bias. More specifically, the process of simple random sampling is unbiased; that is, this process does not produce values which are *consistently* too high or low.

There are situations in which a simple random sample does not suffice. Again, consider our School Debt example. The Rose-Hulman student body is predominantly domestic, with only about 3% of the student body being international students. But, suppose we are interested in comparing the average debt carried between international and domestic students. It is very likely, by chance alone, that in a simple random sample of 25 students none will be international. Instead of a simple random sample, we might consider taking a sample of 13 domestic students and a sample of 12 international students; this is an example of a stratified random sample. This approach is useful when there is a natural grouping of interest within the population.

Definition 4.3 (Stratified Random Sample). A sample in which the population is first divided into groups, or strata, based on a characteristic of interest; a simple random sample is then taken within each group.



Warning

Note that a stratified random sample essentially results in a representative sample within each strata. However, the combined sample may not be representative of the population. If there is interest in using the sample in its entirety, instead of comparing the strata in some way, advanced statistical methodology is required. See texts on analyzing "complex survey design" for a more thorough discussion. Our text will not consider such cases.

There are countless sampling techniques used in practice. The two described above can be very useful starting points for developing a custom method suitable for a particular application. Their benefit stems from their use of randomization as it limits researcher influence on the composition of the sample and therefore minimizes bias.

This section is entitled "Preferred Methods" because while these methods are ideal, they are not always practical. Consider the Deepwater Horizon Case Study described in Chapter 2; conceptually, we can take a simple random sample of the volunteers for our study. However, as with any study involving human subjects, researchers would be required to obtain consent from each subject in the study. That is, any individual has the right to refuse to participate in the study. Therefore, it is unlikely that a simple random sample as described above could be obtained. While random selection is a nice tool, the goal is a sample which is representative of the population. While random sampling is helpful for accomplishing this, we may need to appeal to the composition of the sample itself to justify its use. Based on the characteristics of those willing to participate in the study, do we feel the study participants form a representative group of all volunteers? That is the essential question. This is often why studies report a table summarizing participant demographics such as age, gender, etc. It is also why it is extremely

important for researchers to describe how observations were obtained so that readers may make the judgement for themselves whether the sample is representative.

4.4 Two Types of Studies

Thinking about how the data was collected helps us determine how the results generalize beyond the sample itself (to what population the results apply). When our question of interest is about the relationship between two variables (as most questions are), we must also carefully consider the study design. Too often separated from the statistical analysis that follows, keeping the study design in mind should guide the analysis as well as inform us about the conclusions we can draw. While we will discuss study design more fully in Chapter 25, we introduce some of the concepts here.

In order to illustrate how study design can impact the results, consider the following example.

Example 4.2 (Kangaroo Care). At birth, infants have low levels of Vitamin K, a vitamin needed in order to form blood clots. Though rare, without the ability for her blood to clot, an infant could develop a serious bleed. In order to prevent this, the American Academy of Pediatrics recommends that all infants be given a Vitamin K shot shortly after birth in order to raise Vitamin K levels. As with any shot, there is typically discomfort to the infant, which can be very discomforting to new parents.

Kangaroo Care is a method of holding a baby which emphasizes skin-to-skin contact. The child, who is dressed only in a diaper, is placed upright on the parent's bare chest; a light blanket is draped over the child. Suppose we are interested in determining if utilizing the method while giving the child a Vitamin K shot reduces the discomfort in the infant, as measured by the total amount of time the child cries following the shot.

Within this context, contrast the following two potential study designs:

- (A) We allow the attending nurse to determine whether Kangaroo Care is initiated prior to giving the Vitamin K shot. Following the shot, we record the total time (in seconds) the child cries.
- (B) We flip a coin. If it comes up heads, the nurse should have the parents implement Kangaroo Care prior to giving the Vitamin K shot; if it comes up tails, the nurse should give the Vitamin K shot without implementing Kangaroo Care. Following the shot, we record the total time (in seconds) the child cries.

Note, in both study designs (A) and (B), we only consider term births which have no complications to avoid situations that might alter the timing of the Vitamin K shot or the ability to implement Kangaroo Care.

Note that there are some similarities in the two study designs:

- The underlying population is the same for both designs: infants born at term with no complications.
- There are two groups being compared in both designs: the "Kangaroo Care" group and the "no Kangaroo Care" group.
- The response (variable of interest) is the same in both designs: the time (in seconds) the infant cries.
- There is action taken by the researcher in both designs: a Vitamin K shot is given to the child.

There is one prominent difference between the two study designs:

• For design (A), the choice of Kangaroo Care is left up to the nurse (self-selected); for design (B), the choice of Kangaroo is *assigned* to the nurse by the researcher, and this selection is made *at random*.

Design (A) is an example of an **observational study**; design (B) is a an example of a **controlled experiment**.

Definition 4.4 (Observational Study). A study in which each subject "self-selects" into one of groups being compared in the study. The phrase "self-selects" is used very loosely here and can include studies for which the groups are defined by an inherent characteristic or are chosen haphazardly.

Definition 4.5 (Controlled Experiment). A study in which each subject is *randomly* assigned to one of the groups being compared in the study.

It is common to think that anytime the environment is "controlled" by the researcher that a controlled experiment is taking place, but the defining characteristic is the random assignment to groups (sometimes referred to as the *factor* under study or *treatment* groups). In the example above, both study designs involved a controlled setting (the delivery room of a hospital) in which trained staff (the nurse) deliver the shot. However, only design (B) is a controlled experiment because the researchers randomly determined which treatment the infant would receive.

To understand the impact of random allocation, suppose that we had conducted a study using design (A); further, the results of the study suggest that those infants who were given a shot while using Kangaroo Care cried for a shorter time period, on average. Can we conclude that it was the Kangaroo Care that led to the shorter crying time? Maybe. Consider the following two potential explanations for the resulting data:

(1) Kangaroo Care is very effective; as a result, those children who are given Kangaroo Care cry for less time, on average, following the Vitamin K shot.

(2) It turns out that those nurses who choose to implement Kangaroo Care (remember, they have a choice under design (A) whether they implement the method) are also the nurses with a gentler bedside manner. Therefore, these nurses tend to be very gentle when giving the Vitamin K shot whereas the nurses who choose not to implement Kangaroo Care tend to jab the needle when giving the shot. The reduced crying time is not a result of the Kangaroo Care but the manner in which the shot was given.

The problem is that we are unable to determine which of the explanations is correct under study design (A). Given the data we have collected, we are unable to tease out the effect of the Kangaroo Care from that of the nurse's bedside manner. As a result, we are able to say we observed an association between the use of Kangaroo Care and reduced crying time, but we are unable to conclude that Kangaroo Care caused a reduction in the crying time (that is, the reduced crying time may be due to something else, like the bedside manner of the nurse). In this hypothetical scenario, the nurse's bedside manner is called a **confounder**.

Definition 4.6 (Confounding). When the effect of a variable on the response is misrepresented due to the presence of a third, potentially unobserved, variable known as a confounder.

Note

While both result in estimates we may not trust, confounding is not equivalent to a biased sample.

Confounders can mask the relationship between the factor under study and the response. Did you know there is a documented association between ice cream sales and the risk of shark attacks? As ice cream sales increase, the risk of a shark attack also tends to increase. This does not mean that if a small city in the Midwest increases its ice cream sales that the citizens are at higher risk of being attacked by a shark. As Figure 4.4 illustrates, there is a confounder — temperature. As the temperatures increase, people tend to buy more ice cream; as the temperature increases, people tend to go to the beach, thereby increasing the risk of a shark attack. The two variables, ice cream sales and shark attacks, appear to be related as a result of the confounder of temperature.

Big Idea

Confounders are variables that influence *both* the factor of interest and the response.

Observational studies are subject to confounding; thus, controlled experiments are often considered the gold standard in research because they allow us to infer cause-and-effect relationships from the data. Controlled experiments allow for causal interpretations because the random allocation to the levels of the factor removes the impact of confounders. Let's return to the hypothetical Vitamin-K study in Example 4.2. Suppose there are nurses with a gentle bedside

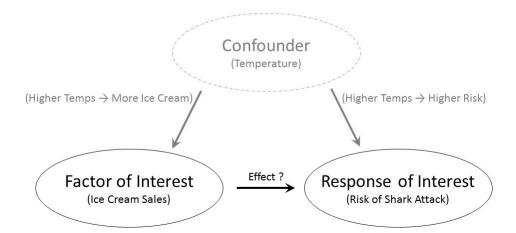


Figure 4.4: Illustration of a confounding variable. The confounder, related to both the response and the factor of interest (or treatment) can make it appear as though there is a causal relationship when none exists.

manner and those who are a little less gentle. If we randomly determine which infants receive Kangaroo Care, then for every gentle nurse who is told to implement Kangaroo Care while giving the shot, there tends to be a gentle nurse who is told to not implement Kangaroo Care. Similarly, for every less-gentle nurse who is told to implement Kangaroo Care while giving a shot, there tends to be a less-gentle nurse who is told to not implement Kangaroo Care. This is illustrated in Figure 4.5. For an observational study, the treatment groups can be unbalanced; for example, Figure 4.5 illustrates a case in which there is a higher fraction (11/12 compared to 1/4) of friendly nurses in the Kangaroo Care group compared to the No Kangaroo Care group. For the controlled experiment however, the treatment groups tend to be balanced with respect to this confounder; there is approximately the same fraction of friendly nurses in both groups. Random assignment is the great equalizer. It tends to result in groups which are similar in all respects; therefore, since we have eliminated all other differences between the groups (other than the treatment they receive), any differences we observe between the groups must be due to the grouping and not an underlying confounding variable.

Big Idea

Randomly assigning subjects to groups balances the groups with respect to any confounders; that is, the groups being compared are similar. Therefore, any differences between the two groups can be attributed to the grouping itself, leading to cause-and-effect conclusions.

While controlled experiments are a fantastic study design, we should not discount the use of observational studies. Consider the Deepwater Horizon Case Study described in Chapter 2;

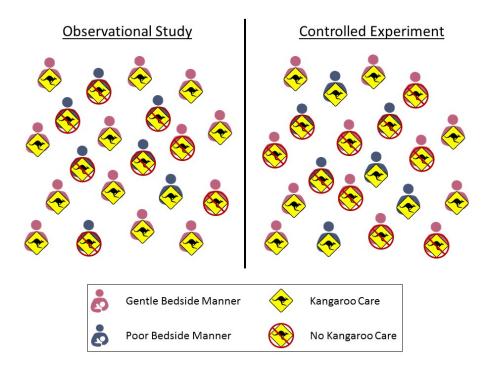


Figure 4.5: Illustration of the impact of random assignment in study design. For the observational study, the treatment groups are unbalanced. For the controlled experiment, the treatment groups are balanced.

suppose we are interested in the following question:

Is there evidence that volunteers who are directly exposed to oil have an increased risk of developing adverse respiratory symptoms compared to those who are not directly exposed to oil?

The response is whether a volunteer develops adverse respiratory symptoms; the factor of interest is whether the volunteer has direct exposure to oil. We could conduct a controlled experiment by randomly determining which volunteers are assigned to wildlife clean up and which are assigned to administrative tasks, for example. However, it may be that volunteer tasks need to be determined by skillset or by greatest need at the time the person volunteers. It may not be feasible to randomly assign volunteers to specific positions. Or, it could be that the data was obtained after the fact; that is, the data is not the result of a planned study in which case random assignment is not possible because volunteers self-selected into positions in the past. If random assignment is not possible, it does not mean the data is useless. But, it does mean we will need to be sure we acknowledge, and potentially address, the potential confounding when performing the analysis and discussing the results.

The big idea is that in order to make causal conclusions, we must be able to state that the groups being compared are balanced with respect to any potential confounders; random assignment is one technique for accomplishing this.

5 Presenting the Evidence (Summarizing Data)

If you open any search engine and look up "data visualization," you will quickly be overwhelmed by a host of pages, texts, and software filled with tools for summarizing your data. Here is the bottom line: a good visualization is one that helps you answer your question of interest. It is both that simple and that complicated.

Fundamental Idea III

The use of data for decision making requires that the data be summarized and presented in ways that address the question of interest and represent the variability present.

Whether simple or complex, all graphical and numerical summaries should help turn the data into usable information. Pretty pictures for the sake of pretty pictures are not helpful. In this chapter, we will consider various simple graphical and numerical summaries to help build a case for addressing the question of interest. The majority of the chapter is focused on summarizing a single variable; more complex graphics are presented in future chapters within a context that requires them.

5.1 Characteristics of a Distribution (Summarizing a Single Variable)

Remember that because of variability, the key to asking good questions is to not ask questions about individual values but to characterize the underlying distribution (see Definition 3.3). Therefore, characterizing the underlying distribution is also the key to a good visualization or numeric summary. For the Deepwater Horizon Case Study described in Chapter 2, the response (whether a volunteer experienced adverse respiratory symptoms) is categorical. As we stated previously, summarizing the distribution of a categorical variable reduces to showing the proportion of individual subjects that fall into each of the various groups defined by the categorical variable. Figure 5.1) displays a bar chart summarizing the rate of respiratory symptoms for volunteers cleaning wildlife.

In general, it does not matter whether the frequency or the relative frequencies are reported; however, if the relative frequencies are plotted, some indication of the sample size should




Figure 5.1: Frequency of adverse respiratory symptoms for volunteers cleaning wildlife following the Deepwater Horizon oil spill.

be provided with the figure, either as an annotation or within the caption. From the above graphic, we see that nearly 28% of volunteers assigned to wildlife experienced adverse respiratory symptoms; the graphic helps address our question, even if not definitively.

Note

When you are summarizing only categorical variables, a bar chart is sufficient. Statisticians tend to agree that bar charts are preferable to pie charts (see this whitepaper and this blog for further explanation).

While a single type of graphic (bar charts) are helpful for looking at categorical data, summarizing the distribution of a numeric variable requires a bit more thought. Consider the following example.

Example 5.1 (Paper Strength). While electronic records have become the predominant means of storing information, we do not yet live in a paperless society. Paper products are still used in a variety of applications ranging from printing reports and photography to packaging and bathroom tissue. In manufacturing paper for a particular application, the strength of the resulting paper product is a key characteristic.

There are several metrics for the strength of paper. A conventional metric for assessing the inherent (not dependent upon the physical characteristics, such as the weight of the paper,

Table 5.1: Breaking length (km) for first 5 specimens in the Paper Strength study.

Specimen	Breaking Length
1	21.312
2	21.206
3	20.709
4	19.542
5	20.449

which might have an effect) strength of paper is the *breaking length*. This is the length of a paper strip, if suspended vertically from one end, that would break under its own weight. Typically reported in kilometers, the breaking length is computed from other common measurements. For more information on paper strength measurements and standards, see the following website: http://www.paperonweb.com

A study was conducted at the University of Toronto to investigate the relationship between pulp fiber properties and the resulting paper properties (Lee 1992). The breaking length was obtained for each of the 62 paper specimens, the first 5 measurements of which are shown in Table 5.1. The complete dataset is available online at the following website: https://vincentarelbundock.github.io/Rdatasets/doc/robustbase/pulpfiber.html

While there are several questions one might ask with the available data, here we are primarily interested in characterizing the breaking length of these paper specimens.

Figure 5.2 presents the breaking length for all 62 paper specimens in the sample through a *dot* plot in which the breaking length for each observed specimen is represented on a number line using a single dot.

With any graphic, we tend to be drawn to three components:

- where the values tend to be,
- how tightly the values tend to be clustered there, and
- the way the values tend to cluster.

Notice that about half of the paper specimens in the sample had a breaking length longer than 21.26 km. Only about 25% of paper specimens had a breaking length less than 19.33 km. These are measures of *location*. In particular, these are known as **percentiles**, of which the **median**, **first quartile** and **third quartile** are commonly used examples.

Definition 5.1 (Percentile). The k-th percentile is the value q such that k% of the values in the distribution are less than or equal to q. For example,

• 25% of values in a distribution are less than or equal to the 25-th percentile (known as the "first quartile" and denoted Q_1).

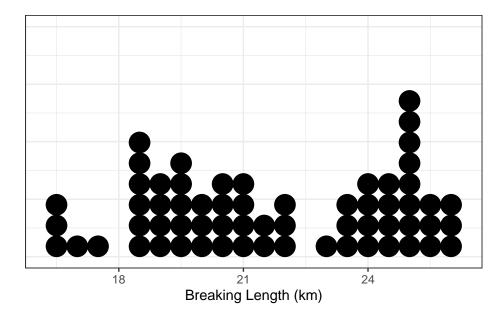


Figure 5.2: Breaking Length (km) for 62 paper specimens.

- 50% of values in a distribution are less than or equal to the 50-th percentile (known as the "median").
- 75% of values in a distribution are less than or equal to the 75-th percentile (known as the "third quartile" and denoted Q_3).

The average is also a common measure of location. The breaking length of a paper specimen is 21.72 km, on average. In this case, the average breaking length and median breaking length are very close; this need not be the case. The average is not describing the "center" of the data in the same way as the median; they capture different properties.

Definition 5.2 (Average). Also known as the "mean," this measure of location represents the balance point for the distribution. If x_i represents the *i*-th value of the variable x in the sample, the sample mean is typically denoted by \bar{x} .

For a sample of size n, it is computed by

$$\bar{x} = \frac{1}{n} \sum_{i=1}^{n} x_i.$$

When referencing the average for a population, the mean is also called the "Expected Value," and is often denoted by μ .

Clearly, the breaking length is not equivalent for all paper specimens; that is, there is variability in the measurements. Measures of *spread* quantify the variability of values within a distribution. Common examples include the **standard deviation** (related to **variance**) and **interquartile range**. For the Paper Strength example, the breaking length varies with a standard deviation of 2.88 km; the interquartile range for the breaking length is 5.2 km.

The standard deviation is often reported more often than the variance since it is on the same scale as the original data; however, as we will see later, the variance is useful from a mathematical perspective for derivations. Neither of these values has a natural interpretation; instead, larger values of these measures simply indicate a higher degree of variability in the data.

Definition 5.3 (Variance). A measure of spread, this roughly captures the average distance values in the distribution are from the mean.

For a sample of size n, it is computed by

$$s^{2} = \frac{1}{n-1} \sum_{i=1}^{n} (x_{i} - \bar{x})^{2}$$

where \bar{x} is the sample mean and x_i is the *i*-th value in the sample. The division by n-1 instead of n removes bias in the statistic.

The symbol σ^2 is often used to denote the variance in the population.

Definition 5.4 (Standard Deviation). A measure of spread, this is the square root of the variance.

Definition 5.5 (Interquartile Range). Often abbreviated as IQR, this is the distance between the first and third quartiles. This measure of spread indicates the range over which the middle 50% of the data is spread.

Note

The IQR is often incorrectly reported as the interval (Q_1, Q_3) . The IQR is actually the width of this interval, not the interval itself.

The measures we have discussed so far are illustrated in Figure 5.3. While some authors suggest the summaries you choose to report depend on the shape of the distribution, we argue that it is best to report the values that align with the question of interest. It is the question that should be shaped by the beliefs about the underlying distribution.

Finally, consider the *shape* of the distribution of breaking length we have observed. The breaking length tends to be clustered in two locations; we call this *bimodal* (each mode is a

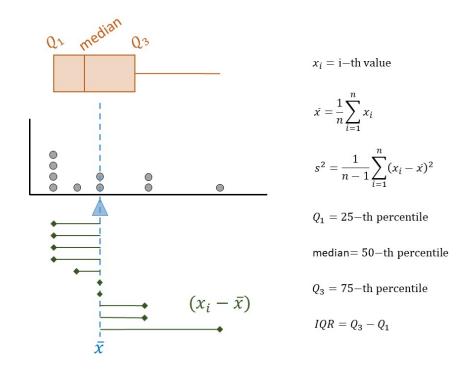


Figure 5.3: Illustration of measures of location and spread for a distribution of values.

"hump" in the distribution). Other terms used to describe the shape of a distribution are *symmetric* and *skewed*. Symmetry refers to cutting a distribution in half (at the median) and the lower half being a mirror image of the upper half; skewed distributions are those which are not symmetric.

Observe that the dot plot above gives us some idea of the location, spread, and shape of the distribution, in a way that the table of values could not. This makes it a useful graphic as it is characterizing the **distribution of the sample** we have observed. This is one of the four components of what we call the *Distributional Quartet*.

Definition 5.6 (Distribution of the Sample). The pattern of variability in the observed values of a variable.

When the sample is not large, a dot plot is reasonable. Other common visualizations for a single numeric variable include:

- jitter plot: similar to a dot plot, each value observed is represented by a dot; the dots are "jittered" (shifted randomly) in order to avoid over-plotting when many subjects share the same value of the response.
- box plot: a visual depiction of five key percentiles; the plot includes the minimum, first quartile, median, third quartile, and maximum value observed. The quartiles are connected with a box, the median cuts the box into two components. Occasionally, outliers are denoted on the graphic.
- *histogram*: can be thought of as a grouped dot plot in which subjects are "binned" into groups of similar values. The height of each bin represents the number of subjects falling into that bin.
- density plot: a smoothed histogram in which the y-axis has been standardized so that the area under the curve has value 1. The y-axis is not interpretable directly, but higher values along the y-axis indicate that the corresponding values on along the x-axis are more likely to occur.

Definition 5.7 (Outlier). An individual observation which is so extreme, relative to the rest of the observations in the sample, that it does not appear to conform to the same distribution.

To illustrate these graphics, the breaking length for the Paper Strength example is summarized using various methods in Figure 5.4. The latter three visualizations are more helpful when the dataset is very large and plotting the raw values actually hides the distribution. There is no right or wrong graphic; it is about choosing the graphic which addresses the question and adequately portrays the distribution.

The numeric summaries of a distribution are known as **statistics**. While parameters characterize a variable at the population level, statistics characterize a variable at the sample level.




Figure 5.4: Four graphical summaries of the breaking length for the Paper Strength example.

Definition 5.8 (Statistic). Numeric quantity which summarizes the distribution of a variable within a *sample*.

Why would we compute numerical summaries in the sample if we are interested in the population? Remember the goal of this discipline is to use the sample to say something about the underlying population. As long as the sample is representative, the distribution of the sample should reflect the **distribution of the population**; therefore, summaries of the sample should be close to the analogous summaries of the population (statistics estimate their corresponding parameters). Now we see the real importance of having a representative sample; it allows us to say that what we observe in the sample is a good proxy for what is happening in the population.

Definition 5.9 (Distribution of the Population). The pattern of variability in values of a variable at the population level. Generally, this is impossible to know, but we might model it.

Statistics being a proxy for the corresponding parameter implies the mean in the sample should approximate (estimate) the mean in the population; the standard deviation of the sample should estimate the standard deviation in the population; and, the shape of the sample should approximate the shape of the population, etc. The sample is acting as a representation in all possible ways of the population.

🕊 Big Idea

A representative sample reflects the population; therefore, we can use statistics as estimates of the population parameters.

Note

Notation in any discipline is both important and somewhat arbitrary. We can choose any symbol we want to represent the sample mean. However, it is convention that we never use \bar{x} to represent a parameter like the mean of the population. The symbol \bar{x} (or \bar{y} , etc.) represents observed values being averaged together. Since the values are observed, we must be talking about the sample, and therefore \bar{x} represents a statistic. A similar statement could be made for s^2 (sample variance) compared to σ^2 (population variance). Again, in reality, the symbols themselves are not important. The importance is on their representation. Statistics are observed while parameters are not.

5.2 Summarizing Relationships

The summaries discussed above are nice for examining a single variable. In general, however, research questions of interest typically involve the relationship between two or more variables. Most graphics are two-dimensional (though 3-dimensional graphics and even virtual reality are being utilized now); therefore, summarizing a rich set of relationships may require the use of both axes as well as color, shape, size, and even multiple plots in order to tell the right story. We will explore these various features in upcoming units of the text. Here, we focus on the need to tell a story that answers the question of interest instead of getting lost in making a graphic. Consider the following question from the Deepwater Horizon Case Study described in #sec-caseDeepwater:

What is the increased risk of developing adverse respiratory symptoms for volunteers cleaning wildlife compared to those volunteers who do not have direct exposure to oil?

Consider the graphic in Figure 5.5; this is *not* a useful graphic. While it compares the number of volunteers with symptoms in each group, we cannot adequately address the question because the research question involves comparing the rates for the two groups; that is, we are lacking a sense of how many volunteers in each group did not report symptoms.

Instead, Figure 5.6 compares the rates within each group. Note that the graphic is still reporting frequency along the y-axis; that was not the primary problem with Figure 5.5. However, by reporting frequencies for both those with respiratory symptoms and those without, we get a sense of the relative frequency with which respiratory symptoms occur.




Figure 5.5: Illustration of a poor graphic; the graphic does not give us a sense of the rate within each group at which volunteers reported symptoms.




Figure 5.6: Comparison of the rate of adverse respiratory symptoms among volunteers assigned to different tasks.

From the graphic, it becomes clear that within the sample a higher fraction of volunteers cleaning wildlife experienced adverse symptoms compared with those without oil exposure. In fact, volunteers cleaning wildlife were 1.79 times more likely to experience adverse respiratory symptoms.

The key to a good summary is understanding the question of interest and addressing this question through a useful characterization of the variability.

6 Assessing the Evidence (Quantifying the Variability in Estimates)

Again, the goal of statistical inference is to use the sample as a snapshot of the underlying population (Figure 1.2). There are generally three reasons people distrust this process:

- 1. Fear that the sample does not represent what is going on in the population.
- 2. Fear that we cannot make a conclusion with a sample of size n (wanting more data).
- 3. Fear that one study is not enough to make a conclusion.

We have already tackled the first fear in Chapter 4; if we are to trust statistical results, we must collect data that is representative of the underlying population. The second and third fears above are tied together, though maybe not obviously. Before launching into a slightly more formal discussion, consider the following thought experiment.

Example 6.1 (Free Throws). Your friend Dave lives for his Wednesday "pick-up" basketball games at the gym. One afternoon, while waiting for a few more players to arrive Dave shoots 10 free throws, of which he makes 3.

While Dave only exhibited a 30% success rate from the free throw line in this sample, we imagine no one is ready to claim *definitively* that Dave has a 30% success rate from the free throw line overall. So, what can we say? Well, if this set of 10 free throws is representative of Dave's free throw performance, then we would say that 30% is an *estimate* for his success rate; that is, the statistic 30% is a good guess at the unknown parameter (overall success rate). There are two ways we might improve our "trust" in this estimate. First, we might consider a larger sample size (make Dave shoot more free throws); let's continue along these lines for a moment.

Example 6.2 (Free Throws (Cont.)). Joe has also been waiting for a few more players to arrive; however, Joe shoots 100 free throws (clearly he has more time on his hands) of which he makes 30.

Again, we probably wouldn't claim *definitively* that Joe has a 30% success rate from the free throw line overall. And again, assuming this set of 100 free throws is representative of his overall performance, we would say 30% is an *estimate* for his success rate. But, we might also say we have more "trust" in our guess for Joe's overall performance compared with our guess

for Dave's. The more shots we observe, the more we seem to "trust" our estimate. This idea is known as the Law of Large Numbers.

Definition 6.1 (Law of Large Numbers). For our purposes, the Law of Large Numbers essentially says that as a sample size gets infinitely large, a statistic will become arbitrarily close (extremely good approximation) of the parameter it estimates.

There are two drawbacks with the Law of Large Numbers. First, it does not tell us how "close" a statistic is to the parameter for any specific sample size; second, we cannot take an infinitely large sample. For our thought experiment, it is probably not feasible to have Dave or Joe shoot thousands of free throws, for example. Our goal then becomes to somehow quantify the "trust" we have in our estimates given the sample size we have available. That is, given that we only saw Dave shoot 10 free throws, can we quantify our "trust" in that 30% estimate of his free throw success? We need some way of measuring "trust," and we do that through a notion of statistical "confidence."

Big Idea

In statistics, our "trust" is tied to the estimate's repeatability; we ask the question "if we were to repeat the study, how much would we expect our estimate to change?"

We will formalize the notion of statistical confidence shortly, but for now, linking our trust in an estimate to its repeatability gets at the last fear. We know that if we repeat a study, the results will change; our job is to quantify (keeping the sample size in mind) the degree to which the results will change. That is, we need to quantify the variability in the estimate across repeated studies (known as sampling variability; we told you statistics was all about variability). This is characterized by the **sampling distribution**.

Definition 6.2 (Sampling Distribution). The distribution of a *statistic* across repeated samples (of the same size) from the population.

This is perhaps the most important of the *Distributional Quartet*; it is the holy grail of statistical inference. Once we have the sampling distribution, inference is straight-forward.

I Fundamental Idea IV

Variability is inherent in any process, and as a result, our estimates are subject to sampling variability. However, these estimates often vary across samples in a predictable way; that is, they have a distribution that can be modeled.

6.1 Conceptualizing the Sampling Distribution

The sampling distribution of a statistic is one of the most fundamental, and yet one of the most abstract, concepts in statistics. Its name is even confusing; the "distribution of the sample" (Definition 5.6) and the "sampling distribution" (Definition 6.2) use similar words but represent two different things. Before we discuss the utility of the sampling distribution, we first focus on making it a bit more tangible.

For the Deepwater Horizon Case Study discussed in Chapter 2, consider the following question:

What proportion of volunteers assigned to clean wildlife develop adverse respiratory symptoms?

In the sample, we observed 15 out of 54 such volunteers (27.8% or a proportion of 0.278). This proportion is a good estimate of the rate of adverse symptoms in the population (assuming the sample is representative, of course).

Now, imagine randomly selecting 54 *new* volunteers from the population (repeating the study). For this new sample, it would be possible to determine the fraction of volunteers that experienced adverse symptoms; we would expect this value to be a bit different than what we obtained in the first sample since the two samples consist of different subjects. Since this second sample is also representative, however, it also provides a good estimate of the parameter. That is, we now have two good estimates of the same parameter.

Now, we could take a third random sample of 54 volunteers and compute the fraction in this third sample which experienced adverse symptoms. This third sample also provides a good (and potentially unique) estimate of the parameter. In fact, we could continue this process m times, for some large number m, as illustrated in Figure 6.1).

Consider what we are describing. With each representative sample, we have constructed an estimate of the parameter. What we have kept from each replicate sample is *not* the values of the variables themselves (whether the volunteers experienced adverse respiratory symptoms); instead, we have retained the *statistic* from each of *m* completely different studies. So, which of these *m* estimates do we trust? All of them. Since each sample is representative of the population, each estimate is a good (not perfect) estimate of the parameter. Since we have all these estimates, we should think about what information they provide. In fact, there is information not only in what these estimates are but in how different they are from one another. Describing the way in which these estimates change from one sample to another is the sampling distribution.

Notice that the sampling distribution is not describing a variable from our study; it is describing a *statistic*. In order to construct a sampling distribution, we go through the following steps:

- 1. Take a sample; record variables of interest.
- 2. Compute the statistic which estimates the parameter and retain this value.

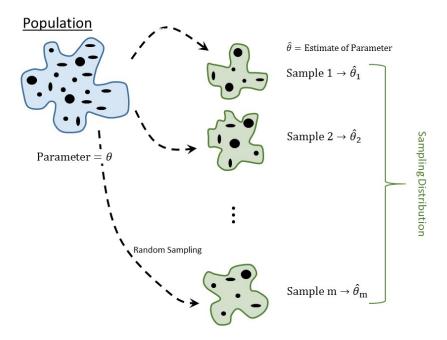


Figure 6.1: Illustration of repeatedly sampling from a population.

- 3. Repeat steps 1 and 2 a large number of times.
- 4. Examine the statistics collected.

So, the sampling distribution is not a plot of the raw values of a variable on individual subjects but a plot of statistics which summarize entire samples. That is, the unit of observation has changed in this distribution. While a sample consists of individual subjects from the population, the sampling distribution consists of individual samples from the population.



Big Idea

Re-read the description of a sampling distribution several times, and return to it often as you read through the text. It takes a while for this to sink in, but if you truly grasp this one concept, the remainder of statistical inference becomes much more accessible.

6.2 Example of a Sampling Distribution

Since this idea is so critical to grasping statistical inference, we are going to walk through the process of generating a sampling distribution for a known data generating process.

Example 6.3 (Dice Experiment). Consider an ordinary six-sided die; we are interested in the proportion of times that rolling the die will result in a 1. Putting this in the language of the statistics, we have the following:

- The population of interest is all rolls of the die. Notice that this population is infinitely large as we could roll the die forever.
- The variable is the resulting value from the roll. Since this can take on only one of six values, this is a categorical variable.
- The parameter of interest is the proportion of rolls that result in a 1.

Our goal is to construct the sampling distribution of the sample proportion of rolls that result in a 1 when the die is rolled 20 times.

What makes this example unique is that we know the value of the parameter. Because of the physical properties of a die, we know that the probability a roll results in a 1 is $\theta = 1/6$. So, statistical inference is not needed here. This example simply provides a vehicle for studying sampling distributions. Also, before going on, notice that the sampling distribution is for the statistic (the sample proportion) and not the parameter; and, it is constructed for a fixed sample size (in this case, n=20 rolls). Going back to the steps for creating a sampling distribution described in the previous section, we have the following steps:

1. Roll a die 20 times, each time recording the resulting value.

- 2. Compute the proportion of times (out of the 20) the resulting value was a 1 and retain this value.
- 3. Repeat steps 1 and 2 a large number of times (let's say 500).
- 4. Plot the resulting values; there should be 500 proportions that we are keeping.

Notice that we are actually rolling a die 10000 times (20 rolls repeated 500 times); we only keep 500 values (one proportion for each set of 20 rolls). This is something you could physically do at home. For example, the first sample might look like that in Figure 6.2.




Figure 6.2: Potential sample from rolling a die 20 times.

For this particular sample, the proportion in the sample (our statistic of interest) would be 0.25 (5/20). That is the statistic we would record. We then repeat this 499 more times. You could try a few out yourself using an online simulator. Figure 6.3 shows the resulting proportions for when we simulated this process with 500 samples, each sample consisting of 20 rolls.




Figure 6.3: Sampling distribution for the proportion of 20 rolls of a die which result in a 1. The distribution is based on repeating the sampling process 500 times.

With modern computing power, there is no need to restrain ourselves to repeating the study 500 times. A simple computer program could replicate rolling the study (20 rolls of a die)

thousands of times. Figure 6.4 is the sampling distribution for the proportion of rolls that result in a 1 based on a sample of size 20, repeating the study 50000 times.




Figure 6.4: Sampling distribution for the proportion of 20 rolls of a die which result in a 1. The distribution is based on repeating the sampling process 50000 times.

⚠ Warning

When looking at a graphical representation of a sampling distribution (such as Figure 6.4), there is a tendency to describe the "data" in the graphic. However, the use of the word "data" here is inappropriate. We restrict the word "data" to refer to information/measurements observed on each subject; this includes numeric and categorical variables directly measured and any transformations of these variables.

For example, the total number of credit hours a student has completed is part of the data observed (quantitative variable). Similarly, using the total number of credit hours to determine class standing (Freshman, Sophomore, Junior, Senior) would also be a part of the data observed (categorical variable).

However, the word "data" should not be used to describe the statistics computed from the sample, or from repeated sampling such as when constructing a sampling distribution, as these quantities are not characterizing the individual units under study.

Notice that the sampling distribution is centered around the true value of the parameter $(\theta = 1/6)$. In general, the sampling distribution of a statistic, when taken from a random sample, is centered on the true value of the parameter. This is the unbiased nature of the data coming out; random samples are representative of the population. Similarly, note that

while no one sample (remember, each value in the distribution represents a statistic from a sample of 20 values) is perfect, none of the samples produced a statistic which is really far from the true parameter. That is, a representative sample may not be perfect, but it will give a reasonable estimate of the parameter. Notice that these properties hold even though we had a relatively small sample size (only rolling the die n=20 times).

Big Idea

The size of the sample is not as important as whether it is representative. A small representative sample is better for making inference than a large sample which is biased.

One of the most useful things about the sampling distribution is that it gives us an idea of how much we might expect our statistic to change from one sample to another. Based on Figure 6.4, we could say that if we roll a die 20 times, the proportion of rolls which result in a 1 is most likely to be between 0.05 and 0.30 (so somewhere between 1 and 6 ones out of the 20 rolls). It would be extremely rare to have 12 of the 20 rolls result in a 1 (notice how small the bar is for a proportion of 0.6). The sampling distribution is therefore giving us an idea of the variability in our statistic.

Remember, our goal was to account for the variability in the statistic (how much it changes from one sample to another) while accounting for the sample size. How is this done? When forming the sampling distribution, we repeated the study, which included ensuring that for each replication, we obtained a new sample that had the same size as the original sample. So, the sample size is baked into the sampling distribution. To see the impact of taking a larger sample, consider rolling a six-sided die 60 times instead of 20 times. When we build the sampling distribution, each replication will then involve repeating the process with 60 new rolls. Figure 6.5 shows the sampling distribution of the proportion of 60 rolls which result in a 1, using 50000 replications. Notice that the distribution is still centered on the true parameter $\theta = 1/6$. The primary difference between this figure and the last is that when we increased the sample size, the sampling distribution narrowed.

We all have this intuition that "more data is better." In truth, we should say "more good data is better" since we understand that having a representative sample is extremely important. With that in mind, the distinction between Figures 6.4 and 6.5 illustrate what we really mean when we say "better" — the statistic is less variable. We have to be careful here. We are **not** saying that the *sample* has less variability; we are saying the *statistic* has less variability. That is, we do not expect our estimate to change as much from one sample to the next. From Figure 6.5, we have that if we roll the die 60 times, we expect the proportion of 1's to be somewhere between 0.1 and 0.25 (somewhere between 6 and 15 ones out of the 60 show up). The proportion is varying much less from one sample to the next compared to when we rolled the die only 20 times.




Figure 6.5: Sampling distribution for the proportion of 60 rolls of a die which result in a 1. The distribution is based on repeating the sampling process 50000 times.

Pig Idea

Larger samples result in statistics which are less variable. That is, as the sample size increases, the sampling distribution of a statistic becomes narrower.

Note

Students often believe that a large sample reduces the variability in the data. That is not true; a large sample reduces the variability in the *statistic*.

⚠ Warning

We repeat a warning here that we stated when introducing the concept of bias. There is a difference between accuracy and precision. Generally, accuracy refers to location (and therefore bias); we say an estimate is accurate when it is unbiased. Precision refers to the variability; we say an estimate is more precise when it has less variability. With regard to sampling distributions, accuracy refers to the center of the sampling distribution while precision refers to its spread.

We have briefly discussed the center and spread of the sampling distribution we constructed above. Before leaving this section, a note on the shape of the sampling distribution is also in

order. Note that for this example, the resulting observations can take on one of only two values (heads or tails). A bar chart summarizing this distribution (that of the variable within the *population*) would consist of only two bars. However, the distribution of the statistic that we record (that is, the *sampling distribution*) is bell-shaped (fig-samplingdistns-dice-histogram2). While the sampling distribution will not always be bell-shaped, it is often the case for the statistics discussed in this text. The key thing to recognize is that the distribution of the variable and the distribution of the statistic are not the same.

6.3 Modeling the Sampling Distribution

Let's return to the Deepwater Horizon Case Study. In particular, suppose we are trying to address the following question:

What proportion of volunteers assigned to clean wildlife develop adverse respiratory symptoms?

We have an estimate for this proportion ($\hat{p} = 0.278$) based on the observed sample. Based on the discussion in the previous section, we know the sampling distribution of this proportion can help us quantify the variability in the estimate. Figure 6.6 represents the sampling distribution of this proportion. From the graphic, we would not expect the proportion of volunteers who experience adverse respiratory symptoms to move much beyond 0.15 and 0.4 if we were to repeat the study; it would almost certainly not move beyond 0.1 and 0.5 if we were to repeat the study.

Now, you might ask "wait, where did this sampling distribution come from? There is no way you actually repeated the study 50000 times, right?" And, you would be right.

In the previous section, we described building the sampling distribution through repeated sampling. In practice, this is never practical. Generally, cost is the limiting factor when collecting data; as a result, we get a single sample to work with. We have essentially argued that the sampling distribution is critical to making inference, but we cannot take multiple samples to make it. Where does that leave us? The answer...modeling. Our goal is to construct a model of the sampling distribution that we can use to make inference.

There are three general techniques for modeling the sampling distribution of a statistic:

- 1. Build an empirical model.
- 2. Build an exact analytical model using results from probability theory.
- 3. Build an approximate analytical model using results from theorems about limits in probability.

We will focus on the first approach; the latter two approaches are discussed in Appendix A. Our emphasis in this chapter is on the conceptual understanding of a sampling distribution and its model. While these latter two approaches differ in their technique, the use of the resulting




Figure 6.6: Sampling distribution for the proportion of volunteers assigned to wildlife who develop adverse symptoms based on a sample of 54 volunteers.

model is the same. We choose to focus on the first approach because it requires less mathematical background and reinforces the conceptual understanding of a sampling distribution discussed above. The idea in constructing an empirical model is to mimic the discussion above regarding the construction of a sampling distribution. Our description references Figure 6.7 often.

We are limited by our resources; because of time and money constraints, we cannot resample from the population (crossed off resamples in Figure 6.7). So, we pretend for a moment that our original sample (colored in green in Figure 6.7) is the population. Our idea is to randomly sample from this original data, creating a resample (colored in orange in Figure 6.7). Forgive the non-technical terms here, but since the orange "blob" is a random sample from the green "blob," then it is representative of the green blob. Therefore, if we construct an estimate $\hat{\theta}^*$ from the orange blob (the star denotes a statistic from a resample), then it should be close to the statistic $\hat{\theta}$ from the green blob; but, since this green blob is representative of the population, $\hat{\theta}$ should be a good estimate of θ . Therefore, we have that

$$\hat{\theta}^* \approx \hat{\theta} \approx \theta \Rightarrow \hat{\theta}^* \approx \theta$$

That is, each resample produces a statistic which is a good estimate of the parameter from the underlying population. The benefit here is that the resamples are taken from the original sample, not the population, and can therefore be constructed in the computer. And, given today's computing power, we are not limited by time or money (10000 resamples can often be taken in a matter of seconds). If you want to see this process in action, we encourage you to check out the free online app located at http://www.lock5stat.com/StatKey/bootstrap_1_cat.html.




Figure 6.7: Illustration of modeling the sampling distribution via bootstrapping.

Again, the idea is to mimic in the computer the resampling that we were unable to do in real life. This process is known as **bootstrapping**.

Definition 6.3 (Bootstrapping). A method of modeling the sampling distribution by repeatedly resampling from the original data.

There are actually several variations of bootstrapping; however, for our purposes currently, we can keep the following details regarding the implementation in mind:

- 1. Each resample (known as a bootstrap resample) is the same size as the original sample.
- 2. Each resample is taken *with replacement*; that means the values from the original sample can show up multiple times. Think of "catch and release" fishing.
- 3. Typically, between 3000 and 10000 bootstrap resamples are taken.

We will avoid actual computation throughout the text, but a quick online search would provide several resources for implementing the process we have described (and its many variants) in various computer programming languages and software packages.

Big Idea

Students often believe that bootstrapping "creates more data." This is not true. Instead, boostrapping resamples from the existing data. By its very nature, it takes the limited information in the sample into account. This highlights the need to have a representative sample when performing analysis.

For the Deepwater Horizon Case Study discussed in Chapter 2, we performed the following steps to create Figure 6.6:

- 1. Select 54 volunteers at random (with replacement) from the original sample of 54 volunteers who had been assigned to clean wildlife.
- 2. For our bootstrap resample, we compute the proportion of those individuals who had experienced adverse respiratory symptoms; this is our *bootstrap statistic*.
- 3. We repeated steps 1 and 2 several thousand times, retaining the bootstrap statistics from each bootstrap resample.
- 4. We plotted the distribution of the bootstrap statistics.

6.4 Using a Model for the Sampling Distribution (Confidence Intervals)

We began this chapter by arguing that quantifying the variability in our estimates was crucial to making inference on the parameters. A model for the sampling distribution of a statistic allows us to visualize the variability in our estimates, and we will capitalize on that in this section. A stepping stone in that direction is simply estimating the variability in the sampling distribution. In Chapter 5, we discussed various metrics for quantifying variability; one such metric was the standard deviation. While we could rely on any metric, the standard deviation is the most common metric used when quantifying the variability of a statistic; to distinguish that we are quantifying the variability of a statistic instead of a variable, we refer to this as the standard error.

Definition 6.4 (Standard Error). The standard error is the estimated standard deviation of a *statistic*; that is, it is the standard deviation from a *model* for the sampling distribution of a statistic. It quantifies the variability in the statistic across repeated samples.

We will see the usefulness of the standard error in Chapter 18. Until then, we simply note that we are able to quantify the variability in our statistics.

Returning to our question for the Deepwater Horizon Case Study — "What proportion of volunteers assigned to clean wildlife develop adverse respiratory symptoms?" — we have an estimate for this parameter: $\hat{p} = 0.278$. However, there is something unsatisfying about this

estimate...it fails to acknowledge the variability in the statistic which we know exists. That is, from the above discussion, we have seen that repeating the study would lead to a different estimate of the parameter. We would like to leverage the information contained in our model for the sampling distribution to provide an estimate which incorporates the variability in this statistic. To do this, we somewhat "reverse engineer" the information we need from the sampling distribution.

Consider the model for the sampling distribution of the sample proportion we constructed in Figure 6.6. From the model, we see that repeatedly resampling from our data, we would not expect to obtain a proportion (of volunteers who experience adverse symptoms) to move much lower than 0.15 or much higher than 0.4. How does this help us in performing inference? Remember that each value in the bootstrap model for the sampling distribution is an estimate of the underlying parameter. So, we can think of the above model as showing us what good estimates of the parameter look like. Another way of saying it: the model for the sampling distribution shows us the reasonable (or plausbile) values of the parameter. Here, by "reasonable," we mean values of the parameter for which the data is consistent. Consider the following statements (which are equivalent):

- Based on our sample of 54 volunteers, it is reasonable that the proportion of volunteers assigned to clean wildlife who would experience adverse respiratory symptoms is between 0.15 and 0.4.
- Our sample of 54 volunteers is consistent with between 15% and 40% of all volunteers assigned to clean wildlife experiencing adverse respiratory symptoms.

There is another way of thinking about how we move from a model for the sampling distribution to a range of plausible values for the parameter. Again, we observed that repeatedly resampling from our data, we would not expect to obtain a proportion (of volunteers who experience adverse symptoms) to move much lower than 0.15 or much higher than 0.4. We admit these are not actual statistics, but they are bootstrap statistics. So, we conclude that bootstrap statistics tend not to move more than approximately 0.125 units from the actual statistic we observed in our sample. Remember, bootstrapping is mimicking the process of a sampling distribution. Therefore, if bootstrap statistics only move about 0.125 units from the actual statistic, then we can conclude that statistics computed from resampling from the population would only move about 0.125 units from the actual parameter. As a result, our statistic must only be about 0.125 units from the true value of the parameter. This leads us to believe that the data is consistent with between 15% and 40% of all volunteers assigned to clean wildlife experiencing adverse respiratory symptoms. Notice we are led the same conclusion.



The model for the sampling distribution of a statistic allows us to determine the reasonable values for the corresponding parameter.

We have just conducted inference for "estimation" type questions. We are able to provide an estimate for the parameter which acknowledges that the data is not perfect and there is variability in sampling procedures. That variability incorporated itself into constructing an estimate that is an interval instead of a single point.

The above interval was chosen arbitrarily by just looking at the sampling distribution and capturing the peak of the distribution. If we want to be more formal, we might try to capture the middle 95% of values. This is known as a **confidence interval**.

Definition 6.5 (Confidence Interval). An interval (range of values) estimate of a parameter that incorporates the variability in the statistic. The process of constructing a k% confidence interval results in these intervals containing the parameter of interest in k% of repeated studies. The value of k is called the *confidence level*.

We have now formally defined "confidence," linking it to the behavior of a statistic across repeated samples. We have "higher confidence" when our process results in capturing the true parameter in a higher percentage of repeated studies.

If we were to capture the middle 95% of statistics in our model of the sampling distribution, a 95% confidence interval, we would obtain an interval of (0.167, 0.407), as shown in Figure Figure 6.8.




Figure 6.8: Construction of a confidence interval via bootstrapping for the proportion of volunteers assigned to wildlife who develop adverse symptoms based on a sample of 54 volunteers.

Process for Constructing a Confidence Interval

The following is a general procedure for constructing confidence intervals:

- 1. Choose a confidence level k (a decimal between 0 and 1, for example 0.95).
- 2. Construct a model for the sampling distribution of the statistic.
- 3. Grab the middle 100k% of values from the model in step (2).

Notice that the definition of a confidence interval, and this general procedure, apply regardless of the technique used for constructing the model of the sampling distribution.

Confidence intervals are often misinterpreted; this comes from their dependence on repeated sampling. When thinking about confidence intervals, think about playing a game of ring toss: you toss a ring in hopes of landing on top of a target. The target is the parameter characterizing the population. The confidence interval is like a ring. Since the confidence interval is constructed from a model of the sampling distribution, it changes with each sample; that is, the confidence interval itself is a statistic. Just like in ring toss where the ring moves with each toss, the confidence interval moves with each sample. However, the target (the parameter of interest) stays fixed. Because of this, there are many incorrect interpretations.

6.5 Incorrect Interpretations of a Confidence Interval

Suppose we have a k% confidence interval; the following are incorrect interpretations of the interval:

- There is a k% chance that individuals in the population have a value of the variable within the confidence interval.
- There is a k% chance (or we are k% sure) that the parameter of interest is inside the confidence interval.
- If we were to repeat the study, there is a k% chance we would see a statistic inside this confidence interval.

The first statement above is incorrect because it neglects that a confidence interval comes from a model for the *sampling distribution* not the *distribution of the sample*. And, sampling distributions describe the variability in *statistics*, not the variability of individuals in the population or sample. Therefore, a confidence interval cannot make statements about where individual observations will fall.

The second statement above is incorrect because it treats the parameter as the thing that is moving. Remember, the target stays fixed in ring toss; so, we can't say there is a probability the target will move into the ring. It is tempting to say "but since the ring is moving, we can have a chance of catching the target." That is a true statement before you throw the ring.

However, once the ring is tossed, you either captured the target or you did not; it no longer makes sense to say "I captured the target with 95% probability." The same holds for confidence intervals. Once the data has been collected, the confidence interval is a fixed quantity. At this point, neither the estimate nor the parameter is moving; so, there is no probability left (it either captured the parameter or it did not). So, the second statement is really ignoring the dependence of confidence interval interpretations on repeated sampling.

The third statement above is incorrect because it neglects the idea that the confidence interval changes with each new sample. Just as the ring changes location with each toss, so does the confidence interval.

For the Deepwater Horizon Case Study, our 95% confidence interval was (0.167, 0.407). Applying the above, the following are **incorrect** interpretations:

- There is a 95% chance that the proportion of volunteers assigned to clean wildlife who will experience adverse symptoms is between 0.167 and 0.407.
- 95% of volunteers assigned to clean wildlife in our sample (or population) had a value between 0.167 and 0.407.

Returning to how we motivated confidence intervals, appropriate interpretations rely on the idea of their capturing reasonable values of the parameter:

• The data is consistent with between 16.7% and 40.7% of volunteers assigned to clean wildlife experiencing adverse symptoms.

We recommend sticking to interpreting a confidence interval as specifying reasonable values for the parameter.



Big Idea

Confidence intervals specify reasonable values of the parameter based on the data observed.

Note

Other texts will interpret confidence intervals with statements like "we are 95% confident that the parameter is inside the interval." We do not recommend this language because the word "confident" is interpreted as "sure" in practice, and that leads to an incorrect interpretation.

This is a difficult concept to wrap our heads around; it seems natural to associate the percentage with the values we have obtained. However, our confidence is in the process, not the resulting interval itself (Figure 6.9). That is, 95% confidence intervals work 95% of the time; however, this statement is about the process of constructing confidence intervals. Once we have computed a confidence interval, it has either worked or not; the problem is of course,

that since we do not know the parameter, we will never know if it worked or not. For this reason, we prefer the interpretation of a confidence interval which avoids these subtleties: a confidence interval specifies the reasonable values of the parameter. The percentage (95% vs 99% for example) then just specifies what we mean by "reasonable."

It may seem like a good idea to make a 100% confidence interval to be sure we always capture the parameter. But, such intervals are not helpful in practice. For example, a 100% confidence interval for the proportion of volunteers experiencing adverse symptoms would be (0, 1). But, this is useless; it essentially says that the proportion has to be a number between 0 and 1, but we already knew that since all proportions are between 0 and 1 by definition. Therefore, we must balance the confidence we desire with the amount of information the interval conveys.

Big Idea

If you want both a high level of confidence but also a narrow interval, increase the sample size. As the sample size increases, the variability in the statistic decreases leading to a narrower interval.

Note

95% confidence intervals are the most common in practice; however, 90%, 98%, and 99% intervals are also used. It is extremely rare to use less than a 90% confidence interval.

6.6 Bringing it All Together

Consider the following question:

Does the study provide evidence that more than 1 in 5 volunteers assigned to clean wildlife develop adverse respiratory symptoms?

Let's answer this question using a confidence interval. Based on the data obtained, we found that the 95% confidence interval (CI) for the proportion of volunteers experiencing adverse symptoms to be (0.167, 0.407). Is this data consistent with more than 1 in 5 volunteers developing adverse symptoms? Yes, since there are proportions within this interval which are larger than 0.2. But, consistency is not the same as evidence; remember, evidence is the idea of "beyond a reasonable doubt." After all, is this data consistent with less than 1 in 5 volunteers developing adverse symptoms? Yes, since there are proportions within this interval which are less than 0.2.

Confidence intervals specify reasonable values — those values of the parameter which are consistent with the data. This data is then consistent with proportions that are both less than 0.2 and greater than 0.2. So, what can we say then? We can say that the study does not




Figure 6.9: Performance of the confidence interval process (for the sample mean) across repeated samples. Our idea of "confidence" requires repeated sampling. The process of constructing a 95% confidence on the captures the parameter in 95% of repeated samples; it is not the specific interval, but the process that works 95% of the time.

provide evidence that more than 1 in 5 volunteers assigned to clean wildlife develop adverse respiratory symptoms, but the data is consistent with this claim.

We can say that the study provides evidence the proportion of volunteers who develop symptoms is less than 0.5; the study provides evidence the proportion of volunteers who develop symptoms is larger than 0.1. That is, the study provides evidence that more than 10% of volunteers develop adverse symptoms and provides evidence this percentage is not larger than 50%. How do we know? Because values less than 10% are not reasonable values of the parameter based on the 95% confidence interval. Values like 0.1 are outside of the confidence interval and are therefore not reasonable. Similarly, values above 0.5 are outside the confidence interval and are therefore not reasonable.

The power of a model for the sampling distribution is that it allows us to determine which values of a parameter are reasonable and which values are not.

7 Quantifying the Evidence (Rejecting Bad Models)

Again, the goal of statistical inference is to use the sample as a snapshot of the underlying population (Figure 1.2). Recall that there are essentially two categories of questions we ask when trying to perform inference:

- Estimation: for example, what *proportion* of volunteers who clean wildlife following an oil spill experience adverse respiratory symptoms?
- Hypothesis Testing: is it reasonable no more than 1 in 5 volunteers who clean wildlife following an oil spill will experience adverse respiratory symptoms; or, is there evidence more than 1 in 5 volunteers who clean wildlife following an oil spill will experience adverse respiratory symptoms?

In Chapter 6, we addressed these questions through the use of confidence intervals — by specifying reasonable values of the parameters through a model of the sampling distribution. However, when our goal is testing a specific hypothesis, there is a second approach; this latter approach is useful when the research question involves multiple parameters and creating a confidence interval is challenging (see Unit III).

In Chapter 3, we described hypothesis testing as being similar to performing a trial in a course of law. Once the prosecution and defense have each presented their case, the jury deliberates and makes one of two decisions:

- 1. Vote "guilty." This happens when the jury believes the facts of the case are not consistent with an innocent defendant; therefore, the prosecution has convinced the jury ("beyond reasonable doubt") that their initial assumption of innocence is not warranted.
- 2. Vote "not guilty." This happens when the jury believes the facts of the case are consistent with an innocent defendant; that is, while the prosecution presented a case that may have been consistent with a guilty defendant, it is still "reasonable" that the defendant is innocent. That is, the prosecution has not convinced the jury that their initial assumption of innocence is unwarranted.

Similar to a jury, we have a working hypothesis (the null hypothesis); we must quantify the evidence from the sample against this hypothesis. If we are convinced that the data is not consistent with this working hypothesis, we will declare we have evidence for the alternative hypothesis. Our goal in this section is to quantify that evidence.

7.1 Some Subtleties

We have described hypothesis testing as being similar to a U.S. trial. That analogy is also helpful in pointing out some subtleties in how we interpret our results. We take a moment to discuss those subtleties before discussing the process itself in order to avoid erroneous interpretations.

The jury weighs the case under the assumption of innocence. That is, they first develop a working hypothesis (the defendant is innocent). Then, the likelihood of the case against the defendant under this assumption is determined. For example, if a defendant were innocent of murder, it is unlikely to have five eye witnesses stating the defendant was seen standing over the victim, holding the murder weapon, and screaming "I killed him!" Since that case against the defendant does not jive with innocence, the jury convicts. If, however, the only case presented is that five eye witnesses place the defendant in the same city as the victim and the defendant matches the description of someone seen fleeing the crime scene, then the jury would not convict. Why? Because the case presented, while pointing toward guilt, is not overwhelming; these things could have happened by chance alone. Therefore, the case, while consistent with guilt does not provide evidence for guilt.

As in Chapter 6, we are making a distinction between "evidence for" a hypothesis and the data being "consistent with" a hypothesis. Evidence for a particular claim is only established by providing evidence against the opposite statement. However, consistency can be established without disqualifying any other statement; that is, data can be consistent with two opposing claims, but data cannot provide evidence for two opposing claims.



Big Idea

Data can be *consistent* with two opposing claims, but data cannot provide *evidence* for two opposing claims.

Also notice that a jury saying "not guilty" is not the same as saying "innocent." That is, a lack of evidence to convict does not imply the defendant is innocent. A lack of evidence is simply a lack of evidence. The defendant may still be guilty, but the evidence has just not proven it.



Big Idea

A lack of evidence for a signal is not evidence for a lack of a signal.

Similarly, when performing a hypothesis test, we will weigh the data under the null hypothesis (our working assumption). Then, the likelihood of our data occurring by chance alone under this hypothesis is determined. If that likelihood is small (data is not consistent with the null hypothesis), we can conclude the data supports the alternative hypothesis (guilty). If, however, that likelihood is large (data is consistent with the null hypothesis), we can only conclude that the data is consistent with the hypotheses. We are *not* able to say "supports the null" because that would be like saying a defendant is innocent. We can't prove innocence because we started by assuming it!

7.2 Assuming the Null Hypothesis

Consider the question we have been asking regarding the Deepwater Horizon Case Study discussed in Chapter 2:

Is there evidence that more than 1 in 5 volunteers assigned to clean wildlife develop adverse respiratory conditions?

Remember, we framed this question through statements about a parameter in Chapter 3:

 H_0 : the proportion of volunteers assigned to clean wildlife that develop adverse respiratory symptoms is no more than 0.20.

 H_1 : the proportion of volunteers assigned to clean wildlife that develop adverse respiratory symptoms exceeds 0.20.

Within the sample we observed that 27.8% of volunteers experienced adverse symptoms, which is certainly more than the 0.20; drawing the connection to the courtroom, the 27.8% is the case presented by the prosecution. From this, we see that the data is at least trending toward the alternative hypothesis. Just as a jury has to ask whether the case is overwhelming (no longer consistent with an innocent defendant) or whether the case is consistent with an innocent defendant, we must ask whether the 27.8% is overwhelming (no longer consistent with 1 in 5 volunteers developing adverse respiratory conditions) or consistent with the null hypothesis. After all, it is possible that we just have a strange sample; that is, it is possible our data is a fluke, resulting in an estimate larger than 0.2 by chance alone.

As we discussed in the previous chapter, we expect our estimate to vary to some degree from one sample to another. Essentially, we need to know if 27.8% of volunteers experiencing symptoms is a strong signal that the rate within the population is larger than 0.2 (1 in 5) or whether 27.8% is simply a fluke that might happen due to sampling variability. While we are going to be attacking the question differently in this chapter than the previous, we see that the key is still variability in the estimate. That is, we are back to the Fourth Fundamental Idea of Inference. As stated above, in order to determine evidence for one statement (captured by the alternative hypothesis), we begin by assuming the opposite statement (captured by the null hypothesis) as our working assumption. That is, if we want to know if 27.8% of volunteers experiencing adverse symptoms is "evidence," we need to figure out what we expect to happen if only 1 in 5 volunteers actually develop adverse respiratory symptoms (the statement represented by the equality portion of the null hypothesis).

Consider this last statement. It is equivalent to saying "what type of actions would we expect of an innocent person?" Only when we know what to expect can we determine if the case in front of us is extreme enough to convict. Only when we know what to expect can we determine if the observed sample provides *evidence* in favor of the alternative. Our strategy is to enter a fake world…a world in which exactly 1 in 5 volunteers actually develop respiratory symptoms. That is, we enter a world in which the null hypothesis is true. Now, in this world, how do we know what to expect? We construct the sampling distribution for the proportion under this assumption that the null hypothesis is true; this is known as the **null distribution**.

Definition 7.1 (Null Distribution). The sampling distribution of a statistic *when* the null hypothesis is true.

The null distribution, the last in our *Distributional Quartet*, is a sampling distribution; it is just a sampling distribution for a world in which the null hypothesis is true. As a result, the process for constructing the null distribution is very similar to the process for constructing the sampling distribution (illustrated in Figure 7.1):

- 1. Sample randomly from a fake population where the null hypothesis is true.
- 2. For each sample, compute the statistic of interest.
- 3. Repeat steps 1 and 2 several thousand times.
- 4. Plot the statistics retained from each sample.

Notice that these are the same steps as constructing a sampling distribution with the exception that instead of sampling from the population of interest, we sample from a hypothetical population in which the null distribution is true.

Since the null distribution is a sampling distribution when a particular hypothesis is true, we are constrained by the same limitations as before. Namely, we are generally unable to construct the actual the null distribution; instead, we must construct a model for it. More, since the null distribution is a sampling distribution, the same techniques we use for modeling the sampling distribution can be modified to model the null distribution. The key is to enforce the null hypothesis to be true. As with sampling distributions, we emphasize the empirical model approach.

Using the computer, we first create a virtual world in which the null hypothesis is true. This often involves adjusting the original sample in order to make it consistent with having been drawn from a population in which the null hypothesis were true. The augmented data becomes the null world. We are then able to bootstrap from the augmented data to simulate repeated sampling when the null hypothesis is true. The details of this procedure are beyond the scope of our current discussion; it is more important to understand the conceptual idea of a null distribution at this point.

Figure 7.2 represents a model for the null distribution of the proportion of volunteers in a sample of 54 assigned to clean wildlife which would develop adverse symptoms when the we assume the actual proportion is 0.20.

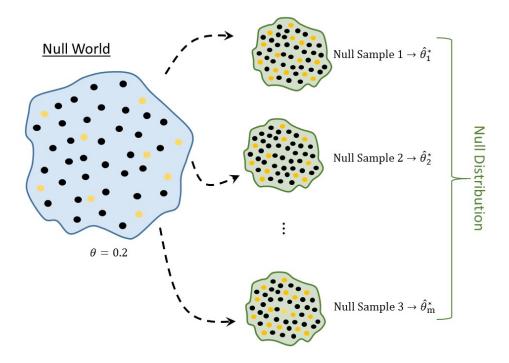


Figure 7.1: Illustration of constructing a null distribution. Notice the similarity to constructing the sampling distributon.

Note

A null distribution is tied to a specific null hypothesis. A sampling distribution does not require a hypothesis to construct. So, while a sampling distribution could be used to address a variety of null hypotheses, a null distribution can only be used to address the corresponding set of hypotheses for which it was developed.

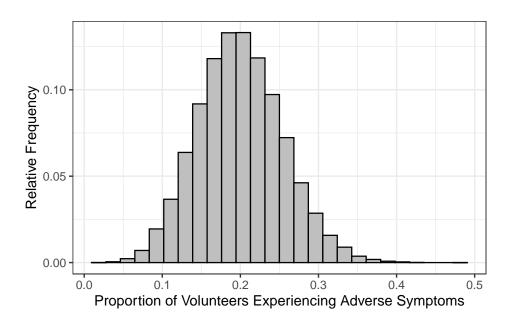


Figure 7.2: Null distribution for the proportion of volunteers assigned to clean wildlife experiencing adverse respiratory symptoms. The null hypothesis is that the proportion is 0.20; this model is based on a sample size of 54.

As we are introducing the concept of a null distribution, we will stick to modeling the null distribution of a statistic. Often times in a statistical analysis, the null distribution is computed for a numerical quantity known as a standardized statistic (see Chapter 12).

7.3 Using the Null Distribution

From Figure 7.2, we see that if the null hypothesis were true — if only 1 in 5 volunteers assigned to clean wildlife experienced symptoms — then in a sample of 54 individuals, we would expect the proportion who experienced symptoms to be somewhere between 0.1 and 0.3. So, we can think of the null distribution as setting up our expectations of what we should have observed if the null hypothesis were true. For example, if the null hypothesis were true, it would be nearly impossible that half of the individuals experienced symptoms. We are able to say that since 0.5 is way off in the tail region of the distribution, and we know that values in

the tail region are more extreme. The question is then how extreme is *our* sample? Again, the null distribution is just setting up expectations; to determine if our sample provides evidence against the null hypothesis, we have to weigh the data against those expectations.

In our sample, we observed 27.8% of volunteers experience symptoms. Since 0.278 is towards the center of the distribution, we would say that it is not an extreme sample. In order to quantify how extreme (or not extreme) it is, we compute the fraction of values which are more extreme (larger in this case) than the value observed; that is, we compute the fraction of values that appear further to the right than 0.278 in the null distribution. Figure 7.3 illustrates this computation. Based on the null distribution, there is a 10.6% chance that if the null hypothesis were true — only 1 in 5 volunteers actually experienced symptoms — that in a random sample of 54 volunteers we would obtain data this extreme or more so (an observed proportion of 0.278 or larger) by chance alone (due to sampling variability). This area is known as the **p-value**.

Essentially, this tail area is quantifying the strength of the evidence. The smaller this area, the further in the tail region our observed statistic is; that is, the smaller this area, the more unexpected our data. Therefore, small areas indicate that the data (our case) does not jive with our expectations under the null hypothesis (innocence), forcing us to conclude the data provides evidence against the null hypothesis (guilty verdict). On the other hand, the larger this area, the closer toward the center our observed statistic is; that is, the larger this area, the less unexpected our data. Therefore, large areas indicate that the data (our case) is consistent with our expectations under the null hypothesis (innocence), forcing us to conclude the data is consistent with the null hypothesis (not guilty verdict).

Definition 7.2 (P-Value). The probability, assuming the null hypothesis is true, that we would observe a statistic, from sampling variability alone, as extreme or more so as that observed in our sample. The p-value quantifies the strength of evidence against the null hypothesis, with smaller values indicating stronger evidence.

For the Deepwater Horizon Case Study, we observed 27.8%, which is towards the center of the null distribution, and our p-value is quite large. Therefore, our data is consistent with what we might expect if the null hypothesis were true, and we conclude that there is no evidence that the rate of those experiencing symptoms exceeds 1 in 5.

It is natural to ask "how small does the p-value need to be to prove a statement?" Like a trial, the strength of the case presented depends on the context. In some studies, a p-value less than 0.01 may be strong evidence while in other studies a p-value less than 10^{-6} is required. And, as in a trial, it is not only the strength of the case but the type of information presented (the presence of DNA may carry more weight than the presence of fingerprints). In statistics, it is important to consider the *effect size* (some measure of the signal in the data) as well as the p-value. That is, consider whether the difference between the estimate and the null value is actually large for the specific context; this is always based on subject-matter expertise. It is often helpful, whenever possible, to report a confidence interval alongside a p-value.




Figure 7.3: Likelihood of obtaining a statistic as extreme or moreso as that of the original sample when the parameter of interest is the proportion of volunteers assigned to clean wildlife experiencing adverse respiratory symptoms. The null hypothesis is that the proportion is 0.20; this is based on a sample of size 54.

Gradient of Evidence

While what constitutes evidence may vary from discipline to discipline, the list below is a good rule of thumb:

- p > 0.1: no evidence against the null hypothesis.
- $0.05 \le p < 0.1$: weak evidence against the null hypothesis.
- 0.01 : some evidence against the null hypothesis.
- $0.001 \le p < 0.01$: evidence against the null hypothesis.
- p < 0.001: strong evidence against the null hypothesis.

Of course, evidence "against the null hypothesis" is equivalent to evidence "for the alternative hypothesis." As with any rule of thumb, the above gradient should not be considered binding and may vary depending on the application.

We can think of the p-value as a measure of whether the data is able to discern the difference between a hypothesized value (the null value) and that observed in the sample. When the p-value is sufficiently low, we might say there is a "statistically discernible" difference.

Warning

You may see some authors refer to a "significance level," often denoted by the Greek letter alpha. The significance level is a pre-determined threshold for declaring a p-value to be "small." That is, p-values below the significance level indicate there is evidence against the null hypothesis, and p-values above the significance level indicate there is no evidence against the null hypothesis.

This has led to authors stating that a study shows a "statistically significant" difference. We dislike this language because the term "significant" suggests the difference is important. However, we may be able to detect a difference that is not meaningful in real life. Determining importance is a matter for discipline experts.

We prefer the term "statistically discernible" and the gradient of evidence provided above instead of the binary decision that results from defining a significance level.

As Regina Nuzzo once quipped¹, we can think of the p-value as an "index of holy shitness." Essentially, small p-values tell you the data was really unexpected under the null hypothesis ("wow, that was unexpected"); and, larger p-values tell you the data was about what you would expect under the null hypothesis ("yeah, that was about what we expected"). When your data surprises you (from the perspective of the null hypothesis), then we start to believe that our underlying assumption about that data (the null hypothesis) is not appropriate.

¹Stated in a keynote presentation at USCOTS23 (https://www.causeweb.org/cause/uscots/uscots23/ keynotes/1).

Big Idea

A p-value should never be reported in isolation. It should always be accompanied by a confidence interval, a numerical summary of the data, or a graphical summary of the data — something which indicates the effect size and variability in the data.

Note

Here, we refer to the "effect size" as the impact of the signal in the data. However, some disciplines rely on standardized effect sizes which are unit-less measurements. These measurements are typically a ratio of the difference between the observed statistic and the null value to a measure of the variability in the data.

Regardless of whether you use the raw effect size or a standardized effect size, the idea is the same — you should understand whether the signal you are detecting in the data is of practical relevance in your discipline.

Process for Computing a P-Value

The following is a general procedure for computing a p-value:

- 1. Define the null and alternative hypotheses.
- 2. Construct a model for the null distribution of the desired statistic.
- 3. Compute the desired statistic for the original sample.
- 4. Overlay the statistic from step (3) on the model developed in step (2), and then compute the area under the curve for values *more extreme* than that observed in step (3).

Notice that the definition of a p-value, and this general procedure, apply regardless of the technique used for constructing the model of the null distribution.

Like confidence intervals, p-values are often misinterpreted. In fact, they have become so abused that some researchers argue against their use. It is our opinion that the p-value can be a useful tool once it is appropriately understood; so, let's dispel some of the misconceptions.

7.4 Incorrect Interpretations of a P-Value

Suppose we have a p-value of p. The following are incorrect interpretations of that p-value:

- There is a 100p% chance that the null hypothesis is correct.
- When p is large, we have evidence (or "the data supports") the null hypothesis is correct.

The first statement is incorrect because it assumes the parameter is moving. Hypotheses are statements about parameters, and parameters are fixed, unknown values. One of the hypotheses is true, and the other is not. Our ignorance does not change this. Since the parameter is a fixed value, it does not make sense to think about the likelihood that either hypothesis is correct. Instead, the p-value quantifies the likelihood of observing a particular statistic assuming the null is true.

The second statement has circular logic. Again, when computing the p-value, we assumed the null hypothesis was correct; so, we cannot prove what we start off assuming. It mistakenly assumes that a lack of evidence for the alternative hypothesis is evidence for the null hypothesis, but that is not true. A large p-value simply indicates the data is consistent with the null hypothesis, but a large p-value cannot provide evidence to support the null hypothesis. This is similar to the fact that in a trial, the jury can declare "not guilty" (you did not convince us of guilt), but they do not declare "innocent."

For the Deepwater Horizon Case Study, we obtained a p-value of 0.106. Considering the above, the following are **incorrect** interpretations:

- There is a 10.6% chance that only 1 in 5 volunteers assigned to clean wildlife will experience adverse symptoms.
- Since the p-value is large, there is evidence (or the data supports the claim) that 1 in 5 volunteers assigned to clean wildlife will experience adverse symptoms.

Thinking about the p-value correctly, we are led to the following conclusion:

• With a p-value of 0.106, our data is *consistent with* only 1 in 5 volunteers assigned to clean wildlife experiencing adverse respiratory symptoms.

Unfortunately, our data is also *consistent with* more than 1 in 5 volunteers assigned to clean wildlife experiencing adverse respiratory symptoms. The fact that our data is consistent with both hypotheses may be an unsatisfying conclusion, but it is still a conclusion nonetheless.

This process of using a p-value to quantify the evidence against a particular model is captured in our last of the *Five Fundamental Ideas of Inference*.

I Fundamental Idea V

With a model for the distribution of a statistic under a proposed model, we can quantify the the likelihood of an observed sample under that proposed model. This allows us to draw conclusions about the corresponding parameter, and therefore the population, of interest.

7.5 Sampling Distributions vs. Null Distributions

Clearly the sampling distribution and null distribution of a statistic are closely related. The difference is that the null distribution is created under a proposed model while the sampling distribution lets the data speak for itself. It is worth taking just a moment to highlight the differences in the use of these two components of the Distributional Quartet.

The sampling distribution is centered on the true value of the parameter; the null distribution is centered on the null value. Once we assume the null hypothesis is true, we have a value for the parameter; as a result, we expect the sampling distribution under this assumption (that is, the null distribution) to be centered on this hypothesized value. So, null distributions of a statistic are always centered on the null value.

When we model the sampling distribution, this leads to a confidence interval. That confidence interval specifies the reasonable value of the parameter based on the observed data.

When we model the null distribution, we are able to compute a p-value. The p-value quantifies how likely our observed statistic is under the proposed null hypothesis.



Big Idea

Model the sampling distribution to construct a confidence interval; to assess a hypothesis the null value is overlaid on the model for the sampling distribution. Extreme values of the model for the sampling distribution are unreasonable values for the parameter. Model the null distribution to compute a p-value; to assess a hypothesis, the statistic from the sample is overlaid on the model for the null distribution. Extreme values of the model for the distribution are values which provide evidence against the null hypothesis.

8 Using the Tools Together

In this unit, we have introduced the key components in both the language and logic of statistical inference. In fact, with a firm grasp of the concepts in this unit, you should be able to read and interpret key statistical findings. All statistical analyses make use of the *Five Fundamental Ideas of Inference* and alternate between the components of the *Distributional Quartet*. The context of each problem differs, but the logic remains the same. In this chapter, we present another analysis based on the Deepwater Horizon Case Study of Chapter 2, annotating it along the way to see how these elements work together fluidly to reach a conclusion. Specifically, we are interested in the following question:

Are volunteers assigned to clean wildlife at higher risk of developing adverse respiratory symptoms compared to those volunteers who do not come into direct contact with oil? If so, estimate the increased risk.

8.1 Framing the Question (Fundamental Idea I)

We are really interested in whether the rate of respiratory symptoms in one group of volunteers is larger than that in a second group. Therefore, our working assumption is that the rate of respiratory symptoms for those assigned to clean wildlife is no more than that for those assigned to tasks which do not involve direct exposure to oil. That is, we have

 H_0 : the rate of adverse respiratory symptoms for volunteers assigned to clean wildlife is no greater than that for those assigned to tasks which do not involve direct exposure to oil.

 H_1 : the rate of adverse respiratory symptoms is greater for volunteers assigned to clean wildlife compared to those assigned to tasks which do not involve direct exposure to oil.

We can also state this more formally with mathematical notation as follows:

Let θ_1 be the rate of developing adverse respiratory symptoms for volunteers assigned to clean wildlife.

Let θ_2 be the rate of developing adverse respiratory symptoms for volunteers assigned to tasks without direct exposure to oil.

 $H_0: \theta_1/\theta_2 \le 1$ $H_1: \theta_1/\theta_2 > 1$ The ratio θ_1/θ_2 is known as the *relative risk* as it captures the increased risk for one group compared to another.

Notice that this is a well-posed question as it centers on parameters which characterize the population. Therefore, it can be answered with appropriate data.

Distribution of the Population: Our questions of interest are about the population and therefore focus on characterizing this distribution.

8.2 Getting Good Data (Fundamental Idea II)

As we are working with previously collected data, we are unable to design a good sampling scheme. The only thing we can do at this point is critique the sample we have. The key question to ask ourselves is whether there is any reason that this group of volunteers differs systematically from other volunteers working oil spills. For example, this oil spill occurred in the Gulf of Mexico; the majority of volunteers were then naturally residents of Gulf states. It is possible that these residents are somehow fundamentally different with respect to their risk of developing adverse respiratory symptoms compared to the remainder of the United States. If that is the case, the results of this study would not generalize to oil spills occurring in the Atlantic. However, it is probably reasonable to say that these results would apply to future oil spills in the Gulf. If, on the other hand, we believe this group of volunteers is representative of volunteers for other oil spills, regardless of location, our results could generalize more broadly.

Also note that this was not a controlled experiment. Volunteers were not randomly allocated to their assignments that we know of. Therefore, our results could be somewhat limited. The two groups should be compared regarding other attributes (this data is unavailable to us currently) in order to determine if they are similar with respect to other variables which may potentially confound the results. If confounding is a concern, we would not be able to conclude that any observed differences were *caused* by the exposure to oil; it could be that volunteers who choose assignments which bring them into contact with oil also share some trait which puts them at higher risk of respiratory symptoms.

8.3 Presenting the Data (Fundamental Idea III)

The heart of this question is comparing the rate of adverse events in each group. Figure 8.1 makes this comparison.

As seen in Figure 8.1, the rate of adverse respiratory symptoms was larger in the group of volunteers assigned to wildlife cleanup. Specifically, the rate of respiratory symptoms was 1.79 times higher in the volunteers assigned to clean wildlife compared to those assigned to tasks with no direct oil exposure.




Figure 8.1: The risk of developing adverse respiratory symptoms for volunteers assigned to clean wildlife is higher than that for those volunteers assigned to tasks which do not have direct exposure to oil.

Notice that we reported the relative risk comparing the two groups as it is directly tied to how we specified the hypotheses above. That is, the statistic we report is governed by the parameter of interest; we compute a value in the sample to estimate the corresponding value in the population.

Distribution of the Sample: graphics and numerical summaries characterize this distribution, informing us about the underlying population. This is possible as long as the sample is representative of the population.

8.4 Quantifying the Variability in the Estimate (Fundamental Idea IV)

While we have an estimate for the increased risk of adverse respiratory symptoms for those volunteers assigned to clean wildlife, the estimate has not taken into account sampling variability. In order to quantify this variability, we use a bootstrap procedure to model the sampling distribution of the relative risk. Observe that we focus on the sampling distribution of the statistic that estimates the parameter of interest.

Recall that bootstrapping mimics the process for generating a sampling distribution. In this case, "repeating the study" involves collecting data from not one, but two groups. So, we

must resample both from the 54 volunteers who were assigned to clean wildlife and the 103 volunteers assigned to tasks not involving direct oil exposure. Each time we resample, we ensure that we select 54 volunteers who clean wildlife and 103 who do not (mimicking the original study). We need the process of the original study to be maintained. Each time we resample from these groups, we compute the relative risk and retain this value. Figure 8.2 shows the model for the sampling distribution for the relative risk comparing these two groups. Again, it is important to note that we are not generating new data; we are resampling (or reusing) the original sample.




Figure 8.2: Model of the sampling distribution for the relative risk comparing volunteers assigned to clean wildlife to volunteers assigned to tasks not involving oil exposure. The model was developed via bootstrapping using 50000 replications.

The study suggests that volunteers assigned to clean wildlife are 1.79 times (90% CI = (1.04, 3.12)) more likely to experience adverse respiratory symptoms compared to those volunteers assigned to tasks not requiring direct exposure to oil. Our data is consistent with volunteers assigned to clean wildlife being at increased risk compared to those who do not have direct exposure to oil.

Sampling Distribution: allows us to quantify the variability in the statistic and provide an interval estimate for the parameter which incorporates this variability.

8.5 Quantifying the Evidence (Fundamental Idea V)

In order to quantify the departure of the data from our working assumption that the risk is for those assigned to clean wildlife is no more than that for those assigned to tasks without direct oil exposure, we rely on a model for the null distribution and compute a p-value.




Figure 8.3: Model of the null distribution for the relative risk comparing volunteers assigned to clean wildlife to volunteers assigned to tasks not involving oil exposure. The null hypothesis assumed the two groups of volunteers had the same risk. The null distribution was developed via bootstrapping using 50000 replications.

There is some (borderline weak) evidence (p = 0.04) to suggest that volunteers exposed to oil have an increased risk of developing adverse respiratory symptoms. Given the estimated level of this increased risk (see the previous section for the confidence interval), we believe this is something health officials should investigate further. It would be worth investigating what aspects of the oil exposure may have led to the increased risk to determine if it can be avoided in the future.

Note we are careful to not claim that the assignments have caused an increase in the risk as this data is not from a controlled experiment. This is one of the limitations of this analysis. However, if we are able to assume the two groups are fairly similar with respect to other attributes — that is, there is no reason why people prone to respiratory symptoms would be more likely to be assigned to wildlife cleaning — then we may have some reason to believe the results are causal. We will wrestle more with these types of conclusions in future units.

Null Distribution: allows us to quantify the level of evidence against a particular claim or working hypothesis.

8.6 Summary

Notice that our analysis moved through the *Five Fundamental Ideas*, and in doing so made use or referenced each of the four components of the *Distributional Quartet*. As we move through the remainder of the text, we will explore how these frameworks are used in various other analysis scenarios. As we do, we reveal additional concepts involved in statistical modeling.

We admit that there are several other questions that may be raised by the above analysis. This unit is meant to introduce the big concepts of inference. We will concern ourselves more with the details as we progress through the text.

Part II

Unit II: Inference on the Overall Mean Response

In Unit I, we developed the language and logic of statistical inference. In this unit, we apply those concepts toward making statements about the mean response of a variable. Throughout, we will take a modeling perspective, using this specific context to develop key concepts around statistical models. This will serve as a bridge between the big ideas in Unit I and the more interesting models in the remainder of the text.

9 Case Study: Birth Weights of Babies

The Centers for Disease Control and Prevention (CDC) — using data provided by the U.S. Department of Health and Human Services, National Center for Health Statistics, the Division of Vital Statistics and the CDC — maintains a database on all babies born in a given year¹. This database contains key metrics on each child born, including the weight of the infant. Low birth weight can be indicative of poor health or illness in children; and, high birth weight can be indicative of obesity later in life. Researchers, therefore, use the database to examine links between lifestyle choices of the parents (such as whether the mother consumed alcohol during pregnancy) and birth weight of the infant.

Chihara and Hesterberg (2011) describe a random sample from this database; specifically, the sample consists of 1009 babies born in North Carolina during 2004. The babies each had a gestation period of at least 37 weeks (full term) and were single births (no twins, triplets, etc.). For each birth in the sample, we have the following information:

- Age of the mother (in years).
- Whether the mother is a current smoker.
- Whether the mother used to bacco during the pregnancy.
- Whether the mother used alcohol during the pregnancy.
- Sex assigned to the child at birth (male or female).
- Weight of the child at birth (grams).
- Length of gestation (length of pregnancy, weeks).

A subset of the collected data is shown in Table 9.1.

We might be interested in using this data to estimate the average birth weight of an infant (carried to full term) born in North Carolina.

¹http://wonder.cdc.gov/natality-current.html

Table 9.1: Subset of a sample of 1009 babies born in North Carolina during 2004.

Subject ID	Age Range (years)	Assigned Sex of Baby	Weight of Baby (g)	Gestation (weeks)
1	30-34	Male	3827	40
2	30-34	Male	3629	38
3	35-39	Female	3062	37
4	20-24	Female	3430	39
5	25-29	Male	3827	38
6	35-39	Female	3119	39
7	20-24	Female	3260	40
8	20-24	Male	3969	40
9	20-24	Male	3175	39
10	25-29	Female	3005	39

10 Model for the Data Generating Process

The numerical summaries of any study are subject to sampling variability. That is, if we were to repeat the study with new subjects, the statistics we compute would almost certainly change to some degree. The key to feeling confident in our results is to quantify the variability in our estimates; this was the rationale for developing a model of the sampling distribution of a statistic (Chapter 6) and a model for the null distribution of a statistic under a specified hypothesis (Chapter 7). Often times, constructing such models requires modeling the data-generating process as a precursor. As in any other discipline, statistical models simplify the actual process being modeled by making certain assumptions. In this chapter, we develop a model for the data generating process that will help us make inference about the mean of a single population. This will also serve as the foundation for how we approach problems in other contexts moving forward in the text.

i Note

The word "model" is generic, and it applies to any representation of a complex process. We will either explicitly describe the model we are discussing (e.g., the model for the sampling distribution of a statistic), or it should be clear from context. Regardless, it is important to know which model is being referenced.

10.1 General Formulation

Consider dropping a tennis ball from the top of a 50-meter building and recording the time required before the ball hits the ground. Applying the principles learned in a first course in physics, we would be able to compute the time precisely using the formula

time =
$$\sqrt{\frac{2(\text{distance})}{9.8}}$$

where $9.8m/s^2$ is the acceleration due to gravity; further, this formula works regardless of the mass of the object. Plugging 50 meters into the equation yields a time of 10.2 seconds. If we were to drop a second tennis ball from the same building, the formula tells us that it will also take 10.2 seconds to hit the ground below. This is known as a **deterministic** system since entering the same input results in the same output each time.

Definition 10.1 (Deterministic Process). A process for which the output is completely determined by the input(s). That is, the output can be determined with certainty.

While we often think of the above formula as a "law," we should recognize that it is just a model. It simplifies extremely complex processes involving the gravitational pull between objects; it makes simplifying assumptions, such as the gravitational pull of your hand when you release the ball is negligible compared to the gravitational pull of the earth. More, the model works really well. Despite how well it works, it does not always match reality. If we were to repeatedly drop tennis balls from the same 50-meter building and record the time before hitting the ground, we might find that the time differs slightly from one ball to the next (it is true that these differences may be negligible, but they would exist nonetheless).

There are several reasons why our observed responses do not line up directly with those predicted by the above equation; for example, our device for measuring time may be subject to some measurement error, a strong gust of wind could alter the results (while the above equation assumes no air resistance), or the person dropping the ball may have inadvertently increased the initial velocity of the ball. These reasons, and others, contribute to the observations not lining up exactly with the model. That is, there is associated noise in the resulting measurements. A model which incorporates this noise might be written as

$$time = \sqrt{\frac{2(distance)}{9.8}} + noise$$

where the noise is not a known quantity. In fact, the noise is not a constant as it varies from one observation to the next! As a result, the output of the model (time) is no longer fully determined by the input (distance); the output also depends on a random component, meaning the same input could result in different outputs. This is known as a **stochastic** model.

Definition 10.2 (Stochastic Process). A process for which the output cannot be predicted with certainty.

In the above example, the deterministic model was extended to include a stochastic component. This extension leads to our general formulation for a statistical model used in this text:

Response = function(predictor variables, parameters) + noise.
$$(10.1)$$

The response we observe is the result of two components:

- A deterministic component that is a function of predictor variables and unknown parameters. It is often this component on which we would like to make inference.
- A stochastic component that captures the unexplained variability in the data generating process.

Since the noise is a random element which varies from one observation to another, it has a distribution. We often place conditions on the structure of this distribution to enable inference on the deterministic component of the model. We discuss this later in the chapter.

This general model introduces what we will see as a theme in statistical modeling — partitioning the variability in the response. The model says that part of the reason the response differs across units (or subjects) is because those units have different values of the predictor, and part of the reason the response differs across units is unexplained random noise. Notice that the response is also governed by the value of the parameters; however, as discussed in Chapter 3, parameters are numeric constants, so do they not vary from one unit to another. The above model just makes explicit how the parameters characterize the population — through the model for the data generating process. As before, while the parameters are constants, they are unknown. Our goal is to use data to make inference on the parameters and therefore make inference on how the data is generated.

The overall goal of a statistical model is to give an explanation for why the value of the response is what it is. How did it come to be? What process generated the values we have observed? Our statistical model says that these values have some deterministic component plus some additional random noise we cannot explain.

We now simplify this general formulation for the specific case of making inference on the population mean.

Note

We note that the above formulation assumes a quantitative response variable, which is the focus of this text. If your response is categorical, you can think of the response as simply the output of a lottery where the number of outcomes and the likelihood of each outcome is determined by a function of the predictor variables and parameters.

That is, the process is still stochastic, but the predictor variables and parameters impact the likelihood of a particular value being observed:

Probability(Response is y) = function(predictor variables, parameters, y)

10.2 Statistical Model for a Quantitative Response with No Predictors

Consider the Birth Weights Case Study described in Chapter 9.. Suppose we are interested in estimating the average birth weight of infants (carried to full term) born in North Carolina, the population from which our sample was taken. Our response variable is the birth weight of the infant. Our question of interest is not about the relationship of the birth weight to any other variable; that is, there are no predictor variables being considered. But, that does

not mean the deterministic portion of our model is empty. We have a parameter of interest: the average birth weight. This parameter lives in the deterministic portion of the model. In particular, consider the following model for the data generating process:

$$(Birth Weight)_i = \mu + \varepsilon_i$$

where μ represents the average birth weight of infants born in North Carolina. In this model for the data generating process, the function that represents deterministic component of Equation 10.1 takes the value μ , a constant. The term ε_i is used to capture the random noise in the *i*-th measurement; the subscript indexes the individual infants in the sample, indicating that the noise in the birth weight for each infant potentially differs. This ε_i captures the difference between the birth weight for the i-th infant and the overall mean birth weight of all infants. This model says that the birth weight for the i-th infant is shifted (as a result of the noise term) from the overall average birth weight μ . Notice that if there were no noise in the system, the data generating process would say that all infants have the same birth weight μ . However, due to genetic variability, differences in the lifestyle of each mother, and measurement error, ε_i is not a constant (noise does exist), resulting in each subject having a different response.

Notice that the deterministic portion of the model describes the mean response through the parameter. This will be a running theme in the models we consider in this text.



🅊 Big Idea

The deterministic portion of the data generating process for Equation 10.1 describes the mean response; naturally, it is often referred to as the mean response function. Note that the mean response function is governed by parameters.

Certainly, this model for the data generating process is a simplified version of reality. Instead of capturing the impact of genetics or the impacts of lifestyle choices made by the infant's mother on the infant's birth weight, we simply allow those impacts to enter the random noise component of the model. This does not make the model for the data generating process incorrect; instead, it simply means our model will not explain the variability in the response. In later chapters we will begin to address this limitation. However, it is worth noting now that all models for the data generating process are simplified versions of reality, but that does not mean they are not helpful.

When the model for the data generating process does not contain a predictor variable, we are saying that the only source of variability in the response is random noise.



Big Idea

The stochastic component of a statistical model captures the unexplained variability due to natural variability in the population or measurement error in the response.

Data Generating Process for Single Mean Response

In general, given a quantitative response variable and no predictors, our model for the data generating process is

$$(Response)_i = \mu + \varepsilon_i \tag{10.2}$$

where μ represents the average response in the population, the parameter of interest.

It is worth pointing out that we have two "models" at this point: a model for the data generating process and a model for the sampling distribution of a statistic. The model for the data generating process is used to develop a model for the sampling distribution (or null distribution) of a statistic. It is the second model that is actually necessary in order to conduct inference; the model for the data generating process is simply a stepping stone to the model of interest.

10.3 Conditions on the Error Distribution

In our model for the data generating process we incorporated a component ε to capture the noise observed in the response. Since the error is a random variable (stochastic element), we know it has a distribution. We typically impose a certain structure to this distribution through the assumption of specific conditions. The more conditions we impose, the easier it is to construct an analytical model for the sampling distribution of the corresponding statistic. However, the more conditions we impose, the less applicable our model is in a general setting. More importantly for our discussion, the conditions we impose dictate how we conduct inference (the computation of a confidence interval or p-value).

Note

Why we need conditions on the stochastic portion of a model can be confusing at first. Think of it this way: saying a term is "random" is just too broad. It is like saying "I am thinking of a number. What number?" There are too many choices to even have a hope of getting it correct. We need structure (boundaries, conditions) on the problem. Placing conditions on what we mean by "random" is like saying "I am thinking of a whole number between 1 and 10." Now, we have a problem we can at least attack with some confidence.

The first condition we consider is that the noise attributed to one observed unit is **independent** of the noise attributed to any other unit observed. That is, the amount of error ε in any one response is unrelated to the error in any other response observed. In context, the error in the birth weight of one infant is unrelated to the error in the birth weight of any other infant.

It is easiest to understand this condition by examining a case when the condition would not hold.

Definition 10.3 (Independence). Two random variables are said to be independent when the likelihood that one random variable takes on a particular value does not depend on the value of the other random variable.

Similarly, two observations are said to be independent when the likelihood that one observation takes on a particular value does not depend on the value of the other observation.

Example 10.1 (Tire Rotation Timing). Suppose we are conducting a study to estimate the amount of time, on average, required for a novice technician to complete a tire rotation on an automobile. We gather a sample of 25 novice technicians. We have the first technician complete a tire rotation, allowing the other technicians to watch, and record the time required to complete the task. We then have the second technician complete a tire rotation, again allowing the other technicians to watch, and record the time required to complete the task. We continue in this way until we have recorded the time required to complete a tire rotation for each of the 25 technicians.

We could model the data generating process for this as

$$(Time)_i = \mu + \varepsilon_i$$

where μ is the average time required to complete a tire rotation for a novice technician. We might estimate the parameter μ with $\frac{1}{25} \sum_{i=1}^{25} (\text{Time})_i$, the sample mean time for the 25 observed technicians.

In Example 10.1, it would *not* be reasonable to assume the errors are independent of one another. Since each technician is able to observe the prior technicians, it is plausible that each technician's performance is dependent on their observations (they are learning tricks by watching others go before them). As a result, the amount of noise in one technician's time is related to the amount of noise in the next technician's time. Requiring independence among observations prohibits these types of situations.

A second condition we consider is that the error for each subject is **identically distributed**. This ensures that every student essentially belongs to the same population.

Definition 10.4 (Identically Distributed). A set of random variables is said to be identically distributed if they are from the same population.

Similarly, a set of observations is said to be identically distributed if they share the same data generating process.

Practically, this condition means that we do not have a systematic component which is causing our population to be different from what we expected. As an example, let's return to Example 10.1.

Example 10.2 (Tire Rotation Timing, Continued). Suppose we are conducting a study to estimate the amount of time, on average, required for a novice technician to complete a tire rotation on an automobile. We gather a sample of 25 novice technicians. All technicians are permitted to begin training in the facility. During the first week of training, we have the first technician complete a tire rotation and record the time required to complete the task. During the second week of training, we have the second technician complete a tire rotation and record the time required to complete the task. We continue in this way until we have recorded the time required to complete a tire rotation for each of the 25 technicians.

Again, we might posit the same model for the data generating process as in Example 10.1, and we might estimate the average time in the same way.

It would *not* be reasonable to assume the errors across all observations are from the same underlying population in Example 10.2. As the technicians are continuing their training over time, the group under study is changing over time. A technician with 10 weeks of training should not be expected to perform in the same way as a technician with 1 week of training. The observations are being drawn from a population that is changing over time. Therefore, we could not say that the underlying distribution is identical for each observation.

While there are many other conditions we might consider imposing on the data generating process, these two conditions (independence and identically distributed) are sufficient for modeling the sampling distribution of the sample mean (our statistic of interest here) using the bootstrap process we described in Chapter 6. In future chapters, we will consider the impact of adding additional conditions or relaxing these conditions.

Note

While not all authors make the distinction, we distinguish between a "condition" and an "assumption." A *condition* is a mathematical requirement for justifying the statistical theory on which we are relying. However, in practice, we are never able to guarantee a condition is satisfied.

As a result, we will need to make certain *assumptions* about the data generating process. In future units, we will see how we determine which assumptions are reasonable. At this point, we simply emphasize that the assumptions we make govern the process by which we model the sampling distribution.

Throughout this section, we have been discussing the conditions placed on distribution of the error term — the stochastic portion of the model. We emphasize that these conditions are placed on the *distribution*, not the values themselves. Further, notice that Equation 10.2 states the response is formed by taking the error and shifting it by the constant μ ; as a result,

the distribution of the responses is simply a shifted version of the distribution of the error terms. Figure 10.1 illustrates this for a hypothetical population. Since the distribution of the response is just a shifted version of the distribution on the errors, any conditions placed on the distribution of the errors can be translated to statements on the distribution of the response.




Figure 10.1: Relationship between the distribution of the response and the distribution of the error term in the model for the data generating process involving a single mean response.

Note

For the model described in Equation 10.2, stating

- the errors in the response for one observation is independent of the error in the response for all other observations, and
- the errors in the response are identically distributed

is equivalent to stating

- the response for one observation is independent of the response for all other observations, and
- the responses are identically distributed.

Before leaving this chapter, it is worth noting that the introduction of a model for the data generating process does not change any of the fundamentals of inference we have previously discussed. This chapter simply introduces a framework that allows us to unify all the methods we discuss in this text.

11 Estimating with Confidence a Single Mean

Consider the Birth Weight Case Study described in Chapter 9. In the previous chapter, we introduced the following model for how birth weights for infants are generated:

$$(Birth Weight)_i = \mu + \varepsilon_i$$

where μ represents the average birth weight of infants born in North Carolina. We also discussed imposing two conditions on the distribution of the random error terms:

- The error in the birth weight for one infant is independent of the error in the birth weight for any other infant.
- The errors in the birth weight for infants are identically distributed.

Within this model for the data generating process, μ is an unknown parameter of interest. Consider the following research goal associated with this parameter:

On average, what is the birth weight of an infant born in North Carolina?

We can construct a point estimate of the parameter μ with the average birth weight of infants in our sample: 3448.26 g. The data is graphically summarized in Figure 11.1.

In order to construct an estimate of μ which incorporates the variability in the sample mean, we must model the sampling distribution of our estimate. The bootstrap procedure for this case would be

- 1. Randomly sample, with replacement, 1009 records from the original sample.
- 2. For this bootstrap resample, compute the mean birth weight and retain this value.
- 3. Repeat steps 1 and 2 many (say 5000) times.

This process is illustrated in Table 11.1. Each row represents the birth weights for a single resample taken with replacement from the original data. The final column is the computed (and retained), sample mean from each resample (the bootstrap statistic).

A plot of the resulting bootstrap sample means is shown in Figure 11.2. Notice that the x-axis is different from that of Figure 11.1. While a graphical summary of the raw data is summarizing the weight of individual infants, the model for the sampling distribution is summarizing the




Figure 11.1: Weight of a sample of full-term infants born in North Carolina.

Table 11.1: Partial print out of first 10 bootstrap resamples and the resulting bootstrap statistic.

Value 1	Value 2	Value 3	Value 1007	Value 1008	Value 1009	Boostrap Mean
3345	3572	3572	 3827	3827	3119	3461.89
3629	2892	3827	 4111	3374	2948	3476.69
2495	3686	3827	 3289	3544	3487	3428.90
3856	3430	3771	 3487	3742	2665	3436.20
3430	3119	4479	 3686	3090	3005	3451.09
3289	3459	3827	 3600	3856	3260	3473.89
2863	3345	3232	 3345	3544	2948	3427.89
3289	4026	3856	 4338	3771	3714	3435.78
3175	3544	3771	 3572	3515	3005	3419.37
3260	3771	3742	 3572	4054	3033	3447.77




Figure 11.2: Bootstrap model for the sampling distribution of the average birth weight for a sample of 1009 infants born in North Carolina.

statistic we compute in various resamples of the same size. In Figure 11.2, we are not keeping track of individual infant weights but average weights for collections of 1009 infants.

Using this model for the sampling distribution, we can then grab the middle 95% of values in order to construct a confidence interval for the parameter of interest. This results in a 95% confidence interval of (3418.73, 3479). Based on this confidence interval, the data is consistent with the birth weight of infants in North Carolina, on average, being between 3418.73 and 3479; that is, these are the reasonable values of the mean birth weight.

Notice that we are able to narrow down the reasonable values of the parameter to a relatively small interval (a difference of about 60 grams). This is not because all babies in North Carolina have an extremely similar birth weight. It is because we have a relatively large sample, allowing us to have high confidence in our estimate of the *average* birth weight of an infant. The confidence interval does not tell us where we expect an individual infant's birth weight to fall; it only communicates what we are estimating the average birth weight of all infants to be based on our observed sample.

Also, notice how much narrower the model for the sampling distribution is compared to the distribution of the variable in the sample. Remember, statistics have less variability than individual values. This also illustrates why a confidence interval could never describe the fraction of values in the population which fall within a certain range — the variability is not comparable because a sampling distribution has a different x-axis than the distribution of the population or sample.

12 Quantifying the Evidence for a Single Mean

Consider the Birth Weight Case Study described in Chapter 9. In the previous chapter, we introduced focused on estimating the mean response of a variable; in this chapter, we consider testing a set of hypotheses involving the mean response.

In 2004, when the data for the Birth Weight Case Study was collected, an infant was considered "full term" if it was born anytime between 37 and 42 weeks. However, in 2013 the American College of Obstetricians and Gynecologists redefined "full term" to mean an infant born anytime between 39 and 40 weeks. Consider the following research question:

Does our study provide evidence that the average gestation time for infants born in North Carolina exceeds 38 weeks (so that, on average, babies are born full term by the new definition)?

This question is captured by the following set of hypotheses:

$$H_0: \theta \le 38$$
 vs. $H_1: \theta > 38$

where θ is the average gestation period (in weeks) of an infant born in North Carolina. This parameter characterizes the following model for the data generating process:

(Gestation Period)_i =
$$\theta + \varepsilon_i$$
.

We will assume that the gestation period for one infant is independent of the gestation period for any other infant and that this data is representative of all infants born in North Carolina; this implies we can assume the errors are independent and identically distributed.

We can estimate the parameter θ with the average gestation period for the babies in our sample: 39.11 weeks. We seek to quantify the evidence against the null hypothesis summarized by this data.

In Chapter 7, we discussed the concept of the null distribution of the statistic. Here, we are interested in modeling the null distribution of the sample mean gestation for a sample of 1009 infants. Notice, just like when modeling the sampling distribution, incorporating the sample size is an important aspect of our model. Assuming the data is consistent with the above

¹https://journals.lww.com/greenjournal/Fulltext/2013/11000/Committee_Opinion_No_579____ Definition_of_Term.39.aspx

two conditions on the distribution of the errors in the model for the data generating process, we can construct a model of the null distribution using a bootstrap procedure. Figure 12.1 illustrates this model based on 5000 bootstrap replications.




Figure 12.1: Model of the null distribution of the sample mean gestation period for a sample of 1009 infants. The model is based on 5000 bootstrap replications under the null hypothesis that the average gestation period is 38 weeks.

Figure 12.1 was constructed using the following procedure:

- 1. Alter the sample to be representative of having come from a population in which the null hypothesis is true; in this case, the data was recentered to have a sample mean of 38.
- 2. Randomly sample, with replacement, 1009 records from the altered original sample.
- 3. For this bootstrap resample, compute the mean gestation period and retain this value.
- 4. Repeat steps 2 and 3 many (say 5000) times.

The resulting simulated data is illustrated in Table 12.1. Each row represents the gestation periods for a single resample taken with replacement from the altered original data. The final column is the computed (and retained), sample mean from each resample. This is the bootstrap statistic *under the null hypothesis*.

Table 12.1: Partial printout of first 10 bootstrap resamples under the null hypothesis and the resulting bootstrap statistic.

Value 1	Value 2	Value 3	Value 1007	Value 1008	Value 1009	Boostrap Mean
---------	---------	---------	------------	------------	------------	---------------

36.88503	36.88503	39.88503	 36.88503	36.88503	37.88503	37.96234
36.88503	37.88503	37.88503	 39.88503	36.88503	35.88503	37.91278
37.88503	37.88503	37.88503	 38.88503	37.88503	39.88503	37.94747
36.88503	37.88503	35.88503	 36.88503	38.88503	37.88503	38.02577
36.88503	36.88503	37.88503	 38.88503	37.88503	40.88503	38.00198
39.88503	38.88503	36.88503	 36.88503	37.88503	37.88503	38.03072
38.88503	38.88503	36.88503	 37.88503	36.88503	38.88503	38.06938
36.88503	35.88503	39.88503	 36.88503	38.88503	37.88503	38.00396
37.88503	36.88503	37.88503	 37.88503	35.88503	35.88503	37.96234
39.88503	37.88503	35.88503	 36.88503	35.88503	36.88503	37.99802

Notice that the null distribution is centered on 38; this is not an accident. Recall that sampling distributions are centered on the true value of the parameter; since a null distribution is just the sampling distribution of a statistic when the parameter is equal to the null value (in this case 38), the null distribution should be centered on the null value. That is, the null distribution is designed to be centered on the null value in the hypotheses. In order to determine if our sample is consistent with our expectations, we overlay our observed sample mean ($\hat{\theta} = 39.11$) on our model for the null distribution. Since this value is in the far right tail of the null distribution (off the edge of the graph in this case), our sample is *inconsistent* with the null distribution. This tells us that if the average gestation period of an infant was only 38 weeks, we would not have expected to observe a sample mean of $\hat{\theta} = 39.11$ or higher in a sample of 1009 infants. Since our data is not consistent with our expectations, our study provides evidence that the population from which our sample was drawn has an average gestation period larger than 38 weeks.

12.1 Standardized Statistics

In the above discussion, we compared the observed sample mean to our distribution of expected sample means if the null hypothesis were true. We were essentially comparing $\hat{\theta}$ to 38 while accounting for the sampling variability of our estimate $\hat{\theta}$, the sample mean. This is a completely valid approach to inference. In this section, we consider an equivalent (conceptually), though alternative, approach which will provide a more general framework for inference.

At its heart, hypothesis testing is about comparing two models for the data generating process. So far, we have stated one of those models:

Model 1:
$$(Gestation Period)_i = \theta + \varepsilon_i$$
.

This is the data generating process under the alternative hypothesis in which no restrictions are placed on the value of θ . However, if the null hypothesis is true, then the model for the data generating process simplifies to

Model 0: (Gestation Period)_i = $38 + \epsilon_i$.

This may not seem like a simpler model, but it is because there are less unknown parameters; specifically, Model 0 has no parameters. A null hypothesis essentially places further restrictions on the data generating process. A hypothesis test is then about comparing these two models.

Big Idea

Hypothesis testing is about comparing two potential models for the data generating process. Specifically, we are asking whether a reduced model is sufficient for explaining the data or if there is evidence a more complex model is necessary.

Note

You may be wondering why we chose the value 38 when writing Model 0 above. After all, the null hypothesis was $H_0: \theta \leq 38$; so, it is natural to wonder why we did not choose 37 or one of the other infinite values below 38. In hypothesis testing, we choose to be as conservative as possible. If we are to declare that the data is not consistent with the null hypothesis, we want to be sure we are able to say the data is not consistent with any part of the null hypothesis.

Recall that our observed statistic was $\hat{\theta} = 39.11$; so, if we are able to state that we would not expect a sample mean gestation period of 39.11 or larger if the mean gestation period in the population is 38, then surely we are able to state that we would not expect a sample mean gestation period of 39.11 or larger if the mean gestation period in the population is 37. That is, if 39.11 is statistically discernible from 38, surely it is statistically discernible from 37 (or any other value less than 38).

That is, when setting up our model for the data generating process under the null hypothesis — when modeling the null distribution of our statistic — we choose the null value because it is the hardest to establish evidence against in a one-sided hypothesis test. If we can statistically discern our statistic differs from the null value, we can statistically discern our statistic differs from any value specified in the null hypothesis. That allows us to confidently provide evidence against the entire null hypothesis.

Essentially, the null hypothesis we have been considering $(H_0:\theta\leq 38)$ is stating that Model 0 is sufficient for explaining the data observed. And, the alternative hypothesis $(H_1:\theta>38)$ is stating that Model 0 is not sufficient, and a more complex model (Model 1) is necessary for explaining the data observed. This is why we refer to hypothesis testing as assessing model consistency. We are determining if there is evidence that the data is inconsistent with a proposed model for the data generating process.

Intuitively, the two proposed models would be equivalent (Model 0 would be sufficient for explaining the data) if they both performed similarly in predicting a response. Model 1 would

be preferred (Model 0 would not be sufficient for explaining the data) if it performs better in predicting the response. While this idea is intuitive, the process for comparing the models is not — we can assess "prediction" by the amount of variability in the data. If Model 0 captures a similar amount of variability in the response as Model 1, then it is reasonable that Model 0 is sufficient.

For Model 0, the amount of variability can be quantified by

$$SS_0 = \sum_{i=1}^n \left[(\text{Gestation Period})_i - 38 \right]^2.$$

That is, we are computing the total amount (summation) that the observed responses deviate (the squared difference) from the proposed mean of 38. For Model 1, the amount of variability can be quantified by

$$SS_1 = \sum_{i=1}^n \left[(\text{Gestation Period})_i - \hat{\theta} \right]^2$$

where $\hat{\theta} = 39.11$, the observed sample mean. Here, we are computing the total amount (summation) that the observed responses deviate (the squared difference) from the "best" estimate for the unknown mean response. Since Model 1 does not place any constraints on the value of the parameter, we use our estimate from the sample.

Notice that these sums of squared (SS) terms are similar to the definition of sample variance discussed in Chapter 5, without the scaling factor. The rationale for using these to assess predictive ability of the model will be further discussed in Chapter 19. Here, we simply note that they are measuring a distance the observed data is from a mean; the difference is whether that mean is unrestricted (and therefore estimated from the data, Model 1) or restricted under the null hypothesis (Model 0). If SS_0 and SS_1 were similar, then it would suggest that $\hat{\theta}$ is close to the null value of 38, and $\hat{\theta}$ differs from the null value only due to sampling variability, which would be in line with the null hypothesis. If, on the other hand, SS_0 and SS_1 differ substantially from one another, it suggests $\hat{\theta}$ differs from the null value more than we would expect due to variability alone. That is, if SS_0 and SS_1 differ substantially from one another, it suggests that our data is not something we would expect to observe under the null hypothesis. Therefore, the difference in these two sums of squares gives us a measure of the signal in the data against the null hypothesis. The larger this difference, the stronger the signal.

However, the same signal can be more difficult to detect (or discern) in the presence of a lot of background noise. Think about having a radio on in the background of a party, and suppose the radio is set to a specific volume. If there are not many people talking at the party, it is easy to hear the radio; the signal is strong relative to the background noise. However, if there are a lot of people talking at the party, the radio is difficult to hear even though its volume

hasn't changed; the signal is weak relative to the background noise. A signal is more difficult to locate if the background noise is elevated. The same principle holds in data analysis.

Consider Figure 12.2. Suppose we want to use each of these datasets (both containing a sample of size n = 20) to test the hypotheses:

$$H_0: \mu = 0$$
 vs. $H_1: \mu \neq 0$

where μ is the population mean. Both datasets have exactly the same observed sample mean response (the black diamond in the figure). Therefore, it can be shown that the difference between SS_0 and SS_1 is exactly the same for both datasets. However, just visually, it should be clear that Dataset A provides stronger evidence against the null hypothesis than Dataset B; that is, Dataset A is more inconsistent with a mean of 0. The difference is the variability — the background noise.




Figure 12.2: Illustration of the need to compare a signal to the noise in the data to assess its true strength.

Therefore, when quantifying the strength of a signal in a statistical analysis, it is common to measure the signal relative to the background noise. Returning to our example for the Birth Weight Case Study, we have that $SS_0 - SS_1$ quantifies our signal. The noise is the variability in the sample given by

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n \left[(\text{Gestation Period})_i - \hat{\theta} \right]^2,$$

the sample variance. We can examine our signal relative to the noise using a signal-to-noise ratio,

$$T^* = \frac{SS_0 - SS_1}{s^2} = 963.2$$

for our example. Such signal to noise ratios are known as standardized statistics.

Definition 12.1 (Standardized (Test) Statistic). Also, known as a test statistic, a standardized statistic is a ratio of the signal in the sample to the noise in the sample. The larger the standardized statistic, the stronger the evidence of a signal; said another way, the larger the standardized statistic, the stronger the evidence against the null hypothesis.

i Note

A standardized statistic is often referred to as a "test statistic," or a "standardized test statistic," because they are heavily used in hypothesis testing.

Of course, the natural question is "when does a standardized statistic become large enough?" Just as we constructed the null distribution for the observed sample mean in order to construct a distribution of our expectations under the null hypothesis, we can construct a null distribution of the standardized statistic to determine our expectations of this ratio under the null hypothesis. Figure 12.3 provides a model for the null distribution of our standardized statistic for the Birth Weight Case Study. This model for the null distribution is constructed in the same way we constructed the model for the null distribution of the sample mean except that instead of retaining the sample mean from each resample, we compute and retain the standardized statistic from each resample.

Notice that using Figure 12.3 we reach the same conclusions as when we used Figure 12.1. In Figure 12.3, we see that our observed standardized statistic of 963.2 is in the far right tail of the null distribution. Therefore, our data is inconsistent with the null hypothesis. That is, if the null hypothesis were true, it would be very unlikely to obtain a sample which produced a standardized statistic this extreme or more so due to sampling variability alone.

If our conclusions do not change, why the two different approaches? It turns out there is some theory that says bootstrapping standardized statistics tends to be a bit more stable computationally, and these standardized statistics are a bit easier to model analytically using probability theory (as we will later see). However, we introduce them because it again provides a nice overarching framework that unifies several of the approaches discussed in the text.



Quantifying evidence to compare two models for the data generating process can be done by comparing the signal in the data to the background noise.




Figure 12.3: Model of the null distribution of the standardized statistic for a sample of 1009 infants. The model is based on 5000 bootstrap replications under the null hypothesis that the average gestation period is 38 weeks.

Before moving on, we should note that there is not a unique standardized statistic. Other standardized statistics are often reported; for example, for hypothesis testing of a single mean response, the ratio

$$\frac{\sqrt{n}\left(\hat{\theta} - \theta_0\right)}{s}$$

is often reported, where θ_0 is the value of the mean response assumed under the null hypothesis.

It can be shown that many standardized statistics are related to one another (for example, the one given here is the square root of the standardized statistic reported previously). When the same conditions are applied to the data generating process, various standardized statistics yield the same conclusions. Again, we opt for the one described earlier because it will provide continuity in the text.

12.2 Computing the P-value

Now that we have a model for the null distribution of the standardized statistic, we can compute a p-value. The p-value is the probability of observing a sample as or more extreme

due only to sampling variability. Our standardized statistic has the form

$$T^* = \frac{SS_0 - SS_1}{s^2}$$

If the sample were more extreme — that is, if it produced a larger signal — then we would expect the difference between SS_0 and SS_1 to be even larger. Therefore, larger values of the standardized statistic present stronger evidence against the null hypothesis. When looking at the null distribution of the standardized statistic, computing the p-value corresponds to computing the area to the right of the observed standardized statistic.

Looking back at Figure 12.3, our observed standardized statistic is not even on the graphic, meaning the p-value (the tail area) is essentially 0. The data therefore provides strong evidence that the average gestation period of infants born in North Carolina exceeds 38 weeks.

As stated in Chapter 7, we should never report a p-value in isolation. We estimated the average gestation period of infants born in North Carolina to be 39.11 weeks. Our p-value tells us that we are able to statistically discern a difference between 39.11 and the 38 weeks from the null hypothesis. This ability to statistically discern this difference is in part due to the sample size of 1009 infants. To know whether this difference is meaningful, we would want to discuss the results with an obstetrician. However, informally, the increase of one week in the gestation period of an infant has a significant impact on both the health of the infant and the life of the mother (ask anyone who is 9 months pregnant!).

Finally, we take this opportunity to remind you that our conclusion is only about the *average* gestation period of infants. We make no claim about how long any individual infant will be carried prior to labor.

Part III

Unit III: Modeling the Average Response as a Function of a Continuous Predictor

Unit I explored the language and logic of statistical inference, and Unit II applied those concepts toward making inference on the mean response of a single variable. Within this context, we developed the idea of a statistical model for the data generating process, and introduced the concept of a standardized statistic. We now refine these ideas as we introduce the use of predictors in the data generating process. Specifically, we model the average response as a linear function of a single predictor.

13 Case Study: Seismic Activity in Greece

At the intersection of the African plate, the Eurasia plate, and the smaller Aegean plate, Greece is one of the most earthquake-prone regions in the world. Between July 2016 and July 2017, Greece experienced 179 earthquakes; by contrast, the state of Texas experienced 28 over the same span of time. In a region with such seismic activity, careful consideration must be given to municipal construction. Understanding how the motion experienced in a location is related to the soil properties in the area or the magnitude and distance of an earthquake is important.

An article in the Journal of Earthquake Engineering (Koutrakis et al. 2002) examined seismic events in Greece occurring between 1978 and 1997. Of interest for construction is characterizing the "strong ground motion," when the earth shakes with enough force to cause damage to infrastructure, with respect to the properties of a location. The study recorded several measurements from 121 stations (representing 93 distinct seismic events)¹. The primary variable of interest is the bracketed duration, "the time interval [in seconds] between the first and last excursion of the peak ground acceleration beyond a certain predefined level." For our purposes, we only consider the data corresponding to a threshold of 2% of the acceleration due to gravity. A longer bracketed duration would correspond to more time an area is exposed to violent shaking (and therefore more time in which damage may occur). In addition to the bracketed duration, the following measurements were available for each observation:

- Moment Magnitude: a measure of the size of the earthquake; larger values indicate more severe earthquakes.
- Epicentral Distance: distance (kilometers) from the epicenter of the earthquake to the location at which the measurement was taken.
- Soil Condition: indicator of the type of soil present at the measurement site. Soil was categorized as one of three types alluvium (soft, fine particles of clay, silt, sand, and gravel), intermediate soil conditions, or tertiary or older rock (those older than 2.58 million years).

The first 5 observations in the dataset are shown in Table 13.1. In this unit, we are particularly interested in characterizing the relationship between the bracketed duration at a location and

¹The original article presented repeated measurements at each location. We present here only the first measurement from each location to simplify any analyses. Repeated measurements are discussed briefly later in the text; for a more thorough treatment of the subject, we recommend a course in Designed Experiments or Biostatistics. The dataset presented here corresponds to that presented in Navidi's "Statistics for Engineers and Scientists" (Chapter 8, Supplementary Exercise 22).

Table 13.1: Data for the first 5 observations from a study characterizing seismic activity in Greece.

Magnitude	Distance from Epicenter (km)	Bracketed Duration (s)	Soil Conditions
6.4	30	8.82	Soft
5.3	6	4.31	Intermediate
5.6	15	5.74	Intermediate
5.2	7	4.08	Intermediate
6.6	31	28.27	Soft

the magnitude of the corresponding earthquake.

14 Myriad of Potential Questions

For the Seismic Activity Case Study, we are primarily interested in characterizing the relationship between the bracketed duration at a location and the magnitude of the corresponding earthquake. First, note that this question is about the relationship between a quantitative response (bracketed duration; see Definition 3.2) and a quantitative predictor (magnitude). Also note that the question is quite broad. We might actually have one of the following more specific questions in mind:

- In general, does the bracketed duration change as the magnitude changes?
- If two earthquakes with different magnitudes occur in the same location, would we expect the same bracketed duration provided the locations have the same soil conditions?
- Is the relationship between the bracketed duration and the magnitude different depending on the soil condition of where the measurement is taken?

These illustrate an array of potential questions we could address with the data. Each represents a different emphasis that we might have in a research question:

- Marginal Relationship: overall, do two variables tend to move together (are they correlated)?
- Isolation of Effect: does a relationship exist after accounting for the effect of additional variables? Or, what is the effect "above and beyond" the effect of additional variables?
- Interplay: how does the relationship between two variables change as a result of a third variable?

There is no right question to ask; each question examines a different facet of the relationship between two quantitative variables. In this unit, we will focus on questions of the first type. However, the framework we introduce is broad enough to be extended to address each of these types of questions. This may sound daunting, but keep in mind that the fundamental ideas we discussed in Unit I and applied in Unit II will continue to form the foundation of the analyses discussed in this unit; namely,

- We are using a sample to say something about the underlying population.
- In order to make inference, we will need a model for the sampling (or null) distribution of our statistic; as a stepping stone, we model the data generating process.
- In order to form a standardized statistic of interest which measures the strength of the signal in the dataset, we think about variability.

The ideas remain the same; the context has changed.

There is one more thing we want to point out before moving on: any relationships we observe are overall trends, not guaranteed to hold for any single individual. Recall that in Unit II we emphasized that our conclusions were about the mean response (the parameter of interest). Specifically, even if we know the average response within the population, due to variability, we do not expect every individual to have that specific value for the response. This will continue in this unit. If we observe, for example, that an increase in the magnitude is associated with an increase in the bracketed duration, we are describing an overall trend. It is highly likely there is some location for which this trend does not hold, simply due to variability.

15 Nature of Collecting Multivariable Data

For the Seismic Activity Case Study, we are primarily interested in characterizing the relationship between bracketed duration and the magnitude of the earthquake. As we discussed in the previous chapter, this general goal might be refined into one of many specific questions:

- In general, does the bracketed duration increase as the magnitude increases?
- If two earthquakes with different magnitudes occur, would we expect the same bracketed duration provided the two locations of interest are the same distance from the epicenter?
- Is the relationship between the bracketed distance and the magnitude different depending on the soil condition of where the measurement is taken?

Notice that these last two questions actually require knowledge of more than just the bracketed duration and the magnitude of each seismic event. In order to address the second question, we would also need the distance from the center of the earthquake; in order to address the third question, we also need the soil conditions of where the measurement is taken. Often, research questions require knowledge of more than just a single variable; such questions are multivariable.

Definition 15.1 (Multivariable). This term refers to questions of interest which involve more than a single variable. Often, these questions involve many variables. Multivariable models typically refer to a model with two or more predictors.

Consider going to the doctor because you are feeling ill. The doctor does not have you simply enter your most prominent symptom (fever, for example) into a computer and then prescribe a medication based solely on that single symptom. Instead, a good physician will review all symptoms you are experiencing, as well as your medical history, other medications, allergies, etc. The physician operates in a multivariable world in which there are many contributing factors to a response. Therefore, when you arrive for this hypothetical visit, they record several variables that may be of interest.

Studies which collect several variables can be observational studies or controlled experiments. If an observational study, we want to ensure the sample of subjects is representative of the target population. Then, for each individual, we simply record several variables. If a controlled experiment, we randomly assign subjects to a particular "treatment" group; afterwards, we would measure the response in addition to other variables. Notice that with the latter, subjects are randomly assigned to only one of the variables; the remaining variables are simply observed; while this could of course be extended, this is the most common implementation.

i Note

When a study is primarily interested in characterizing the relationship between two or more quantitative variables, the data is typically from an observational study.

What we want to emphasize here is that how we collect the data has not really changed from what we have discussed in previous units. The primary difference is that we are very aware that we are collecting several measurements on each subject. The critical element is that our sample be representative of the target population if we want to apply any findings to that population.

16 Summarizing Multivariable Data

For the Seismic Activity Case Study described in Chapter 13, we are primarily interested in characterizing the relationship between the bracketed duration and the magnitude of the corresponding earthquake. As we discussed in the previous chapters, this broad question could be refined into a question falling into one of three categories:

- Marginal Relationship: overall, does the bracketed duration and the magnitude of the corresponding earthquake tend to move together (are these two variables correlated)?
- Isolation of Effect: is there an association between the bracketed duration and the magnitude of the corresponding earthquake after accounting for the impact of additional variables? That is, is there an association between the bracketed duration and the magnitude of the corresponding earthquake "above and beyond" the association with additional variables?
- Interplay: how does the relationship between the bracketed duration and the magnitude of the corresponding earthquake depend on the type of soil?

As always, regardless of the question we ask, the key is developing summaries which help to address that question.

16.1 Characterizing the Marginal Relationship of Two Quantitative Variables

Suppose we are interested in the following question:

In general, does the bracketed duration tend to change as the magnitude of the corresponding earthquake changes?

This question is about the overall (marginal) relationship between these two quantitative variables. Graphically, we can examine the relationship between these two variables using a scatter plot. The response is placed on the y-axis and the predictor along the x-axis. Figure 16.1 illustrates the relationship between the bracketed duration at a location and the magnitude of the corresponding earthquake.

Figure 16.1 highlights several facets of the relationship. First, we note that as the magnitude of the event increases, the bracketed duration also tends to increase. This is intuitive — as the size of the earthquake increases, the length of time the ground shakes with extreme force increases.




Figure 16.1: Relationship between the bracketed duration at a location in Greece and the magnitude of the corresponding earthquake.

This is a *trend*; it is not a universal truth. For example, note the two cases highlighted in blue; the observation with the larger magnitude has a smaller bracketed duration, which is counter to the overall trend in the graphic. The research objective was to characterize the overall trend, not make global statements about all units in the population. Figure 16.1 also reveals that as the magnitude increases, the variability in the bracketed duration also tends to increase. That is, for earthquakes of small magnitudes, it seems fairly easy to anticipate the bracketed duration; however, the bracketed duration is much more difficult to anticipate for larger magnitudes.

A nice visual tool when exploring the relationship between two quantitative variables is a *smoothing spline*. The details of its construction are beyond the scope of this text, but we can think of it as representing where the average observed response is located for a particular value of the predictor; when computing the average response, information is borrowed form nearby points smoothing out the relationship (hence the name). We do want to point out that this is an exploratory device; we should be cautious about over-emphasizing relationships we observe from the smoother. Figure 16.2 illustrates a smoother relating the length of the bracketed duration with the magnitude of the earthquake. The addition of the spline confirms what we had previously stated about the relationship appearing fairly linear (as the magnitude of the earthquake increases so does the bracketed duration at a location). In addition to the spline, there is a confidence band (generalization of a confidence interval) around the line in order to convey the variability in the estimated smoother.

As we have seen, supplementing graphical summaries with numerical summaries can help

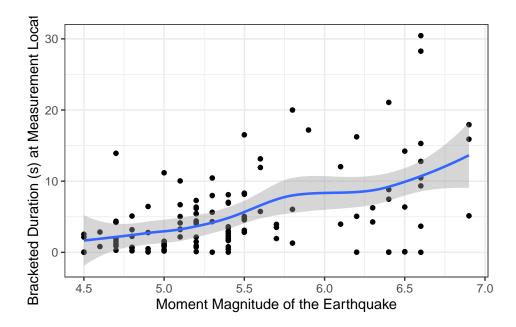


Figure 16.2: Illustrating the use of a smoothing spline to explore the relationship between the bracketed duration and the magnitude of an earthquake for locations Greece.

convey our message. As an example, Figure 16.2 suggests there is a positive, linear relationship between the bracketed duration and the magnitude of the corresponding earthquake. But, can we quantify that relationship? Consider Figure 16.3, which consists of two hypothetical datasets. Both datasets illustrated exhibit a positive, linear relationship between the response and predictor; however, that relationship is much stronger (or more apparent) for Dataset A compared to Dataset B. It would be nice to have a numeric summary which captured this; such a metric is known as the **correlation coefficient**.

Definition 16.1 (Correlation Coefficient). A numerical measure of the *strength* and *direction* of the *linear* relationship between two quantitative variables.

The classical Pearson Correlation Coefficient r is given by the following formula:

$$r = \frac{\sum_{i=1}^{n} \left(x_{i} - \bar{x} \right) \left(y_{i} - \bar{y} \right)}{\sqrt{\sum_{i=1}^{n} \left(x_{i} - \bar{x} \right)^{2} \sum_{i=1}^{n} \left(y_{i} - \bar{y} \right)^{2}}}$$

where \bar{x} and \bar{y} represent the sample means of the predictor and response, respectively.

The correlation between the bracketed duration and the magnitude of an earthquake in our sample is 0.497, indicating the two variables are positively linearly related, though perhaps the relationship is not strong.

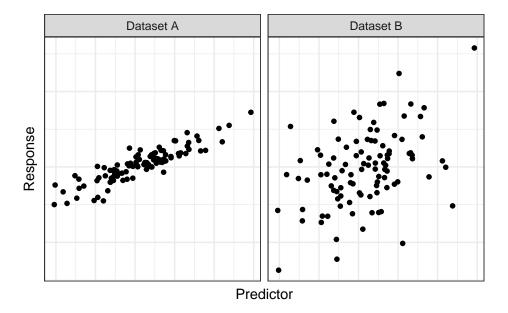


Figure 16.3: The relationship between a response and predictor for two hypothetical datasets;

Dataset A exhibits a stronger correlation between the response and predictor than

Dataset B.

i Properties of the Correlation Coefficient

The Pearson Correlation Coefficient has the following key properties:

- It takes a value between -1 and 1.
- Negative values mean that the variables tend to move in opposite directions.
- Positive values mean that the variables tend to move in the same direction.
- It is unitless and therefore unaffected by unit changes in the variables.

The biggest thing to remember is that a correlation coefficient measures the strength of a *linear* relationship. A correlation of 0 does not mean that two variables are unrelated. It simply means they are not linearly related.

16.2 Visualizing the Impact of a Third Variable on the Marginal Relationship

In the previous section, we stated that in our sample, the bracketed duration tends to increase as the magnitude of the corresponding earthquake increases. It is reasonable to ask the following question:

Is the relationship between the bracketed duration and the magnitude different depending on the soil condition of where the measurement is taken?

That is, we want to determine the impact that a third variable (soil condition) has on the relationship we have observed. While the bulk of this unit will focus on inferential methods for the marginal relationship, graphically assessing questions isolating a single predictor or about the interplay of two predictors is fairly intuitive. In order to add more depth to our graphical representations, we make use of various features of the graphic — the color, shape, and/or size of the points used in plotting — as well as facets (multiple graphics each with a different subset of the data). Figure 16.4 uses color to distinguish between the three possible types of soil conditions at each measurement location. Notice the graphic allows us to both visualize the relationship between the bracketed duration and the magnitude for each soil condition but also facilitates our comparing these relationships across soil conditions.




Figure 16.4: Relationship of the bracketed duration and the magnitude of the corresponding earthquake across across various soil conditions.

Figure 16.4 illustrates that the relationship between the bracketed duration at a location and the magnitude of the corresponding earthquake is similar for both locations that have soft or intermediate soil conditions. However, for locations with rocky conditions, increases in the magnitude of the earthquake are associated with less pronounced increases in the bracketed duration. This potentially suggests that foundations on rocky soils are less subject to the effects of large earthquakes, at least with respect to the amount of time those areas are subject to extreme motion.

While our focus in this chapter has been on the scatter plot, our emphasis remains the same as

when we used simpler graphics in the first unit — summaries need to be constructed to address the question of interest. And, regardless of the type of graphic, it communicates information about how the distribution of the response.

17 Bulding our Statistical Model

In Chapter 10, we introduced the statistical modeling framework. In particular, our general model (see Equation 10.1) for the data generating process was

Response = function(predictor variables, parameters) + noise.

Recall that this model has two components:

- A deterministic component which takes the form of a function of variables and unknown parameters. It is often this component on which we would like to make inference.
- A stochastic component which captures the unexplained variability in the data generating process.

In the previous unit, we made use of this model, but we only scratched the surface of its potential applications. In this unit, we explore some of the capabilities of such a model. Specifically, we consider a model for the data generating process such that the deterministic component is a smooth function (specifically, a line) of potentially several variables. In general, this model building process is known as **regression**.

Definition 17.1 (Regression). Used broadly, this refers to the process of fitting a statistical model for the data generating process to observed data. More specifically, it is a process of estimating the parameters in a data generating process using observed data.

17.1 Statistical Model for A Quantitative Response and Quantitative Predictor

Recall that in Chapter 10, we described a general model for the data generating process of a quantitative response as a function of a single parameter:

$$(Response)_i = \mu + \varepsilon_i$$

where μ represented the average response. We noted that this model is fairly limited as it does not allow the response to depend on additional information collected on each unit. In particular, we might consider a model of the form

(Bracketed Duration)_i =
$$\mu + \varepsilon_i$$
 (17.1)

for the data generating process of the bracketed duration. However, this model does not allow the bracketed duration at a location to depend on the magnitude of the corresponding earthquake. We would like to extend it so that we can account for this additional information.

It is often easier to discuss modeling in the context of the graphics used to visualize them. Consider the Seismic Activity Case Study described in Chapter 13. Let's begin with a broad question:

In general, does the bracketed duration tend to change as the magnitude of the corresponding earthquake changes?

As we are interested in predicting the bracketed duration, we will treat it as the response. In order to imagine what an appropriate model for how the bracketed duration is generated as a function of the magnitude of an earthquake, consider the graphical summary of how these variables are related. As discussed in Chapter 16, we can use a scatter plot to visualize the relationship between our response and predictor. Figure 17.1 overlays a smoothing spline (grey) on the observed data; however, it also includes a proposed model (blue) for the data generating process for which the bracketed duration is linearly related to the magnitude of the corresponding earthquake.

Suppose we feel that this line is a good model for the data generating process. Before proceeding, consider what this statement says. We are not trying to say that the relationship explains every response we observe. Instead, the relationship explains the underlying trend — what happens on average. While not perfect, this linear relationship at least appears plausible. Therefore, we replace μ in Equation 17.1 with the expression for a line; this results in

$$(Bracketed Duration)_i = \beta_0 + \beta_1 (Magnitude)_i + \varepsilon_i$$
 (17.2)

where β_0 represents the intercept of the line and β_1 the slope. Both the bracketed duration and the magnitude of the corresponding earthquake are variables that are measures for each observation. The terms β_0 and β_1 are unknown constants (parameters) governing the model for the data generating process.

Observe that very few points in Figure 17.1 actually fall on the proposed line in the graphic, which is to be expected. This emphasizes the idea that the deterministic portion of the model is not meant to fully capture a data generating process since variability is inherent in any process. This is why statistical models embed a deterministic component alongside a stochastic component — to capture the variability due to error or noise in the data generating process.

The model suggests that the bracketed duration at a location is primarily determined by the magnitude of the corresponding earthquake; however, there is a component we cannot explain.




Figure 17.1: Relationship between the bracketed duration at locations across Greece and the magnitude of the corresponding earthquake. A smoothing spline illustrates the overall relationship, and a line overlayed represents a potential model for the data generating process.

That is, the model does not explain why, for example, when an earthquake with a magnitude of 5.5 hits, all locations do not have the same bracketed duration. This noise is picked up by the ε_i term in the model (as illustrated for a single observation by the red line in Figure 17.1). The model is only capturing the general trend. As in Chapter 10, we refine this model for the data generating process further by placing additional conditions on the random noise term in order to aid in conducting inference.

Simple Linear Regression Model

For a quantitative response and a quantitative predictor, the general form of the simple linear regression model is

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i$$
 (17.3)

where β_0 and β_1 are parameters governing the model for the data generating process.

17.2 Estimating the Parameters

Recall the goal of statistics — to use a sample to say something about the underlying population. There is something intuitive about using the sample mean as an estimate of the population mean; however, now we have a model for the data generating process that has two parameters (the intercept and the slope). We want a method that jointly estimates these parameters.

Before discussing a specific estimation procedure, it is worth emphasizing that our model for the data generating process has two parameters. In equation Equation 17.3, both β_0 and β_1 are unknown constants governing the data generating process; they are parameters. As discussed in Chapter 3, parameters are unknown values, and they will always be unknown. We can use available data to estimate these parameters (either point estimates for interval estimates), and we can use available data to determine if there is evidence the parameter falls outside of a specific region (hypothesis testing), but we will never be able to definitely state the value of these parameters from a sample. As a result, we do not "compute" the values of β_0 and β_1 ; we can only estimate them. Indeed, nearly every scientist and engineer is undoubtedly familiar with a line of "best fit." We are emphasizing here that any such line simply estimates the parameters in our model for the data generating process.

We now turn to how such point estimates are constructed. Our model in Equation 17.3 posits a linear deterministic portion for the data generating process, and we want to use the data observed to estimate what that relationship looks like. That is, we want to draw a line through the points that gives the "best fit." Figure 17.2 illustrates this for a hypothetical dataset; it compares two *estimated relationships*. Again, these are just estimates; neither represents the actual line from which the data was generated. Examining the two estimated

models, something inside us knows that the blue line is preferred to the orange line. The orange line does not seem to represent the pattern in the data because it strays from the cloud of observed points. Intuitively, a model that adequately describes the relationship between the two variables should go through the observed points. Trying to formalize this, we are saying we want a line that is somehow simultaneously as close to all the data as possible.




Figure 17.2: Illustration of two competing estimates of a line which runs through the data.

The most widely used method for accomplishing this goal, for estimating the parameters in Equation 17.3, is known as "the method of least squares." For this reason, the resulting parameter estimates are often referred to as **least squares estimates**. This method essentially chooses estimates for the parameters that minimize the amount of error within the dataset.

Definition 17.2 (Least Squares Estimates). Often called the "best fit line," these are the estimates of the parameters in a regression model chosen to minimize the sum of squared errors. Formally, for Equation 17.3, they are the values of β_0 and β_1 which minimize the quantity

$$\sum_{i=1}^n \left[(\text{Response})_i - \beta_0 - \beta_1 (\text{Predictor})_i \right]^2.$$

The resulting estimates are often denoted by $\hat{\beta}_0$ and $\hat{\beta}_1$.

For those familiar with calculus, we can imagine finding the values of the parameters which minimize the above quantity by taking partial derivatives, setting those partial derivatives equal to 0, and solving the resulting system of equations. In practice, the estimation is carried out numerically using statistical software.

Estimation is often associated with statistics. However, the least squares estimates are actually the result of a mathematical minimization process. In order to make inference on the parameters, we need to quantify the variability in those estimates — that is the heart of a statistical analysis. That is, the statistical aspect is moving into one of the components of the *Distributional Quartet*. In order to construct a model for the sampling distribution of these statistics, we place additional conditions on the stochastic portion of the model for the data generating process. That is the focus of the next chapter.

18 Conditions on the Error Term of a Regression Model

In the previous chapter we developed a general model for the data generating process of a quantitative response as a linear function of a quantitative predictor:

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i.$$

We also discussed a common method for estimating the parameters of this model from a sample — the method of least squares. However, if we are to construct a model for the sampling distribution of these estimates we must add some structure to the stochastic component ε in the model. In this chapter, we focus on the most common conditions we might impose and how those conditions impact the model for the sampling and null distributions (and therefore the computation of a confidence interval or p-value).

18.1 Correctly Specified Model

The first condition we consider is the most important. It states that for every value of the predictor, the average error is 0.

i Mean-0 Condition

The "mean-0 condition" states that for every value of the predictor, the average error is 0.

In practice, this condition implies that the deterministic portion of the model for the data generating process that we have posited is accurate; that is, it implies the *form* of the model is appropriate. For Equation 17.3, this condition states that the response is linearly related to the predictor.

There is a subtlety to this condition; the phrase "for every value of the predictor" is crucial. Without this phrase, the condition would simply state that the model may overestimate for some values of the predictor, but that is balanced out by underestimating for other values of the predictor. That is, the model for the deterministic portion can be wrong everywhere, but those "wrongs" cancel each other out (consider the orange line again in Figure 17.2). Instead,

we want the deterministic portion of the model to be correct everywhere, for all values of the predictor (consider the blue line again in Figure 17.2).

There are two reasons we say that this is the most important condition:

- 1. If this condition is violated, it says your model for the data generating process is fundamentally incorrect in its form. Generally, this is the result of ignoring some curvature in the relationship between the response and predictor or some important additional feature. If the functional form of the model for the deterministic portion is incorrect, you have to begin your modeling process over again.
- 2. This condition allows us to interpret the parameters of the model.

That is, instead of β_0 and β_1 representing solely the intercept and slope, we can interpret these parameters in the context of the research objective.

18.1.1 Interpreting the Parameters

In Chapter 10, we posited a model of the form

$$(Response)_i = \mu + \varepsilon_i$$

where we said that μ represented the average response. When we impose the mean-0 condition in Equation 17.3, we are stating that given a value of the predictor, the average response is given by

$$\beta_0 + \beta_1$$
 (Predictor).

That is, instead of considering the overall average response, we are acknowledging that the average response may depend on the value of the predictor. The mean-0 condition essentially tells us that the deterministic portion of the model for the data generating process represents the average response for a specified value of the predictor.



Big Idea

The deterministic portion of a simple linear regression model specifies the average value of the response given the value of the predictor.

As an example, consider our model for the Seismic Activity Case Study developed in Chapter 17, which considered the bracketed duration at a location as a function of the magnitude of the earthquake:

(Bracketed Duration)_i =
$$\beta_0 + \beta_1$$
(Magnitude)_i + ε_i .

If we are willing to assume that for earthquakes of any magnitude, the average error in the bracketed duration is 0 (the mean-0 condition), then earthquakes with a 5.0 magnitude have an *average* bracketed duration of

$$\beta_0 + \beta_1(5)$$
.

Similarly, earthquakes with a 6.0 magnitude have an average bracketed duration of

$$\beta_0 + \beta_1(6)$$
.

As we have mentioned, the deterministic portion of the model does not specify the exact response for any individual unit; the deterministic portion specifies the trend. We are now able to say that the "trend" we are modeling is the average response. Further, we can estimate this average response by plugging in the least squares estimates $\hat{\beta}_0$ and $\hat{\beta}_1$. Specifically, using the method of least squares, the line of best fit was estimated as

$$(Bracketed Duration) = -19.19 + 4.48(Magnitude)$$

Therefore, we estimate the average bracketed duration for locations with 5.0 magnitude earthquakes to be 3.22 seconds, and we estimate the average bracketed duration for locations with 6.0 magnitude earthquakes to be 7.71 seconds. While we do not expect every location which has a 5.0 magnitude earthquake to have a bracketed duration of 3.22 seconds, we expect the bracketed duration to vary about this length of time. This is huge; it says that when we use a regression model to predict a response, we are actually predicting the *average* response. And, it provides a direct interpretation of the parameters themselves.

Let's begin with the intercept term, β_0 . Notice that in our model above, if we try to predict the bracketed duration for a location with an earthquake which has a magnitude of 0, then our model returns β_0 . In fact, for any regression model, the intercept β_0 is the value of the deterministic portion of the model when the predictor in the model is set to 0, and we know that deterministic portion is the average response.

Interpretation of the Intercept

Consider a regression model of the form Equation 17.3; if we impose the mean-0 condition, the intercept β_0 represents the *average* response when the predictor is equal to 0. Note that the resulting estimate may not be feasible; indeed, considering a predictor of 0 may not even make sense in every context (e.g., an earthquake with a magnitude of 0 is not an earthquake).

For our particular example, our intercept is estimated to be -19.19; this estimates the average bracketed duration for an earthquake with a magnitude of 0. An earthquake with no magnitude

is not an earthquake at all; and, if there is no earthquake, we would not expect the ground to undergo any duration of extreme motion. Further, even if a magnitude of 0 made sense, a negative duration does not.

When predictions for a model do not make logical sense, it is often the result of **extrapolation**. We do not have any data on the bracketed duration for locations which experienced an earthquake with a magnitude less than 4.5. Therefore, we are using a model to predict for a region over which the model was not constructed to operate. This is a lot like using a screw driver to hammer a nail — we are using a tool to accomplish a task for which it was not designed. We should therefore not be surprised when the tool fails. The primary reason extrapolation is dangerous is that without data in a particular region, we have nothing supporting that the posited model for the data generating process will continue to hold in that region. We have illustrated this when discussing the intercept, but extrapolation can occur in any region for which there is no data. For this reason, unless you have strong scientific justification for why a model will hold over all values of the predictor, extrapolation should be avoided.

Definition 18.1 (Extrapolation). Using a model to predict outside of the region for which data is available.

With an interpretation of the intercept β_0 , we now turn our attention to the slope β_1 . Notice that based on our estimates above, the average bracketed duration is 4.48 seconds longer for those locations which experience a 6.0 magnitude earthquake compared to those which experience a 5.0 magnitude earthquake, and this difference is the value of the estimated slope. This is not a coincidence; 4.48 seconds is the change in the average bracketed duration that is associated with a 1-unit increase in the magnitude of an earthquake. The slope is essentially comparing the average response for two values of the predictor that differ by 1 unit.

i Interpretation of the Slope

Consider a regression model of the form Equation 17.3; if we impose the mean-0 condition, the slope β_1 represents the *average* change in the response associated with a 1 unit *increase* in the predictor.

18.1.2 Embedding our Question in a Statistical Framework

Our first fundamental idea centers on the idea that the majority of research questions can be framed in terms of a parameter within the population. This seemed somewhat intuitive when the parameter was simply the mean response. With parameters which are the slope and intercept of a line, this seems less clear. However, the mean-0 condition provides interpretations of the parameters, and this in turn ensures that our questions of interest can be framed in terms of those parameters. Consider the following question:

On average, are changes in the magnitude of an earthquake associated with changes in the bracketed duration observed?

Let's consider how we might write this in terms of a null and alternative hypotheses. The mean-0 condition states that our form for the deterministic portion of the model is correct. That is, imposing the mean-0 condition means we believe the following model for the data generating process is accurate:

$$(Bracketed Duration)_i = \beta_0 + \beta_1 (Magnitude)_i + \varepsilon_i$$

The research question suggests we are looking for a relationship (which we have posited to be linear *if* it exists). Since a relationship between the bracketed duration, on average, and the magnitude is what we are seeking evidence for, the presence of such a relationship should be captured by the alternative hypothesis. In turn, the null hypothesis would capture the idea that the average bracketed duration does not depend on the magnitude. That is, under the null hypothesis, magnitude should not be present in the model at all; that is, our model would reduce to

(Bracketed Duration)_i =
$$\beta_0 + \varepsilon_i$$
.

This reduced model suggests that the bracketed duration does not change at all as the magnitude changes; instead, the average bracketed duration remains β_0 for all values of the magnitude. This is essentially a flat line, which comes from a slope of 0. That is, if the null hypothesis is true — that changes in the magnitude are not associated with any changes in the bracketed duration — then, under our proposed model for the data generating process (which we believe has the correct *form* since we have imposed the mean-0 condition), then $\beta_1 = 0$. Said another way, under the null hypothesis, our model for the data generating process of the bracketed duration should not depend on the magnitude of the earthquake.

Therefore, our null and alternative hypotheses for the above research question can be written as

$$\begin{aligned} H_0: \beta_1 &= 0 \\ H_1: \beta_1 \neq 0 \end{aligned}$$

where β_1 is the parameter linearly relating the bracketed duration to the magnitude — it is the average change in the bracketed duration associated with a 1 unit increase in the magnitude.



Setting the slope parameter to 0 in Equation 17.3 is associated with saying that the predictor is not associated with the average response in a linear fashion — that it does not belong in the model.

The interpretation of our parameters allows us to see that our research questions are characterizing the relationship between the response and the predictor, on average. As in the previous unit, our questions are about the average response; instead of looking at the overall average, however, we are allowing it to depend upon a predictor.

This first condition on the error term — holding the average error to be 0 for all values of the predictor — gives our parameters meaning.

18.2 Independent Errors

The second condition we consider is that the noise attributed to one observed individual is independent of the noise attributed to any other individual observed. That is, the amount of error in any one individual's response is unrelated to the error in any other response observed. This is the same condition we introduced in Chapter 10.

i Independence Condition

The independence condition states that the error in one observation is independent (see Definition 10.3) of the error in all other observations.

With just these first two conditions (that the average error is 0 for all values of the predictors and the errors are independent of one another), we can use a bootstrap algorithm in order to model the sampling distribution of the least squares estimates of our parameters (see Appendix A). However, additional conditions are often considered.

18.3 Same Degree of Precision

The third condition that is typically placed on the distribution of the errors is that the variability of the errors is the same for all values of the predictor. Practically, if this condition does not hold, our response will be more precise in one region than in another. The technical term for this condition is *homoskedasticity*.

i Constant Variance

Also called homoskedasticity, the constant variance condition states that the variability of the errors is the same for all values of the predictor.

With this additional condition imposed, we are able to modify our bootstrap algorithm when constructing a model for the sampling distribution of the least squares estimates.

18.4 Specific form of the Error Distribution

Up to this point in the text, we have discussed bootstrapping as the tool for modeling the sampling distribution (or null distribution) of a statistic. Bootstrapping results in an empirical (data-driven) model for the sampling distribution. And, empirical models for distributions can be unstable in small sample sizes. That is, our model for the sampling distribution may change quite substantially from one sample to another. To be clear, our modeling strategy approach takes into account the sample size and communicates this uncertainty in the form of wider confidence intervals. Nevertheless, this instability should be kept in mind. In such cases, alternative bootstrap algorithms should be considered.

An alternative to building an empirical model is to construct an analytical model. While the term "analytical" as an alternative to empirical might suggest the model is not data-driven, this is a misnomer. The model for the sampling distribution of a statistic is still based on the observed sample; the term "analytical" really means that we can bypass the resampling component of a bootstrap procedure. In particular, the information gained from bootstrapping (the shape and spread of the sampling distribution) is obtained from relying on statistical theory. An analytical approach generally comes from imposing a fourth condition (common in engineering and scientific disciplines) on the distribution of the errors.

18.4.1 Modeling the Population

Before we delve into more detail, let's set the stage for the bigger story being told. Recall that our goal is to say something about the population using a sample. We have developed a process to address this goal:

- 1. Frame our question through a parameter of interest.
- 2. Collect data that allows us to estimate the parameter using the analogous statistic within the sample.
- 3. Summarize the variability in the data graphically.
- 4. Quantify the variability in the statistic through modeling the sampling distribution (or null distribution, whichever is appropriate).
- 5. Using the model for the sampling distribution, determine the reasonable values of the parameter; or, using the model for the null distribution, quantify the evidence in the sample.

This process is presented through our *Five Fundamental Ideas of Inference* and the *Distributional Quartet*. The key step in this process is quantifying the variability by modeling the sampling distribution (or null distribution, whichever is appropriate for our research goal). We have described the construction of these models empirically, through repeating the study by appropriately resampling the data available and performing the analysis on each resample.

Our goal is still to model the sampling distribution; that is the key inferential step. Instead of building an empirical model, we can construct an exact analytical model through an additional step: modeling the population directly.

Big Idea

A model for the sampling distribution of a statistic can often be obtained by modeling the distribution of the population.

Notice yet another usage of the word "model." We have discussed models for the data generating process, and we have discussed models for the sampling (or null) distribution of a statistic. Now, we are discussing a model for the distribution of the population. As we will see, a model for the distribution of the population is closely linked to the model for the data generating process; as such, the model for the distribution of the population is simply a stepping stone to the model for the sampling distribution — the key inferential step. It is important to separate these steps. We are not interested in directly modeling the population; we do it in order to construct a model for the sampling distribution.

Just as with other conditions we have discussed, specifying a particular model for the distribution of the population can allow us to depend on some statistical theory; however, in practice, we are assuming this model for the distribution of the population is appropriate.

18.4.2 The Normal Distribution

Probability, a sub-field of mathematics which is used heavily in statistical theory, is the discipline of modeling randomness. In statistical theory, we make use of probability to model a distribution. In order to get a feel for probability models, consider the following example.

Example 18.1 (Iris Characteristics). The discipline of statistics began in the early 1900's primarily within the context of agricultural research. Edgar Anderson was a researcher investigating the characteristics of the iris. He had collected measurements on over one hundred iris flowers, including their petal length and width and their sepal length and width. The sepal is the area (typically green) beneath the petal of a flower. It offers protection while the flower is budding and then support for the petals after the flower blooms.

Figure 18.1 is a histogram of the sepal width for the iris plants observed by Edgar Anderson; overlayed is the density plot for the same dataset, which we have described as a smoothed histogram. Both the histogram and the density plot are empirical models of the distribution of the sepal width.

Probability models are analytical models for the distribution of a variable. Instead of constructing a density using data alone, we posit a functional form for the density. For example, Figure 18.2 overlays the following function on top of the the iris data:




Figure 18.1: Summary of the distribution of sepal widths for a sample of irises.

$$f(x) = \frac{1}{\sqrt{0.380\pi}} e^{-\frac{1}{0.380}(x-3.057)^2}.$$

A density (whether constructed empirically or posited analytically) is just a model for the distribution of a variable. Further, all density functions share a few basic properties:

- 1. The density is non-negative for all values of the variable.
- 2. The area under the density function must equal 1.

While the value on the y-axis is not directly meaningful, density functions provide a link between the value of the variable and the likelihood of it occurring. Specifically, the probability that a variable falls in a specific range corresponds to the area under the curve in that region. For example, based on the analytical model described above (the blue curve in Figure 18.2), the probability that an iris has a sepal width between 3.5 and 4 centimeters is 0.14, illustrated in Figure 18.3. That is, there is a 14% chance a randomly selected iris will have a sepal width between 3.5 and 4 centimeters based on this model.

While the above model for the population is not a perfect fit to the data, it does capture many of the characteristics present in the sample. Similar to empirical models, analytical models for distributions are just that — models. As stated previously, the term "analytical" does not mean the model does not rely on the data. It is the form of the model that is predetermined; specific components of the model were estimated based on the available data. The particular




Figure 18.2: Summary of the distribution of the sepal widths for a sample of irises with an analytical model of the distribution overlayed.




Figure 18.3: Using the model for a density function to compute a probability.

model illustrated in Figure 18.2, characterized by the bell-shape density, is known as the **Normal Distribution**.

Definition 18.2 (Normal Distribution). Also called the Gaussian Distribution, this probability model is popular for modeling noise within a data generating process. It has the following characteristics:

- It is bell-shaped.
- It is symmetric, meaning the mean is directly at its center, and the lower half of the distribution looks like a mirror image of the upper half of the distribution.
- Often useful for modeling noise due to natural phenomena or sums of measurements.

The functional form of the Normal distribution is

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

where μ is the mean of the distribution and σ^2 is the variance of the distribution.

While there are several nice properties of the Normal Distribution, we are primarily interested in the fact that if the error in a data generating process follows a Normal Distribution (in addition to the other three conditions described above placed on the error term), then the form of the sampling distribution for the least squares estimates of the slope and intercept is known. That is, with all four conditions in place, we have an analytical model for the sampling distribution. This means we avoid simulating in order to build a model for the sampling distribution; so, computationally it is faster. If the errors really are from a Normal Distribution, then we also gain power in our study by imposing this condition.

i Normality

The normality condition states that the distribution of the errors follows the functional form of a Normal distribution (Definition 18.2).

Big Idea

If we are willing to assume the distribution of the errors in Equation 17.3 follows a Normal distribution, then we have an analytical model for the sampling distribution of the least squares estimates.

Let's think about what this condition means for the responses. Given the shape of the Normal distribution, imposing this condition (in addition to the other conditions) implies that some errors are positive and some are negative. This in turn implies that some responses will tend to fall above the line (we will under-predict for these observations), and some response will tend to fall below the line (we will over-predict for these observations).

18.5 Classical Regression Model

We have discussed four conditions we could place on the stochastic portion of the data generating process. Placing all four conditions on the error term is what we refer to as the "Classical Regression Model."

Definition 18.3 (Classical Regression Model). For a quantitative response and single predictor, the classical regression model assumes the following data generating process:

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \epsilon_i$$

where

- 1. The error in the response has a mean of 0 for all values of the predictor.
- 2. The error in the response for one subject is independent of the error in the response for all other subjects.
- 3. The variability in the error of the response is the same for all values of the predictor.
- 4. The errors follow a Normal Distribution.

This is the default "regression" analysis implemented in the majority of statistical packages.



Warning

A "hidden" (typically unstated but should not be ignored) condition is that the sample is representative of the underlying population. In the one sample case (Chapter 10), we referred to this as the errors being "identically distributed." We no longer use the "identically distributed" language for technical reasons; however, we still require that the sample be representative of the underlying population.

We note that "regression" need not require all four conditions imposed in Definition 18.3. Placing all four conditions on the error term results in a specific analytical model for the sampling distribution of the least squares estimates. Changing the conditions changes the way we model the sampling distribution.



Big Idea

The model for the sampling distribution of a statistic is determined by the conditions you place on the stochastic portion of the model for the data generating process.

We have stressed the implications of each condition individually. Figure 18.4 illustrates these conditions working together. The mean-0 condition implies that for a given value of the predictor, the average response is given by the line (shown as the dark green dot in the figure). The Normality condition implies that for a given value of the predictor, the response is distributed evenly about the regression line according to a Normal distribution; further, the shape of the Normal distribution implies that these responses will cluster about the line. The constant variance condition implies that while the responses vary around the line, they do so to the same degree, regardless of the value of the predictor; therefore, the model is just as precise for all values of the predictor. Finally, the independence condition implies the amount one observation deviates from the line is unrelated to the amount any other observation deviates from the line.

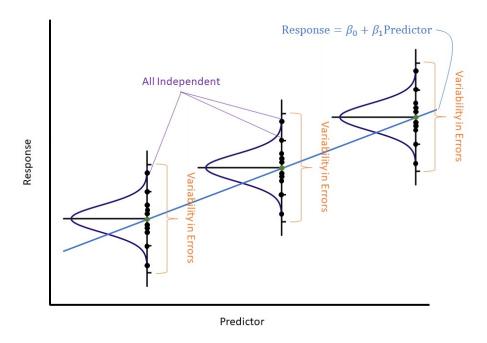


Figure 18.4: Illustration of the conditions on the error term for the classical regression model.

18.6 Imposing the Conditions

Let's return to our model for the data generating process of the bracketed duration as a function of the magnitude of the corresponding earthquake:

$$(\text{Bracketed Duration})_i = \beta_0 + \beta_1 (\text{Magnitude})_i + \varepsilon_i.$$

We were interested in the following research question:

Table 18.1: Summary of the linear model fit relating the bracketed duration at locations in Greece following an earthquake with the magnitude of the event.

Term	Estimate	Standard Error	Lower 95% CI	Upper 95% CI
(Intercept)	-19.194	3.975	-27.066	-11.323
Magnitude	4.484	0.724	3.050	5.917

On average, are changes in the magnitude of an earthquake associated with changes in the bracketed duration observed?

This was captured by the following hypotheses:

$$\begin{aligned} H_0: \beta_1 &= 0 \\ H_1: \beta_1 \neq 0 \end{aligned}$$

Using the method of least squares, we constructed point estimates of the parameters in the model; this leads to the following equation for estimating the average bracketed duration given the magnitude:

$$(Brackted Duration) = -19.19 + 4.48(Magnitude).$$

If we are willing to assume the data is consistent with the conditions for the classical regression model, we are able to model the sampling distribution (see Appendix A) of these estimates analytically and therefore construct confidence intervals. Table 18.1 summarizes the fit for the above model. In addition to the least squares estimates, it also contains the standard error (see Definition 6.4) of each statistic, quantifying the variability in the estimates. Finally, there is a 95% confidence interval for each parameter. Notice that based on the confidence interval for the slope, 0 is not a reasonable value for this parameter. Therefore, we have evidence that the slope coefficient associated with the magnitude differs from 0; that is, the sample provides evidence the average bracketed duration depends on the magnitude of the corresponding earthquake.

Chapter 6 described, in general, how confidence intervals are constructed. Under the classical regression model, there is an analytical model for the sampling distribution, and it is known. As a result, the confidence interval can be computed from a formula.

i Formula for Confidence Interval Under Classical Regression Model

If the classical regression model is assumed, the 95% confidence interval for the parameter β_j can be approximated by

$$\hat{\beta}_j \pm (1.96)$$
 (standard error of $\hat{\beta}_j$)

The confidence interval for the change in the average bracketed duration for each 1-unit increase in the magnitude of an earthquake (the slope β_1) was constructed assuming the classical regression model. Suppose, however, that we are only willing to impose the following conditions:

- The error in the bracketed duration is 0, on average, for earthquakes of any magnitude.
- The error in the bracketed duration for one earthquake is independent of the error in the bracketed duration for any other earthquake.

Since the conditions for the model of the data generating process have been altered, the model for the sampling distribution of the estimates will change, and therefore the corresponding confidence intervals will also change. Under these reduced conditions, we can appeal to a bootstrapping algorithm. Specifically, we could resample (with replacement) 119 earthquakes from the original data; for each resample, we compute the least squares fit (Figure 18.5). Since the observations selected change with each resample, the least squares estimates will also change. By repeating this process over and over again, we can obtain a model for how the estimates would change in repeated sampling.




Figure 18.5: Illustration of a single iteration of a bootstrap procedure to construct an empirical estimate of the sampling distribution for the estimates of the coefficients in a regression model.

Using the empirical model of the sampling distribution for each estimate, we can construct

Table 18.2: Summary of the linear model fit relating the bracketed duration at locations in Greece following an earthquake with the magnitude of the event. This summary only assumes the mean-0 and independence conditions.

Term	Estimate	Standard Error	Lower 95% CI	Upper 95% CI
(Intercept)	-19.194	4.960	-29.258	-9.545
Magnitude	4.484	0.965	2.593	6.437

a confidence interval for each parameter. These updated confidence intervals are shown in Table 18.2.

While the exact interval differs from what we computed previously, our overall conclusion remains the same — the sample provides evidence that the average bracketed duration depends on the magnitude of the corresponding earthquake. It is natural to ask, which confidence interval should we use? That depends on the conditions you are willing to assume, which is an issue we will tackle in Chapter 20.

18.7 Recap

We have covered a lot of ground in this chapter, and it is worth taking a moment to summarize the big ideas. In order to construct a model for the sampling distribution of the least squares estimates, we took a step back and modeled the data generating process. Such a model consists of two components: a deterministic component explaining the response as a function of the predictor, and a stochastic component capturing the noise in the system.

Certain conditions are placed on the distribution of the noise in our model. With a full set of conditions (classical regression model), we are able to model the sampling distribution of the least squares estimates analytically. We can also construct an empirical model for the sampling distribution of the least squares estimates assuming the data is consistent with fewer conditions.

In general, the more conditions we are willing to impose on the data-generating process, the more tractable the analysis; however, the most important aspect is that the data come from a process which is consistent with the conditions we impose, which is discussed in Chapter 20.

19 Quantifying the Quality of a Model Fit

In the previous two chapters, we described a model for describing the data generating process for a quantitative response as a function of a single quantitative predictor:

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i$$

Chapter 17 discussed obtaining estimates of these unknown parameters using the method of least squares. Chapter 18 imposed conditions on the stochastic portion of the model in order to develop a confidence interval for each parameter. In this chapter, we turn to performing inference through the computation of a p-value for a set of hypotheses, and we discuss how to quantify the quality of our model with regard to its utility in making predictions. It turns out these two tasks are very much related and are accomplished through partitioning variability. We will describe what we mean by partitioning variability and how it is used to derive a measure for the overall performance of a model and to develop a standardized statistic for comparing two models. As with previous discussions, these ideas will form another thread of the story that continues throughout the text.

19.1 Partitioning Variability

Let's return to the Seismic Activity Case Study first introduced in Chapter 13. Consider modeling the bracketed duration at a location as a function of the distance the location is from the center of the earthquake using the following model for the data generating process:

(Bracketed Duration)_i =
$$\beta_0 + \beta_1$$
(Epicentral Distance)_i + ε_i .

Using least squares to estimate the parameters, and assuming the data is consistent with the conditions for the classical regression model, the resulting model fit is summarized below in Table 19.1.

Remember, the goal of the model for the data generating process is to explain why the response is the value we see — we are essentially explaining why the values of the response differ from one individual unit to another (the variability in the response). Consider the model for the data generating process summarized above; it includes two reasons why the bracketed duration is not the same value at each measured location:

Table 19.1: Summary of the model fit explaining the bracketed duration as a function of epicentral distance.

Term	Estimate	Standard Error	Lower 95% CI	Upper 95% CI
(Intercept)	4.462	0.726	3.024	5.899
Epicentral Distance	0.029	0.018	-0.007	0.064

- The locations at which the observations are taken are different distances from the epicenter of each earthquake.
- Additional noise due to measurement error in the bracketed duration or additional natural sources we are unable to explain or did not account for in the model.

Looking at the form of the model for the data generating process, it may seem obvious that there are these two sources of variability — two sources for why the bracketed duration differs from one individual observation to another. Our next endeavor is to quantify the amount of variability in the response that can be attributed to each of these components. That is, we move forward with a goal of trying to say something like

$$\begin{pmatrix} \text{Total Variability} \\ \text{in the Bracketed Duration} \end{pmatrix} = \begin{pmatrix} \text{Variability due} \\ \text{to Distance} \end{pmatrix} + \begin{pmatrix} \text{Variability due} \\ \text{to Noise} \end{pmatrix}$$

As we have seen in both Chapter 5 and Chapter 12, variability can be quantified through considering the "total" distance the observations are from a common target (for example, the mean response) where "distance" is captured by squared deviations. That is, the total variability in bracketed duration can be measured by

$$\sum_{i=1}^{n} \left[(\text{Bracketed Duration})_i - (\text{Overall Mean Bracketed Duration}) \right]^2. \tag{19.1}$$

Notice this quantity is related to, but not equivalent to, the sample variance. It measures the distance each response is from the sample mean and then adds these distances up. This is known as the **Total Sum of Squares** since it captures the total variability in the response.

Definition 19.1 (Total Sum of Squares). The Total Sum of Squares, abbreviated SST, is given by

$$SST = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Overall Mean Response}) \right]^2$$

where the overall average response is the sample mean.

We now have a way of quantifying the total variability in the bracketed duration; we now want to partition (or separate) out this variability into its two components: the variability due to the epicentral distance, and the variability due to noise. In order to capture the variability due to epicentral distance, we consider how epicentral distance plays a role in the model for the data generating process: it forms the line which dictates the mean response. That is, the linear portion in the model for the data generating process $\beta_0 + \beta_1$ (Epicentral Distance) is the model's attempt to explain how changes in the epicentral distance explain changes in the bracketed duration; further, this explanation comes in the form of the average response. That is, plugging into the deterministic portion of the model for the data generating process provides a mean response, and if we use the least squares estimates in place of the parameters, we are computing an estimate of the mean response. Finding the variability in the bracketed duration due to the epicentral distance is then equivalent to finding the variability in these estimated (or predicted) mean responses:

$$\sum_{i=1}^{n} \left[(\text{Predicted Bracketed Duration})_i - (\text{Overall Mean Bracketed Duration}) \right]^2. \tag{19.2}$$

This is known as the **Regression Sum of Squares** as it captures the variability explained by the regression line.

Definition 19.2 (Regression Sum of Squares). The Regression Sum of Squares, abbreviated SSR, is given by

$$SSR = \sum_{i=1}^{n} \left[(\text{Predicted Mean Response})_i - (\text{Overall Mean Response}) \right]^2$$

where the predicted mean response is computed using the least squares estimates and the overall mean response is the sample mean.

Finally, the unexplained noise, ε in our model for the data generating process, is the difference between the actual response and the deterministic portion of the model (in our case, the true regression line). This variability in the noise is then the variability in the bracketed duration where the average is *conditional* on the epicentral distance instead of ignoring it (which is what happens when we use the overall sample mean bracketed duration):

$$\sum_{i=1}^{n} \left[(\text{Bracketed Duration})_i - (\text{Predicted Bracketed Duration})_i \right]^2. \tag{19.3}$$

This is known as the **Error Sum of Squares** as it captures the variability not explained by the model but represented by the error term in the model.

Definition 19.3 (Error Sum of Squares). The Error Sum of Squares, abbreviated SSE and sometimes referred to as the Residual Sum of Squares, is given by

$$SSE = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Predicted Mean Response})_i \right]^2$$

where the predicted mean response is computed using the least squares estimates.

With some clever algebra, it can be easily seen that the total variability does in fact partition into these two components. This discussion is represented in Figure 19.1.

Big Idea

The total variability in a response can be partitioned into two components: the variability explained by the predictor and the unexplained variability left in the error term. This is represented in the formula

$$SST = SSR + SSE$$

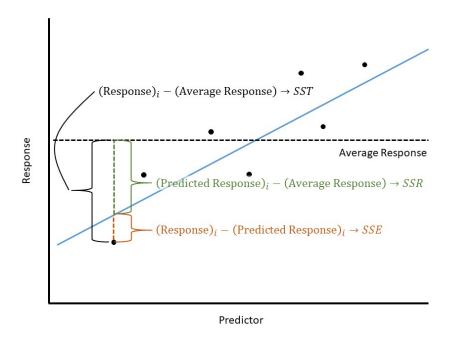


Figure 19.1: Illustration of partitioning the variability of a response using a regression model.

19.2 R-squared

The key to quantifying the quality of a model for the data generating process is to understand that a partition breaks a whole into smaller, distinct components. This means that if you put the components back together, you have the whole. The sums of squares partition the variability in the response into that explained by the deterministic portion of the model for the data generating process and that not explained. We represented this above by the equation

$$SST = SSR + SSE$$

.

The benefit partitioning variability is that it makes clear the breakdown between the variability in the response that the deterministic portion of the model is explaining (SSR) versus the variability in the response that cannot be explained (SSE). We are now in a position to quantify the proportion of the total variability the model is explaining, which is known as the **R-squared** value for the model.

Definition 19.4 (R-Squared). Sometimes reported as a percentage, the R-Squared value measures the proportion of the variability in the response explained by a model. It is given by

R-squared =
$$\frac{SSR}{SST}$$
.

For our model of the bracketed duration as a function of the epicentral distance, the R-squared value turns out to be 0.0216; that is, only 2.16% of the variability in the bracketed duration at a location is explained by its distance from the center of the corresponding earthquake.

As R-squared is a proportion, it must take a value between 0 and 1. If 0, that means our model has no predictive ability within our sample. That is, knowing the predictor does not add to our ability to predict the response any more than guessing. A value of 1 indicates that our model has predicted all the variability in the response; that is, given the predictor, we can perfectly predict the value of the response.

It may appear that obtaining an R-squared value of 1 should be our goal. And, in one sense, it is. We want a model that has strong predictive ability. However, there is a danger in obtaining an R-squared of 1 as well. We must remember that variability is inherent in any process. Therefore, we should never expect to fully explain all the variability in a response. George Box (a renowned statistician) once made the following statement (Box 1979):

"Now it would be very remarkable if any system existing in the real world could be exactly represented by any simple model. However, cunningly chosen parsimonious models often do provide remarkably useful approximations. For example, the law PV = RT relating pressure P, volume V and temperature T of an 'ideal' gas via a

constant R is not exactly true for any real gas, but it frequently provides a useful approximation and furthermore its structure is informative since it springs from a physical view of the behavior of gas molecules.

For such a model there is no need to ask the question 'Is the model true?'. If 'truth' is to be the 'whole truth' the answer must be 'No.' The only question of interest is 'Is the model illuminating and useful?'.

The idea here is that we know the model will not capture the data generating process precisely. Therefore, we should be skeptical of models which claim to be perfect. For example, consider the two models illustrated in Figure 19.2. The model represented by the black line has a perfect fit, but we argue the model represented by the blue line is better. While the black line captures all the variability in the response for this sample, it is certainly trying to do too much. In reality, the blue line captures the underlying relationship while not overcomplicating that relationship. We sacrifice a little quality in the fit for this sample in order to better represent the underlying structure of the population. The black line suffers from what is known as overfitting; the blue line is a more parsimonious (simple) model, balancing complexity with model fit.




Figure 19.2: Illustration of a parsimonious model compared to one which overfits the data.

Students often ask, "if not 1, how high of an R-squared represents a good model?" The answer depends a lot on the discipline. In many engineering applications within a lab setting, we can control much of the external variability leading to extremely high R-squared values (0.95 to 0.99). However, in biological applications, the variability among the population can be quite

large, leading to much smaller R-squared values (0.3 to 0.6). What is considered "good" can depend on the specific application.

Warning

While R-squared is useful for quantifying the quality of a model on a set of data, it should not be used to compare two different models as R-squared always favors more complex models. There are better methods which adjust for the complexity of the model fit.

In addition to the discipline, how you view the R-squared value for a model may depend on the goal of the model. There are generally two broad reasons for developing a statistical model:

- Explain the relationship between a response and one or more predictors. This can involve examining the marginal relationship, isolating the effect, or examining the interplay between predictors.
- Predict a future response given a specific value for the predictors.

If all we are interested in doing is explaining the relationship, we may not be concerned about the predictive ability of the model. That is, since our goal is not to accurately predict a future response, we are primary concerned with whether we have evidence of a relationship. But, if our goal is prediction, we would like that estimate to be precise. In such cases, a high R-squared is required before really relying on the model we have.

What is perhaps counter-intuitive is that, while related, hypothesis testing and the R-squared value may not necessarily yield the same conclusions. That is, it is possible that we have a strong evidence that the average response depends on the predictor (small p-value) and simultaneously conclude that using the predictor would not result in precise predictions (low R-squared value).

19.3 Hypothesis Testing

In addition to quantifying the quality of the model, partitioning the variability in a response into two components is the basis for conducting a hypothesis test to compare two models. In this section, we expand upon the ideas initially presented in Chapter 12, broadening them to add to our unifying framework. Recall that hypothesis testing is really about comparing two models for the data generating process: a more complex model in which the parameters are free to take on any value, and a restricted model in which the parameters are constrained in some way.

When the sample does not provide enough evidence to suggest the more complex model is necessary to explain the variability in the response, we conclude it is reasonable the reduced model for the data generating process is appropriate (some say, we "fail to reject" the null hypothesis). When the sample does provide sufficient evidence to suggest we can discern the difference in the performance of the reduced and complex model, we say the simple model is not sufficient for explaining the variability in the response, and we prefer the more complex model (some say, we "reject" the null hypothesis). Throughout the remainder of this section, we will consider the following research question:

Is there evidence the average bracketed duration for a location following an earthquake is linearly related to the distance the location is from the center of the earthquake?

If we consider the simple linear model for the data generating process described above, this question can be captured using the following set of hypotheses:

$$H_0: \beta_1 = 0 \qquad \text{vs.} \qquad H_1: \beta_1 \neq 0.$$

Again, hypothesis testing is really model comparison; that is, these hypotheses are really suggesting two separate models for the data generating process:

$$\begin{split} \text{Model 1:} \quad & (\text{Bracketed Duration})_i = \beta_0 + \beta_1 (\text{Epicentral Distance})_i + \varepsilon_i \\ \text{Model 0:} \quad & (\text{Bracketed Duration})_i = \beta_0 + \varepsilon_i. \end{split}$$

The model under the null hypothesis (Model 0) has fewer parameters because it is a constrained version of Model 1 resulting from setting $\beta_1 = 0$. In fact, while Model 1 says that there are two components (the epicentral distance and noise) contributing to the variability observed in the bracketed duration, Model 0 says that there is only a single component (noise). So, we can think of our hypotheses as

 H_0 : Model 0 is sufficient for explaining the variability in the response H_1 : Model 0 is not sufficient for explaining the variability in the response.

Regardless of which model we choose, the total variability in the response remains the same. We are simply asking whether the variability explained by the predictor is sufficiently large for us to say it has an impact. In particular, if the null hypothesis were true, we would expect all the variability in the response to be channeled into the noise $(SST \approx SSE)$. In fact, think about computing the error sum of squares for Model 0 above; it would be

$$SSE_0 = \sum_{i=1}^n \left[(\text{Bracketed Duration})_i - (\text{Overall Average Bracketed Duration}) \right]^2$$

since the least squares estimate of $\hat{\beta}_0$ in Model 0 is the sample mean (see Appendix B). But, this is equivalent to the total sum of squares for Model 1 (Equation 19.1). This confirms

our intuition that if the null hypothesis were true, we would expect all the variability in the response to be channeled into the noise $(SST \approx SSE)$.

If, however, the alternative hypothesis is true and the epicentral distance explains some portion of the variability in the bracketed duration, then we would expect some of the variability to be channeled out of the noise term (SSR > 0). Because we have partitioned the variability, we now take a moment to recognize that

$$SSR = SST - SSE$$
,

but we know that the total sum of squares is just the error sum of squares from the reduced model (Model 0) as shown above. Therefore, we can write

$$SSR = SSE_0 - SSE_1, (19.4)$$

where we use the subscripts to denote whether we are discussing the error sum of squares from the reduced model (Model 0) or the full unconstrained model (Model 1). That is, Equation 19.4 reveals that the regression sum of squares is the equivalent of the shift in the error sum of squares as we move from the reduced model under the null hypothesis to the more complex model under the alternative hypothesis.



For a particular dataset, the regression sum of squares quantifies the shift in the error sum of squares as we move from a reduced model to a more complex model. It measures the "signal" in the data represented by the more complex model for the data generating process.

The regression sum of squares represents our signal. The larger the value, the more evidence we have that the data is not consistent with the null hypothesis. However, as we saw in Chapter 12, we should always examine our signal relative to the noise in the data. We already have a measure for the amount of variability due to noise — the error sum of squares! It then seems reasonable to consider the ratio

$$\frac{SSR}{SSE_1} = \frac{SST - SSE_1}{SSE_1} = \frac{SSE_0 - SSE_1}{SSE_1},$$

where again we have added subscripts to emphasize from which model we are computing the sums of squares. While this is a reasonable statistic, it is not yet standardized. Remember that sums of squares capture variability but are themselves not variances, and it turns out a ratio of variances is easier to model analytically. If we take a sum of squares and divide by an appropriate term, known as the **degrees of freedom**, we get a true variance term.

Definition 19.5 (Degrees of Freedom). A measure of the flexibility in a sum of squares term; when a sum of squares is divided by the corresponding degrees of freedom, the result is a variance term.

Rationale for Degrees of Freedom

Degrees of freedom are a very difficult concept to grasp, even for those who have been studying statistics for a while. Here is our way of thinking about them — they are the difference of available terms to work with. For example, think about the total sum of squares associated with a full unconstrained linear regression model described in Equation 17.3:

$$SST = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Overall Mean Response}) \right]^2.$$

The first term of the difference has n different values (one response for each observation). However, the sample mean is just one value. Therefore, there are n-1 degrees of freedom associated with the total sum of squares. This is often described as starting out with n estimates (the data), but needing to estimate one parameter (the mean) along the way, leading to n-1.

Similarly, consider the regression sum of squares for the full unconstrained model:

$$SSR = \sum_{i=1}^n \left[(\text{Predicted Mean Response})_i - (\text{Overall Mean Response}) \right]^2.$$

While there are n predicted values, they are all generated from the same least squares fit $\hat{\beta}_0 + \hat{\beta}_1(\operatorname{Predictor})_i$ which can be computed from two estimates (that for the intercept and slope). Therefore, we begin with only 2 unique values. Again, the sample mean has just one value, leading to 2-1=1 degree of freedom associated with the regression sum of squares.

Finally, consider the error sum of squares for the full unconstrained model:

$$SSE = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Predicted Mean Response})_i \right]^2.$$

We have n initial values (one for each observation). However, as described above, we only need 2 terms to estimate the predicted values. So, we have n-2 degrees of freedom associated with the error sum of squares.

Note that just as the sums of squares formed a partition (SST = SSR + SSE), the corresponding degrees of freedom form a partition ((n-1) = (2-1) + (n-2)).

Again, dividing a sum of squares by its associated degrees of freedom creates a variance term; this term is known as a **mean square**. It is important to note that both sums of squares and mean squares quantify the components of variability in the response, the component explained

by the deterministic portion of the model and the component that is unexplained. However, they serve different purposes.

Definition 19.6 (Mean Square). A mean square is the ratio of a sum of squares and its corresponding degrees of freedom. For a model of the form in Equation 17.3, we have

- Mean Square Total (MST): estimated variance of the responses; this is the same as the sample variance of the response.
- Mean Square for Regression (MSR): estimated variance of the predicted responses.
- Mean Square Error (MSE): estimated variance of the error terms; this is equivalent to the estimated variance of the response for a given value of the predictor (the variance of the response about the regression line).

In each case, the mean square is an estimated variance.

Note

Sums of squares partition the variability of the response into smaller components; mean squares estimate the variance of those smaller components. While SST = SSR + SSE, note that $MST \neq MSR + MSE$.

Since mean squares are proportional to their corresponding sum of squares, an increase in the sum of squares is associated with an increase in the corresponding mean square. We are now ready to define our standardized statistic as the ratio of mean squares. Instead of the ratio

$$\frac{SSR}{SSE_1} = \frac{SST - SSE_1}{SSE_1} = \frac{SSE_0 - SSE_1}{SSE_1},$$

we replace the numerator and denominator with mean squares such that

$$\frac{MSR}{MSE} = \frac{\left(SST - SSE_1\right)/(2-1)}{SSE_1/(n-2)} = \frac{\left(SSE_0 - SSE_1\right)/(2-1)}{SSE_1/(n-2)},$$

where again we have added subscripts to emphasize from which model we are computing the sums of squares. This standardized statistic could be used to quantify the signal-to-noise ratio (the amount of evidence) in the sample for testing the hypotheses

$$\begin{aligned} H_0: \beta_1 &= 0 \\ H_1: \beta_1 \neq 0. \end{aligned}$$

However, we can generalize this for testing a range of hypotheses within our model for the data generating process.

Definition 19.7 (Standardized Statistic for Simple Linear Regression). Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 17.3):

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i$$
.

Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0^1 . A standardized statistic, sometimes called the "standardized F statistic," for testing the hypotheses is given by

$$T^{*}=\frac{\left(SSE_{0}-SSE_{1}\right)/(2-r)}{SSE_{1}/(n-2)},$$

where r is the number of parameters in the reduced model. Defining

$$MSA = \frac{SSE_0 - SSE_1}{2 - r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model.

It can be shown that this standardized statistic is a generalization of the one introduced in Chapter 12. The numerator captures the signal by examining the difference between what we expect the error sum of squares to be under the null hypothesis and what we actually observe; the denominator captures the background noise (relative to the estimated mean response from the full model). Larger values of this standardized statistic indicate more evidence in the sample against the null hypothesis.

We should not lose sight of the fact that our standardized statistic is really a result of partitioning the variability and considering the variability explained by the predictor relative to the noise in the response. Underscoring that the standardized statistic is a result of this partitioning, the analyses of these sources of variability is often summarized in a table similar to that represented in Figure 19.3.

The last entry in the table is the p-value. As with any p-value, it is computed by finding the likelihood, assuming the null hypothesis is true, of obtaining a standardized statistic, by

¹Technically, this standardized statistic only applies to a class of potential hypotheses, but that class is quite broad (see Appendix B).




Figure 19.3: Table summarizing the partitioning of variability in a regression model.

Table 19.2: Analysis of the sources of variability in the bracketed duration as a function of epicentral distance.

Term	DF	Sum of Squares	Mean Square	Standardized Statistic	P-Value
Epicentral Distance	1	85.733	85.733	2.583	0.111
Error	117	3883.708	33.194		

chance alone, as extreme or more so than that observed in our sample. "More extreme" values of the statistic would be larger values; so, the area to the right in the model for the null distribution is needed. The key step is modeling that null distribution. This is where the conditions we place on the error term that were discussed in Chapter 18 come into play. Under the classical regression conditions, we can model the null distribution analytically (see Appendix A); otherwise, we can rely on bootstrapping to model the null distribution.

Let's return to our question that inspired our investigation:

Is there evidence the average bracketed duration for a location following an earthquake is linearly related to the distance the location is from the center of the earthquake?

This corresponds to testing the hypotheses

$$H_0: \beta_1 = 0$$
 vs. $H_1: \beta_1 \neq 0$.

Table 19.2 gives the table summarizing the partitioned sources of variability in the bracketed duration. We have a large p-value (computed assuming the data is consistent with the classical regression model). That is, the sample provides no evidence to suggest that locations further from the center of the earthquake experience a bracketed duration which differs, on average, from those closer to the center of the earthquake.



🅊 Big Idea

Determining if a response is linearly related to a predictor is done by determining if the predictor explains a significant portion of the variability in the response.

In this section, we partitioned variability as a way of evaluating the strength of evidence the predictor plays in determining the response. In the previous section, we used that same partition to quantify the predictive ability of the model for the data generating process. Regardless of our goal, conducting inference or predicting a future response, partitioning the variability is a key step. If inference is our primary aim, this partitioning allows us to determine if a predictor adds to the model above and beyond the error alone. If prediction is our primary aim, the partitioning allows us to quantify the quality of the model's predictive ability.

20 Assessing the Modeling Conditions

We have been considering the simple linear model described in Equation 17.3 for the data generating process of a quantitative response. For example, for the Seismic Activity Case Study, we considered a model that explained the bracketed duration at a location as a function of the magnitude of the earthquake:

(Bracketed Duration)_i =
$$\beta_0 + \beta_1$$
 (Magnitude)_i + ε_i .

Estimates for the unknown parameters in this model were obtained using the method of least squares (Chapter 17). In order to obtain a model for the sampling distribution of these estimates (or the null distribution as appropriate), and thereby conduct inference, we added conditions to the distribution of the error term (Chapter 18). For example, under the "classical regression model" (Definition 18.3) we require the following four conditions:

- 1. The error in the bracketed duration has an average of 0 regardless of the magnitude of the earthquake.
- 2. The error in the bracketed duration for one location is independent of the error in the bracketed duration for any other location.
- 3. The variability of the error in the bracketed duration is the same regardless of the magnitude of the earthquake.
- 4. The errors in the bracketed duration follow a Normal distribution.

We also saw in Chapter 18, that we could develop an empirical model for the sampling distribution of the least squares estimates only enforcing the first two of these conditions on the distribution of the error. Which of these two models for the sampling distribution should be used? Unfortunately, we cannot simply state conditions and then proceed blindly. In order to rely on the p-values and confidence intervals produced from any modeling procedure, the data must be consistent with these conditions.

In this section, we discuss how we assess these conditions qualitatively.

20.1 Residuals

One of the complications we face is that we are imposing conditions on the error term, but we do not observe the error (since the parameters are unknown). However, we are able to

determine the difference between each observation with respect to the estimated model for the data generating process. This difference between each observed response and what we would have predicted for this observation using the least squares estimates, called a residual, is analogous to the error term if the parameters were known. Therefore, residuals should behave like a sample of error terms.

Definition 20.1 (Residual). The difference between the observed response and the predicted response (estimated deterministic portion of the model). Specifically, the residual for the i-th observation is given by

$$(Residual)_i = (Response)_i - (Predicted Mean Response)_i$$

where the "predicted mean response" is often called the predicted, or fitted, value.

Residuals mimic the noise in the data generating process.

For the simple linear regression model, the predicted mean response is computed by plugging into the formula

$$\hat{\beta}_0 + \hat{\beta}_1(\text{Predictor})_i$$

where $\hat{\beta}_0$ and $\hat{\beta}_1$ are the least squares estimates. We can use the residuals to qualitatively assess if the observed data is consistent with each of the four potential conditions we might place on the distribution of the error term.



Big Idea

Residuals, since they mimic the noise in the data generating process, provide a way of assessing the modeling conditions placed on the distribution of the error term. The conditions are placed on the error term, but they are assessed with residuals.

20.2 Assessing the Mean-0 Condition

The error in the bracketed duration has an average of 0 regardless of the magnitude of the earthquake.

It is tempting to read this condition and believe that a rational way to assess this condition is to determine if the average of the residuals is 0. However, while the difference is subtle, the condition is not that the average error is 0; the condition is that the average error is 0 for all values of the predictor. It would seem we need to determine if, for each value of the predictor possible, if the residuals average to 0. This is infeasible numerically because we do not generally have multiple responses for each value of the predictor. We can, however, assess whether the data is consistent with this condition graphically. That is, in order to assess this condition, we need to graphically assess how the average behaves over a range of predictor values. We capture this by looking at the *predicted (or fitted) values*. Figure 20.1 shows the relationship between the residuals and the associated predicted values for the observations in the data set.

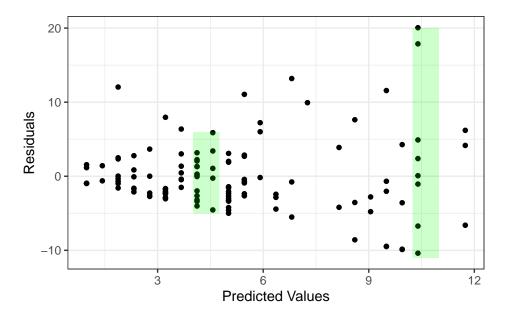


Figure 20.1: Plot of the residuals vs. the predicted values for a model predicting bracketed duration as a function of the magnitude of an earthquake.

If the errors have a mean of 0 for all values of the predictor, then we would expect the residuals to have a mean of 0 for all predicted values. That is, if the data is consistent with the mean-0 condition, then as we move left to right across the plot, the residuals should tend to balance out at 0 everywhere along the x-axis. Imagine a "window" around the residuals (shown in the figure as vertical green rectangles), and imagine moving that window from left to right. If that window has to shift up or down to contain the cloud of residuals (so that the window is no longer centered around 0), that indicates a problem. Any trends in the *location* of this graphic would indicate the data is *not* consistent with the mean-0 condition.

Graphically Assessing the Mean-0 Condition

If the data is consistent with the mean-0 condition, there should be no trends in the *location* of the plot of the residuals against the predicted values.

As we examine Figure 20.1, the residuals do tend to balance out at 0 everywhere along the x-axis. That is, we do not see a trend in the location of the residuals as the predicted values

increase. Therefore, it is reasonable to say the sample is consistent with the mean-0 condition.

20.3 Assessing the Independence Condition

The error in the bracketed duration for one location is independent of the error in the bracketed duration for any other location.

Generally, independence is assessed by considering the method in which the data was collected and considering the context with a discipline expert. By carefully considering the manner in which the data was collected, we can typically determine whether it is reasonable that the errors in the response are independent of one another. Some key things to consider when examining the data collection process:

- Are there repeated observations of the same variable made on the same subject? This often suggests some type of relationship between the observed responses and therefore would not be consistent with errors being independent (see Chapter 33 and Chapter 34).
- Is the response measured over time, such as daily temperature over the course of a month? Data collected over time often exhibits strong period-to-period relationships suggesting the errors are not independent. For example, if the temperature today is above average, it is more likely to be above average tomorrow as well.
- Is there a learning curve in how the data was collected? Learning curves again suggest some dependence from one observation to the next. For example, a new nurse may become better at collecting pulse readings with more practice over time.
- Measurement devices which are failing over time will introduce a dependence from one
 observation to the next. Imagine a bathroom scale that begins to add an additional
 pound each day. Then, being an above average weight one day will most likely lead to
 an above average weight the next, due primarily to the measurement device.

These last three points illustrate a particular deviation from our condition of independence in which two observations collected close together in time are related. When we know the order in which the data was collected, we can assess whether the data tends to deviate from the condition of independence in this manner. This is done graphically through a **time-series plot** of the *residuals*. If two errors were unrelated, then the value of one residual should tell us nothing about the value of the next residual. Therefore, a plot of the residuals over time should look like noise (since residuals are supposed to mimic the noise in the model). If there are any trends, then it suggests the data is not consistent with independence.

Definition 20.2 (Time-Series Plot). A time-series plot of a variable is a line plot with the variable on the y-axis and time on the x-axis.

i Graphically Assessing the Independence Condition

If the data is consistent with the independence condition, we would not expect to see a trend in the *location* or *spread* in a time-series plot of the residuals. Note that this graphic can only be used if the order in which the data was collected is known, and the order indicates some natural timing.

Figure 20.2 shows two hypothetical datasets. In Panel A, the residuals display a trend in the location over time. Knowing that a response was below average suggests the next response will also be below average. In Panel B, the results display a trend in the spread over time. This suggests that measurements taken later in the study were less precise. Both panels are examples of patterns which would suggest the data is not consistent with the condition of independence.




Figure 20.2: Examples of trends in a time-series plot of the residuals. Such trends indicate the data is not consistent with the condition that the errors are independent of one another.

Instead, if the data were consistent with the condition of independence on the error terms, we would expect to see a plot similar to Figure 20.3. Notice there are no trends in the location or spread of the residuals.

For the Seismic Activity Case Study, the data was actually collected over time as earthquakes occurred. More, as technology has changed over time, it is reasonable to fear that the errors in our observations are related over time. In order to assess this, consider the time-series plot of the residuals from our model for the data generating process in which the bracketed duration

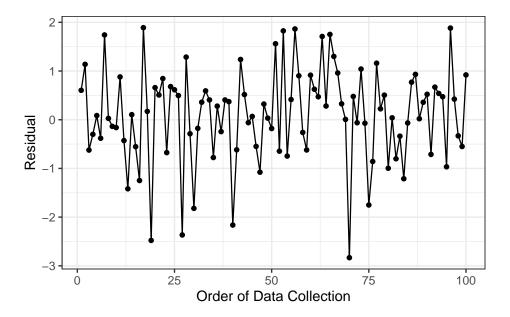


Figure 20.3: Example of a time-series plot of residuals which shows no trends in location or spread. This is consistent with what we would expect if the condition of independence among errors were satisfied.

is modeled as a linear function of the magnitude of the earthquake (Figure 20.4). Based on the figure, there are no clear trends in either the *location* or *spread* of the residuals over time (the figure resembles noise with no patterns). As a result, it is reasonable to assume the data is consistent with the errors being independent of one another.

The condition of independence is another reason we consider randomization when collecting data. Both random sampling and random assignment reduces the likelihood of the errors in two observations being related.

20.4 Assessing Homoskedasticity

The variability of the error in the bracketed duration is the same regardless of the magnitude of the earthquake.

When assessing the "mean-0 condition" above, we saw that the key phrase was "for all values of the predictor." Similarly, homoskedasticity suggests the variability in the errors is consistent for all values of the predictor. Therefore, we rely on the same graphical assessment — a plot of the residuals against the predicted values. However, instead of focusing on a trend in the location of the residuals, we are focused on a trend in the *spread*.




Figure 20.4: Time-series plot of the residuals for a model predicting bracketed duration as a function of the magnitude of an earthquake.

i Graphically Assessing the Constant Variance Condition

If the data is consistent with the constant variance condition, there should be no trends in the *spread* of the plot of the residuals against the predicted values.

Again, imagine a window (illustrated as vertical green rectangles in Figure 20.1) around the residuals. As we move left to right, if the size of the window has to change in order to keep the residuals inside (the window stretches or compresses vertically), then that is an indication that the spread of the residuals is changing. For our example, there is a clear "fan shape" to the residuals as we move left to right suggesting the precision of the model decreases when making larger predictions. This goes back to something we observed in Chapter 16 when examining a plot of the response and predictor (see Figure 16.1). We observed that for large earthquakes (high magnitudes), the bracketed duration was much more variable than for smaller earthquakes. So, our model is not as precise for some values of the predictor. This is evidence that our data is *not* consistent with the condition that the errors have the same variability for all values of the predictor.

This partially explains the differences in the confidence intervals reported in Table 18.1 and Table 18.2. Since there is clear evidence that the data is not consistent with the constant variance condition, then it is not safe to assume the classical regression model. That is, the confidence intervals and p-values, as well as the underlying models for the sampling distribution and null distribution that generated them, constructed assuming the data is consistent

with all four conditions, are suspect. We should instead rely on an empirical model for the sampling distribution of the least squares estimates when constructing confidence intervals or an empirical model for the null distribution of the standardized statistic if computing a p-value.

20.5 Assessing Normality

The errors in the bracketed duration follow a Normal distribution.

Assessing whether observations adhere to a particular distribution is a large area in statistical research. Many methods have been developed for this purpose. We emphasize a single graphical summary known as a **probability plot**. The construction of the plot is beyond the scope of this text, but the concepts underlying its construction actually tie in nicely to the big themes we have been discussing. Recall that if a sample is representative, then it should be a snapshot of the underlying population. Therefore, if we believe the underlying population has some particular distribution, we would expect the properties of this distribution to be apparent in the sample as well.

If we believe the errors follow a Normal distribution, then it is reasonable that the residuals should maintain some of those properties. For example, the 10-th percentile of the residuals should roughly equate to the 10-th percentile expected from a Normal distribution. Mapping the percentiles that we observe to those that we expect is the essence of a probability plot.

Definition 20.3 (Probability Plot). Also called a "Quantile-Quantile Plot", a probability plot is a graphic for comparing the distribution of an observed sample with a theoretical probability model for the distribution of the underlying population. The quantiles observed in the sample are plotted against those expected under the theoretical model.

While a probability plot can be constructed for a host of probability distributions, the most common is the Normal probability plot. The plot compares the observed residuals with those we would expect if the residuals were sampled from a Normal distribution. If the residuals closely aligned with our expectations from a Normal distribution, then we would expect to see a straight line if these are plotted against one another. Trends away from a straight line suggest the proposed Normal distribution is not a reasonable model for the distribution of the errors.

i Graphically Assessing the Normality Condition

If the data is consistent with the normality condition, the Normal probability plot of the residuals should exhibit a straight line with any deviations appearing to be random. Systemic deviations from a straight line indicate the observed distribution does not align with the proposed model.

Figure 20.5 shows the Normal probability plot of the residuals from our model for the data generating process of the bracketed duration.




Figure 20.5: Probability plot of the residuals for a model predicting bracketed duration as a function of the magnitude of an earthquake.

We note there is some systematic departure from a straight line in the graphic. In particular, on the left-hand side of the graphic, from -3 to -1.5 along the x-axis, the points tend to curve upward before changing direction at -1.5 on the x-axis. While we want to avoid over-interpreting small deviations from the linear trend, we should pay attention to clear departures.

We note that of the conditions considered, Normality is probably the least important as the analytic models for the sampling distributions of the least squares estimates are generally fairly robust to this condition. This is especially true in large samples (see Appendix A). However, we can always relax this condition by building an empirical model for the sampling distributions of the least squares estimates. Given the curvature we observed in this graphic, we would consider an empirical model, especially given we have already established the data is not consistent with the condition of homoskedasticity.

For comparison, Figure 20.6 illustrates a hypothetical dataset for which the residuals suggest the Normality condition of the errors is unreasonable as well as an example of when the residuals are consistent with the Normality condition on the errors.




Figure 20.6: Two Normal probability plots of the residuals for two hypothetical datasets. The trend away from a straight line in Panel A suggests assuming the errors follow a Normal distribution would be unreasonable. However, the reasonably straight line in Panel B suggests the data is consistent with the errors following a Normal distribution.

20.6 General Tips for Assessing Assumptions

Each of the methods presented here are qualitative assessments, which means they are subjective. That is okay. As the analyst, in consultation with the discipline expert, it is up to us to determine which conditions we are willing to assume are reasonable to impose. That is, with which conditions do we believe the data are consistent? Here are four overall things to keep in mind.

First, do not spend an extraordinary amount of time examining any one residual plot. If we stare at a plot too long, we can convince ourselves there is a pattern in anything. We are looking for glaring evidence that the data is not consistent with the conditions we have imposed on our model. This is especially true when we have only a few observations. In these settings, reading plots can be very difficult. Again, it is about what we are comfortable assuming.

Second, we have chosen the language carefully throughout this chapter. We have never once stated that a condition "was satisfied." When we perform an analysis, we are making an assumption that the conditions are satisfied. We can never prove that they are; we can only show that the data is consistent with a particular set of conditions. We can, however, provide evidence that a condition is violated. When that is the case, we should be wary of trusting the resulting p-values and confidence intervals which are constructed from imposing that condition. This is not unlike hypothesis testing; just as we can never prove the null hypothesis is true, we cannot prove that a condition is satisfied. The graphics are constructed using residuals, but what we are looking for comes from how we expect the residuals to behave if a particular condition is true. Just because residuals behave a certain way, however, does not guarantee a condition is or is not met. Just like hypothesis testing, we are looking for what we are able to discern given the data.



Big Idea

We cannot prove a condition is satisfied; we can only hope to show the data is consistent with the condition and it is therefore reasonable to assume it is satisfied.

Third, any conditions required for a particular analysis should be assessed. If our sample is not consistent with the necessary conditions, we should choose a different analysis. The inference we obtain from an analysis is only reliable if the data is consistent with any necessary conditions.

Finally, transparency is crucial. Perhaps more important than which conditions we choose to impose is clearly communicating to our audience our decisions. If our analysis is to be replicated or critiqued fairly (and any good scientist or engineer should welcome such critique), we need to be transparent about the decisions we make along the way, including the assumptions we make on the model for the data generating process.

21 Extending the Regression Model

The last several chapters have developed an approach for modeling a quantitative response as a function of a single quantitative predictor using a model of the form:

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i$$
.

This model is well suited for addressing questions about the marginal relationship between two variables. However, as we saw in Chapter 14, not all our questions are about the marginal relationship. The real power of the model in Equation 10.1 is our ability to generalize it to encompass multiple predictors and various types of relationships. In this chapter, we briefly describe how to extend the regression model to address some additional questions of interest.

While this chapter can be considered optional for an introductory treatment to modeling, the concepts discussed here allow us to easily bridge into the remaining topics covered in later units.

21.1 Using a Categorical Predictor

We have described the simple linear model (Equation 17.3) as one which relates two quantitative variables. However, this framework can be extended to make use of a categorical predictor. Specifically, we will consider a binary predictor which categorizes units into one of two groups. More general categorical predictors will be considered in the next unit.

Continuing to work with the data from the Seismic Activity Case Study, Table 21.1 summarizes the bracketed duration at locations based on the soil conditions.

Table 21.1: Numerical summaries of the bracketed duration for locations based on their soil conditions.

Soil Condition	n	Sample Mean	Sample Standard Deviation
Intermediate	43	5.16	4.52
Rocky	24	4.14	5.72
Soft	52	5.86	6.73

We observe that the bracketed duration appears to differ, on average, for locations with rocky soils compared to other types of location. We would like to consider a model for the data generating process of the form

$$(Bracketed Duration)_i = \beta_0 + \beta_1(Soil Condition)_i + \varepsilon_i,$$

but the problem is that we do not know how to multiply a number β_1 and a category like "rocky soil," meaning the above model does not yet even make sense. The key to including a categorical variable in the model for the data generating process is to construct an **indicator** variable.

Definition 21.1 (Indicator Variable). An indicator variable is a binary (takes the value 0 or 1) variable used to represent whether an observation belongs to a specific group defined by a categorical variable.

Indicator variables essentially create a numeric variable which represents a yes/no decision regarding a categorical variable. For example, consider the following variable definition:

$$(\text{Rocky Soil})_i = \begin{cases} 1 & \text{if i-th location has rocky soil} \\ 0 & \text{if i-th location has a different soil type.} \end{cases}$$

Since this variable is numeric, we can use it in our model for the data generating process. Specifically, consider the following model

$$(Bracketed Duration)_i = \beta_0 + \beta_1 (Rocky Soil)_i + \varepsilon_i$$
(21.1)

With a variable capturing when the soil is rocky, we might be tempted to include another indicator variable that takes the value 1 when the soil is not rocky. But, this is not necessary. Think of an indicator variable like a "light switch." The indicator variable turns on when an observation falls into a particular group and turns off otherwise. Just like a single light switch can place a light into the "on" position and into the "off" position, a single indicator variable can capture two groups in the model for the data generating process. If we have a location which has "Intermediate" soil conditions, then that location cannot have "Rocky" soil; therefore, the indicator variable in our model turns off. Setting the indicator to 0 (turning it "off") then leaves only the intercept in the model. The group which has a 0 for the indicator is encoded in the intercept; this is known as the **reference group**.

Definition 21.2 (Reference Group). The group defined by setting all indicator variables in a model for the data generating process equal to 0.

Remember, provided that we have imposed the mean-0 condition, we know that the deterministic portion of the model specifies the mean response for a given value of the predictor. In the case of an indicator variable, the predictor variable can only take on two values (0 or 1). Therefore, it only predicts two mean responses:

- When the indicator variable is 0, the model predicts a mean response of β_0 ; and,
- When the indicator variable is 1, the model predicts a mean response of $\beta_0 + \beta_1$.

So, the "slope" β_1 is really capturing the shift in the deterministic portion of the model from one group to the other. This coincides with the interpretation of the slope discussed in Chapter 18 — β_1 represents the change in the mean response when moving from one group to another. That is, the slope represents the difference in the average response between the two groups.

Interpretation of the Parameters

Suppose we have a simple linear regression model (Equation 17.3) with a quantitative response and a single indicator variable as the sole predictor. Then, the intercept represents the average response for the reference group, and the slope represents the difference in the average response between the two groups.

This shift in the mean response is illustrated in Figure 21.1. The "line" is really connecting the sample mean response of each group.




Figure 21.1: Comparison of the bracketed duration between locations with rocky soil and those with other soil types.

21.2 Including Multiple Precitors

The real power of the model in Equation 10.1 is our ability to generalize it to encompass multiple predictors and various types of relationships. That is, suppose that instead of being interested in the marginal relationship between the bracketed duration and the magnitude of the corresponding earthquake we are interested in *isolating* the effect of the magnitude on the bracketed duration:

On average, does the bracketed duration change as the magnitude of the corresponding earthquake changes if the location remains the same distance from the center of the earthquake?

This question is asking if there is an effect of the magnitude of the earthquake above and beyond the impact due to a locations distance from the epicenter. This question requires a model which has multiple predictors. What bracketed duration would we expect given the magnitude and epicentral distance? We extend the simple linear model to include an additional predictor:

$$(Bracketed Duration)_i = \beta_0 + \beta_1 (Magnitude)_i + \beta_2 (Epicentral Distance)_i + \varepsilon_i.$$
 (21.2)

This more complex model is more difficult to visualize, but conceptually it is similar to the simple linear model in Equation 17.3. Given a value for the magnitude and epicentral distance, we can predict the bracketed duration; the model for the data generating process allows both variables to contribute to the bracketed duration. Our model for the data generating process also allows for random noise to contribute to the bracketed duration; that is, there will still be unexplained variability. One way of envisioning what this model does is to think about taking the linear relationship we previously had and observing that we are now saying that the deterministic portion of the model for the data generating process differs for each group of observations which have a different epicentral distance. For example, consider all locations which were located 10 km away from the center of an earthquake; for this group of earthquakes, Equation 21.2 suggests the bracketed duration is generated by

$$\begin{split} (\text{Bracketed Duration})_i &= \beta_0 + \beta_1 (\text{Magnitude})_i + \beta_2 (10) + \varepsilon_i \\ &= (\beta_0 + 10\beta_2) + \beta_1 (\text{Magnitude})_i + \varepsilon_i. \end{split}$$

Similarly, if we only consider locations which were located 32 km away from the center of an earthquake, then Equation 21.2 suggests the bracketed duration is generate by

$$\begin{aligned} (\text{Bracketed Duration})_i &= \beta_0 + \beta_1 (\text{Magnitude})_i + \beta_2 (32) + \varepsilon_i \\ &= (\beta_0 + 32\beta_2) + \beta_1 (\text{Magnitude})_i + \varepsilon_i. \end{aligned}$$

Figure 21.2 represents this graphically for a range of potential epicentral distances. Essentially, the relationship between the bracketed duration and the magnitude shifts depending on the epicentral distance. The overall trend is similar (the lines are parallel), but where the line is located is really dependent upon the distance of the location from the earthquake. Note that while Figure 21.2 chooses particular values for the epicentral distance, we could have easily made an analogous graphic which chooses values for the magnitude and envisions the relationship between the bracketed duration and the epicentral distance for each of these choices.

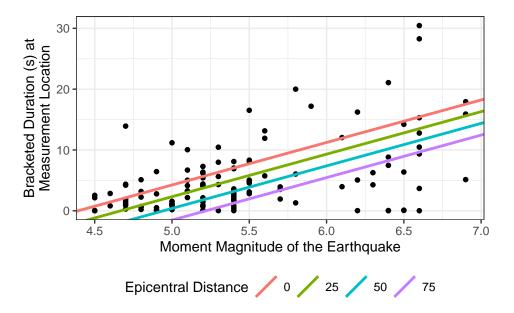


Figure 21.2: Relationship between the bracketed duration and the magnitude of an earthquake after also considering the epicentral distance from an earthquake. Lines estimating this relationship for various values of the epicentral distance are overlayed.

This model has what may appear as an obvious requirement; you cannot use this model to predict the bracketed duration without specifying *both* the magnitude of the earthquake and the epicentral distance of the location. However, it also isolates the effect of the magnitude above and beyond the epicentral distance.

21.2.1 General Model Formulation

Nothing limits us from the inclusion of several predictors. Each predictor is simply added to the model.

General Linear Regression Model

For a quantitative response and one or more predictors, the general form of the linear regression model is

$$\begin{split} (\text{Response})_i &= \beta_0 + \beta_1 (\text{Predictor 1})_i + \beta_2 (\text{Predictor 2})_i + \dots + \beta_p (\text{Predictor } p)_i + \varepsilon_i \\ &= \beta_0 + \sum_{j=1}^p \beta_j (\text{Predictor j})_i + \varepsilon_i \end{split}$$

where β_j for $j=0,1,2,\dots,p$ are the p+1 parameters governing the model for the data generating process.

Note

The predictors in Equation 21.3 can be quantitative variables or indicator variables (Definition 21.1).

The problem, of course, is that the parameters (the β 's in the model) are unknown. However, we can use the method of least squares to estimate each of the parameters simultaneously.

Definition 21.3 (Least Squares Estimates for General Linear Model). The least squares estimates for a general linear model (Equation 21.3) are the values of $\beta_0, \beta_1, \beta_2, \dots, \beta_p$ which minimize the quantity

$$\sum_{i=1}^{n} \left[(\text{Response})_i - \beta_0 - \sum_{j=1}^{p} \beta_j (\text{Predictor } j)_i \right]^2.$$

This minimization procedure is implemented in any statistical software package.

21.2.2 Interpretation of Parameters

The same conditions described in Chapter 18 can be placed on the stochastic portion of the model for the data generating process. Just as with the simple linear model, assuming the model is correctly specified (imposing the mean-0 condition) provides us with an interpretation of each of the parameters.

Consider the model for the data generating process defined in Equation 21.2. If we assume that the error in the bracketed duration has an average of 0 regardless of the magnitude of the corresponding earthquake and the distance of the location to the center of the earthquake, then notice that we are saying the expression

$$\beta_0 + \beta_1$$
 (Magnitude) + β_2 (Epicentral Distance)

defines the average bracketed duration (given the magnitude of the earthquake and epicentral distance of the location).



Big Idea

If we impose the mean-0 condition, the deterministic portion of the general linear regression model specifies the mean response given the value of the predictors.

Therefore, we can interpret the value of β_2 as the change in the average bracketed duration given a 1-kilometer increase in the distance a location is from the center of the earthquake for all locations which experience an earthquake of the same magnitude. This last part is important. In order to interpret one coefficient, we must hold the value of all other predictors fixed.

Interpretation of Parameters

For the general linear model (Equation 21.3), the intercept β_0 is the average response when all predictors are set equal to 0. The j-th coefficient β_i represents the average change in the response associated with a 1-unit increase in the j-th predictor holding the value of all other predictors constant.

This phrase "holding the value of all other predictors constant" has extreme power. It is because of this phrase that we are able to take our first steps toward addressing confounding. For example, consider the model for the data generating process in Equation 21.2. Using the method of least squares, we estimate that for every kilometer further the epicenter of the earthquake is, we can expect the bracketed duration to decrease by 0.08 seconds, on average. Someone might argue as follows: "This is not a controlled experiment; therefore, while there is an association here, it is possible that what is really happening is that earthquakes which were further away tended to also be smaller in magnitude. Therefore, it is not the distance that is driving this observed association but the magnitude of the earthquake." Here, this individual is saying that the earthquake's magnitude is a confounder — related to both the bracketed duration (response) and the variable of interest (distance from the epicenter). If we had fit a marginal model, this would be a valid concern. However, remember our interpretation of β_2 (and our estimate of it) — our fit suggests that for every kilometer further the epicenter of the earthquake is, we can expect the bracketed duration to decrease by 0.08 seconds, on average, holding the magnitude of the earthquake constant. Therefore, since this estimate is comparing two earthquakes of the same magnitude, magnitude cannot be confounding the relationship observed. We have isolated the effect of the epicentral distance from any effect due to the magnitude.

Table 21.2: Summary the linear model fit relating the bracketed duration at locations in Greece following an earthquake with the magnitude of the event and the distance of the location from the center of the earthquake.

Term	Estimate	Standard Error	Lower 95% CI	Upper 95% CI
(Intercept)	-30.715	4.887	-40.395	-21.036
Magnitude	6.991	0.964	5.082	8.900
Epicentral Distance	-0.077	0.021	-0.118	-0.036

Our solution to confounding is to incorporate the relationship between the confounder and the response into our model for the data generating process. Then, any remaining variables cannot be affected by the confounder. Of course this has one major limitation — we cannot account for any variables which are not recorded.

There are entire texts devoted to the topic of addressing confounding. Here, we simply emphasize that regression models allow us to control for the confounders we have observed. The relationships are "adjusted for" these confounders due to the interpretation that a coefficient is the effect "holding all other predictors constant." Regression models allow us to compare similar groups, which are balanced on these confounders after the fact (instead of having addressed confounding through the design of the study).

21.2.3 Inference for the General Formulation

Above, we estimated the change in the average bracketed duration associated with a 1 kilometer increase in the distance from an earthquake while holding the magnitude of the earthquake fixed. However, as we have noted throughout the text, inference requires that we quantify the variability in those estimates.

The same processes that allowed us to model the sampling of the least squares estimates for the simple linear regression model, described in Chapter 18, are applicable for the general linear regression model as well. In particular, under the classical regression model (imposing all four conditions discussed in Chapter 18 on the error term), we can construct an analytical model of the sampling distribution of the least squares estimates (extensions of the results discussed in Appendix A). And, provided we are willing to impose the mean-0 condition and the independence condition, we can construct an empirical model of the sampling distribution using a bootstrapping algorithm. Once we have a model for the sampling distribution of our estimates, we can construct confidence intervals.

Table 21.2 provide a summary of using the data from the Seismic Activity Case Study to fit the model for the data generating process proposed in Equation 21.2; the confidence intervals were computed assuming the data was consistent with all four conditions of the classical regression model.

From the output, we estimate that for every kilometer further the epicenter of the earthquake is, it is reasonable to expect the bracketed duration to decrease between 0.036 and 0.118 seconds, on average, holding the magnitude of the earthquake fixed.

If we are interested in comparing two models, we formulate a standardized statistic and model its null distribution. Suppose for the model we have been considering, we are interested in testing

$$H_0: \beta_1 = \beta_2 = 0$$
 vs. At least one β_j differs from 0,

which corresponds to the following research question:

Does the study provide evidence that the average bracketed duration depends on either the magnitude of the earthquake or the epicentral distance?

Note the complexity of this question and the corresponding hypotheses. We are wanting to drop out two terms from the model for the data generating process. Just as in Chapter 19, the key to testing this hypothesis is to partition out the sources of variability. We can define the error sums of squares for the general linear regression model:

$$SSE = \sum_{i=1}^{n} \left[(\text{Response})_i - \sum_{j=1}^{p} \beta_j (\text{Predictor } j)_i \right]^2.$$

The above hypotheses suggest two alternate models for the data generating process. Under the alternative hypothesis, we have the full unconstrained model presented in Equation 21.2 (call that Model 1). Under the null hypothesis, we have a reduced constrained model for the data generating process (call it Model 0):

(Bracketed Duration)_i =
$$\beta_0 + \varepsilon_i$$
.

If the null hypothesis is true, we would expect the amount of unexplained variability in Model 0 to be the same as the amount of unexplained variability in Model 1 $(SSE_0 \approx SSE_1)$. That is, if the model under the null hypothesis is sufficient for explaining the variability in the bracketed duration, then it should perform as well as the full unconstrained model. If, however, Model 1 explains more of the variability in the bracketed duration (therefore leaving less variability unexplained) than Model 0, then this would indicate the null hypothesis is false. This suggests that a metric for quantifying the signal in the data is given by

$$SSE_0 - SSE_1$$
,

where we use the subscript to denote the model from which the sum of squares was computed. As in Chapter 19, working with variance terms is easier than working with sums of squares

analytically. And, we need to consider the size of the signal relative to the amount of background noise. This leads to a form of our standardized statistic which generalizes our result from Definition 19.7.

Definition 21.4 (Standardized Statistic for General Linear Model). Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 21.3):

$$(\text{Response})_i = \beta_0 + \sum_{j=1}^p \beta_j (\text{Predictor } j)_i + \varepsilon_i.$$

Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0. A standardized statistic, sometimes called the "nested F statistic," for testing the hypotheses is given by

$$T^* = \frac{\left(SSE_0 - SSE_1\right)/(p+1-r)}{SSE_1/(n-p-1)},$$

where p + 1 is the number of parameters in the full unconstrained model (including the intercept) and r is the number of parameters in the reduced model. Defining

$$MSA = \frac{SSE_0 - SSE_1}{p+1-r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model. Just as before, the MSE represents the residual variance — the variance in the response for a particular set of the predictors.

The standardized statistics in Chapter 12 and in Chapter 19 are special cases of the one presented here. The numerator captures the signal by examining the difference between what we expect the error sum of squares to be under the null hypothesis and what we actually observe; the denominator captures the background noise (relative to the estimated mean response from the full model). Larger values of this standardized statistic indicate more evidence in the sample against the null hypothesis.

Table 21.3: Analysis of the sources of variability in the bracketed duration as a function of the magnitude of the earthquake and the epicentral distance.

Term	DF	Sum of Squares	Mean Square	Standardized Statistic	P-Value
Additional Terms	2	1297.653	648.827	28.17	< 0.001
Error	116	2671.787	23.033		

Table 21.3 provides a summary of partitioning the variability to test the hypotheses described above using the data from the Seismic Activity Case Study; the p-value was computed assuming the data was consistent with all four conditions of the classical regression model.

Based on the results, the sample provides strong evidence (p < 0.001) that the average bracketed duration is associated with either (or both) the magnitude of the earthquake and the distance the location is from the center of the earthquake. Note that the p-value does not tell us the form of the relationship, or whether it is only one of the predictors that is important or both. Examining the confidence intervals provides that information. The confidence intervals also allow a discipline expert to determine if the signal we are discerning is meaningful in practice.

21.3 Modifying an Effect

There is one type of question we have not yet addressed — assessing the interplay between two variables on the response.

Consider the following question:

Is the relationship between the average bracketed distance and the magnitude different depending on whether the soil is rocky where the measurement is taken?

This question explains the bracketed duration in terms of both the magnitude as well as whether the soil is rocky. A first pass at such a model might be

$$(Bracketed Duration)_i = \beta_0 + \beta_1 (Magnitude)_i + \beta_2 (Rocky Soil)_i + \varepsilon_i$$

where we use an indicator variable to capture whether the soil is rocky just as before. Just as we did before, we can graphically represent this model (see Figure 21.3), and we find that it is captured by two parallel lines.

The lines are parallel because the coefficient associated with Magnitude (β_1) is the same regardless of the type of soil. That is, regardless of whether the soil is rocky, the change in the bracketed duration, on average, for each 1-unit increase in the magnitude of an earthquake is the same. Our question of interest is essentially, is there evidence that this is not the case?




Figure 21.3: Relationship between the bracketed duration and magnitude for locations with rocky soil and those with other soil types.

So, the above model actually represents the model for the data generating process under the null hypothesis of our current question. Under the null hypothesis, the effect of the magnitude on the bracketed duration (which captures the relationship between these two variables) is the same regardless of soil condition. What we posited above is Model 0 which has an embedded constraint. The question is, how do we form the model for the data generating process which allows the slope to look differently depending on soil type? This new model would serve as our full unconstrained model, Model 1.

Consider adding an additional term to our model above, yielding the following model:

$$\begin{split} (\text{Bracketed Duration})_i &= \beta_0 + \beta_1 (\text{Magnitude})_i + \beta_2 (\text{Rocky Soil})_i \\ &+ \beta_3 (\text{Magnitude})_i (\text{Rocky Soil})_i + \varepsilon_i \end{split}$$

This additional term is formed by taking the product of the indicator variable with the variable magnitude; such a product is known as an **interaction term**.

Definition 21.5 (Interaction Term). A variable resulting from taking the product of two predictors in a regression model. The product allows the effect of one predictor to depend on another predictor, essentially modifying the effect.

In order to really see the impact of the interaction term, observe that if we impose the mean-0 condition, our model for the data generating process says the average bracketed duration for locations with rocky soil is given by

$$(\beta_0 + \beta_2) + (\beta_1 + \beta_3)$$
 (Magnitude)

which comes from plugging in 1 for the indicator variable (representing rocky soils) and then combining terms. The model for the data generating process says the average bracketed duration for locations without rocky soil is given by

$$\beta_0 + \beta_1$$
 (Magnitude)

since we would plug in 0 for the indicator variable (representing soils that are not rocky). Therefore, our revised model for the data generating process that includes the interaction term allows both the intercept and the effect of the magnitude on the bracketed duration to differ for rocky soils and non-rocky soils. That is, the effect of magnitude is allowed to differ across the soil types.



Big Idea

It is common to believe that the interaction term measures the effect between the two variables in the product. However, this is incorrect. The interaction term allows the effect of one predictor to differ across the levels of another predictor.

Visually, this revised model allows two completely different associations — depending on the soil type. This is shown in Figure 21.4. Notice that the lines are no longer parallel. The question of course is which of the two models is more appropriate. Is there actually evidence that the more complex model, which allows the relationship to differ for locations with different soil types, is required? Or, is the more simplistic model, which says the relationship is the same across all locations of different soil types, sufficient?

We can capture our question of interest in the following hypotheses:

$$H_0: \beta_3 = 0$$

 $H_1: \beta_3 \neq 0$

Notice that if the null hypothesis were true, then the slope would be the same regardless of whether the soil was rocky because we resort to Model 0 described above which only allows the intercept to vary across soil types. However, if β_3 is nonzero in Model 1, then the slope will differ for the rocky soil. So, under the null hypothesis, the lines are parallel; under the alternative hypothesis, the lines are not parallel.

Conceptually, comparing these two models is just like any other hypothesis test. By partitioning the variability as described in the previous section, we are able to compute a signal-to-noise




Figure 21.4: Relationship between bracketed duration and the magnitude of an earthquake after also considering the soil conditions of the measurement location. The relationship between the bracketed duration and the magnitude is allowed to differ within each type of soil condition.

Table 21.4: Summary of the model fit explaining the bracketed duration as a function of both magnitude and soil condition at the measurement location. The effect of the magnitude was allowed to differ across soil conditions.

Term	Estimate	Standard Error	95% Lower CI	95% Upper CI
(Intercept)	-23.966	4.531	-32.941	-14.991
Magnitude	5.466	0.834	3.814	7.118
Soil Conditions (Rocky)	11.925	9.181	-6.261	30.112
Interaction: Magnitude & Rocky Soil	-2.614	1.626	-5.835	0.607

Table 21.5: Table partitioning the sources of variability corresponding to testing whether the effect of the magnitude of an earthquake on the resulting bracketed duration differs across soil conditions.

Source	DF	SS	MS	Standardized Statistic	P-Value
Additional Terms	1	62.663	62.663	2.584	0.111
Error	115	2788.908	24.251		

ratio. We can essentially determine how much variability in the response is explained is by the inclusion of the interaction term. That is, we quantify how the unexplained variability decreases when we include the interaction term in the model for the data generating process.

Specifically, Table 21.4 gives the estimates associated with each parameter, and Table 21.5 presents the corresponding partitioning of the sources of variability leading to the computation of the p-value. The confidence intervals and p-value are computed assuming the data is consistent with the conditions of the classical regression model.

From our analysis, our data is consistent with (p = 0.111) the effect of the magnitude being the same across the various soil conditions. The sample provides no evidence that the association between the average bracketed duration and the magnitude of an earthquake depends on the soil conditions.

22 Putting it All Together

For the Seismic Activity Case Study, consider the following question:

Overall, does the average bracketed duration at a location depend on the distance a location is from the center of an earthquake?

While we have touched on this question in previous chapters, we investigate this question more fully in this chapter.

22.1 Graphical Summary

Before developing a statistical model to address our question, we summarize the data graphically. The question suggests the bracketed duration is the response variable and the epicentral distance is the predictor. Figure 22.1 illustrates the relationship between the bracketed duration and the epicentral distance. We note that the axis for the epicentral distance takes logarithmic steps to better illustrate the relationship. That is, moving from 1 to 10 kilometers has roughly the same effect as moving from 10 to 100 kilometers from the earthquake.

22.2 Development of Statistical Model

In order to address our primary question of interest, we must develop a statistical model which explains the data generating process and embeds our question of interest in terms of the parameters of the model. Based on the question of interest, and our observations in the graphical exploration above, our model explaining the generation of the bracketed duration should depend on the epicentral distance. Further, this relationship should be on a logarithmic scale to account for the "stretched" scale in Figure 22.1.

At first, it may seem that the logarithmic scale prevents our use of the "linear" model we discussed in this unit, but the model is flexible enough to accommodate this development. We propose the following model for the data generating process:

$$(Bracketed Duration)_i = \beta_0 + \beta_1 \log_{10}(Epicentral Distance)_i + \varepsilon_i.$$
 (22.1)




Figure 22.1: Relationship between the bracketed duration and the distance from the epicenter of an earthquake for locations measuring seismic activity in Greece.

Notice that instead of just the epicentral distance, our predictor is the base-10 logarithm of the epicentral distance. This transformed variable enters the model just as any other predictor. In addition to modeling the deterministic portion of the data generating process, we must also place conditions on the stochastic portion in order to make inference. We consider the conditions of the classical regression model:

- The error in the bracketed duration is 0, on average for all epicentral distances; that is, the model for the deterministic portion is correctly specified.
- The error in the bracketed duration for one location is independent of the error in the bracketed duration for any other location.
- The variability of the error in the bracketed duration is the same for all locations regardless of the epicentral distance.
- The error in the bracketed duration follows a Normal distribution.

Before imposing these conditions, however, we should assess them.

22.3 Assessment of Conditions

Before making inference regarding our question of interest, we should determine if our data is consistent with the conditions on the error term we have specified. Figure 22.2 is a prob-

ability plot of the residuals used to assess whether the data is consistent with the Normality condition.

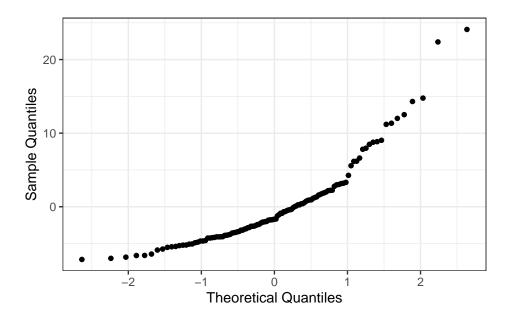


Figure 22.2: Normal probability plot of the residuals corresponding to a model for the bracketed duration of seismic events in Greece as a function of the epicentral distance.

The plot reveals some departure from the linear relationship we would expect if the errors followed a Normal distribution. It seems unreasonable to impose the Normality condition.

Figure 22.3 is a plot of the residuals for the observations in the order in which they were collected. Since the data was collected over time, this plot could reveal potential patterns among the residuals which suggest a departure from independence among the errors. As there are no trends in either the location or spread of the residuals, the data is consistent with the independence conditions.

Figure 22.4 is a plot of the residuals against the predicted values from the deterministic portion of the model. There are no obvious trends in the location of the residuals; that is, the residuals balance out around 0 for all predicted values. Therefore, we are willing to the data is consistent with the condition that the mean of the errors is 0 for each combination of the predictors. We note that while the points balance out around 0 (as indicated by a smoother that remains near 0), they tend to have a larger range on the positive side compared to the negative side; this relates back the observation we made earlier that the residuals do not behave like a sample from a Normal distribution.

We also note that the spread of the residuals is larger on the right side of the graphic compared to the left side. As the predicted values increase, the spread of the residuals also increases.




Figure 22.3: Time-series plot of the residuals corresponding to a model for the bracketed duration of seismic events in Greece as a function of the epicentral distance.

This suggests that for larger bracketed durations, the model is not as precise. Therefore, we are not willing to impose the constant variance condition.

Examining the residuals, we determined that the data is consistent with the following two conditions:

- The error in the bracketed duration for one location is independent of the error in the bracketed duration for any other location.
- The error in the bracketed duration is 0, on average, for all epicentral distances; that is, the deterministic portion of the model for the data generating process is correctly specified.

Since we are only willing to assume these two conditions, we will use an empirical model (bootstrap procedure) for the sampling distribution of the least squares estimates.

22.4 Summary of Model Fit

The parameters in our model are estimated via the method of least squares. The variability in these estimates is quantified using an empirical model of the sampling distribution based on 5000 bootstrap replications. Table 22.1 summarizes the estimates for the parameters in Equation 22.1.

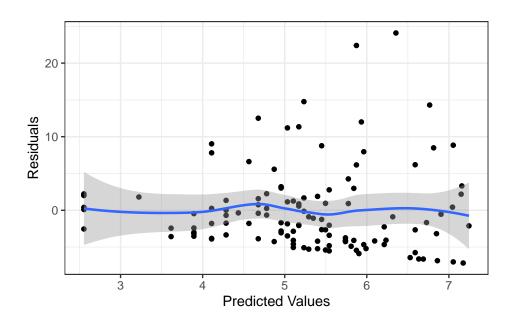


Figure 22.4: Plot of the residuals against the predicted values corresponding to a model for the bracketed duration of seismic events in Greece as a function of the epicentral distance.

Table 22.1: Summary of the model characterizing the bracketed duration of seismic events in Greece as a function of the epicentral distance.

Term	Estimate	Standard Error	95% Lower CI	95% Upper CI
(Intercept)	2.553	1.031	0.464	4.561
$\log 10 (Epicentral Distance)$	2.226	0.985	0.347	4.248

The results suggests that for each 10-fold increase in the number of kilometers a location is from the epicenter of an earthquake (1-unit increase in the base-10 logarithm), the bracketed duration increases between 0.347 and 4.248 seconds, on average. This result is counter-intuitive; it suggests that the further a location is from an earthquake, the longer it is subjected to major ground motion. First, remember that this is data from an observational study; so, the result is not suggesting a causal relationship. Indeed, confounding is playing a role here. Figure 22.5 examines the relationship between the bracketed duration and the epicentral distance in addition to the magnitude of the earthquake and the soil conditions of the location. While it is somewhat difficult to see, note that earthquakes with the largest magnitudes also tended to be recorded further away. While we would need to confirm with a discipline expert, it suggests that the major fault line is approximately 100 kilometers away from the Greek observation stations. As a result, the recording locations in the study experienced the largest amount of major motion when these large quakes occurred. A more robust analysis using the general form of the linear regression model of the previous chapter would reveal that after accounting for these additional terms, the data is consistent with decreases in the bracketed durations, on average, for locations further from the center of an earthquake.




Figure 22.5: Relationship between the bracketed duration and the distance from the epicenter of an earthquake for locations measuring seismic activity in Greece. The relationship is presented for various soil types.

It is important to always place our analyses in the context of the problem and consult with discipline experts. This can sometimes reveal shortcomings in the proposed model for the data generating process. Remember, statistical analyses are not a crystal ball; they reveal potential patterns, but only those that we incorporate into our model for the data generating process.

Part IV

Unit IV: Comparing the Average Response Across Groups

Throughout the text, we have developed the language and logic of statistical inference within the context of a model for the data generating process. In this unit, we extend these ideas to comparing the mean response across groups.

23 Case Study: Organic Foods and Superior Morals

"You are what you eat" is a common phrase, dating back to at least the 1820's¹, used to suggest that if you want to be fit, you must eat healthy foods. However, does the phrase extend to our personality as well as our physique? Recent research has suggested that specific tastes (sweet vs. disgusting, for example) can influence moral processing. That is, certain foods may lead us to be nicer to those around us or lead us to be more judgmental. Organic foods are often marketed using phrases like "pure" or "honest" (Jessica Alba's Honest Company, for example); is there some relationship between the consumption of organic foods and moral behavior?

Dr. Eskine of the Department of Psychological Sciences at Loyola University sought to answer this question (Eskine 2013). He conducted a study to investigate whether exposure to certain types of food had an effect on a person's moral processing. Specifically, he randomized 62 Loyola University undergraduates to one of three food types: organic, comfort, and control. Each participant received a packet containing pictures of four food items from the assigned category:

- Organic Foods: apple, spinach, tomato, carrot
- Comfort Foods: ice cream, cookie, chocolate, brownie
- Control Foods: oatmeal, rice, mustard, beans

The control foods chosen are pre-packaged and are generally considered staple items; organic foods chosen are associated with a healthy diet; and, comfort foods were sweets. After viewing the images for a set period of time, each participant received a packet containing six counterbalanced moral transgressions. An example of such a transgression is produced below:

Bob was at a family gathering when he met Ellen, a second cousin of his that he had seen once or twice before. Bob found Ellen very attractive and he asked her out on a date. Ellen accepted and they began to have a romantic and sexual relationship. They often go on weekend trips to romantic hotels in the mountains.

Participants were then asked to rate the morality of the scenario on a 7-point scale (1 = "not at all morally wrong" to 7 = "very morally wrong"). The average of the morality scores across the six scenarios was used as an overall measure of their moral expectations. A higher

¹https://www.phrases.org.uk/meanings/you-are-what-you-eat.html

Table 23.1: Subset of data from study characterizing moral behavior following exposure to various food categories.

Participant	Food Condition	Response (Avg of Moral Questions)
1	comfort	6.000
2	$\operatorname{comfort}$	3.500
3	$\operatorname{comfort}$	6.167
4	control	5.167
10	control	7.000
12	control	6.833
18	organic	5.500
20	organic	5.500
21	organic	6.333

value indicates high moral expectations (very strict), and a lower value indicates lower moral expectations (very lenient).

Dr. Eskine's analysis revealed the study provided strong evidence (p=0.001) that participants' moral judgments differed, on average, across the various food exposure groups. In particular, those exposed to organic foods had higher moral expectations (an average mean moral judgment of 5.58) compared to those experiencing comfort foods (average mean moral judgment of 4.89) or control foods (average mean moral judgment of 5.08). He therefore concluded that exposure to organic food led to higher moral expectations among the participants.

Understandably, Dr. Eskine's work caught the interest of various media outlets and researchers. Two researchers within the Department of Psychology at Dominican University in Illinois sought to replicate Dr. Eskine's work (Moery and Calin-Jageman 2016). There were several components to their research, but the first phase included a replication of Dr. Eskine's initial study with minor variants. They enrolled 123 college students into their study. The participants were presented with the same food images as in Eskine's study with the exception that celery was used instead of an apple for organic food. The same moral dilemmas were given to participants. As in the original study, the average score from the six moral dilemmas was the primary response for this study. A subset of the collected data, showing three participants from each treatment group (type of food shown), is presented in Table 23.1. The full dataset has been made available by the researchers².

²There were multiple phases to their research. The direct replication of Dr. Eskine's work was Study 1, which is the dataset being considered in this text; it is available at https://osf.io/atkn7/wiki/home/.

24 Framing the Question

"Does exposure to various food types lead to different moral expectations?" The primary question from the Organic Food Case Study (Chapter 23) is about the relationship between two variables: the response (moral expectations; see Definition 3.2) and the **factor** of interest (food type).

Definition 24.1 (Factor). Also referred to as the "treatment" in some settings, a factor is a categorical predictor. The categories represented by this categorical variable are called "levels."

As we saw in Unit III, the majority of interesting research questions involve identifying or quantifying the relationship between two variables. However, the approach to these questions relies on the same language and logic introduced in Unit I. We begin with asking good questions, which involves defining the population of interest and characterizing the variable(s) at the population level through well-defined parameters.

The question posed from the Organic Food Case Study, as stated above, is ill-posed. Almost certainly, there are individuals for which exposure to organic foods may result in higher moral expectations compared to exposure to comfort foods. However, there are almost certainly individuals for which the effect is reversed — higher moral expectations are expected following exposure to comfort foods compared with organic foods. That is, we expect there to be variability in the effect of food types on the resulting moral expectations. The question needs to be refined.

While the study was conducted using college students, the original question seems quite broad (we discuss this discrepancy in more detail in the next chapter). Notice that the original question is not predicated on *consuming* various foods but simply *exposure* to various foods. The question itself is not limited to only those individuals which purchase a specific type of food but concerns all individuals. More, we really see that there are three groups of interest — those which are exposed to organic foods, those exposed to comfort foods, and those exposed to the control foods. We can think of three distinct populations:

- 1. All individuals exposed to organic foods.
- 2. All individuals exposed to comfort foods.
- 3. All individuals exposed to control foods.

We now work to characterize the response within each of these three populations. Since the response of interest is a numeric variable (taking values between 1 and 7 with higher values indicating higher moral expectations), summarizing the variable using the mean is reasonable. That is, we might ask "does exposure to various food types lead to different moral expectations, on average?" Our question is looking for some type of difference in this mean response across the groups; our working hypothesis is then that the groups are all equivalent, on average. This could be framed in the following hypotheses:

 H_0 : the average moral expectations are the same following exposure to each of the three types of food.

 H_1 : the average moral expectations following exposure to food differ for at least one of the three types.

We can represent these hypotheses mathematically as

```
H_0: \mu_{\text{comfort}} = \mu_{\text{control}} = \mu_{\text{organic}}

H_1: \text{At least one } \mu \text{ differs from the others,}
```

where μ_{comfort} is the mean moral expectation for individuals exposed to comfort foods, μ_{control} is the mean moral expectation for individuals exposed to control foods, and μ_{organic} is the mean moral expectation for individuals exposed to organic foods. The question is now well-posed — it is centered on the population and captured through parameters.

For this particular setting, there is an alternative way of thinking about the population. You might argue that there are not three distinct populations; instead, there is only a single population (all individuals) and three different exposures (organic, comfort, and control foods). This is a reasonable way of characterizing the population. The hypotheses remain the same:

```
\begin{split} H_0: \mu_{\text{comfort}} &= \mu_{\text{control}} = \mu_{\text{organic}} \\ H_1: \text{At least one } \mu \text{ differs from the others.} \end{split}
```

The difference is in our interpretation of the parameters. We would describe μ_{comfort} as the mean moral expectation when an individual is exposed to comfort foods, μ_{control} as the mean moral expectation when an individual is exposed to control foods, and μ_{organic} as the mean more expectation when an individual is exposed to organic foods. The distinction, while subtle, is to place emphasis on switching an individual from one group to another instead of the groups being completely distinct. In fact, this latter way of thinking is more in line with how the study was conducted. Individuals were allocated to one of the exposure groups, suggesting that exposure is something that could be changed for an individual.

From an analysis perspective, there is little difference between these two ways of describing the population. The difference is primarily in our interpretation. In many cases, we can envision the population either way; however, there are a few instances where that is not possible. Suppose we were comparing the average number of offspring of mice compared to rats (a lovely thought, we know). It does not make sense to think about changing a mouse into a rat;

here, it only makes sense to think about two distinct populations being compared on some metric. How we describe the population is often related to the question we are asking.

Note

How we describe the population is often connected to the study design we implement. In a controlled experiment, we envision a single population under various conditions. For an observational study, we generally consider distinct populations.

24.1 General Setting

This unit is concerned with comparing the mean response of a numeric variable across k groups. Let $\mu_1, \mu_2, \dots, \mu_k$ represent the mean response for each of the k groups. Then, we are primarily interested in the following hypotheses:

$$H_0: \mu_1 = \mu_2 = \dots = \mu_k$$

 $H_1:$ At least one μ differs

When there are only two groups (k=2), then this can be written as

$$H_0: \mu_1 = \mu_2$$

 $H_1: \mu_1 \neq \mu_2$

Warning

When there are two groups, it makes sense for the alternative hypothesis to include "not equal." While tempting to do something similar when there are more than two groups, it is not possible. The opposite of "all groups are equal, on average" is not "all groups differ, on average." Nor is the opposite "exactly one group differs, on average." The opposite of "all groups equal are equal, on average" is "at least one average differs" which is what we are capturing with the above hypotheses. Note that this is not saying that one group must differ from all other groups, on average. It is saying that at least two of the groups have a different average response. Keep it simple and do not try to get fancy with the notation.

Here we are writing things in the mathematical notation, but let's not forget that every hypothesis has a context. Throughout this unit, we are looking for some signal in the location of the response across the groups. Our working assumption then states that the groups are all similar, on average. This may not be the only comparison of interest to make in practice. For example, it may not be the location that is of interest but the spread of a process. In some applications, managers would prefer to choose the process that is the most precise. These questions are beyond the scope of this unit, but the concepts are similar to what we discuss here.

25 Study Design

Chapter 4 discussed the impact that the design of a study has on interpreting the results. Recall that the goal of any statistical analysis is to use the sample to say something about the underlying population. Observational studies are subject to confounding. In order to use the available data in order to make causal statements that apply within the population, we need to address this confounding. There are two ways of doing this:

- 1. Conduct a controlled experiment. While we do not limit our discussion to controlled experiments in this unit, our discussion will emphasize the elements of a well designed study.
- 2. Use observational data and account for confounders. This can be addressed through regression modeling as suggested in Chapter 21.

As discussed in Chapter 4, controlled experiments (Definition 4.5) balance the groups being compared relative to any potential confounders. As a result, such studies permit causal conclusions to be drawn. While controlled experiments are a useful tool, there are many aspects to consider when designing a study.

25.1 Aspects of a Well-Designed Study

Generally speaking, there are three components to a well-designed study: replication, randomization, and reduction of extraneous noise.



Warning

A study is not poor just because it lacks one of these elements. That is, a study can provide meaningful insights even if it did not make use of each of these elements; every study is unique and should be designed to address the research objective. These elements are simply helpful in creating study designs.

As we have stated repeatedly, variability is inherit in any process. We know there is variability in the population; not every subject will respond exactly the same to each treatment. Therefore, our questions do not seek to answer statements about individuals but about general trends in the population. For example, as discussed in Chapter 24, we may be interested in comparing the mean response across various groups. In order to establish these general trends, we must allow that subject-to-subject variability be present within the study itself. This is accomplished through **replication**, obtaining data on multiple subjects from each group. Each subject's response would be expected to be similar, with variability within the group due to the inherent variability in the data generating process.

Definition 25.1 (Replication). Replication results from taking measurements on different units (or subjects), for which you expect the results to be similar. That is, any variability across the units is due to natural variability within the population.

Warning

The term "replication" is also used in the context of discussing whether the results of a study are replicable. While our use of the term is about replicating a measurement process within a study, this does not downplay the importance of replicating an entire study.

When we talk about gathering "more data," we typically mean obtaining a larger number of replicates. Ideally, replicates will be obtained through random selection from the underlying population to ensure they are representative. The subjects are then randomly allocated to a particular level of the factor under study (randomly allocated to a group) when performing a controlled experiment. This random allocation breaks the link between the factor and any potential confounders, allowing for causal interpretations. However, random allocation preserves any link between the factor and the response, if a link exists. These are the two aspects of randomization.

Definition 25.2 (Randomization). Randomization can refer to random selection or random allocation. Random selection refers to the use of a random mechanism (e.g., a simple random sample, Definition 4.2, or a stratified random sample, Definition 4.3) to select units from the population. Random selection minimizes bias.

Random allocation refers to the use of a random mechanism when assigning units to a specific treatment group in a controlled experiment (Definition 4.5). Random allocation eliminates confounding and permits causal interpretations.

Note

While those new to study design can typically describe random selection and random allocation, they often confuse their purpose. Random selection is to ensure the sample is representative. Random allocation balances the groups with respect to confounders.

It is tempting to manually adjust the treatment groups to achieve what the researcher views as balance between the groups. This temptation should be avoided as balancing one feature of the subjects may lead to an imbalance in other features. Remember, random allocation leads to balance. Of course, random allocation does not guarantee any particular sample is perfectly balanced; however, any differences are due to chance alone. The likelihood of such differences can then be studied and quantified (think about the construction of a model for the null distribution and the definition of a p-value). As the sample size increases, these differences due to chance are minimized.

Even with random allocation providing balance between the groups, there will still be variability within each group. The more variability present, the more difficult it is to detect a signal — to discern a difference in the mean response across groups. The study will have more **power** to detect the signal if the groups are similar. This leads to the third component of a well-designed study — the **reduction of noise**.

Definition 25.3 (Power). In statistics, power refers to the probability that a study will discern a signal when one really exists in the data generating process. More technically, it is the probability a study will provide evidence against the null hypothesis when the null hypothesis is false.

Statistical power is like asking the question "what is the probability a guilty defendant will be found guilty by the jury?"

Definition 25.4 (Reduction of Noise). Reducing extraneous sources of variability can be accomplished by fixing extraneous variables or blocking (Definition 25.5). These actions reduce the number of differences between the units under study.

i Tension between Lab Settings and Reality

Scientists and engineers are trained to control unwanted sources of variability (or sources of error in the data generating process). This creates a tension between what is observed in the study (under "lab" settings) and what is observed in practice (in "real-world" settings). This tension always exists, and the proper balance depends on the goals of the researchers.

Fixing the value of extraneous variables can reduce variability in a study. For example, in the Organic Foods Case Study, the study was only conducted among college students in a psychology class. This choice indirectly impacts the value of an extraneous variable. Depending on the university, for example, this choice could imply the age of participants is most likely between 18 and 22. That reduces the "noise" in how participants respond due to generational differences, which is an extraneous (not directly being considered in the study) variable. However, note that this decision also potentially limits the scope of the study. It may no longer be appropriate to apply these results to the population at large; that is, elderly individuals may respond differently than what we observe in the study.

An additional tool for reducing noise is **blocking**, in which observations which are dependent on one another because of a shared characteristic are grouped together.

Definition 25.5 (Blocking). Blocking is a way of minimizing the variability contributed by an inherent characteristic that results in dependent observations. In some cases, the blocks are the unit of observation which is sampled from a larger population, and multiple observations are taken on each unit. In other cases, the blocks are formed by grouping the units of observations according to an inherent characteristic; in these cases that shared characteristic can be thought of having a value that was sampled from a larger population.

In both cases, the observed blocks can be thought of as a random sample; within each block, we have multiple observations, and the observations from the same block are more similar than observations from different blocks.

Example 25.1 (Overseeding Golf Greens). Golf is a major pastime, especially in southern states. Each winter, the putting greens need to be overseeded with grasses that will thrive in cooler weather. This overseeding can affect how the ball rolls along the green. Dudeck and Peeacock (1981) reports on an experiment that involved comparing the ball roll for greens seeded with one of five varieties of rye grass. Ball roll was measured by the mean distance (in meters) that five balls traveled on the green. In order to induce a constant initial velocity, each ball was rolled down an inclined plane.

Because the distance a ball rolls is influenced by the slope of the green, 20 greens were placed into four groups in such a way that the five greens in the same group had a similar slope. Then, within each of these four groups, each of the five greens was randomly assigned to be overseeded with one of the five types of Rye grass. The average ball roll was recorded for each of the 20 greens.

The data from Example 25.1 are shown in Table 25.1.

It would have been easy to simply assign 4 greens to each of the Rye grass varieties; the random allocation would have balanced any confounders across the five varieties. However, an additional layer was added to the design in order to control some of that additional variability. In particular, greens with similar slopes were grouped together; then, the random allocation to Rye grass varieties happened within the grouped greens. The blocks in this study are the "slope groups." Each block represents greens with a different slope. Certainly, there are more than 4 potential slopes that a green might have; yet, we observed 4 such groups in our study. We can think of these 4 observed slopes as a sample of all slopes that might exist on a putting green.

Within each block, we have five units of observations; these five units were randomized to the five treatment groups (the five Rye grass varieties). Notice the random allocation strategy ensures that each variety appears exactly once within each slope grouping. This study design will allow us to compare the impact of the Rye grass variety while minimizing the extraneous

Table 25.1: Data from the Overseeding Golf Greens example.

Rye Grass Variety	Slope of Green Grouping	Mean Distance Traveled (m)
A	1	2.764
В	1	2.568
C	1	2.506
D	1	2.612
E	1	2.238
A	2	3.043
В	2	2.977
С	2	2.533
D	2	2.675
E	2	2.616
A	3	2.600
В	3	2.183
C	3	2.334
D	3	2.164
E	3	2.127
A	4	3.049
В	4	3.028
С	4	2.895
D	4	2.724
E	4	2.697

variability due to the slope of the green, which is a nuisance characteristic. To see how we capitalize on blocking in the analysis, we refer you to Unit IV of the text.

Note

Blocking is often a way of gaining additional power when limited resources require your study to have a small sample size.

An extreme case of blocking occurs when you repeatedly measure the response on the same subject under different treatment conditions. For example, a pre-test/post-test study is an example of a study which incorporates blocking. In this case, the blocks are the individual subjects, the unit of observation. The response is then observed on the subject both prior to the intervention (the "test") and following the intervention. The rationale here is to use every subject as his or her own "control." This reduces extraneous noise because the two treatment groups (the pre-test group and the post-test group) are identical.

Page Idea

A block is a secondary grouping variable present during the data collection that records a nuisance characteristic. While it reduces extraneous noise in the sample, the block must be accounted for appropriately during the analysis of the data (as described in Unit IV).

25.2 Critiquing the Organic Food Case Study Design

In the previous section, we described various aspects of a well-designed study. We now consider how the Organic Food Case Study incorporated these aspects.

We notice that random allocation was utilized. Each of the 123 participants was randomly assigned to one of three treatment groups (type of food to which the participant was exposed). The random allocation allows us to make causal conclusions from the data as confounders will be balanced across treatment groups. For example, subjects who adhere to a strict diet for religious purposes would naturally tend toward organic foods and higher moral expectations. However, for each subject like this who was assigned to the organic foods group, there is someone like this (on average) who was assigned to the comfort foods group.

We also note that there is replication. Instead of assigning only one subject to each of the three treatment groups, we have several subjects within each group. This allows us to evaluate the degree to which the results vary within a particular treatment group.

As discussed above, the researchers chose only to collect data among college students taking a psychology class. This most likely limits the age of the participants significantly. While this does serve to reduce noise due to generational differences, it also limits the scope of the study. This is a potential limitation of the study.

The study does not make use of blocking. There are a couple of potential reasons for this; first, with such a large sample size, the researchers may not have thought it necessary. Second, it could be that there was a restriction on time. For example, researchers may have considered having students be exposed to each of the three types of food and answering different scenarios after each. However, this would take a longer amount of time to collect data. Third, it could be that researchers were not concerned about any identifiable characteristics that would generate additional variability. Regardless, the study is not worse off because it did not use blocking; it is still a very reliable design.

While it is clear that random allocation was utilized in the design, random selection was not. Students participating in the study are those from a particular lecture hall. As a result, these students were not randomly sampled from all college students (or even from the university student body). As a result, we must really consider whether the conclusions drawn from this study would apply to all college students within the United States. Having additional information on their demographics may help determine this, but in general, this is not something that can be definitively answered. It is an assumption we are either willing to make or not. More, notice that the original question was not focused on college students; however, the sample consists only of college students. This can impact the broader generalization of our results. It is quite possible that we observe an effect in college students that is not present in the larger population. We should always be careful to ensure that the sample we are using adequately represents the population of interest. When the sample does not represent the population, the study may not be useless, but we should critically consider what population the sample represents.

25.3 Collecting Observational Data

An inability to conduct a controlled experiment does not mean we neglect study design. Random sampling is still helpful in ensuring that the data is representative of the population. And, even if random sampling is not feasible, we should still aim to minimize bias and have a sample that is representative of our population. Similarly, ensuring there are a sufficient number of replications to capture the variability within the data is an important aspect of conducting an observational study. When collecting observational data, one of the most important steps is constructing a list of potential confounders and then collecting data on these variables as well. This will allow us to account for these confounders in our analysis (see Chapter 21); we cannot model what we do not collect. Finally, observational studies may still permit the blocking of subjects and accounting for this additional variability in our analysis (Unit IV).

26 Presenting the Data

Chapter 19 introduced the importance of partitioning sources of variability in the response. This theme impacts both the manner in which we summarize data and the analysis we conduct. We keep this idea in mind as we also construct graphical and numerical summaries which address our research objective.

We have already argued that variability makes addressing questions difficult. If every subject in the Organic Food Case Study had the same moral expectation when exposed to the same food type, there would be no need for statistics. We would simply evaluate one subject and determine which treatment to give. Statistics exists because of the ambiguity created by variability in the responses. As a result of this variability, our statistical graphics (and later our model for the data generating process) must distinguish the various sources of variability. That is, with any analysis, we try to answer the question "why aren't all the values the same? What are the reasons for the difference we are observing?"

From the Organic Food Case Study, consider the primary question of interest:

Is there evidence of a relationship between the type of food a person is exposed to and their moral expectations, on average, following exposure?

We are really asking "does the food exposure help explain the differences in the moral expectations of individuals?" We know that there are differences in moral expectations between individuals. But, are these differences solely due to natural variability (some people are just inherently, possibly due to how they were raised, more or less strict in terms of their moral beliefs); or, is there some systematic component that explains at least a portion of the differences between individuals?

A good graphic must then tease out how much of the difference in the moral expectations is from subject-to-subject variability and how much is due to the food exposure. When our response and predictor were both quantitative, a scatterplot was appropriate (Chapter 16). When our response is categorical, we may turn to other options, but the goal is the same—address the research objective and tease out the various sources of variability in the response. Consider a common graphic which is **not** useful in this situation (Figure 26.1).

To determine an appropriate graphic, we need to remember that we want to partition the variability. So, we must not only compare the differences between the groups but also allow the viewer to get a sense of the variability within the group. A common way of doing this




Figure 26.1: Illustration of a poor graphic using the Organic Food Case Study; the graphic does not give us a sense of variability. As a result, it is not clear how different these means really are.




Figure 26.2: Comparison of the moral expectations for college students exposed to different types of food.

within engineering and scientific disciplines is to construct side-by-side boxplots, as illustrated in Figure 26.2.

From the graphic, we see that the moral expectation scores seem to have nearly the same pattern in each of the exposure groups. More, the center of each of the groups is roughly the same. That is, there does not appear to be any evidence that the type of food to which a subject is exposed is associated with moral expectations, on average.

Side-by-side boxplots can be helpful in comparing large samples as they summarize the location and spread of the data. When the sample is smaller, it can be helpful to overlay the raw data on the graphic in addition to the summary provided by the boxplot. We might also consider adding additional information, like the mean within each group. An alternative to boxplots is to use violin plots which emphasize the shape of the distribution instead of summarizing it like boxplots; again, this is most helpful when the sample size is large. Yet another option is to construct density plots which are overlayed on one another; again, this is most appropriate for larger samples. This works when there are only a small number of groups; if the number of groups is large, then placing the distributions side-by-side is much more effective. A comparison of these approaches is in Figure 26.3.




Figure 26.3: Multiple ways to effectively compare a quantitative response (in this case, the moral expectation score) across multiple groups (in this case, the type of food to which the participant was exposed).

⚠ Warning

If you have smaller than 10 observations in a group, it is more appropriate to use an individual value plot (equivalent to a scatterplot when one axis is a categorical variable) or a jitter plot. Having the raw data is important in these cases.

Each of the plots in Figure 26.3 is reasonable. What makes them useful in addressing the research question is that in each plot, we can compare the degree to which the groups differ relative to the variability within a group. That is, we partition the variability. With each plot, we can say that one of the reasons the groups differ is because of exposure to different food types; however, this difference is extremely small relative to the fact that regardless of which food group you were exposed to, the variability in moral expectations with that group is quite large. Since the predominant variability in the moral exposure is the variability within the groups, we would say there is no signal here. That is, the graphics do not indicate there is any evidence that the average moral expectation scores differ across food exposure groups.

The key to a good summary is understanding the question of interest and building a graphic which addresses this question through a useful characterization of the variability.

27 Building the Statistical Model

In Chapter 10, we introduced the statistical modeling framework. In particular, our general model (see Equation 10.1) for the data generating process was

Response = function(predictor variables, parameters) + noise.

Recall that this model has two components:

- A deterministic component which takes the form of a function of variables and unknown parameters. It is often this component on which we would like to make inference.
- A stochastic component which captures the unexplained variability in the data generating process.

In the previous units we made use of this model for making inference on a single mean response and then extended the model so that the mean response of a quantitative variable could depend on a quantitative predictor through a linear function. In this chapter, we discuss how to choose a function for the deterministic portion of the model so that we can compare the mean response of a quantitative variable across the levels of a factor (categorical predictor).

Note

For those that felt comfortable with Chapter 21, you may notice that much of what we discuss is a special case of the concepts discussed for the general linear regression model. However, we have tried to make this unit stand alone; that is, Chapter 21 (indeed, none of Unit III) is required prior to reading the material discussed in Unit IV, though we may occasionally make references to Unit III to draw connections between the material.

27.1 Statistical Model for A Quantitative Response and a Categorical Predictor

For the Organic Food Case Study, we are comparing the moral expectations (quantitative response) for different food exposures (levels of a categorical variable). Our model for the data generating process is best understood in light of the graphic we used to display the data (see Figure 27.1).




Figure 27.1: Moral expectation scores for students following exposure to one of three food types.

Let's consider how the value 3.67, highlighted red in Figure 27.1, was generated. As we mentioned in the previous chapter, there are two sources of variability that contribute to the moral expectation scores (two reasons that the values are not all the same). One reason the moral expectations differ across participants is the fact that some participants were exposed to different food groups. That is, one reason the value 3.67 differs from others observed is because this subject belongs to the organic group and not the comfort or control groups. In order to incorporate the food type into the model for the data generating process, we might initially consider something like the simple linear model (Equation 17.3) discussed in Chapter 17:

(Moral Expectation Score)_i =
$$\beta_0 + \beta_1$$
(Food Exposure Group)_i + ε_i .

However, the food exposure group is not a quantitative variable. That is, we cannot multiply β_1 (a number) and "comfort" (a word) together as the model suggests. Our solution is to consider the deterministic portion of the data generating process to be a piecewise function:

$$\text{function} \left((\text{Food Exposure Group})_i, \text{parameters} \right) = \begin{cases} \mu_1 & \text{if i-th participant exposed to comfort foods} \\ \mu_2 & \text{if i-th participant exposed to control foods} \\ \mu_3 & \text{if i-th participant exposed to organic foods}. \end{cases}$$

This is a completely acceptable function; it involves both the parameters, the mean response μ_1, μ_2, μ_3 for each of the three groups, and the factor of interest (which good group participants are exposed to). However, this function is cumbersome to write, and it is not necessarily easy to communicate to the computer or generalize to more complex settings. In order to create something that is a bit easier to work with, observe how the function works: it receives an input regarding which group a participant belongs to, and it directs you to the appropriate parameter to represent the mean response for their group (the output of the function). We can recreate this through a clever use of indicator variables (Definition 21.1).

27.2 Indicator Variable

An indicator variable is a binary (takes the value 0 or 1) variable used to represent whether an observation belongs to a specific group defined by a categorical variable.

Consider defining the following indicator variables:

$$\begin{aligned} &(\operatorname{Comfort})_i = \begin{cases} 1 & \text{if i-th participant exposed to comfort foods} \\ 0 & \text{otherwise} \end{cases} \\ &(\operatorname{Control})_i = \begin{cases} 1 & \text{if i-th participant exposed to control foods} \\ 0 & \text{otherwise} \end{cases} \\ &(\operatorname{Organic})_i = \begin{cases} 1 & \text{if i-th participant exposed to organic foods} \\ 0 & \text{otherwise}. \end{cases} \end{aligned}$$

Each of these indicator variables acts as a light switch, clicking "on" (taking the value 1) when a participant belongs to the corresponding group, and turning "off" (taking the value 0) when they do not. Further, because of how the categorical variable is defined, each participant must belong to exactly one group; therefore, exactly one light switch is turned on for each participant. These indicator variables allow us to capture the piecewise nature of the above function:

$$(Moral Expectations)_i = \mu_1(Comfort)_i + \mu_2(Control)_i + \mu_3(Organic)_i + \varepsilon_i.$$
 (27.1)

Note

In Chapter 21, a single indicator variable was used to capture two groups; this was possible because the model for the data generating process included an intercept term — a parameter not associated with any predictor variable. In this chapter, there is no intercept term.

Consider the group of individuals exposed to comfort foods; for this group, the "comfort" indicator variable takes the value 1, and the "control" and "organic" indicator variables take the value 0. Therefore, the model suggests the average moral expectations for individuals exposed to comfort foods is

$$\mu_1(1) + \mu_2(0) + \mu_3(0) = \mu_1.$$

Notice that the deterministic portion of our model for the data generating process is specifying the mean response, continuing what we have seen in previous units.



Big Idea

The deterministic component of a statistical model for the data generating process incorporates the parameters which govern the question of interest. It is built to explain differences in the mean response of the units of observation.

The deterministic portion of Equation 27.1 says that every single person exposed to the same food group should have the same moral expectations, and the stochastic portion of Equation 27.1 allows individuals within the same group to vary because of random noise.



Big Idea

The stochastic component of a statistical model for the data generating process captures the unexplained variability due to natural variability in the population or measurement error in the response.

Because this model is very directly partitioning the variability in the response (into the portion explained by the factor of interest and the portion unexplained by the factor of interest), it is often called an ANOVA (ANalysis Of VAriance) model.

ANOVA Model

For a quantitative response and a single categorical predictor (also known as a factor) with k levels, the ANOVA model is

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \varepsilon_i$$
 (27.2)

$$(\text{Group } j)_i = \begin{cases} 1 & \text{i-th unit belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

is an indicator variable capturing whether a unit belongs to the j-th group and $\mu_1, \mu_2, \dots, \mu_k$ are the parameters governing the model for the data generating process.

This is somewhat of a silly name since we have seen that all statistical models partition the variability in the response, and the above model is simply a special case of the general linear regression model. But, the name has stuck.

(Optional) Comparison of ANOVA to General Linear Regression Model

Notice that Equation 27.2 has a very similar form to Equation 21.3. The primary difference is the presence of an intercept term in Equation 21.3. Each indicator is acting as a predictor in the model for the data generating process. If we were to add an intercept to Equation 27.2, and then remove one of the indicator variables to construct a reference group, then we would completely fall under the general linear regression model framework. This means that as we move forward, we can adopt results established for the general linear regression model.

As each parameter in Equation 27.1 represents the mean response of a particular food exposure group, it is intuitive that we would estimate the parameter with the observed sample mean of the corresponding food exposure group. This turns out to be equivalent to estimating the parameters through the method of least squares (Definition 17.2).

Note

The least squares estimate of μ_j in Equation 27.2 is the sample mean response from the j-th group. That is, we use the observed mean response from each group as an estimate of the mean response for each group.

Just as before, while point estimates are helpful, inference requires that we quantify the variability in our estimates. And, just as before, we need to distinguish between the model for the data generating process and the model for the sampling distribution of the parameter estimates and the model for the null distribution of a standardized statistic. And, just as before, to move from the model for the data generating process to a model for the sampling distribution of the parameter estimates, we impose conditions on the stochastic component.

28 Conditions on the Error Term of an ANOVA Model

In the previous chapter we developed a general model for the data generating process of a quantitative response as a function of a categorical predictor:

$$(\text{Response})_i = \sum_{i=1}^k \mu_k(\text{Group } j)_i + \varepsilon_i,$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if the i-th observation is from Group } j \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables to capture the level of the factor to which each observation belongs.

We also discussed a common method for estimating the parameters of this model from a sample — the method of least squares. However, if we are to construct a model for the sampling distribution of these estimates we must add some structure to the stochastic component ε in the model. In this chapter, we focus on the most common conditions we might impose and how those conditions impact the model for the sampling and null distributions (and therefore the computation of a confidence interval or p-value).

28.1 Independent Errors

The first condition we consider is that the noise attributed to one observed individual is independent of the noise attributed to any other individual observed. That is, the amount of error in any one individual's response is unrelated to the error in any other response observed. This is the same condition we introduced in Chapter 10 and Chapter 18.

Independence Condition

The independence condition states that the error in one observation is independent (see Definition 10.3) of the error in all other observations.

With just this condition, we can use a bootstrap algorithm in order to model the sampling distribution of the least squares estimates of our parameters (see Appendix A). However, additional conditions are often considered.

28.2 Same Degree of Precision

The second condition that is typically placed on the distribution of the errors is that the variability of the errors is the same in each group. We introduced this condition in Chapter 18. Practically, this condition implies that the variability of the response within each group is the same for all groups.

i Constant Variance

Also called homoskedasticity, the constant variance condition states that the variability of the errors within each group is the same across all groups.

With this additional condition imposed, we are able to modify our bootstrap algorithm when constructing a model for the sampling distribution of the least squares estimates.

At this point, you might ask "if we assume constant variance in the classical ANOVA model and in the classical regression model, why did we not consider this assumption in the single response case (Chapter 10)?" Remember, the condition is not that the errors are the same; the condition is that the variability of the errors within each group is the same for every group. In Chapter 10, our model for the data generating response did not have a predictor; therefore, there was only a single group to which the response belonged. The variability in the response cannot vary across the predictor if there is no predictor; so, this condition could not be violated in that setting!

28.3 Specific Form of the Error Distribution

The third condition that is typically placed on the distribution of the errors is that the errors follow a Normal distribution, as discussed in Chapter 18. Here, we are assuming a particular structure on the distribution of the error population.

i Normality

The normality condition states that the distribution of the errors follows the functional form of a Normal distribution (Definition 18.2).

Let's think about what this condition means for the responses. Given the shape of the Normal distribution, imposing this condition (in addition to the other conditions) implies that some

errors are positive and some are negative. This in turn implies that some responses will be above average for their group, and some responses will be below average for their group. More, because the distribution of the response within each group is just a shifted version of the distribution of the errors, we know that the distribution of the response variable itself follows a Normal distribution.

With this last condition imposed, we can construct an analytic model for the sampling distribution of the least squares estimates. As in regression modeling, we are not required to impose all three conditions in order to obtain a model for the sampling distribution of the estimates. Historically, however, all three conditions have been routinely imposed in the scientific and engineering literature.

28.4 Classical ANOVA Model

We have discussed three conditions we could place on the stochastic portion of the data generating process. Placing all three conditions on the error term is what we refer to as the "Classical ANOVA Model."

Definition 28.1 (Classical ANOVA Model). For a quantitative response and single categorical predictor with k levels, the classical ANOVA model assumes the following data generating process:

$$(\text{Response})_i = \sum_{j=1}^k \mu_j (\text{Group } j)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if i-th observation belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables and where

- 1. The error in the response for one subject is independent of the error in the response for all other subjects.
- 2. The variability in the error of the response within each group is the same across all groups.
- 3. The errors follow a Normal Distribution.

This is the default "ANOVA" analysis implemented in the majority of statistical packages.

Warning

A "hidden" (typically unstated but should not be ignored) condition is that the sample is representative of the underlying population. In the one sample case (Chapter 10), we referred to this as the errors being "identically distributed." We no longer use the "identically distributed" language for technical reasons; however, we still require that the sample be representative of the underlying population.

We note that "ANOVA" need not require all three conditions imposed in Definition 28.1. Placing all three conditions on the error term results in a specific analytical model for the sampling distribution of the least squares estimates. Changing the conditions changes the way we model the sampling distribution.



Big Idea

The model for the sampling distribution of a statistic is determined by the conditions you place on the stochastic portion of the model for the data generating process.

At this point, you might be wondering what happened to the "mean-0" condition we imposed in regression models. The mean-0 condition stated that the error was 0, on average, for all values of the predictor. Recall that this assumption implied that the model for the mean response was correctly specified — that no curvature was ignored. In our model above, with only a single categorical predictor with k levels (captured through k indicator variables), there is no "trend" being described. That is, instead of saying that the mean response increases, decreases, or has any particular form, the deterministic portion of the model allows the mean response in one group to be completely unrelated to the mean response in any other group. Since there is no "trend" term in the mean response model, we cannot have misspecified that trend. The mean-0 condition is only required when you make a simplifying assumption constraining the mean response to follow a specific functional form.

Note

The lack of the mean-0 condition in ANOVA is why we did not need to consider the mean-0 condition in Unit I when we were focused on inference about a single mean.

28.5 Imposing the Conditions

Let's return to our model for the moral expectation score as a function of the food exposure group given in Equation 27.1:

(Moral Expectations)_i =
$$\mu_1(\text{Comfort})_i + \mu_2(\text{Control})_i + \mu_3(\text{Organic})_i + \varepsilon_i$$
,

Table 28.1: Estimated average moral expectation score for participants exposed to one of three food groups.

Exposure Group	Estimated Mean Moral Expectation Score
Comfort Foods	5.58
Control Foods	5.65
Organic Foods	5.75

Table 28.2: Summary of the model fit relating the moral expectation score of college students to the type of food to which they were exposed.

Term	Estimate	Standard Error	Lower 95% CI	Upper 95% CI
Comfort Foods Group	5.585	0.128	5.331	5.839
Control Foods Group	5.654	0.130	5.397	5.912
Organic Foods Group	5.750	0.131	5.490	6.010

where we use the same indicator variables defined in Chapter 27. We were interested in the following research question:

Does the average moral expectation score differ for at least one of the three food exposure groups?

This was captured by the following hypotheses:

$$\begin{split} H_0: \mu_1 = \mu_2 = \mu_3 \\ H_1: \text{at least one } \mu_j \text{ differs.} \end{split}$$

Using the method of least squares, we constructed point estimates of the parameters in the model; this leads to the following estimates of the average moral expectation score for each exposure group:

If we are willing to assume the data is consistent with the conditions for the classical ANOVA model, we are able to model the sampling distribution of these estimates and therefore construct confidence intervals. Table 28.2 summarizes the results of fitting the model described in Equation 27.1 using the data available from the Organic Foods Case Study. In addition to the least squares estimates, it also contains the standard error (see Definition 6.4) of each statistic, quantifying the variability in the estimates. Finally, there is a 95% confidence interval for each parameter.

Chapter 6 described, in general, how confidence intervals are constructed. Under the classical ANOVA model, there is an analytical model for the sampling distribution, and it is known. As a result, the confidence interval can be computed from a formula.

Formula for Confidence Interval Under Classical ANOVA Model

If the classical ANOVA model is assumed, the 95% confidence interval for the parameter μ_i can be approximated by

$$\hat{\mu}_i \pm (1.96)$$
 (standard error of $\hat{\mu}_i$)

The confidence intervals for the mean moral expectation score within each group were constructed assuming the classical ANOVA model. And, while the confidence intervals are similar for each of the groups, we have not actually addressed the question of interest. We cannot use the confidence intervals given to directly compare the groups. Instead, we must directly attack the hypotheses of interest by computing a p-value. We consider this in the next chapter.

Warning

It is common to try and compare the mean response of several groups by determining if the confidence intervals for the mean response of each group overlap. This is a mistake. If you want to compare groups, you need to do conduct an analysis that directly addresses the hypothesis of interest.

28.6 Recap

We have covered a lot of ground in this chapter, and it is worth taking a moment to summarize the big ideas. In order to compare the mean response in each group, we took a step back and modeled the data generating process. Such a model consists of two components: a deterministic component explaining the response as a function of the predictor, and a stochastic component capturing the noise in the system.

Certain conditions are placed on the distribution of the noise in our model. With a full set of conditions (classical ANOVA model), we are able to model the sampling distribution of the least squares estimates analytically. We can also construct an empirical model for the sampling distribution of the least squares estimates assuming the data is consistent with fewer conditions.

In general, the more conditions we are willing to impose on the data generating process, the more tractable the analysis; however, the most important aspect is that the data come from a process which is consistent with the conditions we impose, which is discussed in Chapter 30.

29 Quantifying the Evidence

In the previous two chapters, we described a model for describing the data generating process for a quantitative response as a function of a single categorical predictor:

$$(\text{Response})_i = \sum_{j=1}^k \mu_j (\text{Group } j)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{i-th unit belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

is an indicator variable.

Chapter 27 discussed obtaining estimates of these unknown parameters using the method of least squares. Chapter 28 imposed conditions on the stochastic portion of the model in order to develop a confidence interval for each parameter. In this chapter, we turn to performing inference through the computation of a p-value for a set of hypotheses. As we saw with regression models in Chapter 19, this is accomplished through partitioning variability.

Figure 29.1 displays a numeric response across three groups for two different datasets. Consider the following question:

For which dataset is there *stronger* evidence that the response is associated with the grouping variable?

Nearly everyone would say that Dataset A provides stronger evidence of a relationship between the grouping variable and the response. We generated these data such that the mean for Groups I, II and III are 5, 6 and 7, respectively, for both Datasets A and B. While there is a difference, on average, in the response across the groups in both cases, it is correct that Dataset A provides stronger evidence for that relationship. The real question is "what is it that leads everyone to make the same conclusion when we have not yet discussed how to analyze this data?" When we ask students why they feel Dataset A provides stronger evidence, we typically hear that it is because the "gaps" between the groups "look bigger." Exactly!




Figure 29.1: Simulated data illustrating that signal strength is determined by partitioning variability. There is a clear signal (difference in the location across groups) for Dataset A but not for Dataset B.

29.1 Partitioning Variability

Subconsciously, when we are deciding whether there is a difference in the average response between the groups, we are partitioning the variability in the response. We are essentially describing two sources of variability: the variability in the response caused by subjects belonging to different groups and the variability in the response within a group (Figure 29.2). In both Datasets A and B from Figure 29.1, the **between-group variability** is the same; the difference in the means from one group to another is the same for both datasets. However, the **within-group variability** is much smaller for Dataset A compared to Dataset B.

Definition 29.1 (Between Group Variability). When comparing a quantitative response across groups, the between group variability is the variability in the average response from one group to another.

Definition 29.2 (Within Group Variability). When comparing a quantitative response across groups, the within group variability is the variability in the response within each group.

The power of Figure 29.1 is that it allows us to examine the between group variability (how the average responses differ from one another) relative to the within group variability (how the responses within a group differ from one another). What we see is that the larger this ratio,




Figure 29.2: Illustration of partitioning the variability in the response to assess the strength of a signal.

the stronger the signal. Quantifying the strength of a signal is then about quantifying the ratio of these two sources of variability. Let this sink in because it is completely counterintuitive. We are saying that in order to determine if there is a difference in the mean response across groups, we have to examine variability. Further, a signal in data is measured by the variability it produces. It is for this reason that comparing a quantitative response across a categorical variable is termed ANalysis Of VAriance (ANOVA).



Big Idea

Consider the ratio of the variability between groups to the variability within groups. The larger this ratio, the stronger the evidence of a signal provided by the data.

This partitioning is a bit easier to visualize here than it was for the simple linear regression model, but the process is actually exactly the same.

29.2 Forming a Standardized Test Statistic

Let's return to our model for the moral expectation score as a function of the food exposure group given in Equation 27.1:

(Moral Expectations)_i =
$$\mu_1(\text{Comfort})_i + \mu_2(\text{Control})_i + \mu_3(\text{Organic})_i + \varepsilon_i$$
,

where we use the same indicator variables defined in Chapter 27. We were interested in the following research question:

Does the average moral expectation score differ for at least one of the three food exposure groups?

This was captured by the following hypotheses:

$$H_0: \mu_1 = \mu_2 = \mu_3$$

 $H_1:$ at least one μ_j differs.

As we stated above, quantifying the strength of a signal is equivalent to quantifying the ratio of two sources of variability. This ratio will form our standardized statistic. Our model acknowledges these two sources of variability; the question we then have before us is the following: how do we measure these sources of variability?

As with the linear regression model, we want to move forward with a goal of trying to say something like

$$\begin{pmatrix} \text{Total Variability} \\ \text{in the Moral Expectations} \end{pmatrix} = \begin{pmatrix} \text{Variability due} \\ \text{to Food Exposure} \end{pmatrix} + \begin{pmatrix} \text{Variability due} \\ \text{to Noise} \end{pmatrix}$$

As we have seen in Chapter 5, Chapter 12, and in Chapter 19, variability can be quantified through considering the "total" distance the observations are from a common target (for example, the mean response) where "distance" is captured by squared deviations. That is, the total variability in the moral expectation score can be measured by

$$\sum_{i=1}^{n} \left[(\text{Moral Expectation})_i - (\text{Overall Mean Moral Expectation}) \right]^2. \tag{29.1}$$

Notice this quantity is related to, but is not equivalent to, the sample variance. It measures the distance each response is from the sample mean and then adds these distances up. This **Total Sum of Squares** is exactly the same as we developed for the regression model (Definition 19.1):

29.3 Total Sum of Squares

The Total Sum of Squares, abbreviated SST, is given by

$$SST = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Overall Mean Response}) \right]^2$$

where the overall average response is the sample mean.

We now have a way of quantifying the total variability in the moral expectation scores we observed; we now want to partition (or separate) out this variability into its two components. In order to capture the variability due to the food exposure groups, we consider how it plays a role in the model for the data generating process: it allows participants from different groups to have a different mean response. That is, the deterministic portion of the model for the data generating process is the model's attempt to explain how changes in the group explain changes in the moral expectation score. Finding the variability in the moral expectations due to the food exposure groups is then equivalent to finding the variability among these estimated (or predicted) mean responses:

$$\sum_{i=1}^{n} \left[(\text{Group Mean Moral Expectation})_i - (\text{Overall Mean Moral Expectation}) \right]^2. \tag{29.2}$$

This term quantifies the variability explained by the groups, and it is called the **Treatment Sum of Squares**, but it is equivalent to the **Regression Sum of Squares** (Definition 19.2):

29.4 Regression Sum of Squares

The Regression Sum of Squares, abbreviated SSR, is given by

$$SSR = \sum_{i=1}^{n} \left[(\text{Predicted Mean Response})_i - (\text{Overall Mean Response}) \right]^2$$

where the predicted mean response is computed using the least squares estimates and the overall mean response is the sample mean.

This is also known as the Treatment Sum of Squares (abbreviated SSTrt) in ANOVA.

We need to be careful here that the overall mean response is the sample mean across all groups, while the "predicted mean response" is the observed sample mean within each group.

Finally, the unexplained noise, ε in our model for the data generating process, is the difference between the actual response and the deterministic portion of the model (in our case, the true mean response in each group). This variability in the noise is then the variability within each group:

$$\sum_{i=1}^n \left[(\text{Moral Expectation})_i - (\text{Group Mean Moral Expectation})_i \right]^2.$$

{eq-anovateststat-sse}

This **Error Sum of Squares** is exactly the same as we developed for the regression model (Definition 19.3):

29.5 Error Sum of Squares

The Error Sum of Squares, abbreviated SSE and sometimes referred to as the Residual Sum of Squares, is given by

$$SSE = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Predicted Mean Response})_i \right]^2$$

where the predicted mean response is computed using the least squares estimates.

Again, the "predicted mean response" for ANOVA is the observed sample mean within each group.

9 I

Big Idea

The total variability in a response can be partitioned into two components: the variability explained by the predictor and the unexplained variability left in the error term. This is represented in the formula

$$SST = SSR + SSE$$

As we have seen repeatedly, hypothesis testing is really model comparison; that is, our hypotheses of interest comparing the means really suggest two separate models for the data generating process:

 $\text{Model 1:} \quad (\text{Moral Expectation})_i = \mu_1(\text{Comfort})_i + \mu_2(\text{Control})_i + \mu_3(\text{Organic})_i + \varepsilon_i$

Model 0: (Moral Expectation)_i = $\mu + \varepsilon_i$,

where μ represents the common value of μ_1, μ_2, μ_3 under the null hypothesis (that is, the shared overall mean moral expectation score). The model under the null hypothesis (Model 0) has fewer parameters because it is a constrained version of Model 1 resulting from setting $\mu_1 = \mu_2 = \mu_3$ (the common value of which we called μ). In fact, while Model 1 says there are two components (food exposure group and noise) contributing to the variability observed in the moral expectations, Model 0 says that there is only a single component (noise).

Regardless of which model we choose, the total variability in the response remains the same. We are simply asking whether the variability explained by the food exposure group is sufficiently large for us to say it has an impact. In particular, if the null hypothesis were true, we would expect all the variability in the response to be channeled into the noise $(SST \approx SSE)$. In fact, think about computing the error sum of squares for Model 0 above; it would be

$$SSE_0 = \sum_{i=1}^n \left[(\text{Moral Expectations})_i - (\text{Overall Mean Moral Expectation}) \right]^2$$

since the least squares estimate of μ in Model 0 is the sample mean (see Appendix B). But, this is equivalent to the total sum of squares for Model 1 (Equation 29.1). This confirms our intuition that if the null hypothesis were true, we would expect all the variability in the response to be channeled into the noise.

If, however, the alternative hypothesis is true and the food exposure group explains some portion of the variability in the moral expectations, then we would expect some of the variability to be channeled out of the noise term (SSR > 0). Because we have partitioned the variability, we now take a moment to recognize that

$$SSR = SST - SSE$$
,

but we know that the total sum of squares is just the error sum of squares from the reduced model (Model 0) as shown above. Therefore, we can write

$$SSR = SSE_0 - SSE_1, \tag{29.3}$$

where we use the subscripts to denote whether we are discussing the error sum of squares from the reduced model (Model 0) or the full unconstrained model (Model 1). That is, Equation 19.4 reveals that the regression sum of squares is the equivalent of the shift in the error sum of squares as we move from the reduced model under the null hypothesis to the more complex model under the alternative hypothesis.



Big Idea

For a particular dataset, the regression sum of squares quantifies the shift in the error sum of squares as we move from a reduced model to a more complex model. It measures the "signal" in the data represented by the more complex model for the data generating process.

Just as we stated in Chapter 19, while sums of squares partition the variability, mean squares quantify the actual variance, and it turns out working with variances is beneficial. To move from a sum of squares to a mean square, we need to divide by the **degrees of freedom** (Definition 19.5).

29.6 Degrees of Freedom

A measure of the flexibility in a sum of squares term; when a sum of squares is divided by the corresponding degrees of freedom, the result is a variance term.

Definition 29.3 (Mean Square (in ANOVA)). A mean square is the ratio of a sum of squares and its corresponding degrees of freedom. For a model of the form in Equation 27.2, we have

- Mean Square Total (MST): estimated variance of the responses; this is the same as the sample variance of the response.
- Mean Square for Regression (MSR): estimated variance of the sample mean responses from each group; this is also called the Mean Square for Treatment (MSTrt) in ANOVA.
- Mean Square Error (MSE): estimated variance of the error terms; this is equivalent to the estimated variance of the response within a group.

In each case, the mean square is an estimated variance. These are equivalent to the MST, MSR, and MSE in the regression model (Definition 19.6).

The MSR quantifies the between group variability and the MSE quantifies the within group variability. The MSR is our signal. The larger this variance, the further apart the observed group sample means are from one another (providing evidence for the alternative hypothesis); the smaller this variance, the closer the observed group sample means are (consistent with the null hypothesis).

Degrees of Freedom for MSR (or MSTrt)

In an ANOVA model, the MSR is capturing the variability among k sample means; therefore, the associated degrees of freedom are k-1.

MSE provides a measure of the noise within the data. Again, in Figure 29.1, the variability between the means is identical for the two datasets; the signal is easier to discern for Dataset A because this variability is larger with respect to the noise.

i Applicability of MSE

The MSE is a pooled estimate of the variance within a group. That is, it is a weighted average of the observed group sample variances. Therefore, interpreting this value only makes sense if we are willing to impose the constant-variance condition. However, regardless of whether we impose the constant-variance condition, the MSE is helpful in computing a standardized statistic.

i Degrees of Freedom for MSE

In an ANOVA model, predicting the mean response requires estimating k parameters (the sample mean within each of the k groups); therefore, the associated degrees of freedom are n-k.

We are now ready to define our standardized statistic as the ratio of mean squares:

$$\frac{MSR}{MSE} = \frac{\left(SST - SSE_1\right)/(k-1)}{SSE_1/(n-k)} = \frac{\left(SSE_0 - SSE_1\right)/(k-1)}{SSE_1/(n-k)},$$

where again we have added subscripts to emphasize from which model we are computing the above sums of squares. This standardized statistic could be used to quantify the signal-to-noise ratio (the amount of evidence) in the sample for testing the hypotheses

$$\begin{split} H_0: \mu_1 &= \mu_2 = \mu_3 \\ H_1: \text{At least one } \mu_j \text{ differs} \end{split}$$

However, we can generalize this for testing a range of hypotheses within our model for the data generating process.

Definition 29.4 (Standardized Statistic for ANOVA). Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 27.2):

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \varepsilon_i,$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{i-th unit belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

is an indicator variable. Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0. A standardized statistic, sometimes called the "standardized F statistic," for testing the hypotheses is given by

$$T^* = \frac{\left(SSE_0 - SSE_1\right)/(k-r)}{SSE_1/(n-k)},$$

where k is the number of parameters in the full unconstrained model and r is the number of parameters in the reduced model. Defining

$$MSA = \frac{SSE_0 - SSE_1}{k - r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model. Just as before, the MSE represents the residual variance — the variance in the response for a particular set of the predictors.

It can be shown that this standardized statistic is a special case of the one discussed in Chapter 21 and is therefore consistent with the ones presented in Chapter 19 and Chapter 12. The numerator captures the signal by examining the difference between what we expect the error sum of squares to be under the null hypothesis and what we actually observe; the denominator captures the background noise (relative to the estimated mean response from the full model). Larger values of this standardized statistic indicate more evidence in the sample against the null hypothesis.

We should not lose sight of the fact that our standardized statistic is really a result of partitioning the variability and considering the variability explained by the factor of interest relative to the noise in the response. Underscoring that the standardized statistic is a result of this partitioning, the analyses of these sources of variability is often summarized in a table similar to that represented in Figure 19.3, which is called an "ANOVA table" (Figure 29.3).

	— Partition	ning Variab	oility			
Source	DF	SS	MS	F	P-Value	
Factor	k-1	SSTrt	MSTrt	$\frac{MSTrt}{MSE}$	р	
Error	n-k	SSE	MSE	1		
Total	n-1	SST				
				St.	andardized	d Test Statistic
				Estima	te of Resid	ual Variability

Figure 29.3: Layout of an ANOVA table which summarizes the analysis conducted. Emphasis is on partitioning the variability.

This table is extremely familiar as we encountered it in Chapter 19. Just as before, the last entry in the table is the p-value. As with any p-value, it is computed by finding the likelihood, assuming the null hypothesis is true, of getting, by chance alone, a standardized statistic as extreme or more so than that observed in our sample. "More extreme" values of the statistic would be larger values; so, the area under the null distribution to the right of the observed statistic is the p-value.

We note that while mathematical formulas have been provided to add some clarity to those who think algebraically, our emphasis is *not* on the computational formulas as much as the idea that we are comparing two sources of variability.

Let's return to the question that inspired our investigation in this chapter:

Does the average moral expectation score differ for at least one of the three food exposure groups?

Table 29.1: Analysis of the sources of variability in the moral expectation sore as a function of the food exposure groups.

Term	DF	Sum of Squares	Mean Square	Standardized Statistic	P-Value
Food Exposure Group	2	0.562	0.281	0.406	0.667
Error	120	82.951	0.691		

This was captured by the following hypotheses:

 $H_0: \mu_1=\mu_2=\mu_3$

 H_1 : at least one μ_i differs.

Table 29.1 gives the ANOVA table summarizing the partitioned sources of variability in the moral expectation score. We have a large p-value (computed assuming the data is consistent with the classical ANOVA model). That is, the sample provides no evidence to suggest the average moral expectation differs across any of the food exposure groups. The study suggests it is reasonable to assume that the foods we are exposed to do not impact our moral expectation, on average.



Big Idea

Determining if a response is related to a categorical predictor is done by determining if the predictor explains a significant portion of the variability in the response.

In this chapter, we partitioned variability as a way of evaluating the strength of evidence the predictor plays in determining the response. As with the linear regression model, partitioning the variability is a key step. By partitioning the variability in the response, we are able to construct a standardized statistic for testing the hypothesis of interest. The model for the null distribution of this statistic depends upon the conditions we are willing to impose on the stochastic portion of the data generating process. Regardless of the conditions we impose, we can interpret the resulting p-value similarly. It provides an indication of whether the data suggests that the average response differs for at least one of the groups.

Of course, the interpretation of the p-value depends on the conditions we impose. We should not choose such conditions without performing some type of assessment to ensure those conditions are reasonable — that the data is consistent with the conditions. That is the focus of the next chapter.

30 Assessing the Modeling Conditions in ANOVA

In this unit we have discussed a model relating a quantitative response to a categorical predictor. For the Organic Food Case Study, our model has had the form

$$(\text{Moral Expectations})_i = \mu_1(\text{Comfort})_i + \mu_2(\text{Control})_i + \mu_3(\text{Organic})_i + \varepsilon_i,$$

where we use the same indicator variables defined in Chapter 27. Further, we considered three conditions on the distribution of the error term:

- 1. The error in the moral expectation score for one individual is independent of the error in the moral expectation score for all other individuals.
- 2. The variability in the error for the moral expectation score within a group is the same for any food exposure group.
- 3. The error in the moral expectation score follows a Normal distribution.

However, while we imposed all three of these conditions in Chapter 29, we could have developed an empirical model for the null distribution of the standardized statistic only enforcing the first of these conditions on the distribution of the error. Unfortunately, we cannot simply state conditions and then proceed blindly. In order to rely on the p-values and confidence intervals produced from any modeling procedure, the data must be consistent with the conditions imposed.

In this section, we discuss how we assess these conditions qualitatively. Just as we saw in Chapter 20, while the conditions are placed on the error terms, they are assessed using **residuals**.

30.1 Residual

The difference between the observed response and the predicted response (estimated deterministic portion of the model). Specifically, the residual for the i-th observation is given by

$$(Residual)_i = (Response)_i - (Predicted Mean Response)_i$$

where the "predicted mean response" is often called the predicted, or fitted, value.

Residuals mimic the noise in the data generating process.

For the ANOVA model, the predicted mean response is the observed sample mean of each group.



Big Idea

The conditions are placed on the error term, but they are assessed with residuals.

30.2 Assessing the Independence Condition

The error in the moral expectation score for one individual is independent of the error in the moral expectation score for all other individuals.

Generally, independence is assessed by considering the method in which the data was collected and considering the context with a discipline expert. By carefully considering the manner in which the data was collected, we can typically determine whether it is reasonable that the errors in the response are independent of one another. Some key things to consider when examining the data collection process:

- Are there repeated observations made on the same subject? This often suggests some type of relationship between the responses and therefore would not be consistent with errors being independent. In particular, look for blocking.
- Is the response measured over time (time-series) such as daily temperature over the course of a month? Time-series data often exhibits strong period-to-period relationships suggesting the errors are not independent. For example, if it is hot today, it will probably be hot tomorrow as well.
- Is there a learning curve in how the data was collected? Learning curves again suggest some dependence from one observation to the next. For example, a new nurse may become better at collecting pulse readings with more practice over time.
- Measurement devices which are failing over time will introduce a dependence from one observation to the next. Imagine a bathroom scale that begins to add an additional pound each day. Then, being above average weight one day will most likely lead to an above average weight the next, due primarily to the measurement device. Generally, independence is assessed through the context of the data collection scheme. By carefully considering the manner in which the data was collected, we can typically determine whether it is reasonable that the errors in the response are independent of one another. Some key things to consider when examining the data collection process:

These last three points illustrate a particular deviation from our condition of independence in which two observations collected close together in time are related. When we know the order in which the data was collected, we can assess whether the data tends to deviate from the condition of independence in this manner. This is done graphically through a **time-series plot** of the *residuals*. If two errors were unrelated, then the value of one residual should tell us nothing about the value of the next residual. Therefore, a plot of the residuals over time should look like noise (since residuals are supposed to mimic the noise in the model). If there are any trends, then it suggests the data is not consistent with independence.

30.3 Time-Series Plot

A time-series plot of a variable is a line plot with the variable on the y-axis and time on the x-axis.

Graphically Assessing the Independence Condition

If the data is consistent with the independence condition, we would not expect to see a trend in the *location* or *spread* in a time-series plot of the residuals. Note that this graphic can only be used if the order in which the data was collected is known, and the order indicates some natural timing.

For the Organic Food Case Study, participants were assessed simultaneously within a large lecture. Therefore, there is no ordering in time to be concerned about, and a time-series plot of the residuals would not be useful here. Since we cannot rely on a graphic, we can only rely on what we know about how the data was collected. Students worked individually on the questionnaire; there was nothing in the data description of how the data was collected that would lead us to believe that one student's response would be impacted by any other student's response. And, no blocking was implemented. Our review of the data collection methods suggests it is reasonable to assume that the errors in the moral expectation score are unrelated to one another.

30.4 Assessing Homoskedasticity

The variability of the error in the moral expectation within each food exposure group is the same across all food exposure groups.

We want the variability in the errors within a group to be the same across the groups. We can do this by examining side-by-side boxplots (or jitter plots, etc.) of the residuals within each of the groups. Figure 30.1 shows the residuals for each individual across the various groups. Notice that the boxes for each group are roughly the same size; that is, the interquartile ranges

are similar. This suggests that the variability within each group is similar from one group to the next. That is, the data is consistent with the constant variance condition.




Figure 30.1: Comparison of the residuals predicting the moral expectation score for college students exposed to different types of food.

There is a second (equivalent) approach to assessing this condition. From the model for the data generating process, we see that the response for any individual is some constant plus noise; therefore, the distribution of the responses for any group is simply a shifted version of the distribution of the errors within the same group. If the variability in the errors for each response is the same, then the variability of the response must be the same for each group. Therefore, we can also examine the side-by-side boxplots (or jitter plots, etc.) of the response instead of the residuals. Figure 30.2 shows the moral expectation score for each individual across the various groups. Just as in the previous graphic, the interquartile ranges are similar for each of the three groups indicating the data is consistent with this condition. The benefit of the first approach is that the residuals will always be centered around 0 within each group; this allows for easy side-by-side comparisons; when looking at the observed response, the data need not be aligned across the groups.

Finally, there is a third (equivalent) approach to assessing this condition — assessing it just as we did for linear regression models. We can create a plot of the residuals against the fitted values. Figure 30.3 shows the residuals plotted against the fitted values. Just as in the previous graphic, the interquartile ranges are similar for each of the three groups indicating the data is consistent with this condition.

This does not have the same pattern as what we might have expected from Chapter 20. Remember, our study only had three groups; therefore, the deterministic portion of the model




Figure 30.2: Comparison of the moral expectation scores for college students exposed to different types of food.

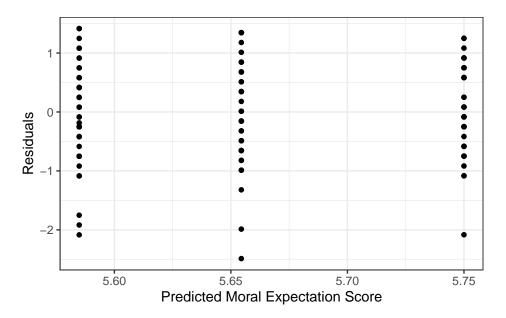


Figure 30.3: Residuals plotted against the predicted the moral expectation scores for college students exposed to different types of food.

for the data generating process is only comparing three groups, and as a result, it can only predict one of three values for the average response. Each vertical "slice" in Figure 30.3 represents the residuals from one of those three predicted responses.



Warning

The order of the groups when plotting the response against the groups need not be the same as the order when plotting the residuals against the fitted values.

⚠ Warning

When plotting the residuals against the fitted values in ANOVA, if two groups are similar, but these differ from the other groups, the vertical "slices" can be nearly on top of one another, making it difficult to assess the constant variance condition. It is for this reason we prefer one of the first two methods discussed in this section.

Graphically Assessing the Constant Variance Condition

If the data is consistent with the constant variance condition, there should be no trends in the *spread* of the residuals (or the response) across each group.

30.5 Assessing Normality

The errors in the moral expectation score follows a Normal distribution.

If the errors follow a Normal distribution, then we would expect the residuals to mimic a sample taken from a Normal distribution. As introduced in Chapter 20, we emphasize the Normal **probability plot** for assessing the Normality condition.

30.6 Probability Plot

Also called a "Quantile-Quantile Plot", a probability plot is a graphic for comparing the distribution of an observed sample with a theoretical probability model for the distribution of the underlying population. The quantiles observed in the sample are plotted against those expected under the theoretical model.

i Graphically Assessing the Normality Condition

If the data is consistent with the normality condition, the Normal probability plot of the residuals should exhibit a straight line with any deviations appearing to be random. Systemic deviations from a straight line indicate the observed distribution does not align with the proposed model.

Figure 30.4 shows the probability plot for the residuals from the Organic Food Case Study.




Figure 30.4: Normal probability plot of the residuals for the Organic Food Case Study.

Overall, the points do tend to follow a straight line. There are some deviations from a linear relationship at each end of the plot, but the deviations are not extreme and do not appear to be systematic. Deviations in the tails are common, especially with larger datasets. And with naturally less data in the tails, it can become more difficult to establish a pattern. We are generally not concerned unless these tails form a part of a larger pattern of deviating from the linear trend. We believe these residuals are consistent with the errors having a Normal distribution.

30.7 General Tips for Assessing Assumptions

First discussed in Chapter 20, we want to remember four things that should be kept in mind when assessing conditions:

- 1. We should not spend an extraordinary amount of time examining any one residual plot; we might convince ourselves of patterns that do not exist. We are looking for major deviations from our expectations.
- 2. We can never prove a condition is satisfied; we can only determine whether the data is consistent with a condition or whether it is not consistent with a condition.
- 3. Any condition required for a particular analysis should be assessed.
- 4. Transparency is crucial.

31 Using the Tools Together

This unit introduced a framework for determining if there is an association between a quantitative response and a categorical predictor. We formed a standardized statistic for measuring the signal, and developed a model for the data generating process which allowed us to model the null distribution of the standardized statistic. In this chapter, we pull these tools together once more to answer a research question.

The primary question we have been addressing in this unit was whether the moral expectations of students were affected by the type of food to which they were exposed. We saw that there was little evidence of a relationship between these two variables. We now use the data from the Organic Food Case Study to answer a related question:

Do the moral expectations of males and females differ?



Warning

This question enforces a gender binary. Whether the questionnaire provided to students during the study only allowed for "male" and "female," or whether the responses from participants only included these two genders, we do not know. Those identifying with a gender other than "male" or "female" are not represented by this study, and that is reflected in the above question.

We note that in general, non-binary genders have been historically overlooked in research design.

31.1 Framing the Question (Fundamental Idea I)

As stated, the above question is ill-posed. We have not identified a variable or parameter of interest. We refine this question to be

Does the average moral expectation score of males differ from that of females?

This question could also be stated as the following set of hypotheses:

Let μ_1 and μ_2 represent the average moral expectation score for males and females, respectively.

 $H_0: \mu_1 = \mu_2$ $H_1: \mu_1 \neq \mu_2$

31.2 Getting Good Data (Fundamental Idea II)

As we are working with previously collected data, our goal in this discussion is not how best to collect the data but making note of the limitations of the data as a result of how it was collected. We previously described the Organic Food Case Study as an example of a controlled experiment. This was true with regard to the primary question of interest (moral expectations and food exposure). However, the subjects were not randomly assigned to gender; therefore, with regard to this question of interest, the data represents an observational study.

It is common for young researchers to believe that if initially a controlled experiment was performed that the data always permits a causal interpretation. However, we must always examine the data collection with respect to the question of interest. Such "secondary analyses" (using data collected from a study to answer a question for which the data was not initially collected) are generally observational studies. As a result, there may be other factors related to gender and moral expectations that drive any associations we observe.



Warning

When answering a question for which a controlled experiment was not originally designed, carefully consider the question as causal interpretations may no longer be appropriate.

31.3 Presenting the Data (Fundamental Idea III)

Our question here is examining the relationship between a quantitative response (moral expectation score) and a categorical predictor (gender). Figure 31.1 compares the distribution of the moral expectation score for the two groups. Note that 3 students did not specify their gender; these subjects are removed from the analysis.



Warning

If the group is unknown, the corresponding observation cannot contributed to the analysis. It is therefore common to remove missing values. In this example, however, we note that not specifying a gender may be informative. These individuals may be indicating that they prefer not to provide their gender or that they do not identify with the two gender options available on the questionnaire.

In this case, given the few number of individuals not indicating their gender, we do not feel confident in using their results to make a claim about students who would not provide their gender on such a questionnaire. It is for this reason they were removed from the analysis.




Figure 31.1: Comparison of the moral expectations of males and females. The average value is added for each group. The three students who did not specify their gender are not represented in this graphic.

We note that there were substantially more females in the study than males. This could be a result of the demographics within the home department of the course or demographics of the university at which the study was conducted. From the graphic, it appears the female participants tended to have higher moral expectations by about 1 point, compared to the male participants.

31.4 Quantifying the Variability in the Estimate (Fundamental Idea IV)

In order to measure the size of the signal, we partition the variability in an ANOVA table, which allows us to compute a standardized statistic. In order to partition the variability, we first consider the following model for the data generating process:

$$(\text{Moral Expectation Score})_i = \mu_1(\text{Male})_i + \mu_2(\text{Female})_i + \varepsilon_i \tag{31.1}$$

Table 31.1: ANOVA table summarizing the comparison of the moral expectation score across gender within the Organic Food Case Study.

Term	DF	Sum of Squares	Mean Square	Standardized Statistic
Gender	1	4.363	4.363	6.517
Error	118	79.008	0.670	

where

$$\begin{aligned} &(\mathrm{Male})_i = \begin{cases} 1 & \text{if i-th participant is male} \\ 0 & \text{otherwise} \end{cases} \\ &(\mathrm{Female})_i = \begin{cases} 1 & \text{if i-th participant is female} \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

are indicator variables capturing the participant's gender. Table 31.1 reports the standardized statistic from our study corresponding to testing the hypotheses

$$H_0: \mu_1 = \mu_2$$
 vs. $H_1: \mu_1 \neq \mu_2$.

Of course, if we were to collect a new sample, we would expect our standardized statistic to change. If we want to conduct inference and determine the strength of evidence in this study, we need a model for its null distribution. In order to construct a model for the null distribution of the standardized statistic, we need to place appropriate conditions on the error term. We have three possibilities:

- 1. The error in the moral expectation score for one individual is independent of the error in the moral expectation score for any other individual.
- 2. The variance of the error in the moral expectation scores for males is the same as the variance of the error in moral expectation scores for females.
- 3. The error in the moral expectation score for individuals follows a Normal Distribution.

Before creating a model for the null distribution and computing a p-value, we need to assess whether the data is consistent with these assumptions. This requires examining the residuals from the model. First, we discuss the assumption of independence. Since the data was collected at a single point in time, known as a *cross-sectional study*, constructing a time-series plot of the residuals would not provide any information regarding this assumption. Instead, we rely on the context of the problem to make some statements regarding whether the data is consistent with this condition (whether making this assumption is reasonable). It is reasonable that the errors are independent. One case in which this might be violated is if students discussed their answers to the questions as they filled out the survey; then, it is plausible that one student

influenced another student's responses. As this is unlikely given the description of the data collection, we feel it is reasonable to assume independence.

Again, note that there is a condition of independence; we are simply saying whether we are willing to assume the condition is satisfied. There is no way to ensure the condition holds.

In order to assess the constant variance condition, let us look back at the boxplots given in Figure 31.1. As the spread of the moral expectation score for each of the two genders is roughly the same, it is reasonable to assume the variability of the errors in each group is the same.

Finally, to assess the Normality condition, we consider a Normal probability plot of the residuals (Figure 31.2). Given that the residuals tend to display a linear relationship, it is reasonable that the errors follow a Normal Distribution.





Figure 31.2: Normal probability plot assessing the assumption that the errors for our model comparing the moral expectation score across gender follow a Normal distribution.

Given that we are comfortable assuming the data is consistent with all conditions from the classical ANOVA model (Definition 28.1), we can make use of an analytical model for the null distribution.

31.5 Quantifying the Evidence (Fundamental Idea V)

Now that we have a model for the null distribution of our standardized statistic, we can determine how extreme our particular sample was by comparing the standardized statistic for our sample with this null distribution (Figure 31.3).



Standardized Statistic Across Repeated Samples Under Null Hypothesis

Figure 31.3: Analytical model of the null distribution of the standardized statistic computed from the Organic Foods Case Study comparing the moral expectation scores across males and females.

Based on the results, the study suggests there is some evidence (p = 0.012) that the average moral expectations of male students differs from that of female students. Looking back at Figure 31.1, females tend to have higher moral expectations on average.

31.6 Conclusion

Throughout this unit, we have examined a framework for examining the association between a quantitative response and a categorical predictor. This reinforces a couple of big ideas we have seen throughout this text:

- The key to measuring a signal is to partition the variability in the response.
- A standardized statistic is a numeric measure of the signal strength in the sample.
- Modeling the data generating process provides us a way of modeling the sampling distribution of the parameter estimates and the null distribution of a standardized statistic when combined with conditions on the stochastic portion of the model for the data generating process.
- Before imposing conditions on the stochastic portion of a data generating process, we should graphically assess whether the data is consistent with these conditions.

Part V

Unit V: Comparing the Average Response Across Correlated Groups

Throughout the text, we have developed the language and logic of statistical inference within the context of a model for the data generating process. In the previous unit, we extended these ideas to comparing the mean response across groups. The key condition required to perform inference was that the error in the observation from one observation was independent of the error in any other observation. A consequence of this condition is that the observations from different groups are independent. In this unit, we consider comparing the mean response across groups under a special case when independence cannot be consumed.

32 Case Study: Paying a Premium for the Experience

Around 2010, Americans were in love with frozen yogurt. Self-serve frozen yogurt shops were popping up in towns across the country. By 2016, the craze had begun to subside¹. What happened to this once booming industry? One possibility is that consumers began to believe that they were paying a premium for the experience; that is, consumers were paying an increased price to purchase a product in a self-serve shop that could be purchased for much less at a local grocery store. Under this theory, the increased price point reflected the experience of eating at the yogurt shop, not an improvement in the quality of the product purchased. While we cannot fully determine why self-serve shops began to disappear, we can assess whether consumers identify a difference in the quality of yogurt from self-serve shops compared with their local grocery retailer.

A group of college students taking a Practice of Science class at a small university conducted a study in which nine consumers were given a cup of vanilla yogurt from each of three different locations:

- East-Side Yogurt: a local frozen yogurt self-serve shop on the east side of town.
- South-Side Yogurt: a local frozen yogurt self-serve shop on the south side of town.
- Name Brand: frozen yogurt from a name-brand regional dairy available at local grocery retailers.

The order in which each participant tasted the yogurt was determined randomly. Each cup was labeled A, B or C to prevent participants from knowing which location each cup represented. The participants were asked to rate the taste, the texture, and the appearance of each cup on a 10-point scale (higher values representing a more appetizing score).

The full dataset is presented below.

 $^{^{1}} https://www.washingtonpost.com/business/economy/baked-goods-coffee-and-cash-rise-from-the-ashes-of-the-frozen-yogurt-craze/2015/12/05/3c1e7d72-99fd-11e5-b499-76cbec161973_story.html?noredirect=on&utm_term=.5012f2fabf08$

Table 32.1: Subset of data from blind taste test comparing the qualities of frozen yogurt across three vendors.

Participant ID	Taste Rating	Texture Rating	Appearance Rating	Type
1	5	7	8	Name Brand
1	7	7	8	South Side
1	8	7	8	East Side
2	8	5	8	Name Brand
2	6	7	9	South Side
2	9	8	10	East Side
3	8	6	6	Name Brand
3	6	9	9	South Side
3	9	5	8	East Side
4	8	7	8	Name Brand
4	6	7	8	South Side
4	9	9	8	East Side
5	8	5	5	Name Brand
5	4	7	8	South Side
5	9	9	8	East Side
6	4	5	8	Name Brand
6	8	9	8	South Side
6	2	7	6	East Side
7	9	7	9	Name Brand
7	9	8	9	South Side
7	10	10	5	East Side
8	5	2	6	Name Brand
8	8	7	7	South Side
8	7	8	7	East Side
9	8	3	8	Name Brand
9	4	8	7	South Side
9	2	8	7	East Side

33 Framing the Question

"Do college students rate (based on taste) yogurts from different vendors differently?" As in the previous two units, the primary question in the Frozen Yogurt Case Study is about the relationship between two variables: the response (taste of the product on scale of 1 to 10; see Definition 3.2) and the factor of interest (vendor of yogurt; see Definition 24.1).

The question, as stated above, is ill-posed. Instead of asking about individual tastes, we want to know if, on average, there is a difference in the taste of yogurt between vendors. That is, we want to test the following hypotheses:

 H_0 : the average taste rating is the same for each of the three vendors. H_1 : the average taste rating differs for at least one of the three vendors.

Mathematically, we write

$$\begin{split} H_0: \theta_1 &= \theta_2 = \theta_3 \\ H_1: \text{At least one } \theta_i \text{ differs from the others} \end{split}$$

where $\theta_1, \theta_2, \theta_3$ represent the average taste ratings for the East Side, Name Brand, and South Side yogurts, respectively. Our question is now centered on the population and captured through parameters, making it a well-posed question. In fact, this seems to be a question we have already addressed. In a way, it is.

In this unit, we will be tackling the same types of questions we did in Chapter 24 — we are comparing the mean of a quantitative response across the levels of a categorical predictor. The difference is that the observations we have observed are not independent of one another, and we must account for this lack of independence in the analysis.

33.1 General Setting

This unit is concerned with comparing the mean response of a numeric variable across k groups. Let $\theta_1, \theta_2, \dots, \theta_k$ represent the mean response for each of the k groups. Then, we are primarily interested in the following hypotheses:

$$\begin{split} H_0: \theta_1 &= \theta_2 = \dots = \theta_k \\ H_1: \text{At least one } \theta_i \text{ differs from the others.} \end{split}$$

When there are only two groups (k = 2), then this can be written as

$$\begin{split} H_0: \theta_1 &= \theta_2 \\ H_1: \theta_1 \neq \theta_2. \end{split}$$

Here we are writing things in the mathematical notation, but let's not forget that every hypothesis has a context. Throughout this unit, we are looking for some signal in the *location* of the response across the groups. Our working assumption then states that the groups are all similar, *on average*.

34 Correlated Data

Chapter 25 discussed three aspects of a well-designed study: replication, randomization, and reduction of noise. Of particular interest for us in this unit is the concept of blocking for creating groups which are as similar as possible to reduce external variability.

34.1 Blocking

Blocking is a way of minimizing the variability contributed by an inherent characteristic that results in dependent observations. In some cases, the blocks are the unit of observation which is sampled from a larger population, and multiple observations are taken on each unit. In other cases, the blocks are formed by grouping the units of observations according to an inherent characteristic; in these cases that shared characteristic can be thought of having a value that was sampled from a larger population.

In both cases, the observed blocks can be thought of as a random sample; within each block, we have multiple observations, and the observations from the same block are more similar than observations from different blocks.

A block is formed by subjects which have some underlying commonality. An extreme case occurs when you repeatedly measure the response on the same unit under different treatment conditions. This is exactly what happened in the Frozen Yogurt Case Study. Each participant was measured three times, once for each of the three vendors. In order to really comprehend the effect of this, consider an alternative study design:

Take each of the nine study participants and randomly assign each to one of the three yogurt vendors (3 participants to each yogurt vendor). Each participant rates only the yogurt assigned.

This alternative study design is completely valid and feasible, but there seems an intuitive advantage to having each participant rate not just one yogurt but instead rate all the yogurts being tested. This is indeed the case; by making the participants taste each of the three vendor yogurts, the three groups are as similar as possible (the same group exposed to each of the three treatments). This reduces the noise across the three levels of the factor being studied. This in turn increases the power of the study (given the same number of participants), but it also has an effect on the analysis we conduct.

Note

Examine the language in the above paragraph; we intentionally use "groups," "treatments," and "levels of the factor" interchangeably. Some disciplines prefer one of these terms to the others.

Suppose one of the participants in the study does not really enjoy frozen yogurt (more of an ice cream fan); then, their rating on the taste will tend to be lower than that given by other participants for each of the vendors. Similarly, a huge frozen yogurt fan will tend to have taste ratings (for each vendor) which are higher than other participants. If you are told the taste rating for the Name Brand vendor given by the anti-frozen-yogurt participant, you have a pretty good guess of what their rating will be for the other vendors. There is a correlation between that participant's responses across the factor levels; that is, those three responses are not independent of one another. This relationship is what provides power in the study, but it is also what prevents us from using the methods previously discussed in this text to analyze the data.

In the Frozen Yogurt Case Study, the order in which the yogurt from each vendor was given to a participant was randomized. That is, participants were not randomly assigned to vendors (randomization at the unit-level); instead, the order of the vendors was assigned within each participant (randomization within the unit). When randomization occurs within the block, and the blocks are balanced with respect to the treatment groups, the study is referred to as a randomized complete block design.

Definition 34.1 (Randomized Complete Block Design). A randomized complete block design is an example of a controlled experiment utilizing blocking. Each treatment is randomized to observations within blocks such that within each block every treatment is present and the same number of observations are assigned to each treatment.

In the Frozen Yogurt Case Study, the blocks are formed by taking repeated measurements (multiple ratings) on each participant. Example 25.1 provides an example of a randomized complete block design where the blocks are not formed by taking repeated measurements on the same units; instead, the blocks were composed of separate units (golf greens) which had a similar characteristic (slope of green) which created the block.

While a randomized complete block design is a controlled experiment that uses blocking, we can also have observational studies that make use of blocks. For example, pre-post tests are an example of a block design in which each participant is tested prior to and following some intervention. The participants act as the blocks; however, the "treatment" groups (pre and post) are not randomized; instead, they occur in a specific order (pre measurement taken, treatment applied, post measurement taken). This would be an example of an observational study with blocks.

Example 34.1 (Jump Height). Biomedical Engineering students conducted a study investigating how vertical jump height changes with age. They hypothesized that jump height would increase with age. The students enrolled 19 participants into their study. Each participant was asked to perform a vertical jump, the height of which was recorded using the Qualisys Motion Capture system. The participant's age and height were also recorded.

Students noted that the 19 participants represented 6 families (2 families of 3 individuals, 2 families of 4 individuals, and 1 family of 5 individuals). The engineering students believed that athletic ability (and therefore jump height) may have both a genetic and environmental component common among families.

The primary goal of the study was to quantify the change in the average vertical jump height as the age of subjects increase while accounting for similarities among members of the same family.

Example 34.1 is another example of an observational study using blocks. Notice that participants are not assigned to their age; so, no random allocation occurs. However, there are still blocks that pertain to an inherent characteristic among units. The jump heights from members of the same family are correlated.

In each of the examples given, we can think of the blocks as being a sample from some larger population. In the Frozen Yogurt Case Study, the participants form the blocks; we can imagine these college students are a sample from among all college students. In Example 25.1, the blocks are formed by the slope of the green; we can imagine these slopes form a sample from among all potential slopes a green could have. And, in Example 34.1, the blocks are created by families; we can imagine the families being a sample from among all families in the world, for example. Blocks are distinguished from secondary factors of interest because comparing the blocks themselves is not relevant to the question; instead, the blocks are a nuisance that result in correlated responses, and that correlation must be accounted for.

Example 34.2 (Greenhouse Study). A greenhouse study was conducted to determine the effects of two types of environmental conditions (amount of light and daily temperature) on the growth of a certain species of plant. Specifically, two amounts of light were considered: 10 hours and 12 hours. And, two temperatures were considered: 70 degrees and 80 degrees Fahrenheit. Researchers ultimately wanted to decide which of the four possible combinations of these conditions resulted in the maximum growth of the plant. A total of twelve plants were randomized to be grown under one of the four combinations of light and temperature (with three plants being randomized to each of the four combinations). After 10 weeks of exposure, the measured dry weight of the plant material for each of the twelve plants was recorded.

For Example 34.2, there are two grouping variables: amount of light and temperature. However, neither variable represents a block. In this example, researchers are particularly interested in comparing the specific levels of each variable. That is, both variables represent factors of interest. Instead of being a "nuisance," both the amount of light and the daily temperature are important variables in examining the response. While the method for addressing a second factor and addressing a block are similar in practice, conceptually, the a secondary factor is very different from a block.

Identifying a Block Variable

It can be helpful to ask two questions when determining if a categorical variable is a factor or a block:

- 1. Are we interested in comparing the levels of the variable to one another?
- 2. If we were to repeat the study, would the exact same levels of the variable be present in the study?

If we answer "yes" to the first question, the variable is almost certainly a factor. This first question is helpful, and is easy to grasp, but it is not fool-proof. It is possible that a variable is a factor, but that factor is not of primary interest in the study (but may be present to adjust for confounding, for example).

If we answer "no" to the second question, the variable is almost certainly a block. Since blocks typically represent a sample from a larger population, if we were to repeat the study, we generally would observed different blocks forming (different individuals, different families, etc.).

35 Presenting Correlated Data

Consider the graphical summary in Figure 35.1 of the data from the Frozen Yogurt Case Study. While this graphic was praised for comparing a quantitative response across multiple groups, when a study makes use of blocking, such graphics are no longer appropriate. In particular, nothing in the graphic suggests correlation among some observations.




Figure 35.1: Illustration of a poor graphic summarizing correlated data; the graphic hides the fact that there were repeated measures on the participants. It is not clear that the responses are in any way correlated.

Summarizing correlated data can be quite difficult. If there are only a few blocks, indicating which observations correspond to the same block using some aesthetic (color, size, shape, etc.) can be helpful. For example, Figure 35.2 uses color to distinguish responses from the same participant. The color allows you to see that one participant (represented by a blue color) does indeed tend to rate all yogurts highly compared to other participants. The color draws out the correlation structure.

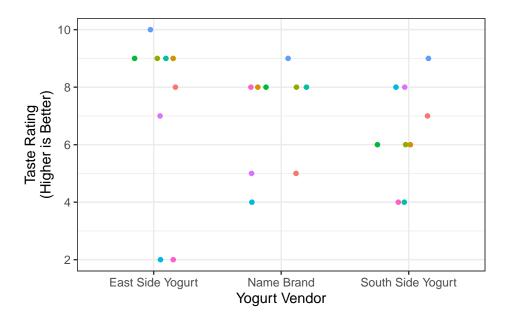


Figure 35.2: Results from a blind taste test comparing how participants rated the taste of yogurt from three different vendors. Ratings from the same participant are displayed using the same color.

Note

When color (or some other aesthetic) is used to denote observations from the same block, we often do not include a corresponding legend since the specific levels are not of interest. That is, which points correspond to Block 1 or Block 2 is not important; it is only important to know which points correspond to the same block.

Even with only nine blocks, it can be difficult to distinguish one participant's response from another in Figure 35.2. Another technique is to connect the responses from a single subject; this is illustrated in Figure 35.3. We note that within participants, there was not a universally preferred yogurt; however, most tended to prefer East Side to the Name Brand.

There is no universally adopted gold standard for summarizing correlated data. The key here is that the correlation in the data should not be ignored and should be illustrated in the summary while still addressing the primary question of interest.

We note that the correlation structure does not impact the average rating for each group. That is, we can still compute and report the average response within each group; however, the correlation structure is what helps us to properly visualize the variability when examining graphical summaries.



Figure 35.3: Results from a blind taste test comparing how participants rated the taste of yogurt from three different vendors. Ratings from the same participant are connected.

Big Idea

A good graphic should aid in partitioning the variability; with correlated responses, this includes indicating values which are related so that we can visually assess the variability between independent groups and within related groups.

36 Modeling Correlated Responses

Our question of interest in this unit is the same as that in our previous unit:

$$H_0: \theta_1 = \theta_2 = \dots = \theta_k$$
 vs. $H_1:$ At least one θ_j differs.

As this is the same question associated with an Analysis of Variance (ANOVA), it seems reasonable to begin with the model for describing a quantitative response as a function of a categorical predictor described in Chapter 27. In this chapter, we extend this model to account for the correlation between responses.

36.1 Statistical Model for Correlated Responses

For the Frozen Yogurt Case Study, we are comparing the average taste rating for different vendors. We might consider the following model for the data generating process (Equation 27.2) introduced in the previous unit:

$$(\text{Taste Rating})_i = \mu_1(\text{East Side})_i + \mu_2(\text{Name Brand})_i + \mu_3(\text{South Side})_i + \varepsilon_i$$

where

$$\begin{aligned} &(\text{East Side})_i = \begin{cases} 1 & \text{if i-th rating associated with east side yogurt vendor} \\ 0 & \text{otherwise} \end{cases} \\ &(\text{Name Brand})_i = \begin{cases} 1 & \text{if i-th rating associated with name brand yogurt vendor} \\ 0 & \text{otherwise} \end{cases} \\ &(\text{South Side})_i = \begin{cases} 1 & \text{if i-th rating associated with south side yogurt vendor} \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

are indicator variables to capture the various factor levels. In order to use this model, the first condition we imposed on the error term was that the error in the rating for one observation is independent of the error in the rating for all other observations. In fact, this condition is required to implement any form of inference (bootstrapping or an analytical approach).

However, for the Frozen Yogurt Case Study, we know this condition is violated. If the errors were independent of one another, it would imply the responses were independent of one another. But, since each participant rated each of the three vendors, the ratings from the same participant are related.

Consider the participant who loves frozen yogurt and tends to always give a higher rating than other participants. This individual would tend to give a higher than average rating regardless of the vendor. That is, the error (which represents the difference between an observed rating and the average rating for that corresponding vendor) for this participant's response for the Name Brand vendor would be a large positive value; however, the error for this participant's response for the East Side vendor would also be a large positive value. That is, knowing the error for one of the participant's responses would help us predict the error for another of their responses. This indicates a dependency. However, knowing this individual's error is large for one vendor tells us nothing about how the next participant's error term will behave.

Note

Violations of the independence condition can occur in clusters, which is what happens when blocks are present. Specifically, while observations from the same block are dependent on one another (correlated), observations from different blocks can remain independent.

At this point in the text, hopefully it is not a surprise that the way to address the correlated error terms is to partition the variability in the response further. Essentially, the blocking in the study informs us of another reason for the variation in the observed taste ratings: observations from the same participant will be similar. We want to tease this out of the variation in ratings among the same individual, and that is done by adding additional terms into the model for the data generating process.

For the Frozen Yogurt Case Study, consider the following model for the data generating process:

$$(\text{Taste Rating})_i = \mu_1(\text{East Side})_i + \mu_2(\text{Name Brand})_i + \mu_3(\text{South Side})_i \\ + \beta_2(\text{Participant 2})_i + \beta_3(\text{Participant 3})_i + \beta_4(\text{Participant 4})_i \\ + \beta_5(\text{Participant 5})_i + \beta_6(\text{Participant 6})_i + \beta_7(\text{Participant 7})_i \\ + \beta_8(\text{Participant 8})_i + \beta_9(\text{Participant 9})_i + \varepsilon_i$$
 (36.1)

where the indicators for the vendors were previously described and

$$(\text{Participant j})_i = \begin{cases} 1 & \text{i-th observation taken from Participant j} \\ 0 & \text{otherwise} \end{cases}$$

is an indicator of whether the observation comes from a particular participant. In this model, the β parameters capture the "bump" in each participant's ratings that is due to the participant's inherent feeling towards frozen yogurt. That is, every observation that is associated with the same participant will share this "bump," capturing the similarity between observations from the same participant.

It may at first appear as if we forgot the indicator for Participant 1; however, it is not needed. Just as with any model, it is often easiest to see what is happening by thinking about the form of the model under specific cases. How do we describe observations (remember there is more than one) for Participant 2? The above model for the data generating process states that the average rating for the East Side vendor from Participant 2 is given by $\mu_1 + \beta_2$. Similarly, the average rating for the East Side vendor from Participant 6 is given by $\mu_1 + \beta_6$. What about Participant 1? Well, Participant 1 would have a 0 for every "participant indicator" variable in the model; therefore, the above model states that the average rating for the East Side vendor from Participant 1 is simply μ_1 .

This affects how we interpret our parameters. In our model μ_1 is no longer the average rating given to East Side Yogurt; it is the average rating given to East Side Yogurt by the first participant. It is the same concept as the "reference group" (see Definition 21.2) discussed in Chapter 21.

Note

If there are b blocks, we need only include b-1 indicator variables and corresponding parameters in the model for the data generating process in order to capture all the blocks. The remaining block is the "reference group" and is captured by the parameters comparing the factor levels under study.

This may seem like it affects our questions of interest. After all, the hypothesis

$$H_0: \mu_1=\mu_2=\mu_3$$

says that the "average taste rating for Participant 1 is the same for all vendors" instead of the "average taste rating across all individuals is the same for all vendors." The latter is the hypothesis we want to test, but we have the parameters specified in terms of the first participant only. This "problem" resolves once we recognize an inherent assumption of our model. Notice the difference between the average ratings for the East Side vendor and Name Brand vendor for Participant 1 is

$$\mu_1 - \mu_2$$
.

And, notice the *difference* between the average ratings for the East Side vendor and the Name Brand vendor for Participant 2 is

$$(\mu_1 + \beta_2) - (\mu_2 + \beta_2) = \mu_1 - \mu_2.$$

In fact, since the "bump" for every observation from Participant j is always β_j , when comparing averages across vendors, these "bumps" cancel out. Therefore, if the mean response for one Participant is the same for all vendors, then it must be that the mean response across vendors is the same for all Participants (see Appendix B)! In context, this means that all individuals must share the same preferences for frozen yogurt vendors. This is a feature of the model, and we will discuss this in the next chapter.

•

Big Idea

The model we introduce for blocking assumes that any difference between the levels of a factor is similar across all blocks.

The model for the data generating process we have been discussing essentially says there are three reasons that the taste ratings differ from one observation to another:

- 1. Ratings applied to different vendors may differ,
- 2. Ratings from different individuals for the same vendor may differ, and
- 3. Even within the same individual, ratings for cups of yogurt from the same vendor may differ due to unexplained variability.

In general, this type of model, often described as a "Repeated Measures ANOVA" model, partitions the variability in the response into three general categories: differences between groups, differences between blocks, differences within blocks.

Repeated Measures ANOVA Model

For a quantitative response and a single categorical predictor (also known as a factor) with k levels in the presence of b blocks, the repeated measures ANOVA model is

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i \tag{36.2}$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{i-th unit belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{i-th unit belongs to block } m \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables capturing whether a unit belongs to the j-th group and m-th block, respectively; and, $\mu_1, \mu_2, \dots, \mu_k$ and $\beta_2, \beta_3, \dots, \beta_b$ are the parameters governing the

model for the data generating process.

This model assumes any differences between groups are similar across all blocks.

In the past, the stochastic portion of the model ε captured the subject-to-subject variability. It no longer has the same role in this case. It now captures the variability in observations within the same block. That is, it captures the fact that if we repeatedly taste the same yogurt, we might rate it differently each time because of our mood or some other external factor that we have not captured. The subject-to-subject variability is captured by the β parameters in the model.

•

Big Idea

In a model without repeated measures (blocks), the error term captures the subject-to-subject variability. In a model with repeated measures, the error term captures the variability between observations within the same block.

There is something else that is unique about the repeated measures ANOVA model. We do not really care about all the parameters in the model. Our question of interest is based on the parameters μ_1, μ_2, μ_3 . We would never be interested in testing something of the form

$$H_0: \beta_2 = \beta_3 \qquad \text{vs.} \qquad H_1: \beta_2 \neq \beta_3$$

as this would be comparing Participant 2 to Participant 3. Such a comparison (does Participant 2 have different yogurt ratings from Participant 3) is not useful. Said another way, we did not put the parameters β_2, \ldots, β_9 into the model because they helped us address a particular research objective; instead, we put them in the model because they captured the observed relationship in the responses. This is the difference between factors and blocks.

Note

The statistical theory underlying models which generalize the repeated measures ANOVA model make use of the terms "fixed effect" and "random effect" instead of factor and blocks. These more technical terms allow the model to generalize to a host of situations not covered by the repeated measures ANOVA model. For our purposes, however, it is sufficient to differentiate between factors and blocks.

Consider applying the questions listed at the end of Chapter 34 for distinguishing between a factor and a block. Notice that if we were to repeat the study, we would use the same three vendors, since they are a fundamental part of the question. However, we would not need to use the same participants in the sample; we would be satisfied with any random sample from the population. So, the values "East Side Yogurt," "South Side Yogurt," and "Name Brand" (at least, the three vendors these represent) are of specific interest. However, we do not care about

"Participant 2" and "Participant 3." These can be any two individuals from the population. Therefore, for the Frozen Yogurt Case Study, the vendor is the factor of interest, while the participant is the block term.

Big Idea

The parameters and corresponding indicator variables capturing the blocking are placed in the model for the data generating process to account for the correlation between responses.

(Optional) Comparison of Repeated Measures ANOVA to General Linear Regression Model

Notice that Equation 36.2 has a very similar form to Equation 21.3. The primary difference is the presence of an intercept term in Equation 21.3. Each indicator is acting as a predictor in the model for the data generating process. If we were to add an intercept to Equation 21.3, and then remove one of the indicator variables used to distinguish the factor of interest, then we would completely fall under the general linear regression model framework. This means that as we move forward, we can adopt results established for the general linear regression model.

The parameters in Equation 36.2 can be estimated using the method of least squares (Definition 17.2). We must keep in mind that these parameters do not correspond directly to the average response observed in each group. As a result, the sample mean from each group is often reported as well.

Just as before, while point estimates are helpful, inference requires that we quantify the variability in our estimates. And, just as before, we need to distinguish between the model for the data generating process and the model for the sampling distribution of the parameter estimates and the model for the null distribution of a standardized statistic. And, just as before, to move from the model for the data generating process to a model for the sampling distribution of the parameter estimates, we impose conditions on the stochastic component.

37 Conditions on the Error Term of the Repeated Measures ANOVA Model

In the previous chapter we developed a general model for the data generating process of a quantitative response as a function of a categorical predictor in the presence of blocking:

$$(\text{Response})_i = \sum_{i=1}^k \mu_k(\text{Group } j)_i + \sum_{m=2}^b (\text{Block } m)_i + \varepsilon_i,$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if the i-th observation is from Group } j \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{if the i-th observation is from Block } m \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables to capture the level of the factor and the block to which each observation belongs.

We also discussed a common method for estimating the parameters of this model from a sample — the method of least squares. However, if we are to construct a model for the sampling distribution of these estimates we must add some structure to the stochastic component ε in the model. In this chapter, we focus on the most common conditions we might impose and how those conditions impact the model for the sampling and null distributions (and therefore the computation of a confidence interval or p-value).

37.1 Conditions on the Error Distribution

In our model for the data generating process, we incorporated a component ε to capture the noise within each block. Since the error is a random variable (stochastic element), we know it has a distribution. We typically assume a certain structure to this distribution.

37.1.1 Correctly Specified Model

The first condition we consider is the most important. It states that for every value of the predictor, the average error is 0. We have actually already mentioned this condition in another form — the inherent assumption we make with the structure of our model. Our model states that any differences in the average response across the levels of the factor are similar across all blocks. This comes from the structure of our model. Therefore, this is equivalent to saying that the deterministic portion of our model for the data generating process is correctly specified. This is the same as the condition we introduced in Chapter 18.

Mean-0 Condition

The mean-0 condition states that the treatment differences are similar across all blocks. Even though we state this as a condition on the error terms, it is equivalent to say that the "the block effect is similar for all observations across treatment groups."

Note

The mean-0 condition can be relaxed through the inclusion of an interaction term (see Definition 21.5), but this is beyond the scope of the text.

37.1.2 Independent Errors

The second condition we consider is that the noise attributed to one observed response for an individual is independent of the noise attributed to the observed response for any other individual. That is, the amount of error in any one observation is unrelated to the error in any other observations. This is the same condition we encountered in Chapter 10, Chapter 18, and Chapter 28.

i Independence Condition

The independence condition states that the error in one observation is independent (see Definition 10.3) of the error in all other observations.

With just these first two conditions, we can use a bootstrap algorithm in order to model the sampling distribution of the least squares estimates of our parameters (see Appendix A). However, additional conditions are often considered.

The idea of assuming independence may seem counterintuitive; this entire unit exists because we felt there was a correlation among the responses. However, this condition is stating that once we account for the correlation induced by the blocks through the incorporation of the block terms in the model for the data generating process, the remaining noise is now independent. We

essentially partitioned out the correlated component, and what remains is now just independent noise.

37.1.3 Same Degree of Precision

The third condition that is typically placed on the distribution of the errors is that the variability of the errors is the same for all combination of the predictors. Again, we encountered this condition in Chapter 18 and Chapter 28. As our model includes both the group comparisons of interest as well as the block terms, violating this condition happens if the response is more precise for one group/block combination than another.

i Constant Variance

Also called homoskedasticity, the constant variance condition states that the variability of the errors within each group is the same across all groups.

With this additional condition imposed, we are able to modify our bootstrap algorithm when constructing a model for the sampling distribution of the least squares estimates.

37.1.4 Specific Form of the Error Distribution

The fourth condition that is typically placed on the distribution of the errors is that the errors follow a Normal distribution, as discussed in Chapter 18. Here, we are assuming a particular structure on the distribution of the error population.

i Normality

The normality condition states that the distribution of the errors follows the functional form of a Normal distribution (Definition 18.2).

Let's think about what this condition means for the responses. Given the shape of the Normal distribution, imposing this condition (in addition to the other conditions) implies that some errors are positive and some are negative. This in turn implies that some responses within a block will be above average for their group, and some responses will be below average for their group. More, because the distribution of the response within each block and group is just a shifted version of the distribution of the errors, we know that the distribution of the response variable itself follows a Normal distribution within a particular block and group. While this is similar to the argument made in Chapter 28 in the context of the ANOVA model, it is often difficult to visualize for a repeated measures ANOVA. Recall that blocking is typically used to increase the power of a study; as a result, it is common that only one observation exists within each group and block combination. This makes using the responses directly to visualize

the shape of the distribution impossible. This is similar then to the linear regression model scenario discussed in Chapter 18.

With this last condition imposed, we can construct an analytic model for the sampling distribution of the least squares estimates. As in regression modeling, we are not required to impose all four conditions in order to obtain a model for the sampling distribution of the estimates. Historically, however, all four conditions have been routinely imposed in the scientific and engineering literature.

37.2 Conditions on the Block Effects

The blocking terms, and the associated β parameters, in the model for the data generating process were used to partition the variability in the response to account for the correlation among responses from the same block. In the previous chapter, we highlighted this when we stated that the β parameters really capture the subject-to-subject variability. Further, remember we stated that blocks should be viewed as a sample from some larger population. As a result, the blocking parameters represent a sample from a population; so, they have a distribution which must be constrained.

Note

The presence of this distribution is why blocks are referred to as "random effects" in the statistical literature when discussing a more general approach to addressing correlated responses.

The easiest way to discuss additional conditions on the blocking parameters is to think about each β as a "bump" attributed to that block. Think about a participant who is not a fan of frozen yogurt; then, regardless of which vendor the yogurt they are tasting originated from, the participant's taste rating will tend to "bump" down compared to others.

37.2.1 Independent Block Effects

The first condition we consider is that the "bump" for one participant is unrelated to the "bump" for any other participant. Practically, one person's taste for frozen yogurt is unaffected by the taste for frozen yogurt of anyone else.

We also impose the condition that the "bump" for a participant is unrelated to the amount of error in the response for that participant. That is, the error term must be independent of the blocking term.

i Independence Condition

The independence condition among the blocks states that the blocks are independent of one another and that the blocks are independent of the error term.

37.2.2 Specific Form of the Distribution

The last condition that is typically placed on the distribution of the "bumps" is that the magnitude of these "bumps" follows a Normal distribution, as discussed in Chapter 18. Here, we are assuming a particular structure on the distribution of the blocks.

i Normality

The normality condition states that the distribution of the block impacts follows the functional form of a Normal distribution (Definition 18.2).

Note

The conditions on the block effects are much more technical than those placed on the error term. The statistical theory for such models is beyond the scope of this text, but they impact the model for the sampling distribution of the estimates in a similar way as the conditions on the error term do.

37.3 Classical Repeated Measures ANOVA Model

We have discussed several conditions we could place on the stochastic portion of the data generating process. Placing all conditions on the error term and blocking effects is what we refer to as the "Classical Repeated Measures ANOVA Model."

Definition 37.1 (Classical Repeated Measures ANOVA Model). For a quantitative response and single categorical predictor with k levels in the presence of b blocks, the classical repeated measures ANOVA model assumes the following data generating process:

$$(\text{Response})_i = \sum_{i=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if i-th observation belongs to group j} \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{if i-th observation belongs to block m} \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables and where

- 1. The error in the response for one subject is independent of the error in the response for all other subjects.
- 2. The variability in the error of the response is the same across all predictors.
- 3. The errors follow a Normal distribution.
- 4. Any differences between the groups are similar across all blocks. This results from the deterministic portion of the model for the data generating process being correctly specified and is equivalent to saying the error in the response, on average, takes a value of 0 for all predictors.
- 5. The effect of a block on the response is independent of the effect of any other block on the response.
- 6. The effect of a block on the response is independent of the error in the response for all subjects.
- 7. The block effects follow a Normal distribution.

This is the default "repeated measures ANOVA" analysis implemented in the majority of statistical packages.



Warning

A "hidden" (typically unstated but should not be ignored) condition is that the sample is representative of the underlying population. In the one sample case (Chapter 10), we referred to this as the errors being "identically distributed." We no longer use the "identically distributed" language for technical reasons; however, we still require that the sample be representative of the underlying population.

We note that "repeated measures ANOVA" need not require all four conditions on the error distribution imposed in Definition 37.1. Placing all four conditions on the error term results in a specific analytical model for the sampling distribution of the least squares estimates. Changing the conditions changes the way we model the sampling distribution.



Big Idea

The model for the sampling distribution of a statistic is determined by the conditions you place on the stochastic portion of the model for the data generating process.

37.4 Imposing the Conditions

Let's return to our model for the yogurt taste ratings as a function of the vendor while accounting for the correlation induced due to the repeated measures across participants given in Equation 36.1:

```
\begin{split} (\text{Taste Rating})_i &= \mu_1(\text{East Side})_i + \mu_2(\text{Name Brand})_i + \mu_3(\text{South Side})_i \\ &+ \beta_2(\text{Participant 2})_i + \beta_3(\text{Participant 3})_i + \beta_4(\text{Participant 4})_i \\ &+ \beta_5(\text{Participant 5})_i + \beta_6(\text{Participant 6})_i + \beta_7(\text{Participant 7})_i \\ &+ \beta_8(\text{Participant 8})_i + \beta_9(\text{Participant 9})_i + \varepsilon_i, \end{split}
```

where we use the same indicator variables defined in Chapter 36. We were interested in the following research question:

Does the average taste rating differ for at least one of the three yogurt vendors?

This was captured by the following hypotheses:

```
\begin{split} H_0: \mu_1 = \mu_2 = \mu_3 \\ H_1: \text{at least one } \mu_j \text{ differs.} \end{split}
```

Using the method of least squares, we constructed point estimates of the parameters in the model. If we are willing to assume the data is consistent with the conditions for the classical repeated measures ANOVA model, we are able to model the sampling distribution of these estimates and therefore construct confidence intervals. Table 37.1 summarizes the results of fitting the model described in Equation 36.1 using the data available from the Frozen Yogurt Case Study. In addition to the least squares estimates, it also contains the standard error (see Definition 6.4) of each statistic, quantifying the variability in the estimates. Finally, there is a 95% confidence interval for each parameter.

We note that while these parameter estimates are somewhat interesting, none of them address our question directly, and none of them estimate the overall average taste rating for a particular vendor. We must remember that the first three estimates in Table 37.1 are really estimating the average rating for only the first participant.

Note

Often the parameter estimates in the repeated measures block design are not of interest.

Table 37.1: Estimated parameters in a model for the taste ratings of yogurt from three vendors using data from a randomized complete block design with 9 blocks.

Term	Estimate	Standard Error	Lower 95% CI	Upper 95% CI
East Side Yogurt	7.000	1.358	4.121	9.879
Name Brand	6.778	1.358	3.899	9.657
South Side Yogurt	6.222	1.358	3.343	9.101
Participant 2	1.000	1.737	-2.683	4.683
Participant 3	1.000	1.737	-2.683	4.683
Participant 4	1.000	1.737	-2.683	4.683
Participant 5	0.333	1.737	-3.350	4.016
Participant 6	-2.000	1.737	-5.683	1.683
Participant 7	2.667	1.737	-1.016	6.350
Participant 8	0.000	1.737	-3.683	3.683
Participant 9	-2.000	1.737	-5.683	1.683

37.5 Recap

We have covered a lot of ground in this chapter, and it is worth taking a moment to summarize the big ideas. In order to compare the mean response in each group in the presence of blocking, we took a step back and modeled the data generating process. Such a model consists of two components: a deterministic component explaining the response as a function of the predictor and the blocks, and a stochastic component capturing the noise in the system.

Certain conditions are placed on the distribution of the noise in our model as well as on the distribution of the block effects. With a full set of conditions (classical repeated measures ANOVA model), we are able to model the sampling distribution of the least squares estimates analytically. We can also construct an empirical model for the sampling distribution of the least squares estimates assuming the data is consistent with fewer conditions.

In general, the more conditions we are willing to impose on the data generating process, the more tractable the analysis; however, the most important aspect is that the data come from a process which is consistent with the conditions we impose, which is discussed in Chapter 39.

38 Quantifying the Evidence

In the previous two chapters, we described a model for describing the data generating process for a quantitative response as a function of a single categorical predictor in the presense of blocking:

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if i-th observation corresponds to group } j \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{if i-th observation corresponds to block } m \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables.

Chapter 36 discussed obtaining estimates of these unknown parameters using the method of least squares (which turned out not to be incredibly helpful). Chapter 37 imposed conditions on the stochastic portion of the model and the block effects in order to develop a confidence interval for each parameter (again, not incredibly helpful). In this chapter, we turn to performing inference through the computation of a p-value for a set of hypotheses. Following the developments in Chapter 19 and Chapter 29, this is accomplished through partitioning variability.

Let's return to our model for the yogurt taste rating as a function of the vendor while accounting for the correlation induced by the repeated measures across participants given in Equation 36.1:

```
\begin{split} (\text{Taste Rating})_i &= \mu_1(\text{East Side})_i + \mu_2(\text{Name Brand})_i + \mu_3(\text{South Side})_i \\ &+ \beta_2(\text{Participant 2})_i + \beta_3(\text{Participant 3})_i + \beta_4(\text{Participant 4})_i \\ &+ \beta_5(\text{Participant 5})_i + \beta_6(\text{Participant 6})_i + \beta_7(\text{Participant 7})_i \\ &+ \beta_8(\text{Participant 8})_i + \beta_9(\text{Participant 9})_i + \varepsilon_i \end{split}
```

where we use the same indicator variables defined in Chapter 36. We were interested in the following research question:

Does the average taste rating differ for at least one of the three yogurt vendors?

This was captured by the following hypotheses:

```
\begin{split} H_0: \mu_1 = \mu_2 = \mu_3 \\ H_1: \text{at least one } \mu_j \text{ differs.} \end{split}
```

Recall that hypothesis testing is really about model comparison. The model for the data generating process described above (and in Equation 36.1) places no constraints on the parameters; this corresponds to the alternative hypothesis. Let's refer to this as Model 1. The null hypothesis places constraints on the parameters; namely, the mean response should be the same across all three groups. Enforcing the constraint suggested by the null hypothesis (with μ representing the common unspecified value), we have a different model for the data generating process, call it Model 0:

$$(\text{Taste Rating})_i = \mu + \beta_2 (\text{Participant 2})_i + \beta_3 (\text{Participant 3})_i + \beta_4 (\text{Participant 4})_i + \beta_5 (\text{Participant 5})_i + \beta_6 (\text{Participant 6})_i + \beta_7 (\text{Participant 7})_i + \beta_8 (\text{Participant 8})_i + \beta_9 (\text{Participant 9})_i + \varepsilon_i$$

where again the blocking terms capture the correlation among the responses.

So, it turns out this set of hypotheses we are considering is more complex than we might first imagine. The null hypothesis suggests replacing the three different terms associated with vendor in the model for the data generating process with a single common parameter while ignoring any parameters associated with the blocking. Just as in Chapter 29, the key to testing this hypothesis is to partition out the sources of variability.

Both Model 1 and Model 0 acknowledge that not all of the variability in the response is fully explained (both models have an error term). As we saw in Chapter 19, we can define the error sum of squares for a model as

$$SSE = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Predicted Response})_i \right]^2$$

where in our case the predicted response is obtained using the parameter estimates obtained by the method of least squares. Specifically, for Model 1 above, the predicted response is given by

$$\begin{split} (\text{Predicted Response})_i &= \hat{\mu}_1(\text{East Side})_i + \hat{\mu}_2(\text{Name Brand})_i + \hat{\mu}_3(\text{South Side})_i \\ &+ \hat{\beta}_2(\text{Participant 2})_i + \hat{\beta}_3(\text{Participant 3})_i + \hat{\beta}_4(\text{Participant 4})_i \\ &+ \hat{\beta}_5(\text{Participant 5})_i + \hat{\beta}_6(\text{Participant 6})_i + \hat{\beta}_7(\text{Participant 7})_i \\ &+ \hat{\beta}_8(\text{Participant 8})_i + \hat{\beta}_9(\text{Participant 9})_i \end{split}$$

where the "hats" denote the use of the corresponding least squares estimate. If the null hypothesis is true, we would expect the amount of unexplained variability in Model 0 to be the same as the amount of unexplained variability in Model 1 $(SSE_0 \approx SSE_1)$. That is, if the model under the null hypothesis is sufficient for explaining the variability in the yogurt test ratings, then it should perform as well as the full unconstrained model. If, however, Model 1 explains more of the variability in the yogurt taste ratings (therefore leaving less variability unexplained) than Model 0, then this would indicate the null hypothesis is false. That is, if incorporating the vendor is informative, then we would expect the amount of unexplained variability in Model 1 to be less than that of Model 0. This suggests that a metric for quantifying the signal in the data is given by

$$SSE_0 - SSE_1$$
,

where we use the subscript to denote the model from which the sum of squares was computed. As in Chapter 29, working with variance terms is easier than working with sums of squares analytically. And, we need to consider the size of the signal relative to the amount of background noise. This leads to a form of our standardized statistic:

$$\frac{(SSE_0 - SSE_1)/(k-1)}{SSE_1/(n-k-b+1)},$$

where again we have added subscripts to emphasize from which model we are computing the above sums of squares. In previous chapters, the degrees of freedom by which we divided were perhaps a bit more intuitive than they are here; so, we take a moment to discuss these further.

Notice that Model 1 has k+b-1 parameters (in our case, 3 parameters to capture the vendors and 9-1=8 parameters to capture the blocks), and Model 0 has b parameters (in our case, 1 parameter to capture the overall average for the first participant and 9-1=8 parameters to capture the blocks). Therefore, Model 1 is using an additional

$$(k+b-1)-(b) = k-1$$

parameters to explain the variability in the response. This gives the degrees of freedom used in the numerator. Now, in the full model, the predicted values require we estimate k + b - 1 parameters using n observations; therefore, the degrees of freedom are

$$n - (k + b - 1) = n - k - b + 1$$

for estimating the error. This gives the degrees of freedom used in the denominator.

Definition 38.1 (Standardized Statistic for Repeated Measures ANOVA). Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 36.2):

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if i-th observation corresponds to group } j \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{if i-th observation corresponds to block } m \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables. Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0. A standardized statistic, sometimes called the "standardized F statistic," for testing the hypotheses is given by

$$T^* = \frac{\left(SSE_0 - SSE_1\right)/(k+b-1-r)}{SSE_1/(n-k-b+1)},$$

where k + b - 1 is the number of parameters in the full unconstrained model and r is the number of parameters in the reduced model. Defining

$$MSA = \frac{SSE_0 - SSE_1}{k + b - 1 - r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model. Just as before, the MSE represents the residual variance — the variance in the response for a particular set of predictors.

It can be shown that this standardized statistic is a special case of the one discussed in Chapter 21 and is therefore consistent with the ones presented in Chapter 19, Chapter 12, and Chapter 29. The numerator captures the signal by examining the difference between what we expect the error sum of squares to be under the null hypothesis and what we actually observe; the denominator captures the background noise (relative to the estimated mean response from the full model). Larger values of this standardized statistic indicate more evidence in the sample against the null hypothesis.

We should not lose sight of the fact that our standardized statistic is really a result of partitioning the variability and considering the variability explained by adding the factor of interest to the blocks, relative to the noise in the response. Underscoring that the standardized statistic is a result of this partitioning, the analyses of these sources of variability is often summarized in an ANOVA table.

Note

Occasionally, an ANOVA table for repeated measures ANOVA will partition the variability in terms of the variability due to the factor, the variability due to the blocks, and the error variability. We have instead presented the ANOVA table as corresponding to a specific set of hypotheses.

Big Idea

Comparing the difference in the unexplained variability between two models allows us to assess if the data is inconsistent with the simpler of the two models.

Just as before, given data, we can compute not only the standardized statistic but a corresponding p-value. As with any p-value, it is computed by finding the likelihood, assuming the null hypothesis is true, of getting, by chance alone, a standardized statistic as extreme or more so than that observed in our sample. "More extreme" values of the statistic would be larger values; so, the area under the null distribution to the right of the observed statistic is the p-value. Of course, the model for the null distribution is developed under the conditions we place on the stochastic portion and the block effects in the model for the data generating process.

Let's return to the question that inspired our investigation in this chapter:

Table 38.1: Analysis of the sources of variability in the yogurt taste scores as a function of the vendors while accounting for the repeated measures on the participants.

Term	DF	Sum of Squares	Mean Square	Standardized Statistic	P-Value
Vendor	2	2.889	1.444	0.319	0.731
Error	16	72.444	4.528		

Does the average taste rating differ for at least one of the three yogurt vendors?

This was captured by the following hypotheses:

$$\begin{split} H_0: \mu_1 = \mu_2 = \mu_3 \\ H_1: \text{at least one } \mu_j \text{ differs.} \end{split}$$

Table 38.1 gives the ANOVA table summarizing the partitioned sources of variability in the yogurt taste ratings. We have a large p-value (computed assuming the data is consistent with the classical repeated measures ANOVA model). That is, the sample provides no evidence to suggest the average yogurt taste rating differs across any of the vendors. While we cannot prove that the average ratings are the same for all vendors, the this study is consistent with the three yogurt vendors having the same taste ratings, on average. This is in line with the idea that consumers were paying a premium for the experience of going to a yogurt-shop; they viewed the product similarly with what could be purchased at a local grocery retailer.



Big Idea

Determining if a response is related to a categorical predictor in the presence of blocking is done by determining if the factor explains a significant portion of the variability in the response above and beyond the blocks.

In this chapter, we partitioned variability as a way of evaluating the strength of evidence the factor of interest plays in determining the response. As with the ANOVA model, partitioning the variability is a key step. By partitioning the variability in the response, we are able to construct a standardized statistic for testing the hypothesis of interest. The model for the null distribution of this statistic depends upon the conditions we are willing to impose on the stochastic portion of the data generating process. Regardless of the conditions we impose, we can interpret the resulting p-value similarly. It provides an indication of whether the data suggests that the average response differs for at least one of the groups.

Of course, the interpretation of the p-value depends on the conditions we impose. We should not choose such conditions without performing some type of assessment to ensure those conditions are reasonable — that the data is consistent with the conditions. That is the focus of the next chapter.

39 Assessing the Modeling Conditions in Repeated Measures ANOVA

In this unit we have discussed a model relating a quantitative response to a categorical predictor in the presence of blocks which induce correlation among the responses. For the Frozen Yogurt Case Study, our model had the form

```
\begin{split} (\text{Taste Rating})_i &= \mu_1(\text{East Side})_i + \mu_2(\text{Name Brand})_i + \mu_3(\text{South Side})_i \\ &+ \beta_2(\text{Participant 2})_i + \beta_3(\text{Participant 3})_i + \beta_4(\text{Participant 4})_i \\ &+ \beta_5(\text{Participant 5})_i + \beta_6(\text{Participant 6})_i + \beta_7(\text{Participant 7})_i \\ &+ \beta_8(\text{Participant 8})_i + \beta_9(\text{Participant 9})_i + \varepsilon_i \end{split}
```

where we use the same indicator variables defined in Chapter 36. Further, we considered several conditions on the distribution of the error term and the distribution of the block effects:

- 1. The error in the taste ratings within one individual is independent of the error in the taste ratings within any other individual.
- 2. The variability in the error in the taste ratings is the same for all vendor and participant combinations.
- 3. The error in the taste ratings follows a Normal distribution.
- 4. The deterministic portion of the model is correctly specified; that is, any differences in the ratings across vendors is the same for all participants.
- 5. One participant's preferences relative to the population is independent of any other participant's preferences; that is, the block effects are independent of one another.
- 6. Each participant's preferences relative to the population is independent of the error in the taste ratings of all individuals; that is, the block effects are independent of the error terms.
- 7. Participants' preferences follow a Normal distribution; that is, the block effects follow a Normal distribution.

While we imposed each of these conditions in Chapter 38, we could have developed an empirical model for the null distribution of the standardized statistic imposing only a subset of these conditions. Unfortunately, we cannot simply state conditions and then proceed blindly. In order to rely on the p-values and confidence intervals produced from any modeling procedure, the data must be consistent with the conditions imposed.

In this section, we discuss how we assess these conditions qualitatively. We note that the last three conditions on the block effects are not easily assessed. These are generally stated and assumed.

Warning

The conditions on the block effects in a repeated measures ANOVA model are not easily assessed and tend to be stated and assumed in practice.

Just as we saw in Chapter 20 and Chapter 30, however, the first four conditions placed on the error distribution can be assessed using **residuals**.

39.1 Residual

The difference between the observed response and the predicted response (estimated deterministic portion of the model). Specifically, the residual for the *i*-th observation is given by

$$(Residual)_i = (Response)_i - (Predicted Mean Response)_i$$

where the "predicted mean response" is often called the predicted, or fitted, value.

Residuals mimic the noise in the data generating process.

For the repeated measures ANOVA model, the predicted mean response is determined by using the least squares estimates from the model:

$$(\text{Predicted Mean Response})_i = \sum_{j=1}^k \hat{\mu}_j (\text{Group } j)_i + \sum_{m=2}^b \hat{\beta}_m (\text{Block } m)_i.$$



Big Idea

The conditions are placed on the error term, but they are assessed with residuals.

39.2 Assessing the Independence Condition

The error in the taste ratings within one individual is independent of the error in the taste ratings within any other individual.

Generally, independence is assessed by considering the method in which the data was collected and considering the context with a discipline expert. By carefully considering the manner in which the data was collected, we can typically determine whether it is reasonable that the errors in the response are independent of one another. When a study design incorporates blocking, it was because researchers had identified a source of correlation prior to conducting the study. Further, the researchers are generally willing to assume that observations within the block are independent of one another and the units in different blocks are independent of one another. That is, the study design considered the possible sources of correlation and accounted for them.

While the incorporation of the blocks into the model for the data generating process should eliminate any correlation in the errors due to observations being from the same block, it is possible that other forms of correlation exist. For example, if the data is collected over time, it is possible that observations collected in proximity to one another exhibit a relationship; and, such a relationship would not be captured by the block terms. When we know the order in which the data was collected, we can assess whether the data tends to deviate from the condition of independence over time. This is done graphically through a **time-series plot** of the *residuals*. If two errors were unrelated, then the value of one residual should tell us nothing about the value of the next residual. Therefore, a plot of the residuals over time should look like noise (since residuals are supposed to mimic the noise in the model). If there are any trends, then it suggests the data is not consistent with independence.

39.3 Time-Series Plot

A time-series plot of a variable is a line plot with the variable on the y-axis and time on the x-axis.

Graphically Assessing the Independence Condition

If the data is consistent with the independence condition, we would not expect to see a trend in the *location* or *spread* in a time-series plot of the residuals. Note that this graphic can only be used if the order in which the data was collected is known, and the order indicates some natural timing. Otherwise, we generally assume the inclusion of the block terms captured the primary source of correlation among responses.

For the Frozen Yogurt Case Study, participants were assessed simultaneously within the class. Therefore, there is no ordering in time to be concerned about. As such, a time-series plot of the residuals would not be useful here. Considering the context, the students were the ones who had designed the study (it was carried out by their instructor); therefore, they were concerned about the quality of the data collection. The students did their best to not influence the ratings of any other participant. It is reasonable to assume the data is consistent with the

independence condition once the repeated measures have been accounted for by the inclusion of the block terms in the model.

39.4 Assessing Homoskedasticity

The variability in the error in the taste ratings is the same for all vendor and participant combinations.

We want the variability of the errors to be the same across all values of the predictors — that is what is being captured with the phrase "all vendor and participant combinations." As a result, unlike Chapter 30, we cannot simply make a boxplot of the residuals across the three vendors. Once there is more than one variable in the deterministic portion of the model for the data generating process, we must resort back to the strategy discussed in Chapter 20. That is, we construct a plot of the residuals against the predicted values.

Warning

Even though our goal is to compare the mean response across groups as in ANOVA, since a repeated measures ANOVA model contains multiple variables in the deterministic portion of the model for the data generating process, we cannot assess the constant variance condition in the same way.

Graphically Assessing the Constant Variance Condition

If the data is consistent with the constant variance condition, there should be no trends in the spread of the residuals when plotted against the predicted mean response (the fitted values).

Figure 39.1 shows the residuals for each individual across the predicted taste rating for that individual observation. Notice that as the predicted taste ratings increase, the variability in the residuals tends to decrease. This is inconsistent with what we would expect if the variability in the errors was constant. That is, if the variability of the errors was constant, we would not expect the spread of the residuals to exhibit this "fan" shape. Therefore, we do not feel it is reasonable to impose the constant variance condition.

39.5 Assessing Normality

The error in the taste ratings follows a Normal distribution.




Figure 39.1: Plot of the residuals against the predicted values from a model comparing the taste ratings across yogurt vendors accounting for subject-variability.

If the errors follow a Normal distribution, then we would expect the residuals to mimic a sample taken from a Normal distribution. As introduced in Chapter 20, we emphasize the Normal **probability plot** for assessing the Normality condition.

39.6 Probability Plot

Also called a "Quantile-Quantile Plot", a probability plot is a graphic for comparing the distribution of an observed sample with a theoretical probability model for the distribution of the underlying population. The quantiles observed in the sample are plotted against those expected under the theoretical model.

i Graphically Assessing the Normality Condition

If the data is consistent with the normality condition, the Normal probability plot of the residuals should exhibit a straight line with any deviations appearing to be random. Systemic deviations from a straight line indicate the observed distribution does not align with the proposed model.

Figure 39.2 shows the probability plot for the residuals from the Frozen Yogurt Case Study.

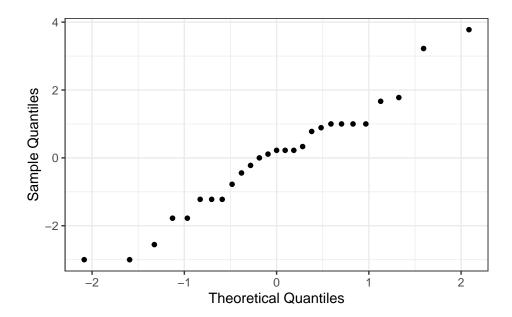


Figure 39.2: Probability plot of the residuals for the Frozen Yogurt Case Study. If the errors follow a Normal distribution, we would expect the residuals to fall along a straight line.

Overall, the points do tend to follow a straight line. It seems reasonable the data is consistent with the errors following a Normal distribution.

39.7 Assessing Whether the Model is Correctly Specified

The deterministic portion of the model is correctly specified; that is, any differences in the ratings across vendors is the same for all participants.

Recall that the *structure* of our repeated measures ANOVA model suggests that any differences between groups are the same across all blocks. This is analogous to assuming a specific functional form for the relationship between two predictors (for example, linear or sinusoidal). Because we are discussing the structure of the deterministic portion of the model for the data generating process, occasionally, we can assess this condition from the context of the problem with the help of a discipline expert. For example, do we believe it is reasonable to believe that every individual will have the same vendor preference for frozen yogurt? Our personal experience alone tells us one person's favorite restaurant is not necessarily everyone's favorite restaurant; so, there is some reason already to doubt this condition is reasonable for the Frozen Yogurt Case Study.

When we (or the discipline experts) do not have enough knowledge to confidently assess this condition, we can examine a graphic of the residuals and predicted values. Recall that in Chapter 18 and Chapter 20, we argued that the deterministic portion of the model being correctly specified was the result of the errors having a value of 0, on average. That is, the mean-0 condition is synonymous with requiring the deterministic portion of the model for the data generating process to have the correctly specified structure.

If the errors have a mean of 0 for all combination of variables in the deterministic portion of the model, then we would expect the residuals to have a mean of 0 for all predicted values. That is, if the data is consistent with the mean-0 condition, then as we move left to right across a plot of the residuals and predicted values, the residuals should tend to balance out at 0 everywhere along the x-axis. Any trends in the *location* of this graphic would indicate the data is *not* consistent with the mean-0 condition.

i Graphically Assessing the Mean-0 Condition

If the data is consistent with the mean-0 condition, there should be no trends in the *location* of the plot of the residuals against the predicted values.

As we examine Figure 39.1, the residuals do tend to balance out at 0 everywhere along the x-axis. That is, we do not see a trend in the location of the residuals as the predicted values increase. Therefore, it is reasonable to say the sample is consistent with the mean-0 condition.

The sample seems to be at odds with our personal experience. Remember, we cannot use the graphic to confirm the condition is met; we can only say the data is *consistent* with the behavior we would expect of the residuals if the the condition were true. It is probably the case that any differences in vendor preferences across individuals is so slight that the structure we have imposed is reasonable in this sample.

39.8 General Tips for Assessing Assumptions

First discussed in Chapter 20, we want to remember four things that should be kept in mind when assessing conditions:

- 1. We should not spend an extraordinary amount of time examining any one residual plot; we might convince ourselves of patterns that do not exist. We are looking for major deviations from our expectations.
- 2. We can never prove a condition is satisfied; we can only determine whether the data is consistent with a condition or whether it is not consistent with a condition.
- 3. Any condition required for a particular analysis should be assessed.
- 4. Transparency is crucial.

In this chapter, we add to these four tips that conditions placed on the block effects cannot be graphically assessed.

40 Using the Tools Together

This unit introduced a framework for determining if there is an association between a quantitative response and a categorical predictor in the presence of blocking. We formed a standardized statistic for measuring the signal, and developed a model for the data generating process incorporating the block effect which allowed us to model the null distribution of the standardized statistic. In this chapter, we pull these tools together once more to answer a research question.

For the Frozen Yogurt Case Study, we found no evidence that college students rated the taste differently, on average, for any of the vendors. However, the students conducting the study also recorded how the participants rated the texture and appearance of the yogurt. In this chapter, we will use the data to answer the following question:

Is there evidence that the appearance of the yogurt differs across vendors?

40.1 Framing the Question (Fundamental Idea I)

As stated, the above question is ill-posed. We have not identified a parameter of interest. We refine this question to be

Is there evidence that, on average, the appearance rating differs for at least one of the vendors?

This question could also be stated as the following set of hypotheses:

Let θ_1 , θ_2 and θ_3 represent the average appearance rating (on a scale of 1-10) of vanilla yogurt from each of the three vendors (East Side, Name Brand, and South Side), respectively.

 $H_0:\theta_1=\theta_2=\theta_3$

 H_1 : At least one θ_i differs

40.2 Getting Good Data (Fundamental Idea II)

As we are working with previously collected data, our goal in this discussion is not how best to collect the data but making note of the limitations of the data as a result of how it was collected. As before, each participant sampled yogurt from each of the three vendors, creating natural blocks. That is, each participant forms a unique block. Since it is quite reasonable that appearance preferences vary substantially between individuals, forming blocks out of the participants should allow us to increase the power of the study because we are accounting for a substantial source of variability in the appearance ratings.

The study was was a controlled experiment since the order in which the samples from each vendor were presented to the participants was randomized. While the study made use of random allocation, it did not make use of random selection. The participants were students taking a particular course; however, it may be reasonable to assume they are representative of college students in the area. The sample size was also limited as only students in this course were included in the study.

If you compare the above paragraph to the corresponding section in Chapter 31, you might be confused because in that section, we described that changing the question resulted in the study no longer being a controlled experiment. However, here, we changed the question of interest and retained the fact that the study was a controlled experiment. The question of whether a study is a controlled experiment is always in regard to whether random allocation was used. In general, if you change the response but keep the primary factor of interest unchanged, the study will remain a controlled experiment. If you change the factor under study, it will become an observational study.

40.3 Presenting the Data (Fundamental Idea III)

Our question here is examining the relationship between a quantitative response (appearance rating) and a categorical predictor (vendor) in the presence of blocks (participants). Figure 40.1 compares the distribution of the appearance rating for the three vendors.

Based on the above graphic, there appears to be less variability among the appearance for the South Side Yogurt vendor, but participants tend to rate the appearance similarly across all three vendors.

40.4 Quantifying the Variability in the Estimate (Fundamental Idea IV)

In order to measure the size of the signal, we partition the variability in an ANOVA table, which allows us to compute a standardized statistic. In order to partition the variability, we



Figure 40.1: Comparison of the appearance ratings for yogurt from three vendors. Color is used to distinguish ratings from the same participant.

first consider the following model for the data generating process:

$$\begin{split} (\text{Appearance Rating})_i &= \mu_1(\text{East Side})_i + \mu_2(\text{Name Brand})_i + \mu_3(\text{South Side})_i \\ &+ \beta_2(\text{Participant 2})_i + \beta_3(\text{Participant 3})_i + \beta_4(\text{Participant 4})_i \\ &+ \beta_5(\text{Participant 5})_i + \beta_6(\text{Participant 6})_i + \beta_7(\text{Participant 7})_i \\ &+ \beta_8(\text{Participant 8})_i + \beta_9(\text{Participant 9})_i + \varepsilon_i \end{split}$$

where

$$(\text{East Side})_i = \begin{cases} 1 & \text{if i-th rating associated with east side yogurt vendor} \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Name Brand})_i = \begin{cases} 1 & \text{if i-th rating associated with name brand yogurt vendor} \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{South Side})_i = \begin{cases} 1 & \text{if i-th rating associated with south side yogurt vendor} \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Participant } m)_i = \begin{cases} 1 & \text{i-th observation taken from Participant } m \\ 0 & \text{otherwise} \end{cases}$$

Table 40.1: ANOVA table summarizing the comparison of the appearance ratings of frozen yogurt across three vendors from the Frozen Yogurt Case Study.

Term	DF	Sum of Squares	Mean Square	Standardized Statistic
Vendor	2	3.185	1.593	1.055
Error	16	24.148	1.509	

are appropriately defined indicator variables. We note here that μ_1 is not the same as θ_1 defined earlier; however, testing

$$H_0: \mu_1 = \mu_2 = \mu_3$$

is equivalent to testing

$$H_0: \theta_1 = \theta_2 = \theta_3$$

as discussed in Chapter 36 (also see Appendix B). Table 40.1 reports the standardized statistic from our study corresponding to testing the hypotheses

$$H_0: \mu_1 = \mu_2 = \mu_3 \qquad \text{vs.} \qquad H_1: \text{at least one μ_j differs.}$$

Of course, if we were to collect a new sample, we would expect our standardized statistic to change. If we want to conduct inference and determine the strength of evidence in this study, we need a model for its null distribution. In order to construct a model for the null distribution of the standardized statistic, we need to place appropriate conditions on the error term, which can include:

- 1. The error in the appearance ratings within one individual is independent of the error in the appearance ratings within any other individual.
- 2. The variability in the error in the appearance ratings is the same for all vendor and participant combinations.
- 3. The error in the appearance ratings follows a Normal distribution.
- 4. The deterministic portion of the model is correctly specified; that is, any differences in the ratings across vendors is the same for all participants.
- 5. One participant's preferences relative to the population is independent of any other participant's preferences; that is, the block effects are independent of one another.
- 6. Each participant's preferences relative to the population is independent of the error in the appearance ratings of all individuals; that is, the block effects are independent of the error terms.
- 7. Participants' preferences follow a Normal distribution; that is, the block effects follow a Normal distribution.

Before creating a model for the null distribution and computing a p-value, we need to assess whether the data is consistent with these assumptions. This requires examining the residuals from the model. First, we discuss the assumption of independence. Since the data was collected at a single point in time, known as a *cross-sectional study*, constructing a time-series plot of the residuals would not provide any information regarding this assumption. Instead, we rely on the context of the problem to make some statements regarding whether the data is consistent with this condition (whether making this assumption is reasonable). It is reasonable that the errors are independent. The primary source of correlation was the repeated measures on each participant, which has been accounted for with the incorporation of the block terms. Additionally, as students were interested in preserving the integrity of the data, they did not influence one another during data collection. We feel it is reasonable to assume independence.

Again, note that there is a condition of independence; we are simply saying whether we are willing to assume the condition is satisfied. There is no way to ensure the condition holds.

In order to assess the constant variance condition, let us examine a plot of the residuals against the predicted values. Figure 40.2 shows a plot of the residuals against the predicted values for our model; the plot does not exhibit any trends in the spread of the residuals as we move across the fitted values. So, it seems reasonable to impose the constant variance condition.




Figure 40.2: Plot of the residuals against the predicted values from a model comparing the appearance ratings across yogurt vendors accounting for subject-variability.

To assess the Normality condition, we consider a Normal probability plot of the residuals (Figure 40.3). Given that the residuals tend to display a linear relationship, it is reasonable that the errors follow a Normal Distribution.

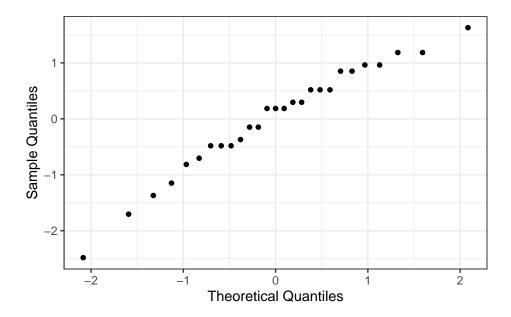


Figure 40.3: Normal probability plot assessing the assumption that the errors for our model comparing the appearance ratings across vendores follow a Normal distribution.

Looking at Figure 40.2, the residuals tend to balance around 0 across all values of the x-axis. This suggests the data is consistent with the deterministic portion of our model for the data generating process being correctly specified. That is, the sample is consistent with the belief that any differences in appearance across vendors is similar for all individuals. Thinking about the context, this seems reasonable; we don't have any reason to believe that preferences in the appearance of frozen yogurt might differ wildly across individuals.

As is common in practice, while we are not able to assess the conditions placed on the block effects, we are willing to assume them. Given the discussion above, we are comfortable assuming the data is consistent with all conditions from the classical repeated measures ANOVA model (Definition 37.1); therefore, we can make use of an analytical model for the null distribution.

40.5 Quantifying the Evidence (Fundamental Idea V)

Now that we have a model for the null distribution of our standardized statistic, we can determine how extreme our particular sample was by comparing the standardized statistic for our sample with this model for the null distribution (Figure 40.4).

Based on the results, the study does not provide any evidence that college students rate the appearance of the vanilla frozen yogurt different, on average, for any of the three vendors. That is, the sample is consistent with the appearance of the yogurt being similar, on average, across




Figure 40.4: Analytical model of the null distribution of the standardized statistic computed from the Frozen Yogurt Case Study comparing the appearance rating of frozen yogurt across three vendors.

all three vendors. We note that while frozen yogurt shops start with a base of vanilla, their attraction is generally the ability to customize your order ("mix-ins," unique flavors, etc.). This study only examined vanilla yogurt because it is readily available from both commercial retailers as well as local grocery retailers.

40.6 Conclusion

Throughout this unit, we have examined a framework for examining the association between a quantitative response and a categorical predictor in the presence of blocking. This reinforces a couple of big ideas we have seen throughout this text:

- The key to measuring a signal is to partition the variability in the response.
- A standardized statistic is a numeric measure of the signal strength in the sample.
- Modeling the data generating process provides us a way of modeling the sampling distribution of the parameter estimates and the null distribution of a standardized statistic when combined with conditions on the stochastic portion of the model for the data generating process.
- Before imposing conditions on the stochastic portion of a data generating process, we should graphically assess whether the data is consistent with these conditions.

References

- Box, George E P. 1979. "Robustness in the Strategy of Scientific Model Building." In *Robustness in Statistics*, edited by R L Launer and G N Wilkinson, 201–36. Academic Press.
- Chihara, Laura, and Tim Hesterberg. 2011. Mathematical Statistics with Resampling and r. Wiley.
- Doyle, Sir Arthur Conan. 1890. The Sign of the Four. Spencer Blackett.
- Dudeck, A E, and C H Peeacock. 1981. "Effects of Several Overseeded Ryegrasses on Turf Quality, Traffic Tolerance and Ball Roll." In *Proceedings of the Fourth International Turfgrass Research Conference*, edited by R W Sheard, 75–81.
- Eskine, Kendall J. 2013. "Wholesome Foods and Wholesome Morals? Organic Foods Reduce Prosocial Behavior and Harshen Moral Judgments." Social Psychological and Personality Science 4 (2): 251–54.
- Goldstein, Bernard D, Howard J Osofsky, and Maureen Y Lichtveld. 2011. "The Gulf Oil Spill." The New England Journal of Medicine 364: 1334–48. https://doi.org/10.1056/NEJMra1007197.
- Johnson, Eric J, and Daniel Goldstein. 2003. "Do Defaults Save Lives?" Science 302: 1338–39. Koutrakis, S I, G P Karakaisis, P M Hatzidimitriou, P K Koliopoulos, and V N Margaris. 2002. "Seismic Hazard in Greece Based on Different Strong Ground Motion Parameters." Journal of Earthquake Engineering 6 (1): 75–109. https://doi.org/10.1080/13632460209350411.
- Lee, J. 1992. "Relationships Between Properties of Pulp-Fibre and Paper."
- Moery, Eileen, and Robert J Calin-Jageman. 2016. "Direct and Conceptual Replications of Eskine (2013): Organic Food Exposure Has Little to No Effect on Moral Judgments and Prosocial Behavior." Social Psychological and Personality Science 7 (4): 312–19.
- Tintle, Nathan, Beth L Chance, A J Rossman, S Roy, T Swanson, and J VanderStoep. 2015. *Introduction to Statistical Investigations*. Wiley.

A Approaches for Modeling Sampling Distributions

Throughout the text, we have alluded to various strategies for modeling the sampling distribution of a statistic. Here, we sketch out these methods in the context of a regression model. As modeling the data generating process provides a unifying framework, the approaches discussed in this appendix generalize to each of the models discussed in this text.

Throughout this appendix, we will assume the simple linear regression model

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i$$

is appropriate for describing the data generating process. Further, all methods we discuss require that we impose the mean-0 and independence conditions. Specifically, we assume

- the average error is 0 for all values of the predictor, and
- the error in one observation is independent of the error for all other observations.

Additional constraints will be discussed as needed. We further assume we have a sample of size n from the population.

We note that all methods described in this section extend directly to the general linear regression model (Equation 21.3) and therefore extend to all models presented in the text.

A.1 Residual Bootstrap

There are several bootstrap algorithms; a very efficient and foundational algorithm for regression models is the "residual bootstrap." In addition to the mean-0 and independence conditions, it also requires the constant variance condition.

For the above simple linear regression model, we begin by obtaining the least squares estimates $\hat{\beta}_0$ and $\hat{\beta}_1$ of the parameters. The residual bootstrap proceeds according to the following steps:

1. Compute the residuals

$$(Residuals)_i = (Response)_i - (Predicted Response)_i$$

- 2. Take a random sample of size n (with replacement) of the residuals; call these values $e_1^*, e_2^*, \dots, e_n^*$.
- 3. Form "new" responses $y_1^*, y_2^*, \dots, y_n^*$ according to the formula

$$y_i^* = \hat{\beta}_0 + \hat{\beta}_1(\text{Predictor})_i + e_i^*.$$

4. Obtain the least squares estimates $\hat{\alpha}_0$ and $\hat{\alpha}_1$ by finding the values of α_0 and α_1 that minimize

$$\sum_{i=1}^{n} \left(y_{i}^{*} - \alpha_{0} - \alpha_{1}(\operatorname{Predictor})_{i}\right)^{2},$$

the result of fitting a regression model of the "new" responses with the predictor from the original sample.

5. Repeat steps 2-4 many (say 5000) times.

Each pass through the algorithm, we retain the least squares estimates from the bootstrap resample, the $\hat{\alpha}_0$ and $\hat{\alpha}_1$. The distribution of these estimates across the resamples is a good empirical model for the sampling distribution of the least squares estimates.

A.1.1 Bootstrap for Hypothesis Testing

The above discussion describes the use of bootstrapping in order to model the sampling distribution of our parameter estimates. If our goal is to model the null distribution of a standardized statistic, we have an additional step. Remember, the null distribution is a sampling distribution when the null hypothesis is true.

Consider testing the hypotheses

$$H_0: \beta_1 = 0$$

$$H_0: \beta_1 \neq 0$$

Then, under the null hypothesis, the model for the data generating process is given by

$$(Response)_i = \gamma_0 + \varepsilon_i.$$

Therefore, in Step 3 of the above algorithm, we generate new data assuming the null hypothesis is true by using the formula

$$y_i^* = \hat{\gamma}_0 + e_i^*$$

where $\hat{\gamma}_0$ is the least squares estimate from fitting the reduced model to the original sample. Notice that this model generates "new" responses that are not dependent upon the predictor. Therefore, when this data is used to fit the model in Step 4 of the above algorithm, we are fitting data under the null hypothesis.

Instead of retaining the parameter estimates at each iteration of the algorithm, we compute and retain the standardized statistic. The distribution of these standardized statistics across the resamples is a good empirical model for the null distribution of the standardized statistic. This algorithm updates that proposed above to ensure that the generation step makes use of the null hypothesis.

A.2 Wild Bootstrap

In the previous section, we introduced the residual bootstrap; the algorithm there required three conditions be imposed. In this section, we discuss an alternate bootstrap algorithm, "the wild bootstrap," which only requires the mean-0 and independence conditions.

Before discussing the algorithm, we note that the wild bootstrap is not technically necessary. The version of bootstrapping illustrated in the text, known as "case-resampling," where we resample observations directly, only requires the mean-0 and independence conditions. However, the performance of the case-resampling algorithm can be quite poor in some settings, particularly when the sample size is small.

The wild bootstrap is an alteration of the residual bootstrap which removes the need to impose the constant-variance condition. For the above simple linear regression model, we begin by obtaining the least squares estimates $\hat{\beta}_0$ and $\hat{\beta}_1$ of the parameters. The wild bootstrap proceeds according to the following steps:

1. Compute the residuals

$$(Residuals)_i = (Response)_i - (Predicted Response)_i$$

2. Construct new pseudo-residuals by multiplying each residual by a random variable U_i with mean 0 and variance 1, for example a sample from a Normal distribution with mean 0 and variance 1:

$$e_i^* = U_i(\text{Residual})_i$$
.

3. Form "new" responses $y_1^*, y_2^*, \dots, y_n^*$ according to the formula

$$y_i^* = \hat{\beta}_0 + \hat{\beta}_1(\text{Predictor})_i + e_i^*.$$

4. Obtain the least squares estimates $\hat{\alpha}_0$ and $\hat{\alpha}_1$ by finding the values of α_0 and α_1 that minimize

$$\sum_{i=1}^{n} \left(y_i^* - \alpha_0 - \alpha_1(\operatorname{Predictor})_i\right)^2,$$

the result of fitting a regression model of the "new" responses with the predictor from the original sample.

5. Repeat steps 2-4 many (say 5000) times.

Each pass through the algorithm, we retain the least squares estimates from the bootstrap resample, the $\hat{\alpha}_0$ and $\hat{\alpha}_1$. The distribution of these estimates across the resamples is a good empirical model for the sampling distribution of the least squares estimates.

If we are interested in using a wild bootstrap to model the null distribution of the standardized statistic, we make the same adjustments to the above algorithm that we did for the residual bootstrap algorithm.

A.3 Classical Theory

In general, classical theory comes from making an additional assumption about the functional form of the distribution of the error terms. In doing so, we are able to rely on statistical theory (or probability theory) in order to determine an analytical model for the sampling distribution. Before stating our result of interest, we first introduce a new analytical model. In Chapter 18, we introduced the idea of specifying the functional form of the density of a distribution, and we gave the Normal distribution (Definition 18.2) as an example. There are countless analytical models we might posit for a distribution; we are particularly interested in two such distributions.

The location-scale t-distribution, also known as the non-standardized t-distribution, is a probability model that has the following functional form:

$$f(x) = \frac{\Gamma\left(\frac{\nu+1}{2}\right)}{\Gamma\left(\frac{\nu}{2}\right)\tau\sqrt{\pi\nu}}\left(1 + \frac{1}{\nu}\left(\frac{x-\theta}{\tau}\right)^2\right)^{-(\nu+1)/2}$$

where ν , τ , and θ are parameters that govern the shape of the distribution, and $\Gamma(\cdot)$ represents the Gamma function¹. The parameter θ represents the mean of the distribution; ν represents the "degrees of freedom," which impacts the shape and spread of the distribution, and τ is the scale parameter that governs the spread of the distribution.

The F-distribution is a probability model that has the following functional form:

$$f(x) = \frac{\Gamma\left(\frac{\nu_1}{2} + \frac{\nu_2}{2}\right)}{\Gamma\left(\frac{\nu_1}{2}\right)\Gamma\left(\frac{\nu_2}{2}\right)} \left(\frac{\nu_1}{\nu_2}\right)^{\nu_1/2} x^{(\nu_1/2)-1} \left(1 + \frac{\nu_1}{\nu_2}x\right)^{-(\nu_1+\nu_2)/2}$$

where ν_1 and ν_2 are parameters that govern the shape of the distribution, and $\Gamma(\cdot)$ represents the Gamma function. The parameter ν_1 represents the "numerator degrees of freedom," and ν_2 represents the "denominator degrees of freedom." These parameters work together to determine the spread of the distribution; the mean of this distribution is essentially 1 for nearly any value of ν_2 seen in practice.

The location-scale t-distribution and F-distribution are graphically illustrated in Figure A.1 for various choices of their parameters.

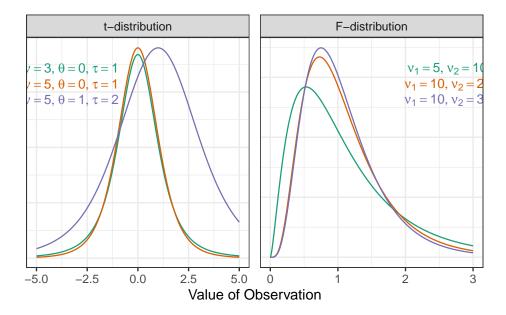


Figure A.1: Illustration of the location-scale t-distribution and F-distribution for various choices of their parameters.

We are now ready to present our two primary results. Suppose that in addition to the mean-0 and independence conditions, we are willing to assume the constant variance condition and

¹The gamma function is an extension of a factorial: https://mathworld.wolfram.com/GammaFunction.html

the Normality condition; that is, we are willing to assume the Classical Regression Model (Definition 18.3). Then, we have that the least squares estimates can be modeled by the location-scale t-distribution with degrees of freedom $\nu=n-2$, mean $\theta=\hat{\beta}_j$, and scale $\tau=\hat{\eta}_j\sqrt{\frac{n-4}{n-2}}$ where $\hat{\eta}_j$ represents the standard error of the estimate $\hat{\beta}_j$. And, we have that the null distribution of the standardized statistic in Definition 19.7 can be modeled by the F-distribution with numerator degrees of freedom $\nu_1=1$ and denominator degrees of freedom $\nu_2=n-2$.

A.4 Asymptotic Theory

Suppose that in addition to the mean-0 and independence conditions, we are willing to assume the constant variance condition. Note that we are *not* willing to assume the Normality condition; that is, we are leaving the specific form of the error distribution unspecified. Without this fourth condition, we are unable to rely on the classical theory developed above. However, it turns out that the classical theory still holds under this more relaxed condition provided the sample size is large. This stems from a result in probability known as the Central Limit Theorem.

Suppose we have a random sample (the observations are independent and identically distributed) of size n from a population with mean μ and variance σ^2 . Then, as the sample size increases, the sampling distribution of the quantity

$$\frac{\sqrt{n}\left(\bar{y}-\mu\right)}{\sigma}$$

can be modeled by a Normal distribution with a mean of 0 and a standard deviation of 1, where \bar{y} represents the sample mean. The term "the" in "the Central Limit Theorem" is misleading; there are actually several different Central Limit Theorems, though the above is the one most commonly presented in a traditional introductory statistics course.

It turns out there is a similar result for the least squares estimates. In particular, if we have data consistent with the model for the simple linear regression model for which the data is consistent with the mean-0 condition, the independence condition, and the constant variance condition, then as the sample size increases, the sampling distribution of the quantity

$$\frac{\hat{\beta}_j - \beta_j}{\hat{\eta}_j}$$

can be modeled by a Normal distribution with a mean of 0 and a standard deviation of 1, where $\hat{\beta}_j$ represents the least squares estimate of the parameter β_j and $\hat{\eta}_j$ represents the standard

error of $\hat{\beta}_j$. These results are known as asymptotic theory because they rely on allowing the sample size to approach infinity.

These results provide nice analytical models for the sampling distributions of the least squares estimates; however, we are left asking how large the sample size needs to be in order to rely on these results. There is no specific criteria.

It can be shown that a t-distribution (with mean 0 and scale parameter 1) is essentially equivalent to a Normal distribution (with mean 0 and standard deviation 1) when the degrees of freedom are large. Specifically, once the degrees of freedom exceed 30, the two distributions are nearly identical. Therefore, many advocate that provided the sample size is large enough so that the degrees of freedom associated with the mean square for error exceeds 30, we can rely on the results suggested by the classical theory even if we are unable to assume the distribution of the error follows a Normal distribution.

B Mathematical Results for Standardized Statistics

Here, we present some mathematical results associated with standardized statistics. The material provided in this appendix is not necessary for understanding the concepts presented in the text. It is provided only for those who are curious about the mathematics underlying some of the claims made throughout.

B.1 Single Mean Response, Same Signal

In Chapter 12, we argued that the level of background noise can make it difficult to detect a signal. That is illustrated in the following theorem.

Theorem B.1. Consider two samples of the same size n. Suppose the sample mean is equivalent in both samples, and consider testing the hypotheses

$$H_0: \mu = \mu_0 \qquad \textit{vs.} \qquad H_1: \mu \neq \mu_0$$

with each sample. The difference in the sum of squares

$$SS_0 - SS_1 = \sum_{i=1}^n \left(y_i - \mu_0\right)^2 - \sum_{i=1}^n \left(y_i - \bar{y}\right)^2$$

will be the same for each sample.

Proof. Observe that the difference in the sums of squares can be written as

$$\begin{split} SS_0 - SS_1 &= \sum_{i=1}^n \left(y_i - \mu_0\right)^2 - \sum_{i=1}^n \left(y_i - \bar{y}\right)^2 \\ &= \sum_{i=1}^n \left[\left(y_i - \mu_0\right)^2 - \left(y_i - \bar{y}\right)^2 \right] \\ &= \sum_{i=1}^n \left[y_i^2 - 2y_i \mu_0 + \mu_0^2 - y_i^2 + 2y_i \bar{y} - \bar{y}^2 \right] \\ &= \sum_{i=1}^n \left[\mu_0^2 - 2y_i \left(\mu_0 - \bar{y}\right) - \bar{y}^2 \right] \\ &= \sum_{i=1}^n \mu_0^2 - 2\left(\mu_0 - \bar{y}\right) \sum_{i=1}^n y_i - \sum_{i=1}^n \bar{y}^2 \\ &= n\mu_0^2 - 2n\mu_0 \bar{y} + 2n\bar{y}^2 - n\bar{y}^2 \\ &= n\mu_0^2 - 2n\mu_0 \bar{y} + n\bar{y}^2 \\ &= n\left(\mu_0 - \bar{y}\right)^2. \end{split}$$

That is, the difference in the sum of squares is a function only of the null value and the sample mean. Therefore, if two samples of the same size have the same sample mean, then the difference in the sums of squares will be equivalent for the two samples. \Box

B.2 Single Mean Response, Two Standardized Statistics

In Chapter 12, we presented the standardized statistic

$$T_1^* = \frac{SS_0 - SS_1}{s^2}.$$

However, another popular standardized statistic for testing hypotheses for a single mean is

$$T_2^* = \frac{\sqrt{n} \left(\bar{y} - \mu_0 \right)}{s}.$$

However, these two standardized statistics are related.

Theorem B.2. Consider a sample of size n collected to test the following hypotheses:

$$H_0: \mu = \mu_0 \qquad \textit{vs.} \qquad H_1: \mu \neq \mu_0.$$

Define

$$T_{1}^{*}=\frac{SS_{0}-SS_{1}}{s^{2}},$$

where

$$SS_0 = \sum_{i=1}^n \left(y_i - \mu_0\right)^2$$

$$SS_1 = \sum_{i=1}^n \left(y_i - \bar{y} \right)^2.$$

And, define

$$T_{2}^{*} = \frac{\sqrt{n}\left(\bar{y} - \mu_{0}\right)}{s}$$

where s is the sample standard deviation. Then, $T_1^* = \left(T_2^*\right)^2$.

Proof. From the proof of Theorem B.1, we have that

$$SS_0 - SS_1 = n \left(\mu_0 - \bar{y} \right)^2$$
.

We can then rewrite T_1^* as

$$T_1^* = \frac{n \left(\mu_0 - \bar{y}\right)^2}{s^2}.$$

This form of T_1^* then makes the result clear.

B.3 Least Squares Estimate for Intercept-Only Model

In Chapter 19, we state that the least squares estimate for the parameter μ in a model of the data generating process that has the form

$$(Response)_i = \mu + \varepsilon_i$$

is the sample mean response. While this is the intuitive estimate for μ we presented in Chapter 10, we had not yet related that to the method of least squares introduced in Chapter 17.

Theorem B.3. Consider a sample of size n; suppose we posit the following model for the data generating process:

$$y_i = \mu + \varepsilon_i$$

where y_i represents the i-th observed value of the response variable and μ is the sole parameter. The least squares estimate of μ is \bar{y} , the sample mean.

Proof. Observe that, by definition (see Definition 17.2), the least squares estimate for this model is the value of μ which minimizes the quantity

$$\sum_{i=1}^{n} \left(y_i - \mu \right)^2.$$

Observe that if we add and subtract the sample mean \bar{y} inside the squared quantity, we can rewrite the above quantity as

$$\begin{split} \sum_{i=1}^{n} \left(y_{i} - \mu\right)^{2} &= \sum_{i=1}^{n} \left(y_{i} - \bar{y} + \bar{y} - \mu\right)^{2} \\ &= \sum_{i=1}^{n} \left[\left(y_{i} - \bar{y}\right)^{2} + 2\left(y_{i} - \bar{y}\right)\left(\bar{y} - \mu\right) + \left(\bar{y} - \mu\right)^{2} \right] \\ &= \sum_{i=1}^{n} \left(y_{i} - \bar{y}\right)^{2} + 2\left(\bar{y} - \mu\right) \sum_{i=1}^{n} \left(y_{i} - \bar{y}\right) + \sum_{i=1}^{n} \left(\bar{y} - \mu\right)^{2} \\ &= \sum_{i=1}^{n} \left(y_{i} - \bar{y}\right)^{2} + 2\left(\bar{y} - \mu\right)\left(n\bar{y} - n\bar{y}\right) + n\left(\bar{y} - \mu\right)^{2}, \end{split}$$

where we make use of the fact that $\bar{y} = \frac{1}{n} \sum_{i=1}^{n} y_i$ implies that $\sum_{i=1}^{n} y_i = n\bar{y}$. Now, we notice that the second term in the above expansion resolves to 0. Therefore, we have that

$$\sum_{i=1}^{n}\left(y_{i}-\mu\right)^{2}=\sum_{i=1}^{n}\left(y_{i}-\bar{y}\right)^{2}+n\left(\bar{y}-\mu\right)^{2}.$$

Notice that the second term is a squared difference; therefore, the second term must always be non-negative. Therefore, the second term can only make the sum larger; further, the second term will reach 0, the smallest value possible, when $\mu = \bar{y}$. Therefore, the sum is minimized when we take $\mu = \bar{y}$. This gives that \bar{y} minimizes the quantity of interest.

This theorem allows us to conclude that the error sum of squares for a model of the form

$$(Response)_i = \mu + \varepsilon_i$$

is given by

$$\sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Overall Mean Response}) \right]^2,$$

where the overall mean response is the sample mean response observed.

B.4 Class of Hypotheses that are Testable

Chapter 19 developed a general standardized statistic for testing hypotheses in a linear model described by Equation 17.3. Chapter 21 extended this standardized statistic to all linear models. This standardized statistic allows us to have a unifying testing framework across the text. However, not all hypotheses can be tested using this standardized statistic. This framework applies to linear hypotheses. For example, consider the model (Equation 17.3):

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i.$$

For this model, we can use the standardized statistic in Definition 19.7 to test any null hypothesis of the form

$$H_0: c_0\beta_0 = c_2, \ c_1\beta_1 = c_3$$

or

$$H_0: c_0\beta_0 + c_1\beta_1 = c_2,$$

where c_0, c_1, c_2, c_3 are known constants. That is, we can test a hypothesis of the form

$$H_0: \beta_1 = 0$$
 vs. $H_1: \beta_1 \neq 0$

because this can be written by setting $c_1 = 1$ and $c_0 = c_2 = c_3 = 0$ in the first form of the hypothesis given above. We can also test a hypothesis of the form

$$H_0: \beta_1 = 0, \beta_0 = 3$$
 vs. $H_1: \text{At least one } \beta_i \text{ differs}$

which corresponds to a reduce model of

$$(Response)_i = 3 + \varepsilon_i$$

because this can be written by setting $c_0 = c_1 = 1$, $c_2 = 3$, and $c_3 = 0$ in the first form of the hypothesis given above.

However, we cannot test a hypothesis of the form $\beta_0\beta_1=1$; while this may be a valid hypothesis, it cannot be written in either of the above forms described above. We would need additional theory to test such a hypothesis.

B.5 Equivalence of Hypotheses with Blocking

Our model for the data generating process comparing a quantitative response across the k levels of a factor in the presence of b blocks is

$$(\text{Response})_i = \sum_{i=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i$$

where (Group j) and (Block m) are appropriately defined indicator variables. Without loss of generality, consider the case where k=2 and b=2. Then, our model reduces to

$$(Response)_i = \mu_1(Group \ 1)_i + \mu_2(Group \ 2)_i + \beta_2(Block \ 2)_i + \varepsilon_i.$$

We argue heuristically in Chapter 36 that testing the hypotheses

$$H_0: \mu_1 = \mu_2$$
 vs. $H_1: \mu_1 \neq \mu_2$

which are hypotheses about the average response in Block 1 is equivalent to testing hypotheses about the overall average response across all blocks. To establish this more formally, note that based on the above model for the data generating process, the overall average response for Group 1, call it θ_1 , is given by

$$\theta_1 = \frac{1}{2}\mu_1 + \frac{1}{2}\left(\mu_1 + \beta_2\right) = \mu_1 + \frac{\beta_2}{2},$$

where we average the average response from Block 1 with that of Block 2. Similarly, the average response for Group 2, call it θ_2 , is given by

$$\theta_2 = \frac{1}{2}\mu_2 + \frac{1}{2}\left(\mu_2 + \beta_2\right) = \mu_2 + \frac{\beta_2}{2}.$$

Now, note that $\theta_1 - \theta_2 = \mu_1 - \mu_2$; therefore, $\theta_1 = \theta_2$ if and only if $\mu_1 = \mu_2$. That is, testing whether the overall average response across all individuals is equal across groups is equivalent to testing whether the average response within Block 1 is equal across groups.

C Glossary

C.1 Five Fundamental Ideas of Inference

This text revolves around five fundamental ideas of inference. These were introduced in the text, and they are provided here for quick reference along with a link to where they were first introduced.

Fundamental Idea I (Chapter 3)

A research question can often be framed in terms of a parameter that characterizes the population. Framing the question should then guide our analysis.

Fundamental Idea II (Chapter 4)

If data is to be useful for making conclusions about the population, a process referred to as drawing inference, proper data collection is crucial. Randomization can play an important role ensuring a sample is representative and that inferential conclusions are appropriate.

Fundamental Idea III (Chapter 5)

The use of data for decision making requires that the data be summarized and presented in ways that address the question of interest and represent the variability present.

Fundamental Idea IV (Chapter 6)

Variability is inherent in any process, and as a result, our estimates are subject to sampling variability. However, these estimates often vary across samples in a predictable way; that is, they have a distribution that can be modeled.

Fundamental Idea V (Chapter 7)

With a model for the distribution of a statistic under a proposed model, we can quantify the the likelihood of an observed sample under that proposed model. This allows us to draw conclusions about the corresponding parameter, and therefore the population, of interest.

C.2 Distributional Quartet

This text refers to what we call the "distributional quartet" — the four key distributions that are central to nearly any analysis. These are introduced early in the text and are

- The distribution of the population; this characterizes the pattern of variability of a variable across individual units in the population. While this is not directly observed, we sometimes posit a model for this distribution.
- The distribution of the sample; this characterizes the pattern of variability of a variable across individual units in the sample. This is what we summarize (graphically/numerically) using the available data.
- The sampling distribution of the statistic; this characterizes the pattern of variability of a statistic across repeated samples. While this is not directly observed, we model it by applying conditions on the stochastic portion of the model for the data generating process.
- The null distribution of a (often standardized) statistic; this is the sampling distribution of a statistic when a specified null hypothesis is enforced.

C.3 Models for the Data Generating Process

This text takes a modeling approach to inference. The following models are introduced in the text; each model is presented with a link to where the model was fully defined in the text.

Data Generating Process for Single Mean Response (Equation 10.2)

In general, given a quantitative response variable and no predictors, our model for the data generating process is

$$(Response)_i = \mu + \epsilon_i$$

where μ represents the average response in the population, the parameter of interest.

Simple Linear Regression Model (Equation 17.3)

For a quantitative response and a quantitative predictor, the general form of the simple linear regression model is

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i$$

where β_0 and β_1 are parameters governing the model for the data generating process.

General Linear Regression Model (Equation 21.3)

For a quantitative response and one or more predictors, the general form of the linear regression model is

$$\begin{split} (\text{Response})_i &= \beta_0 + \beta_1 (\text{Predictor 1})_i + \beta_2 (\text{Predictor 2})_i + \dots + \beta_p (\text{Predictor } p)_i + \varepsilon_i \\ &= \beta_0 + \sum_{i=1}^p \beta_j (\text{Predictor j})_i + \varepsilon_i \end{split}$$

where β_j for $j=0,1,2,\ldots,p$ are the p+1 parameters governing the model for the data generating process.

ANOVA Model (Equation 27.2)

For a quantitative response and a single categorical predictor (also known as a factor) with k levels, the ANOVA model is

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{i-th unit belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

is an indicator variable capturing whether a unit belongs to the j-th group and $\mu_1, \mu_2, \dots, \mu_k$ are the parameters governing the model for the data generating process.

Repeated Measures ANOVA Model (Equation 36.2)

For a quantitative response and a single categorical predictor (also known as a factor) with k levels in the presence of b blocks, the repeated measures ANOVA model is

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{i-th unit belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{i-th unit belongs to block } m \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables capturing whether a unit belongs to the *j*-th group and *m*-th block, respectively; and, $\mu_1, \mu_2, \dots, \mu_k$ and $\beta_2, \beta_3, \dots, \beta_b$ are the parameters governing the model for the data generating process.

This model assumes any differences between groups are similar across all blocks.

C.4 Glossary

The following key terms were defined in the text; each term is presented with a link to where the term was first encountered in the text.

Alternative Hypothesis (Definition 3.10) The statement (or theory) about the parameter capturing what we would like to provide evidence for; this is the opposite of the null hypothesis. This is denoted H_1 or H_a , read "H-one" and "H-A" respectively.

Average (Definition 5.2) Also known as the "mean," this measure of location represents the balance point for the distribution. If x_i represents the *i*-th value of the variable x in the sample, the sample mean is typically denoted by \bar{x} .

For a sample of size n, it is computed by

$$\bar{x} = \frac{1}{n} \sum_{i=1}^{n} x_i.$$

When referencing the average for a population, the mean is also called the "Expected Value," and is often denoted by μ .

- Between Group Variability (Definition 29.1) When comparing a quantitative response across groups, the between group variability is the variability in the average response from one group to another.
- **Bias (Definition 4.1)** A set of measurements is said to be biased if they are *consistently* too high (or too low). Similarly, an estimate of a parameter is said to be biased if it is *consistently* too high (or too low).
- Blocking (Definition 25.5) Blocking is a way of minimizing the variability contributed by an inherent characteristic that results in dependent observations. In some cases, the blocks are the unit of observation which is sampled from a larger population, and multiple observations are taken on each unit. In other cases, the blocks are formed by grouping the units of observations according to an inherent characteristic; in these cases that shared characteristic can be thought of having a value that was sampled from a larger population.

In both cases, the observed blocks can be thought of as a random sample; within each block, we have multiple observations, and the observations from the same block are more similar than observations from different blocks.

Bootstrapping (Definition 6.3) A method of modeling the sampling distribution by repeatedly resampling from the original data.

Categorical Variable (Definition 1.5) Also called a "qualitative variable," a measurement on a subject which denotes a grouping or categorization.

Classical ANOVA Model (Definition 28.1) For a quantitative response and single categorical predictor with k levels, the classical ANOVA model assumes the following data generating process:

$$(\text{Response})_i = \sum_{i=1}^k \mu_j (\text{Group } j)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if i-th observation belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables and where

- 1. The error in the response for one subject is independent of the error in the response for all other subjects.
- 2. The variability in the error of the response within each group is the same across all groups.
- 3. The errors follow a Normal Distribution.

This is the default "ANOVA" analysis implemented in the majority of statistical packages.

Classical Regression Model (Definition 18.3) For a quantitative response and single predictor, the classical regression model assumes the following data generating process:

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \epsilon_i$$

where

- 1. The error in the response has a mean of 0 for all values of the predictor.
- 2. The error in the response for one subject is independent of the error in the response for all other subjects.
- 3. The variability in the error of the response is the same for all values of the predictor.
- 4. The errors follow a Normal Distribution.

This is the default "regression" analysis implemented in the majority of statistical packages.

Classical Repeated Measures ANOVA Model (Definition 37.1) For a quantitative response and single categorical predictor with k levels in the presence of b blocks, the classical repeated measures ANOVA model assumes the following data generating process:

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if i-th observation belongs to group j} \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{if i-th observation belongs to block m} \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables and where

- 1. The error in the response for one subject is independent of the error in the response for all other subjects.
- 2. The variability in the error of the response is the same across all predictors.
- 3. The errors follow a Normal distribution.
- 4. Any differences between the groups are similar across all blocks. This results from the deterministic portion of the model for the data generating process being correctly specified and is equivalent to saying the error in the response, on average, takes a value of 0 for all predictors.

- 5. The effect of a block on the response is independent of the effect of any other block on the response.
- 6. The effect of a block on the response is independent of the error in the response for all subjects.
- 7. The block effects follow a Normal distribution.

This is the default "repeated measures ANOVA" analysis implemented in the majority of statistical packages.

- **Codebook (Definition 1.7)** Also called a "data dictionary," these provide complete information regarding the variables contained within a dataset.
- **Confidence Interval (Definition 6.5)** An interval (range of values) estimate of a parameter that incorporates the variability in the statistic. The process of constructing a k% confidence interval results in these intervals containing the parameter of interest in k% of repeated studies. The value of k is called the *confidence level*.
- **Confounding (Definition 4.6)** When the effect of a variable on the response is mis-represented due to the presence of a third, potentially unobserved, variable known as a confounder.
- **Controlled Experiment (Definition 4.5)** A study in which each subject is *randomly* assigned to one of the groups being compared in the study.
- **Correlation Coefficient (Definition 16.1)** A numerical measure of the *strength* and *direction* of the *linear* relationship between two quantitative variables.

The classical Pearson Correlation Coefficient r is given by the following formula:

$$r = \frac{\sum_{i=1}^{n} \left(x_{i} - \bar{x} \right) \left(y_{i} - \bar{y} \right)}{\sqrt{\sum_{i=1}^{n} \left(x_{i} - \bar{x} \right)^{2} \sum_{i=1}^{n} \left(y_{i} - \bar{y} \right)^{2}}}$$

where \bar{x} and \bar{y} represent the sample means of the predictor and response, respectively.

- **Degrees of Freedom (Definition 19.5)** A measure of the flexibility in a sum of squares term; when a sum of squares is divided by the corresponding degrees of freedom, the result is a variance term.
- **Deterministic Process (Definition 10.1)** A process for which the output is completely determined by the input(s). That is, the output can be determined with certainty.
- **Distribution (Definition 3.3)** The pattern of variability corresponding to a set of values.
- **Distribution of the Population (Definition 5.9)** The pattern of variability in values of a variable at the population level. Generally, this is impossible to know, but we might model it.
- **Distribution of the Sample (Definition 5.6)** The pattern of variability in the observed values of a variable.
- **Error Sum of Squares (Definition 19.3)** The Error Sum of Squares, abbreviated SSE and sometimes referred to as the Residual Sum of Squares, is given by

$$SSE = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Predicted Mean Response})_i \right]^2$$

where the predicted mean response is computed using the least squares estimates.

- **Estimation (Definition 3.7)** Using the sample to approximate the value of a parameter from the underlying population.
- **Extrapolation (Definition 18.1)** Using a model to predict outside of the region for which data is available.
- Factor (Definition 24.1) Also referred to as the "treatment" in some settings, a factor is a categorical predictor. The categories represented by this categorical variable are called "levels."
- **Frequency (Definition 3.4)** The number of observations in a sample falling into a particular group (level) defined by a categorical variable.
- **Hypothesis Testing (Definition 3.8)** Using a sample to determine if the data is consistent with a working theory or if there is evidence to suggest the data is not consistent with the theory.
- **Identically Distributed (Definition 10.4)** A set of random variables is said to be identically distributed if they are from the same population.

Similarly, a set of observations is said to be identically distributed if they share the same data generating process.

Independence (Definition 10.3) Two random variables are said to be independent when the likelihood that one random variable takes on a particular value does not depend on the value of the other random variable.

Similarly, two observations are said to be independent when the likelihood that one observation takes on a particular value does not depend on the value of the other observation.

- Indicator Variable (Definition 21.1) An indicator variable is a binary (takes the value 0 or 1) variable used to represent whether an observation belongs to a specific group defined by a categorical variable.
- **Interaction Term (Definition 21.5)** A variable resulting from taking the product of two predictors in a regression model. The product allows the effect of one predictor to depend on another predictor, essentially modifying the effect.
- Interquartile Range (Definition 5.5) Often abbreviated as IQR, this is the distance between the first and third quartiles. This measure of spread indicates the range over which the middle 50% of the data is spread.
- Law of Large Numbers (Definition 6.1) For our purposes, the Law of Large Numbers essentially says that as a sample size gets infinitely large, a statistic will become arbitrarily close (extremely good approximation) of the parameter it estimates.

Least Squares Estimates (Definition 17.2) Often called the "best fit line," these are the estimates of the parameters in a regression model chosen to minimize the sum of squared errors. Formally, for Equation 17.3, they are the values of β_0 and β_1 which minimize the quantity

$$\sum_{i=1}^{n} \left[(\text{Response})_i - \beta_0 - \beta_1 (\text{Predictor})_i \right]^2.$$

The resulting estimates are often denoted by $\hat{\beta}_0$ and $\hat{\beta}_1$.

Least Squares Estimates for General Linear Model (Definition 21.3) The least squares estimates for a general linear model (Equation 21.3) are the values of $\beta_0, \beta_1, \beta_2, \dots, \beta_p$ which minimize the quantity

$$\sum_{i=1}^{n} \left[(\text{Response})_i - \beta_0 - \sum_{j=1}^{p} \beta_j (\text{Predictor } j)_i \right]^2.$$

Mean Square (Definition 19.6) A mean square is the ratio of a sum of squares and its corresponding degrees of freedom. For a model of the form in Equation 17.3, we have

- Mean Square Total (MST): estimated variance of the responses; this is the same as the sample variance of the response.
- Mean Square for Regression (MSR): estimated variance of the predicted responses.
- Mean Square Error (MSE): estimated variance of the error terms; this is equivalent to the estimated variance of the response for a given value of the predictor (the variance of the response about the regression line).

In each case, the mean square is an estimated variance.

Mean Square (in ANOVA) (Definition 29.3) A mean square is the ratio of a sum of squares and its corresponding degrees of freedom. For a model of the form in Equation 27.2, we have

- Mean Square Total (MST): estimated variance of the responses; this is the same as the sample variance of the response.
- Mean Square for Regression (MSR): estimated variance of the sample mean responses from each group; this is also called the Mean Square for Treatment (MSTrt) in ANOVA.
- Mean Square Error (MSE): estimated variance of the error terms; this is equivalent to the estimated variance of the response within a group.

In each case, the mean square is an estimated variance. These are equivalent to the MST, MSR, and MSE in the regression model (Definition 19.6).

- Multivariable (Definition 15.1) This term refers to questions of interest which involve more than a single variable. Often, these questions involve many variables. Multivariable models typically refer to a model with two or more predictors.
- **Normal Distribution (Definition 18.2)** Also called the Gaussian Distribution, this probability model is popular for modeling noise within a data generating process. It has the following characteristics:
 - It is bell-shaped.
 - It is symmetric, meaning the mean is directly at its center, and the lower half of the distribution looks like a mirror image of the upper half of the distribution.
 - Often useful for modeling noise due to natural phenomena or sums of measurements.

The functional form of the Normal distribution is

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

where μ is the mean of the distribution and σ^2 is the variance of the distribution.

- **Null Distribution (Definition 7.1)** The sampling distribution of a statistic *when* the null hypothesis is true.
- **Null Hypothesis (Definition 3.9)** The statement (or theory) about the parameter that we would like to *disprove*. This is denoted H_0 , read "H-naught" or "H-zero".
- **Null Value (Definition 3.11)** The value associated with the equality component of the null hypothesis; it forms the threshold or boundary between the hypotheses. Note: not all questions of interest require a null value be specified.
- **Numeric Variable (Definition 1.6)** Also called a "quantitative variable," a measurement on a subject which takes on a numeric value *and* for which ordinary arithmetic makes sense.
- **Observational Study (Definition 4.4)** A study in which each subject "self-selects" into one of groups being compared in the study. The phrase "self-selects" is used very loosely here and can include studies for which the groups are defined by an inherent characteristic or are chosen haphazardly.
- **Outlier (Definition 5.7)** An individual observation which is so extreme, relative to the rest of the observations in the sample, that it does not appear to conform to the same distribution.
- **P-Value (Definition 7.2)** The probability, assuming the null hypothesis is true, that we would observe a statistic, from sampling variability alone, as extreme or more so as that observed in our sample. The p-value quantifies the strength of evidence against the null hypothesis, with smaller values indicating stronger evidence.
- **Parameter (Definition 3.6)** Numeric quantity which summarizes the distribution of a variable within the *population* of interest. Generally denoted by Greek letters in statistical formulas.

Percentile (Definition 5.1) The k-th percentile is the value q such that k% of the values in the distribution are less than or equal to q. For example,

- 25% of values in a distribution are less than or equal to the 25-th percentile (known as the "first quartile" and denoted Q_1).
- 50% of values in a distribution are less than or equal to the 50-th percentile (known as the "median").
- 75% of values in a distribution are less than or equal to the 75-th percentile (known as the "third quartile" and denoted Q_3).

Population (Definition 1.1) The collection of subjects we would like to say something about. Power (Definition 25.3) In statistics, power refers to the probability that a study will discern a signal when one really exists in the data generating process. More technically, it is the probability a study will provide evidence against the null hypothesis when the null hypothesis is false.

Probability Plot (Definition 20.3) Also called a "Quantile-Quantile Plot", a probability plot is a graphic for comparing the distribution of an observed sample with a theoretical probability model for the distribution of the underlying population. The quantiles observed in the sample are plotted against those expected under the theoretical model.

R-Squared (Definition 19.4) Sometimes reported as a percentage, the R-Squared value measures the proportion of the variability in the response explained by a model. It is given by

$$R-squared = \frac{SSR}{SST}.$$

Randomization (Definition 25.2) Randomization can refer to random selection or random allocation. Random selection refers to the use of a random mechanism (e.g., a simple random sample, Definition 4.2, or a stratified random sample, Definition 4.3) to select units from the population. Random selection minimizes bias.

Random allocation refers to the use of a random mechanism when assigning units to a specific treatment group in a controlled experiment (Definition 4.5). Random allocation eliminates confounding and permits causal interpretations.

Randomized Complete Block Design (Definition 34.1) A randomized complete block design is an example of a controlled experiment utilizing blocking. Each treatment is randomized to observations within blocks such that within each block every treatment is present and the same number of observations are assigned to each treatment.

Reduction of Noise (Definition 25.4) Reducing extraneous sources of variability can be accomplished by fixing extraneous variables or blocking (Definition 25.5). These actions reduce the number of differences between the units under study.

Reference Group (Definition 21.2) The group defined by setting all indicator variables in a model for the data generating process equal to 0.

- **Regression (Definition 17.1)** Used broadly, this refers to the process of fitting a statistical model for the data generating process to observed data. More specifically, it is a process of estimating the parameters in a data generating process using observed data.
- Regression Sum of Squares (Definition 19.2) The Regression Sum of Squares, abbreviated SSR, is given by

$$SSR = \sum_{i=1}^{n} \left[(\text{Predicted Mean Response})_i - (\text{Overall Mean Response}) \right]^2$$

where the predicted mean response is computed using the least squares estimates and the overall mean response is the sample mean.

- **Relative Frequency (Definition 3.5)** Also called the "proportion," the fraction of observations falling into a particular group (level) of a categorical variable.
- **Replication (Definition 25.1)** Replication results from taking measurements on different units (or subjects), for which you expect the results to be similar. That is, any variability across the units is due to natural variability within the population.
- **Residual (Definition 20.1)** The difference between the observed response and the predicted response (estimated deterministic portion of the model). Specifically, the residual for the i-th observation is given by

$$(Residual)_i = (Response)_i - (Predicted Mean Response)_i$$

where the "predicted mean response" is often called the predicted, or fitted, value.

Residuals mimic the noise in the data generating process.

- **Response (Definition 3.2)** The primary variable of interest within a study. This is the variable you would either like to explain or estimate.
- **Sample (Definition 1.2)** The collection of subjects for which we actually obtain measurements (data).
- **Sampling Distribution (Definition 6.2)** The distribution of a *statistic* across repeated samples (of the same size) from the population.
- Simple Random Sample (Definition 4.2) Often abbreviated SRS, this is a sample of size n such that *every* collection of size n is equally likely to be the resulting sample. This is equivalent to a lottery.
- **Standard Deviation (Definition 5.4)** A measure of spread, this is the square root of the variance.
- **Standard Error (Definition 6.4)** The standard error is the estimated standard deviation of a *statistic*; that is, it is the standard deviation from a *model* for the sampling distribution of a statistic. It quantifies the variability in the statistic across repeated samples.

Standardized (Test) Statistic (Definition 12.1) Also, known as a test statistic, a standardized statistic is a ratio of the signal in the sample to the noise in the sample. The larger the standardized statistic, the stronger the evidence of a signal; said another way, the larger the standardized statistic, the stronger the evidence against the null hypothesis.

Standardized Statistic for ANOVA (Definition 29.4) Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 27.2):

$$(\text{Response})_i = \sum_{j=1}^k \mu_j(\text{Group } j)_i + \varepsilon_i,$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{i-th unit belongs to group } j \\ 0 & \text{otherwise} \end{cases}$$

is an indicator variable. Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0. A standardized statistic, sometimes called the "standardized F statistic," for testing the hypotheses is given by

$$T^* = \frac{\left(SSE_0 - SSE_1\right)/(k-r)}{SSE_1/(n-k)},$$

where k is the number of parameters in the full unconstrained model and r is the number of parameters in the reduced model. Defining

$$MSA = \frac{SSE_0 - SSE_1}{k - r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model. Just as before, the MSE represents the residual variance — the variance in the response for a particular set of the predictors.

Standardized Statistic for General Linear Model (Definition 21.4) Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 21.3):

$$(\text{Response})_i = \beta_0 + \sum_{j=1}^p \beta_j (\text{Predictor } j)_i + \varepsilon_i.$$

Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0. A standardized statistic, sometimes called the "nested F statistic," for testing the hypotheses is given by

$$T^* = \frac{\left(SSE_0 - SSE_1\right)/(p+1-r)}{SSE_1/(n-p-1)},$$

where p+1 is the number of parameters in the full unconstrained model (including the intercept) and r is the number of parameters in the reduced model. Defining

$$MSA = \frac{SSE_0 - SSE_1}{p + 1 - r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model. Just as before, the MSE represents the residual variance — the variance in the response for a particular set of the predictors.

Standardized Statistic for Repeated Measures ANOVA (Definition 38.1) Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 36.2):

$$(\text{Response})_i = \sum_{i=1}^k \mu_j(\text{Group } j)_i + \sum_{m=2}^b \beta_m(\text{Block } m)_i + \varepsilon_i$$

where

$$(\text{Group } j)_i = \begin{cases} 1 & \text{if i-th observation corresponds to group } j \\ 0 & \text{otherwise} \end{cases}$$

$$(\text{Block } m)_i = \begin{cases} 1 & \text{if i-th observation corresponds to block } m \\ 0 & \text{otherwise} \end{cases}$$

are indicator variables. Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0. A standardized statistic, sometimes called the "standardized F statistic," for testing the hypotheses is given by

$$T^* = \frac{\left(SSE_0 - SSE_1\right)/(k+b-1-r)}{SSE_1/(n-k-b+1)},$$

where k + b - 1 is the number of parameters in the full unconstrained model and r is the number of parameters in the reduced model. Defining

$$MSA = \frac{SSE_0 - SSE_1}{k + b - 1 - r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model. Just as before, the MSE represents the residual variance — the variance in the response for a particular set of predictors.

Standardized Statistic for Simple Linear Regression (Definition 19.7) Consider testing a set of hypotheses for a model of the data generating process of the form (Equation 17.3):

$$(Response)_i = \beta_0 + \beta_1 (Predictor)_i + \varepsilon_i.$$

Denote this model as Model 1, and denote the model that results from applying the parameter constraints defined under the null hypothesis as Model 0^1 . A standardized statistic, sometimes called the "standardized F statistic," for testing the hypotheses is given by

$$T^* = \frac{\left(SSE_0 - SSE_1\right)/(2-r)}{SSE_1/(n-2)},$$

where r is the number of parameters in the reduced model. Defining

¹Technically, this standardized statistic only applies to a class of potential hypotheses, but that class is quite broad (see Appendix B).

$$MSA = \frac{SSE_0 - SSE_1}{2 - r}$$

to be the "mean square for additional terms," which captures the shift in the error sum of squares from the reduced model to the full unconstrained model, we can write the standardized statistic as

$$T^* = \frac{MSA}{MSE}$$

where the mean square error in the denominator comes from the full unconstrained model.

Statistic (Definition 5.8) Numeric quantity which summarizes the distribution of a variable within a *sample*.

Statistical Inference (Definition 1.3) The process of using a sample to characterize some aspect of the underlying population.

Stochastic Process (Definition 10.2) A process for which the output cannot be predicted with certainty.

Stratified Random Sample (Definition 4.3) A sample in which the population is first divided into groups, or strata, based on a characteristic of interest; a simple random sample is then taken within each group.

Time-Series Plot (Definition 20.2) A time-series plot of a variable is a line plot with the variable on the y-axis and time on the x-axis.

Total Sum of Squares (Definition 19.1) The Total Sum of Squares, abbreviated SST, is given by

$$SST = \sum_{i=1}^{n} \left[(\text{Response})_i - (\text{Overall Mean Response}) \right]^2$$

where the overall average response is the sample mean.

Variability (Definition 3.1) The notion that measurements differ from one observation to another.

Variable (Definition 1.4) A measurement, or category, describing some aspect of the subject. Variance (Definition 5.3) A measure of spread, this roughly captures the average distance values in the distribution are from the mean.

For a sample of size n, it is computed by

$$s^2 = \frac{1}{n-1} \sum_{i=1}^{n} \left(x_i - \bar{x} \right)^2$$

where \bar{x} is the sample mean and x_i is the *i*-th value in the sample. The division by n-1 instead of n removes bias in the statistic.

The symbol σ^2 is often used to denote the variance in the population.

Within Group Variability (Definition 29.2) When comparing a quantitative response across groups, the within group variability is the variability in the response within each group.